

The Holonet Machine

A Single Finite Geometry as Processor, Network, and Memory

Machine-Verified Computer Architecture on the $W(3,3)$ Substrate

Wil Dahn | github.com/wilcompute/W33-Theory

June 2026

Abstract

We present a complete, machine-verified computer architecture in which the processor, the network, and the memory are not coupled designs but *one finite object*: the symplectic generalized quadrangle $W(3,3) = \text{GQ}(3,3) = \text{SRG}(40, 12, 2, 4)$. On this 40-point geometry, routing a packet is applying a gate is addressing memory — an algebraic identity, not an analogy — and a single automorphism group $W(E_6)$ of order 51840 is simultaneously the processor’s Clifford gate group, the network’s symmetry, and the memory’s code automorphisms. We make each architectural layer a theorem with an executable witness: an enumerated instruction set ($\text{Sp}(4,3)$ of order 51840 plus one cubic gate, universal); a diameter-2, bisection-100, better-than-Ramanujan interconnect that is provably non-planar (hence photonic); a distance-3 qutrit code with a measured $P_L \sim Ap^2$ threshold curve; an I/O boundary saturating the Holevo qutrit capacity $\log_2 3$; a complexity placement (Clifford = P, +cubic = BQP) whose resource is contextuality, with the contextual fraction $1/10$ *derived* from the no-ovoid geometry; and a leaderless consensus layer tolerating five Byzantine faults. The classical layer is polynomial-time, so the architecture runs as software — we ship a runnable virtual machine that routes, corrects errors, teleports state, and self-reproduces, plus an installable `holonet` command and a test suite. The quantum advantage is the one priced resource (classical emulation cost 9^t for t cubic gates), executed here for small t . Every constant is arithmetic in the single integer $q = 3$.




Figure 1: One object, four readings. On $W(3,3)$ the processor, network, memory, and clock/I-O are projections of a single 40-point incidence geometry with a single symmetry group $W(E_6)$.

1 The Substrate and the One Group

$W(3,3)$ is the symplectic generalized quadrangle over $\mathbb{F}_3^4$ with alternating form $\langle u, v \rangle = u_1v_2 - u_2v_1 + u_3v_4 - u_4v_3$. Its collinearity graph is strongly regular, $\text{SRG}(40, 12, 2, 4)$: 40 points of degree 12, $\lambda = 2$, $\mu = 4$; 40 totally-isotropic lines, four points each, four through each point. Every number below is arithmetic in nine integers ($q = 3$, $\lambda = 2$, $\mu = 4$, $k = 12$, $f = 24$, $g = 15$, $\Phi_6 = 7$, $\Phi_4 = 10$, $\Phi_3 = 13$), forced by $q! = 2q$ (unique solution $q = 3$).

Result 1 (One group). *The symplectic group $\text{Sp}(4,3)$, built by closure from the 40 point-transvections, has order exactly $51840 = |W(E_6)| = 3^4(3^2 - 1)(3^4 - 1)$. It acts transitively on the 40 nodes, the 40 line-contexts, and the 160 point-on-line flags, preserving collinearity. Hence one group is at once the processor’s Clifford gate group, the network’s automorphisms, the memory’s code automorphisms, and the readout’s contextuality symmetry (Fig. 1).*

2 Processor: an Enumerated Universal ISA

The core computes in balanced ternary — the most economical radix ($E(b) = b/\ln b$ is minimised at e ; base 3 at 2.731 beats binary and base 4 at 2.885) — with native word $27 = 3^3$. The instruction set is enumerated, not abstract: the degree-2 Clifford layer is the finite group $\text{Sp}(4,3)$ (51840 opcodes = $|W(E_6)|$; the per-qutrit Clifford is $\text{Sp}(2,3)$ of order $24 = f$), classically simulable by Gottesman–Knill; and a single degree-3 macro-op, the E_6 cubic form on the **27**, lifts it to universal (Lloyd–Braunstein). The power source is contextuality.

3 Interconnect and Floorplan

Result 2 (The fabric is diameter 2, maximally fault-tolerant, and non-planar). *$\text{GQ}(3,3)$ has diameter 2 (any-to-any in two hops), vertex/edge connectivity 12 (survives 11 crash faults), second eigenvalue $\lambda_2 = 2$ far below the Ramanujan bound $2\sqrt{11} = 6.63$ (a better-than-Ramanujan expander), and minimum bisection exactly $100 = (n/4)(k - \lambda_2)$, met by an explicit $20|20$ cut. Routing is table-free: two nodes are adjacent iff $\langle x, y \rangle = 0$, one ALU operation — the address is the route. By Thompson’s VLSI theorem a 2-D layout needs area $\gtrsim (\text{bisection}/2)^2 \sim 2500$, so the fabric is intrinsically non-planar and belongs in a photonic (3-D / fiber / wavelength) medium.*

The 240 links partition into 40 four-point cliques (lines), and the graph is 1-factorable into a 12-slot conflict-free link schedule; vertex/edge-transitivity gives perfect load balancing.

4 Memory and the Measured Threshold

The memory is the $[[66, 8, 3]]_3$ qutrit code (distance 3, $n - k = 58$ syndrome qutrits). We verify the mechanism on the runnable $[[5, 1, 3]]_3$ stand-in (same distance-3 correction): all 40 single-qutrit errors carry distinct syndromes and are corrected to fidelity 1. Monte-Carlo over a depolarizing channel gives a *measured* threshold curve $P_L(p) \sim Ap^2$ with $A \approx 8$, crossing break-even $P_L = p$ at a pseudo-threshold $p_{th} = 1/A \approx 0.12$. The substrate’s larger code has $A \sim \binom{66}{2} = 2145$, hence $p_{th} \sim 5 \times 10^{-4}$ — the same p^2 mechanism at a different size. Adding syndrome-measurement noise at the same per-round rate removes the single-round break-even; *repeating* the measurement ($R = 3$ rounds, majority vote) restores it (Fig. 2, left), the standard fault-tolerance gadget.




Figure 2: The machine’s scorecard. *Left*: the measured logical error $P_L(p)$ of the runnable $[[5, 1, 3]]_3$ code vs. the physical depolarizing rate, with the break-even diagonal $P_L = p$; the perfect-syndrome curve bends below it ($\sim Ap^2$, $A \approx 8$), single-round noisy syndromes have no threshold, and three repeated measurements restore it. *Right*: the statistical significance of the contextuality witness ($\chi = 10 > \alpha = 7$) vs. detected photons, crossing 5σ at a few hundred.

5 Energy and I/O

The degree-2 Clifford datapath is unitary, hence reversible and Landauer-free (Bennett); base 3 imposes no per-bit penalty ($kT \ln 3 / \log_2 3 = kT \ln 2$). The only irreducible cost is syndrome export: $(n - k)kT \ln 3 \approx 2.6 \times 10^{-19}$ J/cycle at 300 K — the machine’s metabolism. The I/O boundary is the line geometry: the 40 lines are measurement contexts; the air-gap carrier is a photon’s OAM in $\{\ell = -1, 0, +1\}$ (the balanced-ternary digit), saturating the Holevo capacity $\log_2 3 = 1.585$ bit/photon; the logical port is $k = 8$ qutrits (12.68 bit/block).

6 Complexity and the Contextual Resource

Result 3 (P to BQP, with a derived contextual fraction). *All 12 single-qutrit stabilizer states have non-negative Gross–Wigner function, so the Clifford layer is efficiently classically simulable (P). The degree-3 resource (the Strange state) has Wigner negativity, mana $\ln(5/3)$, and robustness $R = 3$; by Howard et al. this contextuality is the necessary resource, lifting the machine to BQP. The contextual fraction is derived from the geometry: a global Kochen–Specker value-assignment satisfying all 40 contexts would be an ovoid, and $W(q)$ has none for q odd (the maximum partial ovoid is $7 < 10$); by exact integer programming the most contexts any classical assignment can satisfy is 36, so $\text{CF} = (40 - 36)/40 = 1/10 = 1/\Phi_4$. Equivalently, the Cabello–Severini–Winter witness $\chi = \sum_v P(v)$ obeys $\chi \leq \alpha = 7$ (noncontextual) versus $\chi = 10$ (quantum, state-independent), a gap tolerating 30% noise; a simulated heralded-photon run clears it at 5σ with ~ 840 photons. The classical emulation cost of t cubic gates is 9^t : the Pashayan–Wallman–Bartlett quasiprobability estimator (each t -gate sample in $[-3^t, 3^t]$) is unbiased but needs $\sim 9^t$ samples.*

7 The Machine, Executed

Because the Clifford layer is polynomial-time, the architecture runs as software. A single file (`analysis/holonet_node.py`) builds the fabric, routes a packet (address is route, $\mu = 4$ multipath), runs a qutrit register as a stabilizer tableau over $\mathbb{F}_3$, and self-reproduces by splicing. Three programs execute the von Neumann components: it *corrects* errors (the $[[5, 1, 3]]_3$ cycle, fidelity 1), *teleports* state across the fabric (the two classical trits routed in two hops), and *reproduces* (a verified quine fixed point). A leaderless consensus run contracts disagreement by 1/3 per round and outvotes five Byzantine nodes (boundary exactly at $t = 5$, the Dolev bound $2t + 1 \leq 12$). The whole stack ships as the installable `holonet` command (`route`, `teleport`, `correct`, `reproduce`, `verify`, `info`) with a 13-test suite, all passing.

8 Falsifiable Predictions

Claim	Refuted if
P1 radix-12, diameter-2 wiring	a build needs a third hop / radix $\neq 12$
P2 OAM-trit saturates Holevo 1.585 bit/photon	a 3-mode channel misses/beats it
P3 $[[66, 8, 3]]_3$ pseudo-threshold $\sim 5 \times 10^{-4}$	encoding fails to help below it
P4 magic robustness 3 / mana $\ln(5/3)$	a lower-1-norm stabilizer decomposition
P5 5-Byzantine / 11-crash tolerance	6 Byzantine survived, or 5 not
P6 contextual fraction 1/10 ($\chi = 10 > 7$)	the benchtop witness $\hat{\chi} \leq 7$

Theorems (not falsifiable, only verifiable): $|\text{Sp}(4, 3)| = 51840 = |W(E_6)|$; diameter 2, connectivity 12, bisection 100; distance 3; Clifford = P / +cubic = BQP; 1-factorability; $\alpha = 7$, CF = 1/10.

Honest scope

The group orders, the bisection, the contextual fraction, $\alpha = 7$, the robustness $R = 3$, and the threshold coefficient A are computed or measured in the public witnesses. The quantum values (BQP, $\chi = 10$, Holevo saturation) assume the Witting orthonormal realisation in $\mathbb{C}^4$ and the standard quantum-computing theorems (Gottesman–Knill, Lloyd–Braunstein, Howard et al., Pashayan–Wallman–Bartlett). The $[[5, 1, 3]]_3$ code is a runnable stand-in for $[[66, 8, 3]]_3$ (same mechanism, different size). The QEC, teleportation, and threshold are exact / Monte-Carlo simulations (single-round; a full circuit-level treatment is future work). No physical photonic build exists yet; the quantum advantage is the priced 9^t dial, run here only for small t . “Architecture of life” means the von Neumann architecture of self-reproduction (the rigorous computational notion), not biology.

9 Conclusion

The Holonet is a complete finite computer architecture in which one 40-point geometry, one group, and one integer $q = 3$ generate the processor, the network, and the memory together; every layer is a theorem with a runnable witness, the device specification is its own audit, and the whole classical machine runs on any computer today. Its companion papers develop the physics (`photonic_holonet.tex`) and the infrastructural implications (`holonet_practical_implications.tex`); its quickstart is `HOLONET.md`. The first physical milestone — the benchtop contextuality test certifying the 1/10 fraction — is decided by a few hundred photons.

Witnesses: `analysis/w33_isa_encoding.py`, `w33_noc_floorplan.py`, `w33_one_group_machine.py`,
`w33_complexity_advantage.py`, `w33_contextual_fraction.py`, `w33_ks_inequality.py`, `w33_magic_dial.py`,
`holonet_node.py`, `holonet_qec_demo.py`, `holonet_teleport_demo.py`, `holonet_quine.py`, `holonet_`
`quantum_packet.py`, `holonet_consensus_demo.py`, `holonet_threshold_demo.py`, `holonet_ks_experiment.py`,
`w33_architecture_capstone.py`. **Standard results:** Gottesman–Knill; Lloyd–Braunstein; Howard–Wallman–
Veitch–Emerson (2014); Pashayan–Wallman–Bartlett (2015); Cabello–Severini–Winter; Dolev; Lamport–Melliar–
Smith; Thompson (VLSI); Bennett; Landauer.