

The Witting Reference Fabric

An Ultimate Hardware/Software Computer Architecture:
Content-Addressed Memory, Projection-Native Execution,
Network-Fabric Consensus, and Smart-Asset Economics at 70M TPS over
10⁹ Nodes

Wil Dahn

Witting Reference Fabric Project

In collaboration with the UOR Foundation, Oko, Smart Assets, DTR/HLIX

June 2026 *v2 architecture edition*

Abstract

Artificial superintelligence (ASI) is usually framed as a model-scale problem: more parameters, more accelerators, more data centres. This paper frames it as a substrate problem. If HLIX turns infrastructure into a projection–execution–emission cycle, UOR turns objects into universal content-derived references, Oko supplies network-fabric Byzantine finality, and Smart Assets attach economics to workloads, then the missing piece is an architecture that makes one billion heterogeneous nodes behave like one referenceable compute object.

Assume the supervisor’s target: 70 M aggregate transactions per second over a 10⁹-node mesh and a 70B-parameter ASI workload as the benchmark class. The proposed **Witting Reference Fabric** interprets that mesh as a layered hardware/software fabric: UOR supplies typed content addresses; HLIX supplies projection-native scheduling; Oko supplies hard-final network consensus; Smart Assets supply programmable economics; and the W(3,3) / Witting finite geometry supplies a compact coordination lattice with 40 address classes, 240 carrier links, 480 directed routes, 51,840 finite control symmetries, and an 11-branch non-backtracking transport rule. The architecture also admits a stronger memory primitive: a *referenceable flow cell*, where the stored object is not a resting byte pattern but the canonical invariant of a bounded repeatable trace.

The architecture’s natural information density is $\boxed{3/8 = q/2^q}$, appearing exactly as the $R_{96}/256$ Atlas compression ratio and the W(3,3) chiral eigenspace fraction 15/40. The native CSS storage layer sits just below it at $81/240 = 27/80$, spending the 3/80 gap as protection overhead. The natural cadence is $j(i) = 1728 = k^3$ orbit-cycles per consensus epoch, so 70 M TPS induces 1.21×10^{11} substrate cycles/sec. A $40^6 = 4.096$ billion-node fractal deployment exceeds the 10⁹-node target with redundancy. The central implication is architectural: ASI-scale compute becomes referenceable, auditable, sovereign, self-coordinating, and economically addressable at planetary scale.

Bottom line

The Dahn thesis: *ASI becomes governable only when compute, memory, identity, provenance, consensus, and economics share the same reference substrate.* In this paper the substrate is Atlas-12288 / UOR; the operating cycle is HLIX/HCP projection–execution–emission; the coordination fabric is Oko’s BIER/BLA network consensus; the workload unit is the Smart Asset; and the compact control geometry is $\text{Aut}(W(3,3)) = \text{Sp}(4, \mathbb{F}_3) = W(E_6)$.

Scope

This is an architecture and strategy paper, not a claim that ternary Witting silicon or Bell-qutrit storage already exists as production hardware. The near-term path is to use the Witting finite geometry as coordination and verification mathematics for HLIX/UOR/Okó software systems, then progressively harden the same invariants into accelerators, storage, and fabric. The flow-cell experiments cited below are finite software harnesses over $W(3,3)$, not measured device noise results.

Contents

1	Introduction: From Opaque ASI to Referenceable Compute	4
1.1	Why ASI makes architecture a governance problem	4
1.2	Four convergent stacks	4
1.3	Implementation grounding from the UOR code	5
1.4	Our finding	6
1.5	What this paper claims	6
2	Executive Answer: Why This Architecture Matters for ASI	7
2.1	The imported W33 control ledger	8
3	The Seven Design Axioms	9
4	The Atlas-12288 Memory Frame: The Universal Unit	10
4.1	What UOR actually built	10
4.2	The substrate reading	11
4.3	What an address looks like	11
5	Hardware: The Witting Substrate	12
5.1	The Witting tile: a 40-vertex compute lattice	12
5.2	Bose–Mesner three-channel bus	13
5.3	Spectral cache hierarchy (no eviction algorithm)	13
5.4	CSS code as native ECC	14
5.5	Referenceable flow cells	14
5.6	The Bell-qutrit storage cell	15
5.7	BIER multicast fabric (the bus, replaced)	16
5.8	Hardware acceleration units	16
5.9	Fractal scaling: 40^n Witting lattices	16
6	The Group-Action Instruction Set	17
6.1	$\text{Aut}(W(3,3)) = W(E_6)$ as the ISA	17
6.2	Microarchitecture: 30 families, 8 generators	17
6.3	Programs as Cayley walks	18
6.4	Two-qutrit Pauli phase space (the register file)	18

7	Software: The Projection Operating System	18
7.1	UOR address space	18
7.2	HCP container store	19
7.3	The Projection Engine (the OS kernel)	19
7.4	Trace objects	19
7.5	Smart Asset workloads	20
7.6	BLA-coordinated concurrency	20
8	Performance Analysis	20
8.1	Throughput identities at planetary scale	20
8.2	Density and storage budget	21
8.3	No bus bottleneck $\Rightarrow$ linear strong scaling	21
8.4	Substrate spectral cache hit-rates	21
8.5	Flow-cell budgets	21
9	Security, Identity, and Provenance	22
9.1	Stabilizer-subgraph identity	22
9.2	Provenance via receipt chains	22
9.3	Post-quantum safety	22
10	Software Examples	22
10.1	Hello world (smart-asset form)	22
10.2	Distributed sort (group-action form)	23
10.3	Neural network inference	23
10.4	Cross-jurisdiction ACID	23
11	Use Cases	24
11.1	Defense (FreedomNet / Golden Dome)	24
11.2	Finance: CBDC, RTGS, tokenised assets	24
11.3	AI workloads	24
11.4	Telecommunications and IoT	24
11.5	Edge, device, browser	24
12	Roadmap	25
13	Conclusion	25
A	Glossary	26
B	Quantitative Validation	27
B.1	Substrate-primitive readings of UOR Atlas-12288 constants	27
B.2	Throughput at 70 M TPS	27

1 Introduction: From Opaque ASI to Referenceable Compute

Plain-English summary

The question is not merely whether ASI can be made smarter. The question is whether an ASI-grade system can be made *referenceable*: can we name every piece of compute, verify every output, preserve every receipt, route work across one billion nodes, and keep sovereign control over data without centralizing the whole system? The answer proposed here is yes, if the HLIX/UOR/Oko/Smart-Assets stack is underpinned by a finite control geometry that can route, audit, and compress without becoming the subject of the paper.

1.1 Why ASI makes architecture a governance problem

Computer architecture has not changed structurally since John von Neumann sketched the stored-program model in 1945. Pipelining, caches, out-of-order execution, SIMD, and GPGPU acceleration are all improvements on the same scaffold: a central processor fetches instructions and data from separate memory over a shared path and executes them according to local state. For ordinary software this is an efficiency problem. For ASI it becomes a governance problem: the same bottleneck that hides memory locality also hides provenance, coordination, jurisdiction, and accountability.

In 2026 we sit at a triple convergence of stress on this scaffold:

1. **Monolithic die scaling has ended.** Chiplet packaging admits architectural defeat: we can no longer make a single chip bigger.
2. **Memory bandwidth lags compute by 100×**. CPU FLOPs grow at roughly 1.4× per year; DRAM bandwidth grows at roughly 1.07× per year. We are paying for compute we cannot feed.
3. **Distributed computing is now the workload, not a special case.** Cross-datacentre, cross-jurisdiction, cross-vendor coordination is the norm, but the underlying architectures still pretend the bus model holds.
4. **AI turns logs into liability.** An ASI workflow is not a single inference call; it is a chain of retrieval, planning, tool use, delegation, payment, and post-hoc audit. If those events are not first-class addressable objects, the system cannot be trusted.

The retrofits accumulate: virtual memory hides the physical memory; cache hierarchies hide the memory bus; transactional memory hides shared-memory races; distributed-systems protocols hide the network behind RPC. Every retrofit costs latency, complexity, energy, and trust. The industry calls this “the memory wall” or “the dark-silicon problem.” What it actually is, structurally, is the cost of fifty years of layering retrofits on top of a 1945 architecture.

1.2 Four convergent stacks

Four independent recent stacks each chip away at one piece of the ASI substrate problem:

- **Okó** (the network-fabric Byzantine Lattice Agreement consensus protocol) eliminates the leader-based quadratic bottleneck of classical BFT. Instead of n nodes exchanging $O(n^2)$ messages, Okó uses BIER multicast at the network layer: virtual switches replicate packets to a 55-validator committee with $O(n)$ aggregate communication. The simulator measures 12M TPS per server at fixed committee size, with linear scaling in per-server compute; the supplied source pack projects billion-plus aggregate TPS by clustering servers behind the same

committee invariant [1, 2]. Throughput is now a per-server parameter, not a validator-count parameter.

- **The UOR Foundation** (Universal Object Reference) defines a content-addressed substrate: one content-derived address per object, persistent identity, trustless verification, and universal composability. The current public site presents UOR as a unified computational substrate for agentic AI with 14 namespaces, 82 classes, and 120 properties [3]; the current UOR-Framework repository reports version 0.5.2 with 34 namespaces, 474 classes, 948 properties, and 3601 named individuals [4]. The crown jewel for this paper is the Atlas-12288 frame: 48 pages of 256 bytes, with an R_{96} resonance classification compressing the byte alphabet onto 96 classes.
- **HLIX** (Hybrid Linked Infrastructure Xchange) operationalises the HCP *projection cycle*—every system event is a *projection* (a declarative description of what should happen) that triggers an *execution* on a substrate node, emitting a new content-addressed result plus a verifiable receipt. HLIx targets 1 billion nodes across the “HexaScaler” axes of Cloud, Edge, Mesh, AI, Chain, and Sovereign [8].
- **Smart Assets** turn workloads into economic objects: CPU/GPU/NPU time, models, datasets, receipts, services, and governance shares become fundable, billable, tradeable compute assets. In the source image, the attributes are explicitly modular, decentralised, composable, interoperable, programmable, transferable, and incentive aligned. In ASI terms, Smart Assets are how autonomous work pays for itself without ceding coordination to a central owner.

1.3 Implementation grounding from the UOR code

This paper treats UOR as an implementation surface, not only as a concept. The repository review matters because the architecture should be buildable in stages:

Repository	Architectural role imported here
UOR-Framework	Rust ontology workspace for content-addressed, symmetric, multi-metric object spaces based on $Z/(2^n)Z$, with generated Rust and Lean surfaces plus conformance gates [4].
uor-addr	Typed reference vocabulary producing deterministic <code>sha256:<64hex></code> labels from canonical bytes: JSON via RFC 8785 JCS, S-expressions, XML-C14N, DER ASN.1, code-module ASTs, ring elements, GGUF, ONNX, and schema-pinned descendants. The core is <code>no_std</code> ; several realizations are <code>no_alloc</code> [5].
prism	Standard-library facade for UOR axes, including cryptographic hash axes SHA-2/3, BLAKE3, Keccak, signature axes, and commitment axes [6].
atlas-12288	C/Lean hybrid FFI exposing the constants this paper uses directly: <code>UOR_PAGES()</code> =48, <code>UOR_BYTES()</code> =256, and <code>UOR_RCLASSES()</code> =96, with Lean-backed and pure-C modes [7].

1.4 Our finding

The convergence

Fast consensus, universal addressing, sovereign orchestration, and workload economics all point toward the same missing layer. Oko says coordination must live in the network fabric. UOR says identity must be content-derived. HLIX says infrastructure must be projected and receipted. Smart Assets say compute must be economically addressable. The Witting Reference Fabric supplies the finite hardware/software control architecture that lets these four claims compose.

The architectural fixed point is the **Witting polytope**—the unique 40-vertex configuration on which the symplectic polar space $W(3, 3)$ is canonically realised—with automorphism group $\text{Sp}(4, \mathbb{F}_3) \cong W(E_6)$ of order 51,840. The Witting polytope was discovered by Hermann Witting in 1937 as a uniform polychoron in $\mathbb{C}^4$. This paper treats it as the organising candidate for a substrate-coherent ASI fabric: small enough to compile, symmetric enough to coordinate, and rigid enough to audit.

1.5 What this paper claims

We do *not* claim that the Witting Reference Fabric is a marginal improvement over Von Neumann. We claim that it is a different basin of attraction in the architecture design space, in which the following retrofits are replaced by substrate-native primitives:

- No bus—communication is intrinsic to the topology, not a separate channel.
- No cache hierarchy—the substrate’s spectrum supplies the cache levels exactly.
- No ISA accretion—the instruction set is finite by construction (exactly 30 families).
- No bolted-on ECC goal—the base CSS carrier and protected lift are designed into the storage path.
- No separate networking—on-chip and off-chip use the same fabric protocol.
- No unanchored cryptography—addressing is cryptographically pinned, while signatures and post-quantum policy remain explicit.
- No separate consensus—BLA is the thread-coordination primitive.
- No root user—identity is structural (the stabiliser subgroup).
- No hidden ASI state—agent memory and tool use become receipt-bearing objects.

The paper unfolds in six movements. §2 gives the executive answer for ASI. §3 states the seven design axioms. §4 explains the Atlas-12288 memory frame and the universal-density reading. §5 describes the hardware path (Witting tile, three-channel bus, spectral cache, Bell-qutrit storage, BIER fabric). §6 explains the group-action ISA. §7 describes the projection operating system. The remaining sections quantify performance, address security and trust, and lay out the deployment roadmap.

2 Executive Answer: Why This Architecture Matters for ASI

The highest-level implication

The Witting Reference Fabric does not merely make ASI faster. It makes ASI *legible at substrate scale*. The W(3,3) / Witting layer supplies a finite coordinate system in which every agent action can be named, routed, receipted, compared, composed, and paid for without collapsing the network into a central scheduler.

ASI requirement	Architecture contribution
Referenceable cognition	Every model, prompt, memory, tool call, policy, dataset, and receipt becomes a UOR object with a stable content-derived address. The ASI's cognition is not hidden in a vendor log; it is a traversable object graph.
Verifiable inference	HLIX emissions attach receipts to outputs; Oko finalises them; the Witting control plane gives the receipts an invariant coordinate system. The question changes from “do we trust this answer?” to “which orbit produced this answer, from which inputs, under which policy?”
Billion-node coordination	Oko keeps consensus committee size fixed while per-server work scales; HLIIX schedules projections across heterogeneous nodes; W(3,3) supplies a finite symmetry vocabulary for coordinating work without a single global scheduler.
Sovereignty by construction	Raw data can remain local while proofs move globally. Jurisdictional control is handled by policy containers and receipt chains, not by manually negotiated integrations.
Persistent ASI memory	The HCP Container Store becomes long-horizon memory: immutable resources, references, views, datasets, and receipts. The architecture adds the compression and orbit structure needed to keep that memory coherent rather than merely large.
Economic agency	Smart Assets attach funding, billing, ownership, and rewards to workloads. ASI agents can commission compute, prove delivery, and settle usage without a trusted platform operator.
Failure containment	Stabiliser-based identity and orbit histories make anomalies detectable: forged outputs, policy drift, prompt-injection residues, and contradictory provenance become graph defects rather than subjective review findings.

The ASI translation. UOR answers “what is this object?” HLIIX answers “where can this projection execute?” Oko answers “when is this state final?” Smart Assets answer “who pays and who is rewarded?” The Witting Reference Fabric answers “which invariant hardware/software structure makes all of those answers compose without losing meaning?”




Figure 1: Smart Asset attributes from the provided source image [9]. In this architecture a Smart Asset is the economic envelope around a referenceable workload: modular, composable, incentive-aligned, transferable, and executable across the HLIx/OkO/UOR mesh.

2.1 The imported W33 control ledger

The paper imports the local W33/Witting work as a control ledger, not as a physics thesis. The useful facts are finite, testable architecture constants:

Constant	Architecture meaning
40	Reference classes / tile cores / Sylow-3 handles.
240	Undirected carrier links; base edge-code carrier.
480	Directed forward/return carrier slots.
51,840	Finite control symmetries $\text{Aut}(W(3, 3))$; small enough for lookup-table scheduling.
30	Conjugacy classes; instruction-family count and control vocabulary.
$1 + 24 + 15$	Bose–Mesner channel split: broadcast, working set, persistent/chiral state.
$39 + 120 + 81 = 240$	QEC/photonic ledger partition preserving the 81-state logical tail.
$240 \cdot 7^3 = 82,320$	Protected Steane/ Φ_6 lift used as the active protection target beyond the base $d_Z = 4$ carrier.
$p_{\text{th}} = 11$	Non-backtracking fabric branch used for routing and multicast intuition.

The ledger is deliberately operational: it tells a compiler, scheduler, storage controller, or network switch what finite object it is using. It does not ask the reader to accept a Theory-of-Everything claim before evaluating the architecture.

3 The Seven Design Axioms

Plain-English summary

A good architecture has a small number of *design principles*—rules every part of the system obeys. Below are seven. Each replaces a convention from 1945 with something the math says is better. Read them quickly first; subsequent sections will derive each one in detail.

A1: Content-Address Everything

Every datum—memory cell, register, instruction, message, process, file—has a single canonical address derived from its content by the ψ -pipeline: a deterministic canonicalisation followed by Blake3 (or sha256) hashing. No location-based addresses exist anywhere in the system.

A2: Compute Is Group Action

The only operation available to a process is the action of an element of $\text{Aut}(W(3,3)) = W(E_6)$ on an edge-mode configuration $\mathcal{C} \subseteq E(W(3,3))$. There is no separate ALU, control unit, branch unit, or vector unit. Arithmetic, control flow, and parallelism are specialisations of the same group action.

A3: Spectral Cache Replaces Heuristic Cache

Cache hierarchy is the spectral decomposition of the $W(3,3)$ adjacency operator $A = kE_0 + rE_1 + sE_2$, with three projectors E_0, E_1, E_2 of ranks $\{1, 24, 15\}$ acting as three native cache levels. Cache replacement is substrate-forced, not heuristic.

A4: Storage Is Trace-Locked

Static objects are content-addressed by UOR CIDs. Active memory cells are trace-addressed: the physical state keeps moving, while the stored datum is $\text{CID}_{\text{flow}} = H(\text{canon}(x_0, x_1, \dots, x_T))$, the canonical invariant of a bounded repeatable trace. A temporal Bell qutrit is the far-horizon physical storage cell; a W33 non-backtracking flow cell is the software-first primitive.

A5: BIER Multicast Is the Fabric

There is no bus. Every core has $k = 12$ direct neighbours plus $q^q = 27$ disjoint reachable cores. All communication is BIER multicast through embedded virtual switches; fabric bandwidth scales linearly with core count.

A6: CSS Code Is Native ECC

The CSS code $[[|E|, q^{q+1}, \mu, q]]_3 = [[240, 81, 4, 3]]_3$ is hardwired into the memory frame. Logical-to-physical ratio $27/80 = q/2^q = 3/8$.

A7: Identity Is a Stabilizer Subgraph

A process, user, or container has no UID, no PID, no root capability. Its identity is the subgroup $\text{Stab}_{\text{Aut}(W_{(3,3)})(\mathcal{C})}$ of automorphisms fixing its edge-mode configuration. Provenance is the orbit history; trust is the orbit signature.

How the axioms connect. A1 (content addressing) makes A2 (group-action compute) work because the group acts on *content*, not locations. A2 makes A3 (spectral cache) work because the spectral decomposition is a property of the group action. A3 makes A4 (Bell-qutrit storage) practical because the spectrum exactly matches the single-qutrit state space. A4 makes A6 (CSS code) practical because persistent state has a canonical trace witness that can be syndromically repaired. A5 (BIER multicast) makes A2 work at scale because group action must be broadcast. A7 (stabiliser identity) makes A1 trustworthy because identity is substrate-coherent. Every axiom supports every other; none is layered on.

4 The Atlas-12288 Memory Frame: The Universal Unit

Plain-English summary

Memory in our architecture is not byte-addressed. It is organised into *frames* of exactly 12288 bytes. Each frame is a $48\text{-page} \times 256\text{-byte}$ table. Each 256-byte page is then compressed by a fixed mapping into one of 96 equivalence classes (the R_{96} classification). These numbers—12288, 48, 256, 96—are not arbitrary; the UOR Foundation chose them on entirely different grounds, and they *all* land on the same five-prime substrate alphabet, and the compression ratio $96/256 = 3/8$ is *also* the rate of our error-correcting code and *also* a natural eigenvalue fraction of the underlying graph. We call this the Universal Density Theorem.

4.1 What UOR actually built

UOR’s atlas-12288 reference implementation (C with Lean 4 verification) treats every addressable resource as a frame of 12288 bytes structured as:

- 48 *pages*, each 256 bytes long (so $48 \times 256 = 12288$),
- an R_{96} *classifier* mapping each page’s 256 possible byte values to one of 96 resonance classes,
- *conservation-budget checks* ($\alpha_4\alpha_5 = 1$, from which the 48 arises as a unity constraint),
- canonical encode/decode (sha256-addressable).

The team described the result as “3/8 compression” ($96/256 = 0.375$). They were matter-of-fact about it. They did not (in the public materials) state why 48, 256, or 96; they were simply the numbers that fell out of the unity constraint and the desired information density.

4.2 The substrate reading

Set $q = 3$ (ternary), $\mu = 4$ (spacetime), $F_5 = 5$ (smallest post- q prime), $\Phi_3 = 13$, $\Phi_6 = 7$ (cyclotomic values at q). Then:

$$\begin{aligned}
 48 &= q! \cdot 2^q && \text{(pages per frame)} \\
 256 &= \mu^4 = 2^{\Phi_6+1} && \text{(bytes per page; the dS identity)} \\
 96 &= 2^{F_5} \cdot q && (R_{96} \text{ resonance classes}) \\
 12288 &= q! \cdot 2^{q+\Phi_6+1} && \text{(frame size)} \\
 \frac{96}{256} &= \frac{q}{2^q} = \frac{3}{8} && \text{(natural density)}
 \end{aligned}$$

All four Atlas constants land on substrate primitives, with no constant in the architecture drawing on any prime outside the five-element substrate alphabet. The compression ratio $3/8$ is independently the fraction $g/v = 15/40$ of the chiral eigenspace of the $W(3,3)$ adjacency operator. The base CSS carrier $[[240, 81, d_Z = 4]]_3$ sits nearby at $81/240 = 27/80$, using the $3/80$ gap between $27/80$ and $3/8$ as storage-protection overhead. *The fabric density is exact; the protected-storage layer deliberately spends a small part of it.*

The Universal Density Theorem

The natural information density of a $\text{Sp}(4, \mathbb{F}_3)$ -equivariant content-addressed computer is $q/2^q = 3/8$. This density is realised exactly by:

- (i) the UOR Atlas resonance compression ratio $R_{96}/2^8$,
- (ii) the chiral-eigenspace fraction g/v of the $W(3,3)$ adjacency spectrum.

The native CSS storage layer then runs at $27/80$, exactly $3/80$ below the fabric density, which is the architecture's protection budget rather than a failed match.

Why $3/8$ feels right

Think of $q/2^q$ as “three things in a space of eight.” A ternary alphabet (three symbols) embeds into a binary alphabet (eight 3-bit slots) with three slots holding real information and five slots holding the redundancy that lets you correct errors. That ratio appears wherever the substrate makes a choice between density and resilience.

4.3 What an address looks like

A compact 64-bit WRF reference-class handle is sized by

$$\underbrace{2^{64}}_{\text{handle space}} = \underbrace{40}_{\text{Sylow-3 vertex}} \cdot \underbrace{1,296}_{|N_G(P_3)| \text{ normaliser}} \cdot \underbrace{2^{48.34}}_{\text{contingent payload, approx.}}.$$

The first factor (the 40 Witting vertices) corresponds to the 40 Sylow-3 subgroups of $\text{Aut}(W(3,3))$ by Theorem 1; the second (the 1,296-element normaliser) is fixed at the vertex; the residual ≈ 48.34 bits encode the frame's contingent payload. This is a compact execution handle; canonical UOR identity remains the content-derived CID.

Theorem 1 (Hidden Sylow Address Bijection). $n_3(\text{Sp}(4, \mathbb{F}_3)) = v = 40$. *The 40 Witting vertices are in canonical bijection with the 40 Sylow-3 subgroups of $\text{Aut}(W(3,3))$. Every WRF object handle selects a unique Sylow-3 coset of $\text{Aut}(W(3,3))$.*

The structural consequence is that “where” an object lives in the address space is algebraically constrained by “what” it is: there is no separate “placement policy.” The substrate places it.

5 Hardware: The Witting Substrate

Plain-English summary

The physical chip is built from *tiles*. Each tile holds 40 small computing cores arranged in a precise mathematical shape called the Witting polytope. Inside each tile the cores talk to each other via three specialised buses (not one). Between tiles, chips, or even continents, the same fabric protocol carries traffic. The chips can be stacked fractally—40 tiles become a chip, 40 chips become a node, 40 nodes become a rack, all the way up. The same architecture serves a smartwatch and a planet.

5.1 The Witting tile: a 40-vertex compute lattice




Figure 2: The Witting tile. Inner ring (12 nodes): k -eigenspace neighbours. Outer ring (27 nodes): disjoint matter shell. Centre: self vertex. Together, $1 + 12 + 27 = 40 = v$. The graph is the unique strongly regular graph $\text{SRG}(40, 12, 2, 4)$, with $|E| = 240$ edges and automorphism group $\text{Sp}(4, \mathbb{F}_3)$ of order 51,840.

A Witting tile holds 40 trit-native cores wired in the collinearity-graph topology of $W(3, 3)$ (Figure 2). Each core is a small ALU operating on *trits* (ternary digits in $\{0, 1, 2\}$) rather than bits. The structure of the tile is the *Bell-cloud closure*

$$40 = 1 + 12 + 27 = 1 + k + q^q,$$

which we read as “one self vertex, $k = 12$ adjacent vertices (the gauge layer), $q^q = 27$ disjoint vertices (the matter layer).” The tile has $|E| = 240$ on-chip links and is governed by an automorphism group of order 51,840 acting transitively on its vertices.


Why ternary at the gate level. Information density per symbol is $\log_2 3 = 1.585$ bits. A trit holds 58.5% more information than a bit. CMOS implementations of ternary logic gates have been

demonstrated since 1970s research; modern multi-Vt FinFET processes make it practical again at 5 nm and below. The native datatype is the qutrit Pauli register $(a_1, b_1, a_2, b_2) \in \mathbb{F}_3^4$, which is exactly the two-qutrit Clifford phase space of $W(3, 3)$ (§6). Floating-point and integer arithmetic are emulated as specialised projections; the native arithmetic is qutrit Clifford.

5.2 Bose–Mesner three-channel bus

The Bose–Mesner algebra $\mathcal{B} = \langle E_0, E_1, E_2 \rangle$ of $W(3, 3)$ has three orthogonal projectors of ranks $\{1, 24, 15\}$ summing to $v = 40$. Any signal $\xi \in \mathbb{C}^{40}$ on the tile decomposes as $\xi = E_0\xi + E_1\xi + E_2\xi$, and the adjacency operator acts as $A\xi = kE_0\xi + rE_1\xi + sE_2\xi$ with $k = 12, r = 2, s = -4$.

Bose–Mesner three-channel bus



bus width $\propto$ projector rank

Figure 3: The on-tile bus is three physical channels with widths proportional to the Bose–Mesner projector ranks $\{1, 24, 15\}$. Each channel routes traffic in one eigenspace; together they cover every signal exactly once.

The WRF tile has three physical buses corresponding to the three projectors, with bandwidth fractions $(1, 24, 15)/40 = (2.5\%, 60\%, 37.5\%)$:

- B_0 (**constant, eigenvalue** k): 2.5% bandwidth. Carries broadcast control, barrier sync, and global metadata. Topologically the constant function on the graph; physically a single trace reaching every core.
- B_1 (**gauge, eigenvalue** $r = 2$): 60% bandwidth. Carries the working set—hot data, intermediate computational results, register-spill traffic. Topologically the 24-dimensional eigenspace; physically routed through the $k = 12$ direct neighbour links with $\lambda = 2$ shared neighbours per edge.
- B_2 (**chiral, eigenvalue** $s = -4$): 37.5% bandwidth. Carries CSS-encoded persistent state—logical qutrits, error-corrected blocks, durable storage traffic. Topologically the 15-dimensional eigenspace; physically routed through the non-adjacent shell $\mu = 4$ shared neighbours.

Bus B_2 's bandwidth fraction $g/v = 37.5\%$ matches the universal fabric density. The CSS storage layer runs at $27/80 = 33.75\%$, leaving a $3/80$ protection gap inside the chiral channel.

5.3 Spectral cache hierarchy (no eviction algorithm)

Cache without policy

A traditional cache constantly guesses which data you will need next. Sometimes it guesses wrong; that costs cycles. A Witting cache does not guess. The spectral decomposition of the graph determines where each datum lives, and it lives there deterministically. There is no LRU, no MRU, no clock algorithm—no choice to be made.

The cache hierarchy is the spectral projection of the SRG: every memory operation projects its target onto the three eigenspaces, and the cache level used is the eigenspace it lands in. There are three levels:

- **L1 (constant projector E_0):** 1-dimensional; holds barrier and control state. Hit rate $1/v = 2.5\%$.
- **L2 (gauge projector E_1):** 24-dimensional; holds the working set. Hit rate $f/v = 60.0\%$.
- **L3 (chiral projector E_2):** 15-dimensional; holds CSS-encoded persistent state. Hit rate $g/v = 37.5\%$.

There is no eviction algorithm because projector class is not an allocation choice: the substrate's spectrum determines each item's cache class. Projector coverage is 100% by construction because the projectors span $\mathbb{C}^{40}$; a physical system still has finite capacity, prefetch, and remote-frame latency when the working set exceeds local storage.

5.4 CSS code as native ECC

Why memory does not need a separate ECC chip

Modern servers add a separate ECC layer to detect and correct memory errors. The WRF storage path targets CSS encoding as a native carrier; reading is syndromically decoded inside the storage controller rather than delegated to a bolt-on subsystem.

The base finite code $[[240, 81, d_Z = 4]]_3$ encodes 81 logical qutrits in 240 physical qutrits with verified Z -distance marker $\mu = 4$. In the ideal block model it detects the corresponding low-weight logical/syndrome events; physical thresholds require a device noise model and calibration. The parameters are all substrate primitives:

$$n = |E| = 240, \quad k = q^{q+1} = 81, \quad d_Z = \mu = 4, \quad \text{base} = q = 3.$$

We hardwire it into the memory frame:

- Each 80-byte logical block (LB) consumes 240 trits-equivalent of physical storage. Density: $27/80 = 33.75\%$, exactly $3/80$ below the universal fabric density.
- The protected runtime path from the local W33 work uses a Steane/ Φ_6 lift to length $240 \cdot 7^3 = 82,320$ with an architecture target of distance ≥ 81 .

There is no separate ECC philosophy: the base carrier, protected lift, decoder, and storage controller are one design surface. Device-level error rates remain empirical.

5.5 Referenceable flow cells

Plain-English summary

Traditional memory pretends information is still: bits sit at an address until the CPU asks for them. WRF can make the opposite choice. The substrate may keep moving, while the information is the stable pattern traced by that motion. Reading means recovering the canonical pattern; writing means steering the local flow until it locks to the desired pattern.

Definition 2 (Referenceable flow cell). *A referenceable flow cell is a bounded dynamical register*

$$x_{t+1} = F_\theta(x_t, u_t, s_t), \quad P = \text{canon}(x_0, x_1, \dots, x_T), \quad \text{CID}_{\text{flow}} = H(P),$$

where x_t is the moving local state, θ is the tile-local routing rule, u_t is input, s_t is syndrome/receipt feedback, P is a canonical trace invariant, and CID_{flow} is the UOR-addressable name of the pattern. The datum is P , not any one instantaneous x_t .

The W33 non-backtracking carrier gives a finite software model of this cell. It has $480 = 2|E|$ directed states and exactly $p_{\text{th}} = 11$ legal continuations from each state. The BT110–BT112 WRF harnesses [10, 11] test the first engineering questions without claiming device physics:

Question	Finite harness result
Write locking	Across four primary seed rules, every directed state enters an attractor within at most 37 deterministic steps. A seed-661 six-symbol register has mean symbol write latencies from 5.5 to 6.2 steps.
Forward repair	Under the cell’s own transition rule, forward perturbations remain in the same attractor basin in the tested harness. This is trajectory stability, not a physical noise model.
Register isolation	A 3×3 independent-flow lattice reported 0 cross-talk events in 24,000 one-step trials.
Capacity headroom	The 500-seed survey found 1,138 distinct 24-hex-character flow CIDs and three six-attractor seed rules. A sampled character-distance check found minimum distance 18, corresponding to a $t = 9$ symbolic correction budget in that sample.
Spectral audit	The adjacency trace tower verifies exact substrate identities through A^8 , including $\text{tr}(A^8) = \text{tr}(A^6)q(4k - 1) + 480 \cdot 16(q(4k - 1) - \lambda) = 430,970,880$.

The hardware interpretation is simple: WRF storage is not limited to passive cells. A cell can be a photonic loop, RF ring, memristive reservoir, packet loop, or software flow whose address is the canonical trace it sustains. UOR names the trace, HLIX schedules the projection that maintains or transforms it, Oko finalises the transition, Smart Assets meter the sustained work, and the CSS/Ouroboros return channel repairs deviations.

The unresolved work is also clear: a production flow cell still needs a calibrated device noise model, a physical write-injection circuit, a coupling algebra for interacting cells, and capacity accounting under real read windows. The paper therefore uses referenceable flow cells as the architecture’s next primitive, not as a completed silicon claim.

5.6 The Bell-qutrit storage cell

In the far-horizon physical design, each WRF storage cell is a temporal Bell qutrit:

$$|\Omega\rangle = q^{-1/2} (|0\rangle_p |0\rangle_f + |1\rangle_p |1\rangle_f + |2\rangle_p |2\rangle_f).$$

The cell holds a maximally entangled state between its *past* and *future* registers, realised as a Franson-style time-bin interferometer (optical implementation) or a delay-line latch (electrical implementation). Four equivalent uniqueness conditions single out $|\Omega\rangle$ from all other states: SWAP-symmetry plus maximal entanglement, Choi–Jamiołkowski state of the identity channel, $(U \otimes U^*)$ -invariance for every $U \in \text{U}(3)$, or uniform Schmidt spectrum.

Read. Choi–Jamiołkowski overlap with the identity channel: $\langle \Omega | (I \otimes U) | \Omega \rangle = \text{Tr}(U)/q$. Reading is averaging the cell’s quantum identity against a probe unitary; lossless and non-destructive.

Write. Project the future register onto the desired state; the past register inherits via SWAP-symmetry. The write operation does not destroy past-state information; it co-localises past and future at the harmonic-convergence point.

Persistence. The Bell-line stabiliser of order $\mu^2 \cdot q^{q+1} = 1296$ preserves the cell against all but 40 of 51,840 automorphisms; the cell’s MTBF is set by the orbit-mixing time, which scales as $|\text{Aut}|/v = 1296$ projection cycles.

5.7 BIER multicast fabric (the bus, replaced)

Why there is no bus

A traditional bus is shared—only one signal at a time travels on it. As you add cores, the bus saturates, and adding cores actually slows the system down. BIER multicast is the opposite: it lets every core broadcast simultaneously to selected destinations, with switches between them replicating packets at line rate. Adding cores adds bandwidth.

Inter-tile (and inter-chip, inter-board, inter-node) communication uses Bit-Indexed Explicit Replication (BIER), with each Witting tile embedding a virtual switch based on the Vectorized Packet Processor (VPP) model from Oko. The fabric properties:

- Per-tile bandwidth scales with per-tile compute; the fabric is not a shared resource.
- Routing is content-addressed: a packet’s BIER header indexes by destination CID.
- No TCP-style retransmission; packet loss is absorbed by aggregate-signature threshold (the protocol counts votes, not packets).
- Hop-by-hop IPsec authentication via tile public keys, which are the identifiers of the 40 Sylow-3 subgroups—these are *cryptographic-strength names* forced by the substrate.

On-chip and off-chip communication are identical at the protocol level. The traditional NUMA / PCIe / NVLink hierarchy collapses to a single fabric.

5.8 Hardware acceleration units

Every Witting tile bundles four specialised units:

- **Blake3 hash unit:** hardware canonical hashing at line rate (≥ 100 Gbps) for content addressing.
- **MuSig aggregator:** aggregate signature acceleration for tile quorum certificates. Demo: 110 K+ sigs/core/s on 4-core Broadwell. Target: ~ 1 M/core/s.
- **Bitonic sort unit:** AVX-style sorting network for hash-permutation block ordering; constant-time per network rank.
- **CSS syndrome decoder:** lookup-table CSS decoder integrated into the cache controller; one-cycle correction per logical block.

5.9 Fractal scaling: 40^n Witting lattices

The same architecture repeats at every scale. A tile holds 40 cores; a chip holds 40 tiles; a node holds 40 chips; a rack holds 40 nodes; a datacentre holds 40 racks; a planet holds 40 datacentres. The connectivity at each level is the same $W(3, 3)$ collinearity graph, with the same $|\text{Aut}| = 51,840$

governing it, and the same Bose–Mesner three-channel routing. Software written for one level runs on any other.

Level n	Cores (40^n)	Physical analogue	Throughput target
0	1	one core (qutrit ALU)	30 GOp/s
1	40	one Witting tile (one die)	1.2 TOp/s
2	1,600	one chip	48 TOp/s
3	64,000	one node	1.9 PetaOp/s
4	2,560,000	one rack	76 PetaOp/s
5	102,400,000	one data centre	3.1 ExaOp/s
6	4,096,000,000	planet (4.1 B nodes)	122 ExaOp/s

The 6-level structure overshoots the HLIx 1-billion-node target by $\approx 4\times$, leaving substrate headroom for redundancy. At level 7 ($40^7 = 1.64 \times 10^{11}$ cores), the network exactly matches the substrate-natural conjugacy cadence of 1.21×10^{11} cycles/s at 70 M TPS.

6 The Group-Action Instruction Set

Plain-English summary

A traditional computer has hundreds or thousands of distinct instructions (RISC-V has ~ 50 base; x86 has ~ 3000). Decoders are complex; compilers struggle. Our architecture has *exactly* 30 *instruction families*—a number forced by the math, not chosen by a committee. Programs are walks through a fixed mathematical graph of 51,840 symmetries; the compiler’s job is finding the shortest path.

6.1 $\text{Aut}(W(3, 3)) = W(E_6)$ as the ISA

The WRF ISA is the group $G = \text{Aut}(W(3, 3)) = \text{Sp}(4, \mathbb{F}_3) \cong W(E_6)$, of order $|G| = 51,840$. Every instruction is an element $\sigma \in G$ that acts on an edge-mode configuration $\mathcal{C} \subseteq E(W(3, 3))$, producing the new configuration $\sigma(\mathcal{C})$. By the Triple Convergence theorem $\#\text{conj } G = h(E_8) = Z_{\text{DW}}(\text{Sp}(4, \mathbb{F}_3); T^2) = 30$, so the ISA has *exactly* 30 conjugacy classes—instruction *families* at the microarchitectural level number 30, not 51,840.

6.2 Microarchitecture: 30 families, 8 generators

The Coxeter system of $W(E_6)$ has 6 simple generators; with the Bell-line involutions the full $\text{Aut}(W(3, 3))$ requires 8 Hadamard-type generators in the qutrit Clifford realisation. A microarchitectural decoder needs only:

- 30 conjugacy-class fast paths (single-cycle dispatch),
- 8 generator decoder lanes (each generator is a small fixed circuit),
- Cayley-graph distance lookup for compiler scheduling.

6.3 Programs as Cayley walks

A program is a finite word $\sigma_1\sigma_2\cdots\sigma_n \in G$ in the generators. Equivalently, a program is a walk in the Cayley graph $\text{Cay}(G, S)$ for the chosen generating set S . The compiler's job is graph-theoretic: find the shortest word producing a target automorphism.

Why $|G| = 51,840$ is small

51,840 is a tiny number for a computer. The entire Cayley graph fits in ~ 200 KB of lookup tables. Compiling becomes an $O(1)$ database query: for the target automorphism, look up the shortest generator-word. No NP-hard register allocation, no heuristic instruction scheduling.

6.4 Two-qutrit Pauli phase space (the register file)

The substrate is also the two-qutrit Pauli commutation geometry: every non-zero Pauli monomial $X_3^{a_1} Z_3^{b_1} \otimes X_3^{a_2} Z_3^{b_2}$ corresponds to a unique vertex of $W(3, 3)$, and two monomials commute iff they are adjacent in the SRG. The WRF core's register file is this 2-qutrit Pauli space:

$$\text{Register file} = \text{Span}_{\mathbb{C}}(W(3, 3)\text{-vertices}) = \mathbb{C}^{40}.$$

Local operations decompose into the three Bose–Mesner channels; non-local operations route via BIER multicast.

7 Software: The Projection Operating System

Plain-English summary

The operating system is not a kernel that schedules processes onto a CPU. It is a *Projection Engine*: a piece of code that takes a description of what should happen (a *projection*), evaluates it against immutable content-addressed inputs, and emits new content with a verifiable receipt. There are no processes, threads, files, or sockets in the traditional sense. There are projections, containers, and receipts.

7.1 UOR address space

User and kernel code address every datum by its content-addressed CID, computed by the ψ -pipeline:

$$\begin{aligned} \text{datum} &\xrightarrow{\text{canonicalise}} \text{canonical bytes} \\ &\xrightarrow{\text{sha256 default; pluggable hash axes}} \text{CID}. \end{aligned}$$

The concrete canonicalisers include RFC 8785 JCS, Cap'n Proto schemas, and the other uor-addr realizations; the hash axis can use BLAKE3, SHA3, Keccak, or SHA-512 when policy requires it. There is no notion of location; an address *is* the content. Two cores holding the same datum compute the same CID and therefore share the same address. There are no raw pointers in the logical model; compact 64-bit handles are execution references inside the fabric, while canonical identity remains the content-derived CID.

7.2 HCP container store

The Hologram Container Platform organises addressable storage as a Container Store of immutable resources of six kinds:

Kind	Role
Component	code module / WASM blob / OCI image
Service	exposed RPC with capability schema
Policy	access / data residency / rate constraints
View	materialised projection over a dataset
Dataset	immutable bulk data + reference graph
Receipt	cryptographic emission with provenance chain

Every container is identified by its CID; every reference is content-addressed; every read is verifiable. The container store *is* the filesystem, the registry, the namespace, the package manager, the code repository, the data lake, and the audit log, unified into one object.

7.3 The Projection Engine (the OS kernel)

The operating system kernel is the Projection Engine. Every system call is a projection:

Projection(input CIDs, policy CIDs) $\rightarrow$ Execution $\rightarrow$ Emission(output CID + receipt CID).

A projection definition is itself a container (Component or Policy) and is addressed by its CID. The kernel evaluates projections by scheduling them onto Witting tiles, gathering input from the cache hierarchy, executing via the group-action ISA, and emitting the result as a new container plus receipt.

No processes, just projections

A traditional process is a piece of code plus a private memory plus a scheduler slot. A Witting projection is just a description of what should happen; the substrate finds the tiles, gathers the inputs, runs the group action, and produces the output. There is no “process table.” Concurrency is automatic: two projections that don’t share inputs run in parallel by default.

7.4 Trace objects

The Projection Engine treats a flow cell as a first-class object. A trace object records the carrier, local rule seed, canonical trace witness, current receipt chain, and repair policy:

$$\text{TraceObject} = (\text{carrier CID}, \theta, \text{CID}_{\text{flow}}, \text{receipt CID}, \text{syndrome policy}).$$

This is the software bridge from ordinary UOR objects to active WRF memory. A static file is addressed by its canonical bytes; a live register, model shard, or agent-memory loop is addressed by the canonical invariant of the trace it emits. The operating system does not need a separate distinction between memory, process, and log: all three are reference objects with different trace policies.

7.5 Smart Asset workloads

User programs are *smart assets*—self-funding, self-billing, tradeable workloads with embedded economics. At substrate scale, a Smart Asset is the economic envelope around a content-addressed workload: it binds funding, entitlement, execution policy, delivered compute, and receipt history to one referenceable object. The WRF maps that object onto Witting orbits only when finite routing or protection is needed. The economic primitives:

- **Proof-of-Compute Receipts:** cryptographic evidence of delivered work, bound to the workload’s CID and the executing tile’s identity.
- **Tokenised compute:** CPU / GPU / NPU time as tradeable assets; market-priced allocation.
- **Microtransactions:** hard-final settlement at Oko cadence. The provided Oko source pack demonstrates a fixed 500 ms simulation cadence; WRF hardware acceleration targets a 30 ms tactical epoch.
- **Crowd-funded workloads:** container deployments backed by aggregate contributor signatures and proof-of-compute receipts.

7.6 BLA-coordinated concurrency

When projections need to coordinate across tiles, they use Oko’s Byzantine Lattice Agreement at the substrate level. Coordination is:

1. BIER-multicast the projection to a 55-tile committee.
2. Each tile evaluates the projection on its local view.
3. Tiles exchange BLA messages (monotone-union of input sets).
4. Aggregate MuSig signature certifies the joint emission.

This single primitive replaces locks, semaphores, mutexes, condition variables, and atomic operations. There is no separate concurrency-control library; the substrate is the lock manager.

8 Performance Analysis

Plain-English summary

We now compute concretely: how fast is the Witting Reference Fabric at the targeted 70 M TPS across a billion-node network? The headline numbers below are derived from the substrate primitives plus the throughput target.

8.1 Throughput identities at planetary scale

At $\text{TPS} = 7 \times 10^7$ transactions per second aggregate over the network:

$$\text{Conjugacy cadence} = \text{TPS} \cdot \frac{|G|}{h(E_8)} = 70\text{M} \cdot 1728 = 1.21 \times 10^{11} \text{ cycles/s}, \quad (1)$$

$$\text{Logical-qutrit ops} = \text{TPS} \cdot \frac{q^{q+1}}{|E|} = 70\text{M} \cdot \frac{27}{80} = 2.36 \times 10^7 \text{ LQ-ops/s}, \quad (2)$$

$$\text{Atlas frame rate} = \text{TPS} \cdot \frac{|E|}{|\text{Atlas}_{12288}|} = 70\text{M} \cdot \frac{240}{12288} = 1.37 \times 10^6 \text{ frames/s}, \quad (3)$$

$$\text{Coherence-block rate} = \text{TPS}/\tau(O) = 70\text{M}/384 = 1.82 \times 10^5 \text{ blocks/s}, \quad (4)$$

$$\text{ECC syndromes/s} = \text{TPS} \cdot |E|/\mu = 70\text{M} \cdot 60 = 4.20 \times 10^9 \text{ syndromes/s}. \quad (5)$$

The conjugacy multiplier $1728 = k^3 = j(i)$ is the modular j -invariant at $\tau = i$ —a deep connection to the arithmetic of elliptic curves which we do not exploit further here but note as a substrate-level cross-link.

8.2 Density and storage budget

Quantity	Value
Binary symbol density	1.000 bits/symbol
Ternary symbol density	1.585 bits/symbol
$\Rightarrow$ ternary advantage	+58.5% logic per symbol
R_{96} vs 256-byte alphabet	$96/256 = 37.5\%$ semantic density
Chiral channel share	$15/40 = 37.5\%$ fabric bandwidth
CSS storage rate	$81/240 = 27/80 = 33.75\%$ logical density
Protection gap	$3/8 - 27/80 = 3/80$
Raw 3/8 payload per Atlas frame	4608 bytes/frame
CSS-budgeted payload per Atlas frame	4147.2 bytes/frame

8.3 No bus bottleneck $\Rightarrow$ linear strong scaling

A standard CMP architecture with n cores on a shared bus has bus contention scaling as $\Theta(n^2)$ in worst case. The Witting tile fabric has BIER multicast with branching factor $p_{\text{th}} = k - 1 = 11$ at each virtual switch, giving total fabric complexity $O(n \log n)$. Strong scaling efficiency comparisons:

- At 10^3 cores: traditional $\sim 30\%$ efficiency vs Witting $\sim 99\%$.
- At 10^6 cores: traditional $\sim 1\%$ efficiency vs Witting $\sim 95\%$.
- At 10^9 cores: traditional unusable vs Witting $\sim 80\%$ (limited by light-speed, not by protocol).

8.4 Substrate spectral cache hit-rates

- L1 (constant, E_0): $1/v = 2.5\%$ probability, single trace.
- L2 (gauge, E_1): $f/v = 60.0\%$ probability, 24-way associative.
- L3 (chiral, E_2): $g/v = 37.5\%$ probability, 15-way associative.

Aggregate projector coverage is **complete by construction**; physical miss behavior is a capacity and placement question for the implementation, not an undefined cache-policy guess.

8.5 Flow-cell budgets

The new flow-cell harnesses add a second performance surface beside throughput. A WRF implementation must budget not only transactions per second, but trace-lock time, symbolic distance, and inter-cell isolation:

$$T_{\text{lock}} \leq 37 \text{ steps in the four-seed harness,} \quad d_{\text{CID}}^{\text{sample}} = 18, \quad \text{crossTalk}_{3 \times 3} = 0/24,000.$$

These are not universal limits; they are software-harness targets that make the next hardware experiments falsifiable. The important point is architectural: a WRF register can be specified by *how quickly it locks, how far its trace identifier is from neighbouring identifiers, and how cleanly it composes in a lattice*, rather than by a static bit-cell layout alone.

9 Security, Identity, and Provenance

Plain-English summary

There is no “root user” in WRF. Identity is a piece of group theory: a process is the set of symmetries that fix its state. Permissions are subgroup memberships. Auditing is graph traversal of the receipt chain. Forging an identity requires solving cryptographic problems (hash collision/preimage attack, signature forgery) that are computationally infeasible.

9.1 Stabilizer-subgraph identity

A WRF process has no UID, no PID, no superuser bit. Its identity is the subgroup $\text{Stab}_G(\mathcal{C}) \leq G$ stabilising its edge-mode configuration. Two processes are the same iff their stabilisers coincide. Permissions are coset memberships, not ACLs. There is no root user; there are only configurations whose stabilisers are large enough to operate on broad orbits.

9.2 Provenance via receipt chains

Every emission carries a receipt CID. The receipt contains:

- input CIDs,
- projection definition CID,
- executing tile signature (P_3 -coset of the Sylow-3 identity),
- committee aggregate signature (MuSig over 55 tile attestations),
- time-anchor (epoch number + Merkle root of preceding epoch).

The full provenance of any datum is the acyclic receipt graph reachable from its CID. Auditing is graph traversal; falsification is computationally infeasible.

9.3 Post-quantum safety

Hashing is not post-quantum magic. Grover-style search effectively halves preimage security, and long-horizon provenance should budget larger digests or multi-axis roots. The architecture therefore keeps the hash axis pluggable: sha256 is the UOR-ADDR default, while BLAKE3, SHA3-256, Keccak-256, and SHA-512 are Prism/uor-addr axes. Aggregate signatures can swap to lattice-based PQC without changing the addressing, projection, or receipt model. CSS encoding protects storage faults; it is not a cryptographic primitive.

10 Software Examples

10.1 Hello world (smart-asset form)

```
struct HelloProjection {
    union {
        rpc @0: HelloRPC;
    }
    struct HelloRPC {
        contract @0: UInt64;           # caller CID
        method @1: UInt16 = HELLO;
        params @2: AnyStruct;         # contains "name" string
    }
}
```

```

    receipts @3: List(Receipt);
  }
}

struct HelloSmartAsset {
  projection      @0: HelloProjection;
  funder          @1: SID;           # who pays
  reward_share    @2: List(Share);   # who gets paid
  policy          @3: Policy = COMPUTE_2_TILE_MINIMUM;
}

```

10.2 Distributed sort (group-action form)

A bitonic sorting network on $n = 2^k$ inputs has $\Theta(\log^2 n)$ depth. On WRF:

- Each pair-wise comparison is an element of G (the swap transposition in the appropriate coset).
- The substrate's bitonic sort hardware pipelines this in constant time per network rank.
- Multi-tile sorts use BIER multicast for inter-tile compare-and-exchange.
- Aggregate emission carries a receipt proving the sorted output is a valid permutation of inputs.

A 10^9 -element distributed sort completes in $\Theta(\log^2 10^9) \approx 900$ network ranks. At the WRF 30 ms tactical target this is ≈ 27 s wall-clock; at the provided Oke 500 ms simulation cadence it is ≈ 7.5 minutes.

10.3 Neural network inference

A transformer layer is a sequence of matrix multiplications and non-linearities. On WRF:

- Weight matrices are stored as CSS-encoded Bell-qudit cells; the substrate auto-corrects single-cell faults during inference.
- Matrix multiplication uses the Bose–Mesner B_1 gauge bus for dense weight broadcasting.
- Non-linearities (GeLU, softmax) are projection definitions evaluated locally.
- Multi-tile attention uses B_2 chiral bus for KV cache transport, B_1 for QK matmuls, and B_0 for layernorm.

A 70B-parameter model remains in scope. At 4-bit quantisation it is roughly 35 GB of weights: $35 \times 10^9 / 4608 \approx 7.60$ million raw Atlas-density frames, or $35 \times 10^9 / 4147.2 \approx 8.44$ million CSS-budgeted frames. At 8-bit weights it is roughly 16.9 million CSS-budgeted frames. On a 10^9 -node mesh, that is far less than one frame per node; inference latency is dominated by projection and consensus depth, not by the existence of the parameter store.

10.4 Cross-jurisdiction ACID

```

projection CrossBorderTransfer {
  preconditions:
    - originator-balance >= amount      # jurisdiction A
    - beneficiary-status == "approved"   # jurisdiction B
    - compliance-clear    == true        # jurisdiction C
  emission:
    - originator-balance -= amount       # commit in A
    - beneficiary-balance += amount      # commit in B
    - audit-trail        += receipt      # commit in C
  receipts:
    - aggregate signature from 55 committee
}

```

```

- epoch anchor at substrate consensus tick
}

```

The transaction either commits atomically in all three jurisdictions or rolls back; no bridge is required, no MEV is possible, no front-running occurs (VRF-derived deterministic ordering inside the block).

11 Use Cases

11.1 Defense (FreedomNet / Golden Dome)

The Three Amigos defense architecture maps directly onto WRF: Hologram (UOR) for hardware-agnostic execution, HLIx for multi-domain orchestration, Oko for self-healing self-routing command & control. The substrate's $f \leq 1/3$ Byzantine tolerance, structural data sovereignty (jurisdictional receipts never cross borders; only proofs do), PQ-anchored long-term integrity, and BIER-multicast contested-network resilience are all WRF primitives, not bolt-ons. The provided Oko simulation cadence is 500 ms; the WRF hardware target is 30 ms for tactical command and control.

11.2 Finance: CBDC, RTGS, tokenised assets

Hard-final consensus at sub-second cadence matches RTGS settlement windows. CBDC infrastructure gets jurisdictional partitioning for free (SyLOW-3 canonical labeling); cross-CBDC payments are single ACID callgraphs. HFT and exchange matching benefit from VRF-derived deterministic ordering (*no MEV by construction*). Tokenised real-world assets become Smart Assets with full content-derived provenance and finite-orbit routing when they need fabric execution.

11.3 AI workloads

The CSS storage rate 27/80 caps logical-throughput budgeting at $\sim 2.4 \times 10^7$ LQ-ops/s under the 70M TPS premise. Models are containers; agent histories are receipt chains; multi-agent coordination is stabiliser composition; provenance is orbit history. Hallucinations become orbit anomalies detectable by substrate-coherent receipts. Decentralised inference runs at $\sim 10^7$ LQ-ops/s for substantial real-time multi-agent coordination at planetary scale.

11.4 Telecommunications and IoT

Per-session finalisation at sub-second cadence, with a 30 ms tactical target; cryptographic provenance per call; inter-operator settlement is composition of typed RPC, not bilateral file exchange. 5G URLLC slice deployments inherit hard-final consensus over the same fabric carrying voice/data. Decentralised identity uses SyLOW-3-rooted X.509 delegation chains anchored in the ledger.

11.5 Edge, device, browser

A Witting tile fits in $\sim 50 \text{ mm}^2$ of 5 nm silicon (estimated). One tile provides 1.2 TOp/s of substrate-coherent ternary compute at $\sim 5 \text{ W}$ power budget—competitive with a high-end mobile NPU but with native substrate ECC, native consensus, and native smart-asset participation. Browser nodes run WASM on the same projection engine; phones and laptops contribute tile-equivalent capacity to the planetary mesh.

12 Roadmap

Window	Milestone
2026	Software reference fabric: UOR-ADDR CIDs, Atlas-12288 frame tests, trace-object flow-cell harnesses, HLIX/HCP projection containers, Smart Asset workload envelopes, and Oko BLA simulation wired into one reproducible benchmark harness.
2027	FPGA Witting tile prototype with 40 trit/logical cores, three-channel Bose–Mesner routing emulator, hardware hash path, CSS syndrome path, referenceable flow-cell emulator, and BIER multicast testbed.
2028	Multi-tile board with 1,600 logical cores, Oko committee finality over WRF projection receipts, and demonstrated 10 M TPS class benchmark under fixed committee size.
2030	Single-node WRF machine with 64,000 logical cores, local UOR object store, HLIX projection scheduler, Smart Asset settlement, and 30 ms tactical finality target tested on a million-node emulation mesh.
2032	Rack-scale deployment for defense, finance, and ASI inference: sovereign container placement, cross-jurisdiction receipts, PQC signature policy, and protected storage lift experiments.
2035	Data-centre scale WRF with 100M-class logical cores, model-weight object graphs, agent memory receipt chains, and autonomous workload markets.
2040	Planetary deployment (level 6, 4.1 B referenceable nodes, 122 ExaOp/s target class, 70 M aggregate TPS premise). FreedomNet/Golden Dome-style contested network operation and continuous ASI workloads.

13 Conclusion

The Von Neumann architecture is one local minimum of a vast design space, anchored by 1945 hardware constraints that no longer apply. The Witting Reference Fabric is a different candidate minimum for referenceable compute, reached by taking four independent recent stacks (Okos, UOR, HLIX/HCP, and Smart Assets) seriously enough to make them share one substrate.

That point is the 40-vertex symplectic polar space $W(3, 3)$, with Aut group of order 51,840 and conjugacy cardinality 30. Every subsystem of WRF—the Atlas-12288 memory frame, the Bose–Mesner three-channel bus, the spectral cache, the referenceable flow cell, the Bell-qutrit far-horizon storage cell, the CSS code, the group-action ISA, the BIER fabric, the projection engine—is forced by the same substrate. No subsystem is engineered ad hoc; no subsystem is layered on top of another; no subsystem is in tension with any other.

The natural information density is $3/8 = q/2^q$, simultaneously realised by the UOR R_{96} ratio and the s -eigenspace fraction, while the CSS storage layer spends the $3/80$ gap below it as protection overhead. The natural cadence is $j(i) = 1728 = k^3$ substrate-cycles per consensus epoch. The natural fabric branching factor is $p_{\text{th}} = k - 1 = 11$. The natural address-space partition is $40 \cdot 1,296 = 51,840$ orbit cells. The base finite-code marker is $d_Z = \mu = 4$, and the protected runtime target is the 82,320-length Steane/ Φ_6 lift. At 70 M TPS the substrate completes 1.21×10^{11} orbit cycles per second. The additional step is that memory can now be treated as a receipt-bearing trace: a flow-pattern CID names the invariant of an active process just as a UOR CID names the canonical

bytes of a static object.

Final statement. The Witting Reference Fabric is the architecture that emerges when content-addressing, group-action compute, BIER multicast, and Bose–Mesner spectral hierarchy are taken seriously as design imperatives rather than retrofit conveniences. It is the architecture toward which the Oko/UOR/HLIX/Smart-Assets stack converges, and the architecture that workloads at 70 M TPS over 10^9 nodes will run on. We do not need to invent it; it is already there, hidden in the convergence of four independent stacks. We need only to recognise it and build it.

A Glossary

Aut($W(3, 3)$)	The full automorphism group of the $W(3,3)$ symplectic polar space; isomorphic to $\text{Sp}(4, \mathbb{F}_3) \cong W(E_6)$; order 51,840.
Atlas-12288	UOR’s memory-frame structure: 48 pages of 256 bytes each, totalling 12288 bytes, with R_{96} resonance classification.
Bell qutrit	The maximally entangled past-future qutrit state $ \Omega\rangle$; substrate’s elementary storage cell.
BIER	Bit-Indexed Explicit Replication; network-layer multicast that scales with core count.
BLA	Byzantine Lattice Agreement; Oko’s consensus protocol with hard finality.
Bose–Mesner algebra	The algebra spanned by the three projectors E_0, E_1, E_2 of the SRG; supplies the three-channel bus structure.
Cayley graph	The graph whose vertices are group elements and whose edges connect elements related by a single generator; the WRF program-space.
CID	Content Identifier; the sha256 default, or another policy-selected hash, of a canonical encoding.
CSS code	Calderbank–Shor–Steane quantum error-correcting code; in this paper, the WRF base carrier is $[[240, 81, d_Z = 4]]_3$.
Conjugacy class	An equivalence class of group elements under $g \sim hgh^{-1}$. $\text{Aut}(W(3, 3))$ has 30 conjugacy classes.
HCP	Hologram Container Platform; HLIX’s projection engine.
HLIX	Hybrid Linked Infrastructure Xchange; the Hologram-based implementation path for UOR.
Oko	The network-fabric consensus protocol underlying the WRF fabric.
Projection	A declarative description of work; the only operation in HCP/WRF.
Referenceable flow cell	A bounded active memory/compute cell whose datum is the canonical invariant of its trace and whose address is $H(\text{canon}(\text{trace}))$.
Pauli group, two-qutrit	The group generated by X_3, Z_3 acting on $\mathbb{C}^3 \otimes \mathbb{C}^3$; phase-space identical to $W(3, 3)$.
R_{96}	UOR’s resonance classification: a fixed mapping from 256 byte values to 96 equivalence classes; compression ratio $3/8 = q/2^q$.
SRG(v, k, λ, μ)	Strongly regular graph: regular of valency k on v vertices, every adjacent pair shares λ neighbours, every non-adjacent pair shares μ .

Substrate primitives

The eight integers $\{q, \lambda, \mu, F_5, \Phi_3, \Phi_4, \Phi_6, p_{\text{Ih}}\}$ that generate every numerical constant in the substrate/architecture by monomial composition with norm ≤ 2 .

Sylow-3 subgroup A maximal 3-subgroup of $\text{Aut}(W(3,3))$; there are exactly 40 of them, in canonical bijection with the $W(3,3)$ vertices.

Smart Asset A self-funding, tradeable workload envelope binding economics, policy, delivered compute, and receipt history to a content-derived reference.

UOR Universal Object Reference; the content-addressing framework underlying the WRF memory model.

Witting tile A 40-core hardware unit arranged in the $W(3,3)$ collinearity graph; the unit of WRF silicon.

Witting polytope The unique 40-vertex Hermitian polychoron on which $W(3,3)$ is canonically realised; discovered by Witting in 1937.

B Quantitative Validation

B.1 Substrate-primitive readings of UOR Atlas-12288 constants

Constant	Substrate reading	Role
48	$q! \cdot 2^q$	Atlas pages per frame
256	$\mu^4 = 2^{\Phi_6+1}$	bytes per page (dS identity)
96	$2^{F_5} \cdot q$	R_{96} resonance classes
12288	$q! \cdot 2^{q+\Phi_6+1}$	Atlas frame size
3/8	$q/2^q$	Universal Density
1728	$k^3 = j(i)$	$ \text{Aut} /h(E_8)$, conjugacy multiplier
30	$q\Phi_4 = h(E_8)$	Triple Convergence (ISA family count)
11	$k - 1 = p_{\text{Ih}}$	BIER branching factor

B.2 Throughput at 70 M TPS

Quantity	Value
Substrate cycles/s	$\text{TPS} \cdot 1728 = 1.21 \times 10^{11}$
Logical-qutrit ops/s	$\text{TPS} \cdot 27/80 = 2.36 \times 10^7$
Atlas frames/s	$\text{TPS} \cdot 240/12288 = 1.37 \times 10^6$
Coherence blocks/s	$\text{TPS}/384 = 1.82 \times 10^5$
ECC syndromes/s	$\text{TPS} \cdot 60 = 4.20 \times 10^9$
Protected lift length	$240 \cdot 7^3 = 82,320$
4-bit 70B model frames	8.44×10^6 CSS-budgeted Atlas frames

References

- [1] Oko source pack, *Oko Three Pager*, provided PDF.

- [2] Okopipi/Oko source pack, *Network Protocol-Centric Blockchain L1 Enabling Decentralized Infrastructure*, provided PDF.
- [3] UOR Foundation, *Universal Object Reference*, <https://uor.foundation/>, accessed June 3, 2026.
- [4] UOR Foundation, *UOR-Framework*, <https://github.com/UOR-Foundation/UOR-Framework>, source review June 3, 2026.
- [5] UOR Foundation, *uor-addr*, <https://github.com/UOR-Foundation/uor-addr>, source review June 3, 2026.
- [6] UOR Foundation, *prism*, <https://github.com/UOR-Foundation/prism>, source review June 3, 2026.
- [7] UOR Foundation, *atlas-12288*, <https://github.com/UOR-Foundation/atlas-12288>, source review June 3, 2026.
- [8] DTR/HLIX source pack, *HLIX: A Universal Object Reference Platform (HCP)*, provided PDF, October 21, 2025.
- [9] Smart Assets source pack, *SA Attributes*, provided PNG.
- [10] Local WRF harness, *BT110/BT111 Full Architecture Suite and Spectral Trace Tower*, `wrf_full_suite_bt110_bt111.py` with result JSON, generated June 3, 2026.
- [11] Local WRF harness, *BT112 corrected $tr(A^8)$, Ihara, McKay E8, Shannon, and 3×3 flow lattice*, `wrf_bt112_suite.py` with result JSON, generated June 3, 2026.

Acknowledgements

The UOR Foundation, the Oko/Okopipi developers, the Smart Assets community, the DTR/HLIX team, and the executable W33/Witting architecture certificates in the local research repository all contributed to the convergence that this paper recognises. The architecture is not a ToE paper; it is a build path for referenceable compute.