

The Photonic Holonet

A Single Self-Entangled Photon as Universal Computer, Universal Network, and Clock
on the $W(3, 3)$ Substrate

Wil Dahn

June 2026 — machine-verified edition

Abstract

We present a complete, machine-verified architecture in which one self-entangled photon, traversing one finite geometry, is simultaneously the hardware, the software, and the network of a universal quantum computer. The geometry is the symplectic generalized quadrangle $W(3, 3)$: forty points, forty lines, automorphism group with faithful projective subgroup $\mathrm{PSp}(4, \mathbb{F}_3)$ of order 25920 and full Weyl extension $W(E_6) \cong \mathrm{PSp}(4, 3):2$ of order 51840. The symplectic double cover $\mathrm{Sp}(4, 3)$ also has order 51840, but its central involution is invisible on projective objects; the two extensions must not be conflated. The photon carries $W(3, 3)$ twice over — once as the forty Witting rays of its path–polarization space $\mathbb{C}^4$ (states), once as the forty Pauli displacement classes of its past–future time-bin space $\mathbb{C}^9$ (operators) — and the two carriers realize the same incidence geometry, making the machine’s states and gates one set of objects. Every architectural layer is a theorem with an executable witness: the preparation circuit (two Clifford gates of tabletop optics), the routing fabric (an atlas of 540 hypercube charts glued along the 1620 apartments of the Tits building), the protected memory (the Steinberg module of dimension 81), the dual-torsor compiler (three regular copies of both the flat $\mathbb{F}_3^3$ program shell and the noncommutative Heisenberg state shell, with canonical $27+27+27$ triality selected by the Heisenberg center), the oriented control plane (top homology $H_2 = 5_\omega \oplus 5_{\omega^2} \oplus 30$, whose conjugate 5-sectors fuse to one Weyl-visible 10-channel), the mirror middleware (a 2160-slot D_{12} transport bus whose full Clifford lift factors as $24 \cdot 45 \cdot 48 = 51840$), the fuel (the matter shell is exactly the magic sector, with exact Kochen–Specker budget $36/40$ and contextual fraction $1/10$), the clock (an internal $\mathbb{Z}_{12}$ gauge clock, two external references $\mathbb{Z}_7, \mathbb{Z}_{13}$, and an irrational Boerdijk–Coxeter drive realizing a discrete time quasicrystal), and universality (the photon’s own tritter, phase plate, and modulator generate the complete Clifford group; the matter shell supplies the magic). We also give the recursive engineering law: a level- n holonet has 40^n photonic leaves, $(40^n - 1)/39$ $W(3, 3)$ instances, and reversible routing diameter bounded by $8n = 8 \log_{40} N$, while durable commits use a separate Császár/tomotope clock $T(n) = 4(7^n - 1)$. The runtime contains an exact parabolic packet decoder: the Bell-line stabilizer splits the 1296 compass incidences into $648 + 324 + 162 + 162$, giving the complete binary prefix code $\{0, 10, 110, 111\}$ for mirror, schedule, and two chiral caches, while the dual point stabilizer supplies two persistent 648-slot address sheets. The two cache branches are now resolved internally: each is the same homogeneous space H/C_4 , each is a two-sheet cover of one canonical 81-route base H/D_8 , and their 324-slot union has full equivariant commutant D_8 acting in 81 four-state fibers. This split address bus is provably distinct from the older non-split ternary $81 \rightarrow 162 \rightarrow 81$ transport memory. At the communication edge, the Witting delayed-query desk is now a frame compiler: 40 tetrads times 4 Alice slots times 4 Bob slots give a 640-entry admission ROM, with 480 off-diagonal data handshakes and 160 diagonal contextual witness apertures. The corresponding analyzer bank is sparse physical optics rather than a generic multiport: the 40 Witting bases split as $1 + 12 + 27$ analyzers with at most 12 nonzero matrix entries and only cube-root phase weights over $\sqrt{3}$. The newest storage layer turns the tomotope cover pathology into an engineering

feature: durable storage is cover-indexed by a $\mathbb{Z}_k^3$ fiber over the stable 48-block packet ABI, sentinel-aware compilation can trade commit-time slack for lower $g = 15$ fault energy, the first exact guard band is the arithmetically forced desynchronization at $n = 5$ with remainder $24 = f$, and the Grünbaum–Coxeter Wythoff operations of the tomotope are exactly the packet-growth ladder 4, 12, 24, 48, 96.

Beyond the architecture, the same substrate is a candidate *physics*. The choice $q = 3$ is triply forced — by the geometric master equation $q! = 2q$, by the minimal magic (contextuality) dimension, and by the holographic central charge $c = 24 = f$ — and from it unfolds an exact exceptional skeleton: the Eisenstein tower $q = 3 \rightarrow \{3, 7\}$ Hurwitz surfaces $\rightarrow$ the Witting polytope $\rightarrow E_6/E_7/E_8 \rightarrow$ the complex Leech $\rightarrow$ the Monster at $c = 24$, threaded by $G_2 = \text{Aut}(\mathbb{O})$ and welded by a single order-3 Eisenstein element. The Standard Model descends from E_6 by trification ($\dim \text{SU}(3)^3 = 24 = f$, $\dim \text{SM} = 12 = k$), with the three generations, gauge unification ($\sin^2 \theta_W = 3/8$), and tri-bimaximal PMNS mixing all one S_3 triality. The falsifiable layer is on a live test sheet: the cosmological predictions ($n_s = 29/30$, $r = 1/300$, $f_{NL} = 1/72$) are consistent, the neutrino sum is resolved by the seesaw ($\sum m_\nu \approx 0.06 \text{ eV}$), and proton decay ($\tau_p \sim 10^{35-36} \text{ yr}$, Hyper-K) together with two demonstrator readouts (contextual fraction 1/10, pump Chern 2) are sharp near-term tests. No fitted parameters appear anywhere: every constant is substrate arithmetic at $q = 3$.

Contents

1	Overture: the machine is the geometry is the carrier	6
2	The substrate	8
2.1	The two carriers and the operator–operand duality	9
2.2	The five vacua and the Schläfli geography	9
2.3	The parabolic router: address and route are dual	9
3	The carrier: how to self-entangle a photon	10
3.1	Stage A: spatial (one element)	11
3.2	Stage B: temporal (two Clifford gates)	11
4	The network: atlas, building, and routing	11
4.1	Hypercube networking, not hypercube decoration	12
5	The middleware: tomotope mirrors and the runtime bus	13
5.1	Product-grid closure ABI v2	14
5.2	ABI v2 consumer, E6 claim table v2, and D4/V4 branch audit	14
5.3	BT1486–BT1488 ABI v2 retwined CSS and splice manifest	15
5.4	BT1489–BT1491 row symmetry and Fano–E6 square	15
5.5	BT1492–BT1494 canonical Fano fiber and pulse release lock	15
6	The software: braids, teleported gates, universality	16
6.1	Exact braid words	16
6.2	The photon is its own gate	16
6.3	Instruction set architecture	16
6.4	The minimal classical shadow	16
6.5	The Universality Theorem	17
7	Fractal scaling: the computer is the network	17

8	Runtime closure: compile, monitor, commit	18
8.1	The infinite-cover warning	19
8.2	Cover-indexed storage, rerouting, and guard bands	19
9	The fuel: matter equals magic	21
10	Memory, protection, and the immune system	23
11	Time: three clocks and a quasicrystal	26
12	Synchronization: beacons and schedules	27
13	Build sheet and falsifiable witnesses	32
14	Implications, corrections, and the ethos	33
14.1	Physics	33
14.2	Computing	37
14.3	Networking	38
14.4	The physics-to-architecture dictionary	39
14.5	Beyond the laboratory	40
14.6	The ethos	41
A	Verification ledger	42
B	Architecture Completeness and the Three Physical Residuals	44
B.1	The architecture is the substrate’s symplectic continuum	45
B.2	The fault-tolerant layer is the substrate’s lattice tower [Face 5]	45
B.3	The universal gate set is degree two plus degree three	46
B.4	The code lattices are the moonshine ladder [Face 5]	47
B.5	What is physically needed: demonstrator versus fault-tolerant machine [Face 6]	47
B.6	The gates are geometric: holonomic Clifford from the gauge connection	48
B.7	The quantum power is the gauge curvature	48
B.8	The fuel is contextuality: the register is a Wigner phase space and magic is its negativity	49
B.9	From the matter-graph lattice to continuum Yang–Mills: the loop closes	49
B.10	Why “holo-net”: the network is a holographic code	50
B.11	The boundary central charge is $f = 24$: the Monster CFT is the holonet’s holographic dual	51
B.12	The falsifiability ledger, and the demonstrator as an experiment on the theory	54
B.13	The geometric origin of the clock: the triangulation ladder $3 \rightarrow 4 \rightarrow 7$	55
B.14	The clock is a mass and the supercycle is the gauge group: running the machine is the physics	56
B.15	Four further faces of the machine: time, topological order, scrambling, and periodic orbits	57
B.16	A near-term falsification test: the demonstrator as an experiment on the substrate	57
B.17	BT1402 runtime-frontier handoff: readout, port, and certificate	58
B.18	BT1403 eraser-lift port: the Hesse boundary is the readout boundary	59
B.19	BT1404 holonet scope: one Hesse grid is one microframe	60
B.20	BT1405 continuous Q6 path router	60

B.21	BT1406 tomatope-body edge pulse scheduler	61
B.22	BT1407 microframe transaction composer	61
B.23	BT1408 Witting contextual communication bridge	61
B.24	BT1409 Witting duplex admission scheduler	62
B.25	BT1410 Witting delayed-query frame compiler	63
B.26	BT1411 Witting basis analyzer unitaries	63
B.27	BT1412 toroidal Q4 oscillator boundary	64
B.28	BT1413 Q4 plaquette-tomatope compiler	65
B.29	BT1414 Csaszar-Szilassi dual port	65
B.30	BT1415 even-projection Steinberg/CSS syndrome ledger	66
B.31	BT1416 sparse CSS intertwiner	67
B.32	BT1417 linear-optical dual-port primitive synthesis	67
B.33	BT1418 finite D4-quartic magic injection	68
B.34	BT1697 typed packet ABI	68
B.35	BT1698–BT1700 execution stack	69
B.36	BT1701–BT1703 runtime safety layer	70
B.37	BT1704–BT1706 replay and retry economics	71
C	The exceptional skeleton and its physics	71
C.1	The one object and its seven faces	72
C.2	BT1707–BT1709 contextuality and Hesse crossover	75
C.3	The $\{3, 7\}$ Hurwitz tower and the genus-7 crossover	76
C.4	Two Schläfli families, a vertex-figure selection, and the Witting body	77
C.5	The register atlas, the explicit E8 weld, and the Hesse engine	78
C.6	The Eisenstein tower qutrit $\rightarrow$ E6 $\rightarrow$ E8, the Golay gap, and the quadrangle ladder	79
C.7	A guardrail on the 27, and the Klein quartic as the E7 rung	79
C.8	The three exceptional rungs made explicit	80
C.9	The magic square, the complex Leech, and the icosian clock	80
C.10	The ceiling, the Suzuki tower, and the G2 thread	81
C.11	The ledger, the Standard Model, and the tower as one figure	81
C.12	The measurable layer: a scorecard and a unification [Faces 2, 3, 7]	82
C.13	Three live tests, scrutinised	82
C.14	Fixing the neutrino, predicting the proton, and the balance sheet	83
C.15	Why $q = 3$, and one triality for the neutrino [Faces 1, 4]	83
C.16	Pass 1 — three closures: the neutrino matrix, the selection census, the threshold number	84
C.17	Pass 2 — sharper: a Majorana invariant, a cyclotomic merge, a code curve	86
C.18	Pass 3 — deeper: a cyclotomic skeleton, a projective Majorana ratio, a decoder correction	87
C.19	Pass 4 — deepest: the cubic form, the Witting degrees, and the measured curve	88
C.20	Self-error-correction and why the primitive is one	88
C.21	Why the primitive is one <i>massless</i> photon	90
C.22	The single carrier, the geon, and the fractal of nested shells	91
C.23	Pass 5 — the capstone: machine = world, gravity to a number, and a fittable sky template	92
C.24	Pass 6 — sharpening the capstone: a fit, a derived frequency, and a running G	93
C.25	Pass 7 — three edges pressed: a resolution reality, the last identification closed, and an honest hierarchy	93

C.26 Pass 8 — the hierarchy’s exponent, a designed closure test, and one clock for the whole sky	94
C.27 Pass 9 — the breakthrough: the complete primordial spectrum, normalisation included	95
C.28 Pass 10 — the scale residue closed to one unit: amplitude as a law, all masses as integers, two exponents as one	96
C.29 Pass 11 — the potential identified (Starobinsky R^2), the entropy derived, the tower audited	96
C.30 Pass 12 — the residue as one number, the scalaron meets the neutrinos, and the tower as one plot	97
C.31 Pass 13 — one particle for inflation and matter, the baryons, and the master ledger	98
C.32 Pass 14 — the tension resolved, three hierarchies reconciled, and dark matter completed	99
C.33 Pass 15 — the dark matter stress-tested, the lightest neutrino, and the mass ladder .	99
C.34 Pass 16 — the lightest neutrino pinned, the $0\nu\beta\beta$ mass, and the descent to dark energy	101
C.35 Pass 17 — the cosmological constant exact, neutrino dark energy, and the final ledger	101
C.36 Pass 18 — a mechanism for the cosmological constant, the joint sky test, and the standalone distillation	102
C.37 Pass 19 — the floor derived from the 600-cell, and the dated decision	103
C.38 Pass 20 — the one residual input, the floor base settled, and the look-elsewhere bound	104
C.39 Pass 21 — the computed look-elsewhere, the pinned Yukawa, and the tower on one axis	105
C.40 Pass 22 — the depth-stable significance, an honest neutrino-mixing failure, and dated discovery years	106
C.41 Pass 23 — the predicted mixing, the derivation ledger, and the separation from $SO(10)$	106
C.42 Pass 24 — machine = world: the code is the matter algebra, and why 3	107
C.43 Pass 25 — the protected 27 is one Standard-Model generation, and charge is $1/3$ because $q = 3$	108
C.44 Pass 26 — why the Standard Model is consistent, and the weak angle as a 27-trace .	109
C.45 Pass 27 — the three forces meet at $M_{\text{Pl}}e^{-\Phi_6}$ with $\alpha_{\text{GUT}}^{-1} = f$, and the proton decays	110
C.46 Pass 28 — the Higgs sector, one structural supersymmetry, and the state of the solution	111
C.47 Pass 29 — gravity is the spectral action, and the continuum gap is two theorems . .	112
C.48 Pass 30 — the emergent spacetime is K3, and (T2) reduces to one convergence theorem	112
C.49 Pass 31 — the convergence mechanism, demonstrated in 4D, and (T2’) narrowed to one input	113
C.50 Pass 32 — cracking the frontier: both theorems collapse to one finite condition . . .	114
C.51 Pass 33 — the last link is a theorem: the gravity derivation closes to two narrow conditions	115
C.52 Pass 34 — the system: the substrate as a computer and network architecture	116
C.53 Pass 35 — the datapath: the router, the clock, and the reliability budget	117
C.54 Pass 36 — the silicon questions: the ISA enumerated, the power budget, the floorplan	118
C.55 Pass 37 — the boundary questions: complexity class, distributed clock, I/O, and one group	119
C.56 Pass 38 — the advantage measured, the scheduler, and the falsifiable datasheet . . .	120
C.57 Pass 39 — the planetary computer: the universal VM, the everyone-node fractal, and the bench test	121
C.58 Pass 40 — the machine made runnable: the emulator and three frontier applications	121
C.59 Pass 41 — proof of life: the machine corrects, teleports, and reproduces, executed . .	122
C.60 Pass 42 — the benchtop number derived, and the quantum packet delivered	123

C.61 Pass 43 — consensus run against an adversary, the KS inequality, and the holonet CLI	123
C.62 Pass 44 — the experiment simulated, the magic dial turned, the tool installed	124
C.63 Pass 45 — the threshold curve measured, the quickstart, and the unified Machine paper	125
C.64 Pass 46 — the minimal machine, the matched efficiency, the circuit-level threshold, and the scorecard	125
C.65 Pass 47 — the router on a 4-bit CPU, the ternary encoding tax, and the live scorecard	126
C.66 Pass 48 — the assembler, packet energy, and the live router	127
C.67 Pass 49 — exported retro targets and golden traces	127
C.68 Pass 50 — firmware-to-fabric accounting	128
C.69 Pass 51 — the machine audits itself, and the live playground	129
C.70 Pass 52 — the parity law, and the performance face	130
C.71 Pass 53 — parity not primality, a falsifiable prediction, and the table-free win	131
C.72 Pass 54 — the control arm: an explicit even- q ovoid, and the scaling win	132
C.73 Pass 55 — the two-arm discriminator, runnable end to end	132
C.74 Pass 56 — two contextualities, separated; and why one photon in $\mathbb{C}^4$	133
C.75 Pass 57 — the contextuality tax: the defect is one movable point-star	134
C.76 Pass 58 — the spread side of the tax, and the anatomy of the optima	135
C.77 Pass 59 — the tax orbits: one group moves the defect's ground states	135
C.78 Pass 60 — the perp states: the canonical orbit named, the spectrum explained . . .	136
C.79 Operator-on-photon self-applied circuit packet	137
C.80 Internal optic algebra, physical dictionary, and falsifier regimes	137
C.81 Internal Clifford census, lens calibration, and passive-active protocol	137
C.82 OAM recentering, exact optical calibration, protocol table, and literature bridge . .	138
C.83 Recentered witnesses, leakage regimes, and appendix splicer	138
C.84 Appendix validation, recentered protocol table, and claim ledger	138
C.85 Full operator/OAM appendix splice and recentered transaction ABI	138
C.86 LG radial covariance, full witness lane sheet, and OAM front end	139
C.87 Lab tomography, LG alphabet, and the Hesse/T witness loop	140
C.88 The Witting reject shell as the universal fuel rail	140
C.89 Accepted control rail and the full Witting transaction cycle	141
C.90 Physical automaton, Fano detector bins, and the universal ABI	142
C.91 Fano charge, time-bin envelope, and guard-page closure	143

1 Overture: the machine is the geometry is the carrier

Classical architecture separates three things: the processor that acts, the program that directs, and the network that connects. The thesis of this paper is that at the level of a single photon on a finite symplectic geometry the three are one object, and that this is not a metaphor but a chain of theorems.


Principle 1.1 (Identity of the three layers). *Transporting the photon through the mesh is applying a gate is routing a packet. One physical process; three descriptions.*

Rosetta stone of the $q = 3$ substrate (read this first). This paper has two halves: the *machine* (the architecture, below) and the *world* it computes (§C). Both are the one $q = 3$ Eisenstein object — the Witting polytope $3\{3\}3\{3\}3\{3\}3$ (240 vertices = E8 roots), its symmetry $ST\#32$ (order $155520 = 3|\text{Aut}(W(3,3))|$), equivalently the degree-two cyclotomic skeleton $\{\Phi_3, \Phi_4, \Phi_6\} = \{13, 10, 7\}$ at $q = 3$ — seen from *seven faces*: **(1) selection** ($q = 3$ forced: $\Phi_4 = q^2 + 1$ factors the de Sitter cubic and $v = (q+1)\Phi_4 = 40$), **(2) constants** (the Witting degrees $\{12, 18, 24, 30\} = \{k, h(E_7), c=f, h(E_8)\}$), **(3) gauge** ($\dim \text{SU}(3)^3 = 24 = c$, $\dim \text{SM} = 12 = k$, $\sin^2 \theta_W = 3/8$), **(4) neutrino** ($\Phi_3/q^2 = 13/9$, PMNS $1/\Phi_3$), **(5) code** (the GKP tower $A_2 < D_4 < E_8$), **(6) demonstrator** (contextual fraction $1/\Phi_4 = 1/10$, pump Chern $\lambda = 2$), and **(7) cosmology** ($N = 2(v - \Phi_4) = 2h(E_8) = 60$ e-folds; $n_s = 29/30$, $r = k/N^2 = 1/300$). The faces share integers — $k=12, \Phi_4=10, \Phi_3=13, c=f=24, v=40$ each serve ≥ 3 faces (Table 1) — which is why they are one object (Fig. 9). *No free integer parameter*: every geometric integer is a cyclotomic value, GQ count, or Witting degree, and even the moonshine ceiling closes — the Leech kissing number factors as $196560 = 6\mu q^2 \Phi_3 \Phi_4 \Phi_6$, so the Monster (196883), $\tau = \mu q^2 \Phi_6 = 252$, and the fine-structure integer $1/\alpha = 137 = \Phi_3 \Phi_4 + \Phi_6$ are all $\{\mu, q, \Phi_3, \Phi_4, \Phi_6\}$ arithmetic; the lone non-integer input is the 0.036 of 137.036 (QED running). *The one experiment to run first*: the benchtop, calibration-free contextual fraction $1/\Phi_4 = 1/10$, which tests three faces at once.

The identification rests on seven exact pillars, each proved by explicit computation in the project repository (every theorem below carries its witness script):

1. **Operator–operand duality** (§2.1): the same $W(3,3)$ is drawn by the photon’s *states* (orthogonality of the 40 Witting rays in $\mathbb{C}^4$) and by its *operators* (commutation of the 40 two-qutrit Pauli classes on $\mathbb{C}^9$).
2. **Routing–gate identity** (§4): the mesh is an atlas of 540 three-dimensional hypercube charts whose native XOR routing moves are exactly the register operations, glued along the 1620 apartments of the Tits building of $\text{Sp}(4, \mathbb{F}_3)$.
3. **Parabolic address–route duality** (§2.3): the point parabolic supplies two 648-slot address sheets, whereas the Bell-line parabolic decodes the compass fabric by the complete prefix code 0, 10, 110, 111; its cache quarter is an exact D_8 -over-81 four-state address fiber. Address and route are the two nodes of the same C_2 building, while the separate ternary transport extension carries temporal state.
4. **Dual-torsor state–program compilation**: the Steinberg memory restricts to both canonical order-27 shell groups as three regular copies. Bell-shell programs use commuting $\mathbb{F}_3^3$ coordinates; point-shell states use Heisenberg 3^{1+2} coordinates; both compile into the same 81 logical qutrits.
5. **Mirror middleware** (§5): the cyclic selector clock and the mirror transport bus are different 2160 worlds. The photon is clocked by the C_{12} side but routed by the D_{12} side, whose full Clifford lift is the exact product $24 \cdot 45 \cdot 48$.
6. **Universality from optics** (§6.5): the photon’s own beam-splitting elements generate the complete Clifford group (symplectic closure = 51840, exact), and its matter shell is precisely its magic supply.
7. **Fractal holonet scaling** (§7): recursively replacing every site by a fresh $W(3,3)$ core gives a universal computer/network with exact 40^n leaves and logarithmic reversible routing.

And the machine is not separate from the world it computes. The second half of the paper (§C) shows that the *same* $q = 3$ Eisenstein structure that fixes this architecture — the Witting polytope, its symmetry ST#32, the cyclotomic skeleton $\{\Phi_3, \Phi_4, \Phi_6\}$ — is also the selection of $q = 3$, the substrate constants, the Standard-Model gauge group, the neutrino masses, the error-correcting code, the benchtop experiment, and inflationary cosmology. These seven faces, and the single object behind them, are laid out in §C.1 (Fig. 9); the reader who wants the unifying picture before the engineering may read that subsection first.



One loop: carrier → data plane → router → building → runtime → tomatope packet → mirror bus → carrier.

Figure 1: The photonic holonet as an engineering stack. The same photon is data, instruction, and packet; the same finite geometry supplies the network, middleware, and runtime.

2 The substrate

Definition 2.1 ($W(3,3)$). *Let $V = \mathbb{F}_3^4$ with the alternating form $\langle u, v \rangle = u_1v_3 - u_3v_1 + u_2v_4 - u_4v_2$. The points of $W(3,3)$ are the 40 projective points of V ; the lines are the 40 totally isotropic projective lines. Each point lies on 4 lines and each line carries 4 points; collinearity defines the strongly regular graph $\text{SRG}(40, 12, 2, 4)$.*

Three related symmetry groups occur, and their roles are different:

$$\underbrace{\text{Sp}(4,3)}_{2\text{-PSp}(4,3), \text{ Clifford double cover}}, \quad \underbrace{\text{PSp}(4,3)}_{25920, \text{ faithful projective action}}, \quad \underbrace{W(E_6) \cong \text{PSp}(4,3):2}_{51840, \text{ outer Weyl extension}}.$$

The first and third groups both have order 51840 but are not the same extension: the center of $\text{Sp}(4,3)$ acts trivially on projective points and lines, while the outer involution in $W(E_6)$ acts nontrivially. The projective group acts with permutation rank 3. The ATLAS index-40 subgroups are the two building parabolics $3^{1+2}:2A_4$ and $3^3:S_4$ [21]. The substrate primitives are

$$q = 3, \quad \lambda = 2, \quad \mu = 4, \quad k = 12, \quad f = 24, \quad g = 15, \quad \Phi_6 = 7, \quad \Phi_4 = 10, \quad \Phi_3 = 13,$$

and every numerical claim in this paper reduces to arithmetic in these.

Formal grounding (the master formalism). In the master manuscript [1] the whole structure is forced by one Diophantine seed: the *master equation* $q! = 2q$ has the unique positive-integer solution $q = 3$, and the SRG quadratic $x^2 - q!x + 2^q = (x-2)(x-4)$ then delivers $\lambda = 2$, $\mu = 4$ and the full parameter closure ($k = 2^q + q + 1$, $v = 40$, multiplicities $f = 24$, $g = 15$, cyclotomic ladder $\Phi_3 = 13$, $\Phi_6 = 7$, $\Phi_{12} = q^4 - q^2 + 1 = 73$, $\Theta = q^2 + 1 = 10$). Two of those closure constants reappear in this paper as *spectral field discriminants*: the chart-web’s irrational eigenvalue pairs $1 \pm \sqrt{10}$ and $(-1 \pm \sqrt{73})/2$ (§4) live precisely in $\mathbb{Q}(\sqrt{\Theta})$ and $\mathbb{Q}(\sqrt{\Phi_{12}})$ — the machine’s transport spectrum speaks the same cyclotomic dialect as the physics tier ladder.

2.1 The two carriers and the operator–operand duality

Theorem 2.2 (Two realizations, one geometry). (i) *The 40 Witting rays in $\mathbb{C}^4$ (the photon’s path $\otimes$ polarization space) have orthogonality graph isomorphic to the collinearity graph of $W(3, 3)$.* (ii) *The 40 projective classes of two-qutrit Pauli displacements $D(v) = X^a Z^b \otimes X^c Z^d$ on $\mathbb{C}^9$ (the photon’s past $\otimes$ future time-bin space) have commutation graph equal to the same $W(3, 3)$, via $D(u)D(v) = \omega^{\langle u, v \rangle} D(v)D(u)$.* Witnesses: `bt817_self_entangled_photon_atlas.py`, `bt821_operator_operand_duality.py`.

Thus a point of the substrate is simultaneously a pure state the photon can occupy and a unitary the photon can enact; a line is simultaneously a measurement context (orthonormal tetrad) and a stabilizer group (maximal commuting Pauli set with its 9-state joint eigenbasis). The $360 = 40 \times 9 = f \cdot g$ two-qutrit stabilizer states are the joint eigenbases of the forty lines.

2.2 The five vacua and the Schläfli geography

Every maximal subgroup class of $\mathrm{PSp}(4, 3)$ decomposes the substrate in its own canonical way; the maximal indices are the classical inventory of the 27 lines on a cubic surface.

index	subgroup	points	lines	vacuum	cubic surface
27	$2^4:A_5$	[40]	[40]	homogeneous (icosahedral)	27 lines
36	S_6	[40]	[10, 30]	spread fibration	36 double-sixes
40	parabolic	[1, 12, 27]	[4, 36]	holonet point vacuum	trihedra I
40	parabolic	[4, 36]	[1, 12, 27]	dual line vacuum	trihedra II
45	$(2T \times 2T):2$	[8, 32]	[16, 24]	binary polar	45 tritangents

The physical decomposition $40 = 1 + 12 + 27$ (self + gauge shell + matter shell) is the *choice of the point-parabolic vacuum* among five inequivalent readings; the icosahedral vacuum sees a featureless substrate. The platonic ladder threads the table: the binary tetrahedral group $2T$ (the 24-cell) lives in every polar-pair stabilizer, the real octahedral group $O_h (= S_4 \times \mathbb{Z}_2$, order 48) is the chart symmetry of §4, and the binary icosahedral group $2I = \mathrm{SL}(2, 5)$ (the 600-cell) is the core of every spread stabilizer, with point orbits [20, 20] and line orbits [10, 30] echoing the Boerdijk–Coxeter decomposition $600 = 20 \times 30$. *Witnesses: bt808–bt813.*

2.3 The parabolic router: address and route are dual

The two index-40 parabolics are not duplicate descriptions of one local neighborhood. They execute different halves of a packet protocol. This becomes visible only after the compass census: there are two 216-element classes of icosahedral A_5 cores, and exactly $1296 = 216 \cdot 6$ cross-class pairs meet in D_{10} .

Theorem 2.3 (Bell–compass parabolic router). *Fix an isotropic Bell line. Its projective Siegel parabolic $3^3:S_4$ has four orbits on the 1296 incident compass pairs:*

$$162_L + 162_R + 324 + 648.$$

They are distinguished objectwise by the Bell line lying in the pentad core’s 5_L , 5_R , 10, or 20 line orbit. The normalized orbit sizes satisfy the Kraft equality

$$2^{-3} + 2^{-3} + 2^{-2} + 2^{-1} = 1,$$

so they admit the complete prefix decoder

$$110 \mapsto \text{left cache}, \quad 111 \mapsto \text{right cache}, \quad 10 \mapsto \text{schedule}, \quad 0 \mapsto \text{dark mirror}.$$

The outer Weyl involution fuses exactly the two 162 cache orbits, leaving $324_{\text{cache}} + 324_{\text{schedule}} + 648_{\text{mirror}}$. By contrast, the dual point parabolic has two regular projective orbits of size 648; the outer extension preserves rather than fuses these address sheets. Witness: `bt859_bell_compass_parabolic_router.py`.

The equality case of Kraft–McMillan means the routing words fill a binary decision tree with no unused branch [22]. Here the code lengths were not optimized from assumed traffic probabilities: they were read from exact group orbits. A packet first asks whether its Bell line is in the dark mirror half; if not, whether it is in the common schedule quarter; only the remaining cache quarter needs the third, chiral bit. The decoder is self-delimiting and constant-depth, with expected word length

$$\frac{1}{2}(1) + \frac{1}{4}(2) + 2 \cdot \frac{1}{8}(3) = \frac{7}{4}$$

under the invariant uniform measure on the 1296 compass pairs.

orbit	Bell-line position	prefix	engineering action
648	dark 20-orbit	0	enter mirror bus
324	common schedule 10-orbit	10	timetable-local route
162	absolute pentad 5_L	110	left cache
162	absolute pentad 5_R	111	right cache




Figure 2: The two C_2 parabolics implement a dual protocol. Point localization selects one of two address torsors; Bell-line localization decodes a complete variable-length route word.

This theorem also closes an honesty boundary. The Bell stabilizer and the compass incidence set both have cardinality 1296, but they are not a regular G -set identification. Projectively, the Bell stabilizer has order 648 and the orbit stabilizers are 1, 2, 4. The useful result is stronger than the count: a hierarchical decoder and a persistent address bit are forced by the two parabolic actions.

3 The carrier: how to self-entangle a photon

Entanglement is a relation between tensor factors; nothing requires the factors to be distinct particles. Self-entanglement is entanglement between one photon's own registers.

3.1 Stage A: spatial (one element)

A diagonal photon meets a polarizing beam splitter: $\frac{1}{\sqrt{2}}(|H\rangle + |V\rangle) \mapsto \frac{1}{\sqrt{2}}(|H, a\rangle + |V, b\rangle)$. Path and polarization are now maximally entangled *within* the photon; this two-qubit $\mathbb{C}^4$ is the Witting carrier of Theorem 2.2(i).

3.2 Stage B: temporal (two Clifford gates)

A tritter (symmetric 3-port coupler) is the qutrit Fourier gate F_3 ; a three-bin delay ladder $(0, \tau, 2\tau)$ defines past and future registers; one electro-optic modulator applies $CX_{p \rightarrow f}$ conditioned on bin index. The temporal Bell qutrit

$$|\Omega\rangle = CX_{p \rightarrow f}(F_3 \otimes I)|0\rangle_p|0\rangle_f = \frac{1}{\sqrt{3}} \sum_{j=0}^2 |j\rangle_p |j\rangle_f$$

is verified exactly; the photon is entangled with its own future. Inserting any U in the future arm, the recombination visibility is the *trace-Choi witness*

$$V(U) = \frac{|\text{Tr } U|}{3}, \quad V(F_3) = \frac{1}{3}, \quad V(X) = V(Z) = 0,$$

so the photon implements the Choi-Jamiołkowski isomorphism on itself: it measures channels against its own past. *Witness: `bt820_self_entanglement_protocol.py`.*




Figure 3: Preparation schematic. One photon, four registers. The PBS entangles path with polarization (Stage A); the tritter–delay–EOM chain prepares the temporal Bell qutrit $|\Omega\rangle$ (Stage B); the dashed loop is the irrational Boerdijk–Coxeter drive of §11.

4 The network: atlas, building, and routing

Theorem 4.1 (The hypercube atlas). *$W(3, 3)$ contains exactly 540 skew line pairs. For each pair the non-collinearity graph on its 8 points is the cube graph Q_3 , with full symmetry group O_h of order 48 (these are the only order-48 stabilizers; the binary $2O$ does not occur). Each chart admits an exact $\mathbb{F}_2^3$ Gray-code addressing in which chart edges are unit XOR moves and the antipode of a vertex is its unique collinear partner. Every $W(3, 3)$ nonedge is a chart edge in exactly $12 = k$ charts; every edge is an antipode pair in exactly $9 = q^2$. Witnesses: `bt773`, `bt777`, `bt778`.*

4.1 Hypercube networking, not hypercube decoration

The cube is not only a picture for one chart. It is the local interconnect primitive. A chart gives eight simultaneously valid addresses

$$000, 001, 010, 011, 100, 101, 110, 111 \in \mathbb{F}_2^3,$$

and the elementary network instruction is the dimension-order hypercube hop

$$a \mapsto a \oplus e_i, \quad i \in \{1, 2, 3\}.$$

This is the standard e-cube rule of hypercube networks: compare source and target, flip one differing coordinate at a time, and finish in at most three local hops [13]. In the holonet that same XOR flip is also a register operation, because the chart address is a finite “which-register” word on the photon’s self-entangled carrier. The local routing primitive is therefore simultaneously

$$\text{bit flip} = \text{qutrit-frame update} = \text{optical path choice}.$$

A packet in the photonic holonet is not merely a payload. The natural packet header is

$$\mathcal{P} = (\text{payload}, \text{Pauli frame}, \text{chart word}, \text{apartment hop}, \text{mirror slot}, \text{clock phase}).$$

The payload is the quantum state or teleported gate; the Pauli frame is the two-trit feedforward record; the chart word is the $\mathbb{F}_2^3$ route inside the current cube; the apartment hop selects the next chart; the mirror slot chooses the D_{12} transport bus of §5; and the clock phase is the cyclic C_{12} interrupt that keeps the selector coherent. This is why the same object can be called a computer architecture and a network architecture. A node never sends data to a processor: the act of transport is the processor’s gate action.

Theorem 4.2 (The fabric is the building). *The chart-adjacency web (540 nodes, two charts adjacent when one’s axis is a disjoint transversal pair of the other) is 6-regular with exactly 1620 edges, and these edges are in canonical bijection with the 1620 apartments of the Tits building of $\text{Sp}(4, \mathbb{F}_3)$. The web has diameter 5; its adjacency module contains the Steinberg representation twice and presses against the Ramanujan bound only on a 15-dimensional sentinel eigenspace. Witnesses: bt777, bt779, bt744.*

Routing a packet means choosing XOR moves inside charts and apartment hops between them; since chart moves *are* register operations (§6), routing is computation. At scale the recursion is the holonet ladder: 40^n cores at level n , the same $\text{SRG}(40, 12, 2, 4)$ at every level, with the level-0 chart’s $1 + 12 + 27$ degree split realized as self pole, gauge shell, and matter shell of each node.

Corollary 4.3 (Two-plane routing bound). *Inside a chart the worst route length is 3, the diameter of Q_3 . Between charts the worst route length is 5, the diameter of the apartment web. One recursive digit of a holonet address therefore costs at most*

$$3 + 5 = 8$$

reversible transport moves. BT827 globalizes this to $8n = 8 \log_{40} N$ for $N = 40^n$ photonic leaves.




Figure 4: Network schematics. Left: a single chart with its $\mathbb{F}_2^3$ addressing (XOR routing native). Center: the 540-chart web whose links are the apartments of the Tits building (sketched as a circulant fragment). Right: the two-level routing stack.

5 The middleware: tomotope mirrors and the runtime bus

Definition 5.1 (Tomotope packet). *In this paper the tomotope is used in the only way the finite calculation currently justifies: as a local middle-layer incidence packet. A tomotope packet has four vertex carriers, twelve local edge axes, sixteen triangular face labels, and forty-eight edge-face incidence blocks. The doubled base/shadow fiber gives 192 flags. Thus the tomotope is the durable packet format sitting between the fast cube router and the global mirror bus; it is not introduced as an ambient continuous torus, and it is not identified with the rectangle clock.*

The atlas contains two different 2160-element boundary worlds. They must not be collapsed. The rectangle side is cyclic: its stabilizer is C_{12} and it supplies the selector's phase clock. The chart-transversal side is dihedral: its stabilizer is D_{12} and it supplies the mirror transport bus. Equal cardinality is the red herring; different stabilizer structure is the theorem.

Theorem 5.2 (The D_{12} mirror bus). *The 540 cube charts carry exactly four common-transversal slots each, so the global chart-transversal space has $540 \cdot 4 = 2160$ slots. The map*

$$(\text{chart}, \text{common transversal line}) \mapsto (\text{chart}, \text{antipode pair cut out by that transversal})$$

is an equivariant bijection onto the atlas antipode-slot space. Its projective stabilizer has order 12 and GAP identifies it as D_{12} , not C_{12} . Witness: `bt815_global_2160_transversal_gset.py`.

Locally, a skew-line chart has four common isotropic transversals. Read those four transversals as the four vertex carriers of the tomotope middle layer. Each transversal is a K_4 , and each K_4 has three opposite-edge axes and four triangular faces. Thus the chart carries

$$4 \text{ vertices}, \quad 4 \cdot 3 = 12 \text{ edge labels}, \quad 4 \cdot 4 = 16 \text{ face labels},$$

and the middle incidence packet has

$$4 \cdot 3 \cdot 4 = 48$$


edge-face blocks, with the doubled base/shadow fiber restoring 192 flags. This is the local $\langle r_0, r_3 \rangle$ tomotope block layer, not an analogy to a toroidal polyhedron. *Witness: `bt814_tomotope_middle_layer_from_residual_tetrahedra.py`.*

Theorem 5.3 (Photonic mirror middleware). *The full optical Clifford runtime factors as*

$$|\mathrm{Sp}(4, \mathbb{F}_3)| = 51840 = 24 \cdot 2160 = 24 \cdot 45 \cdot 48.$$

Here $24 = f$ is the full $\mathrm{Sp}(4, 3)$ lift of the projective D_{12} mirror-slot stabilizer, 45 is the hyperbolic polar-pair/tritangent geography, and 48 is the local tomatope edge-face middle layer (also the order of the chart group O_h). Equivalently, the photon is clocked by the cyclic selector side and routed by the dihedral mirror side; the complete Clifford group is the lift of that mirror bus through the tomatope middle carrier. Witness: `bt826_photonic_mirror_middleware.py`.

This is the architectural closure missing from a bare universality claim. The tritter, phase plate, and controlled delay generate the whole Clifford group, but the group is *organized* as a middleware stack: full slot stabilizer, global mirror bus, polar vacuum geography, and tomatope local block. In software terms, C_{12} is the clock interrupt, D_{12} is the mirror router, 48 is the local packet format, and $24 \cdot 45 \cdot 48$ is the finite runtime.



fast XOR routing is committed through the tomatope middle layer and reflected by the mirror bus.

Figure 5: Middleware schematic. The cube is the reversible local router; the tomatope is the edge-face packet format; the D_{12} bus is the global mirror transport object.

5.1 Product-grid closure ABI v2

The 12 closure strands are modeled as

$$C_3 \times V_4, \quad V_4 = C_2 \times C_2,$$

rather than as a bare C_{12} . The C_3 coordinate is the Szilassi/Fano channel and qutrit phase axis; the V_4 coordinate is the four-triangle E6 gauge-shell/D4 branch axis.

The E6 firewall square is now upstream in the claim DAG through the exact nodes

$$36 \rightarrow 72, \quad 72 + 6 = 78, \quad 72 + 9 = 81.$$

This supports the exact finite ABI/CSS branch while leaving blocked formula and physics claims downstream of the transcription gate.

The closure packet ABI is upgraded to v2. Each packet carries C_3 channel, V_4 branch, active/guard outputs, 72-row sector metadata, the nine-mode firewall gap, and claim-DAG dependencies.

5.2 ABI v2 consumer, E6 claim table v2, and D4/V4 branch audit

The closure ABI v2 is executable. It regenerates 12 packets, 48 events, and 72 rows while preserving the dual-axis metadata. The three C_3 channels each carry 24 rows, and the four V_4 triangles each carry 18 rows.

The E6-merged claim DAG now produces a table with explicit E6 provenance for the 72-sector, the 81 closure, and the $C_3 \times V_4$ grid. Blocked formula and physical claims remain downstream.

On branch classes, D_4 shear fixes each V_4 triangle, while τ_4 acts as an order-two translation preserving the triangle partition as a set. The two operations still fail to commute on full branch/phase states, so retwining remains part of the exact decoder rule.

5.3 BT1486–BT1488 ABI v2 retwined CSS and splice manifest

BT1486 reruns the CSS join from the ABI v2 consumer output. The result is not just the old 72-row count: every row satisfies the retwined X - and Z -syndrome checks, while the C_3 channel profile remains $P_0 = P_1 = P_2 = 24$ rows and the V_4 triangle profile remains $T_0 = T_1 = T_2 = T_3 = 18$ rows.

BT1487 classifies the branch symmetries behind those four triangles. The full triangle-partition stabilizer is S_4 of order 24, giving the Fano point reading $7 \cdot 24 = 168$. The physically used square subgroup is D_4 of order 8, giving the Fano flag reading $21 \cdot 8 = 168$. Thus the active detector-bin count is Fano-native, not merely an optical count.

BT1488 updates the splice cascade: the E6 claim table v2 from BT1484 supersedes the BT1472 table as the preferred paper table, and the preferred exact finite insert packet is now BT1474, BT1480, BT1482, BT1484, BT1486, and BT1487. The Golden Quartic/Moebius-ball equation fill remains blocked pending rendered formula transcription; no public source-code repository is assumed.

5.4 BT1489–BT1491 row symmetry and Fano–E6 square

BT1489 lifts the S_4 , D_4 , and V_4 branch actions from triangle labels to the actual ABI v2 row layer. All 24 elements of S_4 , all 8 elements of the square D_4 subgroup, and all 4 V_4 translations act as honest permutations of the 72 active/guard value rows. The lift preserves the C_3 channel, row kind, qutrit value, guard slot, and column formula, so the branch symmetry is now a decoder-row symmetry.

BT1490 is the new finite-geometry lock:

$$24 = 4 \cdot 6 = 3 \cdot 8, \quad 72 = 3 \cdot 24, \quad 168 = 7 \cdot 24 = 21 \cdot 8, \quad 81 = 72 + 9.$$

The shared 24-state fiber is V_4 branch times six row-value slots on the ABI side, and three local Fano arms times the D_4 flag stabilizer on the Fano side. Thus the same finite fiber feeds the E6/CSS 72-sector, the 81-closure after adjoining the $q^2 = 9$ firewall gap, and the Fano-native 168 active-bin bus.

BT1491 makes the paper splice idempotent. The main holonet paper now imports the BT1480–BT1491 exact finite packet through one controlled block; rerunning the splicer replaces that block rather than duplicating it. The claim firewall remains intact: this is an exact finite ABI and incidence statement, not a detector calibration, optical-layout claim, or imported Golden Quartic/Moebius particle interpretation.

5.5 BT1492–BT1494 canonical Fano fiber and pulse release lock

BT1492 removes the last arbitrary index from the 24-state bridge. In the Fano plane, $GL(3, 2)$ has order 168. Fixing one point leaves a point stabilizer of order 24, and this stabilizer acts as the full S_4 on the four Fano lines not through that point. Fixing one incident flag leaves an order-8 flag

stabilizer with the D_4 order profile. Therefore the shared fiber is canonical:

$$24 = 3 \cdot 8,$$

namely three Fano lines through the anchor point times the D_4 flag stabilizer.

BT1493 compiles those row symmetries down to the existing physical holonet interfaces. A row action first selects a BT1411 detector slot, then hands that slot to the BT1374 rule mirror_slot mod 4, and finally occupies the matching BT1407 Hesse epilogue lane for its C_3 channel and row-value slot. All 24 S_4 actions compile across all 72 rows, while the order-8 D_4 square subgroup is the native square-pulse layer.

BT1494 repairs the broad photonic-QEC release lock. The legacy tests expected selected **PART_*** artifacts at the repository root, while the preserved copies lived under **manuscripts/parts**. The repair restores those root release-lock targets and regenerates the live scheduler, harmonic bus, fusion splice, and projective screen/bulk/QEC artifacts.

6 The software: braids, teleported gates, universality

6.1 Exact braid words

In the anyonic Fibonacci representation $\sigma_1 = \text{diag}(\zeta^4, -\zeta^2)$ ($\zeta = e^{i\pi/5}$) one has *exactly* $\sigma_1^5 = Z$ and $\sigma_1^{10} = I$ — not merely projectively (verified in the cyclotomic ring $\mathbb{Z}[\zeta_{10}]$). In the dual ($\pm$) encoding of the local $K_{3,3}$ cycle code $\mathbb{F}_2^4$ this makes every register bit-flip an exact five-letter braid word, and the flat global register (the unique connected gluing with trivial holonomy across the atlas) carries the action coherently: a zero-Berry-phase memory. *Witnesses: bt740, bt741.*

6.2 The photon is its own gate

A gate U is stored as the self-entangled state $(I \otimes U)|\Omega\rangle$; Bell-projecting an input register against the photon’s *past* register outputs U (times a known Pauli correction) on the *future* register — verified for all nine outcomes. The machine’s gate library is a library of self-entangled states; executing software means interfering the photon with its own history. *Witness: bt821.*

6.3 Instruction set architecture

The practical instruction set is finite:

$$\mathcal{I}_{\text{holo}} = \{F_p, F_f, S_p, S_f, CX_{p \rightarrow f}, \sigma^5 = Z, D_{12}\text{-mirror}, M_{36}\text{-magic}\}.$$

F and S generate single-qutrit Clifford frame changes; CX couples past to future; $\sigma^5 = Z$ gives exact braid-register flips; the D_{12} mirror operation moves a packet through the atlas without confusing it with the cyclic selector clock; and M_{36} denotes injection from the thirty-six magic rays of §9. The Clifford part normalizes Pauli frames and is efficiently classically trackable; it is not universal by itself. Universality starts only when the Clifford runtime can consume the intrinsic magic sector. This is the engineering meaning of “matter equals magic”: the same thirty-six rays that appear as the matter shell are the non-Clifford fuel for the processor.

6.4 The minimal classical shadow

The smallest classical headline compatible with the substrate is the pair

$$(\lambda, q) = (2, 3).$$

This is the same two-state, three-symbol signature made famous by the Wolfram–Smith weakly universal Turing machine. The present paper does not use that result as the proof of quantum universality; the proof is the Clifford closure plus magic/contextuality [17, 16]. The point is sharper: the classical shadow of the photonic holonet lands on the same minimality boundary. The quantum machine has forty rays and an $81 = q^4$ protected register, but its smallest visible finite-state control alphabet is already the substrate pair 2, 3.

6.5 The Universality Theorem

Theorem 6.1 (Clifford completeness from optics). *The symplectic images of the photon’s physical elements — the tritter F_3 on either register, the quadratic phase plate S , and the delay-conditioned CX — generate all of $\text{Sp}(4, \mathbb{F}_3)$ (closure of order 51840, exact). Hence tabletop optics realize the complete two-qutrit Clifford group modulo Paulis and phases. Witness: `bt825_universality_theorem.py`.*

Theorem 6.2 (Universality). *Clifford completeness, together with the machine’s intrinsic magic supply (§9) and the strict positivity of its contextual fraction (1/10, exact), implies universal quantum computation by magic-state injection; for qutrits, contextuality is necessary and sufficient for magic distillation (Howard–Wallman–Veitch–Emerson) [3]. One self-entangled photon on the $W(3, 3)$ mesh is a universal quantum computer whose transport is its gate action.*

7 Fractal scaling: the computer is the network

The single-core theorem is not the end of the architecture. The substrate is self-similar: any node may be replaced by a fresh copy of the same forty-node holonet, with the parent line/point incidence becoming the routing grammar for the child layer.

Definition 7.1 (Recursive holonet). *Let $\mathcal{H}_0$ be one self-entangled photonic qutrit. Define $\mathcal{H}_n$ recursively by replacing each point of one copy of $W(3, 3)$ with a copy of $\mathcal{H}_{n-1}$. Thus $\mathcal{H}_1$ is the single-core machine of this paper and $\mathcal{H}_n$ is a 40-ary fractal computer/network.*

Theorem 7.2 (BT827 holonet scaling law). *A level- n holonet has*

$$N_n = 40^n$$

photonic leaf cores and

$$I_n = 1 + 40 + \cdots + 40^{n-1} = \frac{40^n - 1}{39}$$

total $W(3, 3)$ instances. Its total edge-qutrit slots are $240I_n$, its directed transport slots are $480I_n$, its chart routers are $540I_n$, its apartment links are $1620I_n$, its mirror slots are $2160I_n$, and its local Clifford runtime atoms are $51840I_n$. Worst-case reversible routing across the recursive address uses at most

$$8n = 8 \log_{40} N_n$$

moves, where $8 = 3 + 5$ is one in-chart hypercube diameter plus one chart-web apartment diameter. Durable commits use the distinct Császár/tomotope clock

$$T(0) = 1, \quad T(g) = 4(7^g - 1) \quad (g \geq 1).$$

Witness: `bt827_holonet_fractal_architecture.py`.

level n	leaves 40^n	$W(3, 3)$ instances	mirror slots	runtime atoms	route bound
1	40	1	2160	51840	8
2	1600	41	88560	2125440	16
3	64000	1641	3544560	85069440	24
4	2560000	65641	141784560	3402829440	32




Figure 6: Fractal scaling schematic. The drawn fans show only a few children; the rule is exact substitution of forty child holonets at each substrate point. Routing grows logarithmically in the number of photonic leaves because each recursive digit costs at most eight reversible moves.

This resolves the computer-engineering question. A conventional machine grows by adding processors connected by a separate network. The holonet grows by adding copies of the same object; each copy contains its own processor, router, memory, clock, and packet format. There is no von Neumann bus to saturate. At every scale, the edge set is both communication bandwidth and entangling-gate availability, the chart atlas is both routing table and instruction decoder, and the Steinberg 81-sector is both protected memory and fault-tolerant logical space.

8 Runtime closure: compile, monitor, commit

The architecture now has an executable operating-system loop.

address word $\longrightarrow$ packet route $\longrightarrow$ sentinel monitor $\longrightarrow$ durable tomatope commit.

BT828–BT831 make each arrow finite and testable.

Theorem 8.1 (BT828 packet compiler). *A recursive address word with digits in $\{0, \dots, 39\}$ compiles to a sequence of digit packets*

(Q_3 -XOR hops, apartment hops, D_{12} mirror slot, C_{12} phase, tomatope block).

Each digit uses at most three XOR hops and at most five apartment hops, hence at most eight reversible moves. The stress route

$$(0, 7, 14, 21, 28, 35) \longrightarrow (39, 32, 25, 18, 11, 4)$$

compiles to 47 reversible moves, below the level-six bound 48. Witness: `bt828_holonet_packet_compiler.py`.

Theorem 8.2 (BT829 sentinel monitor). *The projector onto the $g = 15$ adjacency sector of $W(3, 3)$ has exact entry profile*

$$P_{-4}(x, x) = \frac{3}{8}, \quad P_{-4}(x, y) = \begin{cases} -\frac{1}{8}, & x \sim y, \\ \frac{1}{24}, & x \approx y. \end{cases}$$

Consequently a full context line K_4 is sentinel-invisible, while localized, mirror, and shell faults activate the sentinel with exact energies:

<i>fault</i>	<i>energy</i>	<i>normalized fraction</i>
<i>single point</i>	3/8	5/13
<i>adjacent edge</i>	1/2	5/19
<i>nonedge mirror pair</i>	5/6	25/57
<i>gauge shell $N(x)$</i>	6	5/7
<i>matter shell</i>	27/8	5/13

The immune system is therefore not generic noise detection: it is blind to legal context-line traffic and tuned to off-context routing and shell congestion. Witness: `bt829_fault_sentinel_monitor.py`.

Theorem 8.3 (BT830 two-phase commit clock). *The runtime has two clocks. The reversible prepare phase finishes in at most $8n$ moves. The durable commit phase waits for*

$$T(n) = 4(7^n - 1).$$

For $1 \leq n \leq 24$, $T(n)$ is always divisible by 24 and by 8, so it is aligned with both the local runtime stabilizer and the single-digit route window. But it is not always divisible by the full $8n$ route epoch: the first desynchronization is $n = 5$, with remainder $24 = f$. Thus the commit layer is deliberately slower and cover-like; it is not merely the route clock renamed. Witness: `bt830_two_phase_commit_clock.py`.

8.1 The infinite-cover warning

The tomatope contributes a harder boundary than the earlier packet picture showed. Monson, Pellicer, and Williams prove that the tomatope has infinitely many distinct minimal regular covers; for coprime odd $p, q > 1$ there are minimal covers R_p, R_q of the tomatope such that neither covers the other, and neither covers the special R_2 cover [20]. They also record the toroidal cover arithmetic

$$|\text{Mon}(Q_k)| = 36864 k^6, \quad Q_k : (4, 24, 32, 8, 4)k^3$$

for vertices, edges, triangles, tetrahedra, and octahedra.

Architecturally this is a big correction. The 48-block tomatope middle layer is a stable local ABI, but the global regular cover is not unique. Durable storage must therefore carry a *cover index* k as an implementation gauge. Fast routes compile to the same invariant packet, while persistent storage may choose one minimal cover family without invalidating the other non-comparable families. *Witness: `bt831_tomotope_minimal_cover_architecture.py`.*

8.2 Cover-indexed storage, rerouting, and guard bands

The infinite-cover warning would be useless as engineering unless it changed the runtime. BT832–BT834 make it operational. The principle is that the local tomatope packet is the ABI and the chosen regular cover is the storage backend.

Theorem 8.4 (BT832 cover-indexed durable storage). *For every supported cover index k , durable storage is the lifted packet space*

$$\{0, \dots, 47\} \times \mathbb{Z}_k^3, \quad |\text{slots}| = 48k^3,$$

with projection back to the fast-route ABI by forgetting the $\mathbb{Z}_k^3$ coordinate. The kernel to the base tomatope packet has order

$$k^6 = (k^3)^2,$$

matching the Monson–Pellicer–Williams toroidal cover law $|\text{Mon}(Q_k)| = 36864k^6$. Changing k changes durable capacity and kernel size, but not the packet program seen by the compiler. Tested cover gauges $k = 3, 5, 7, 11, 13$ all reduce exactly to the same 48-block ABI, and larger k reduces the load fraction of the same stress program. Witness: `bt832_cover_indexed_durable_storage.py`.

Theorem 8.5 (BT833 sentinel-aware packet rerouting). *The BT829 projector supplies a compiler-visible cost function. Before the durable commit, each digit route may insert up to two W33 waypoint points and minimize*

sentinel energy first, reversible move count second.

The extra reversible moves are charged to the slower BT830 commit phase, not to the BT827 fast $8n$ route bound. In the level-six stress program the direct route uses 47 moves and total sentinel energy 4; the sentinel-aware route uses 58 moves, still far below the 470592-tick commit window, and lowers the total sentinel energy to 2. Each tested digit route strictly lowers its own $g = 15$ energy; adjacent endpoints can be context-completed to zero, and nonedge mirror pairs drop below their direct $5/6$ activation. Witness: `bt833_sentinel_aware_packet_rerouter.py`.

Theorem 8.6 (BT834 desynchronization guard-band arithmetic). *The two-phase clock has an exact number-theoretic boundary. Full route-epoch synchronization is*

$$8n \mid T(n) = 4(7^n - 1),$$

equivalently

$$2n \mid 7^n - 1.$$

For every odd prime power $p^a \mid n$ with $p \neq 7$, this requires $\text{ord}_{p^a}(7) \mid n$; if $7 \mid n$, synchronization is impossible because the base is not coprime. The first failure is $n = 5$: $\text{ord}_5(7) = 4 \nmid 5$, and the remainder is

$$T(5) \bmod 40 = 24 = f.$$

Thus the first cover-index where durable storage and the full route epoch separate is not empirical drift; it is the $f = 24$ guard band of the tomatope-backed runtime. Witness: `bt834_desync_guard_band_arithmetic.py`.

Theorem 8.7 (BT838 tomatope Wythoff runtime ladder). *The Grünbaum–Coxeter/tomatope operation tables give the local tomatope counts*

$$\text{original, rectified, truncated, maximal expanded, omnitruncated} = 4, 12, 24, 48, 96.$$

These are exactly the runtime packet-growth layers:

$$4 = \mu \quad (\text{vertex carriers}), \quad 12 = k \quad (\text{edge axes}), \quad 24 = f \quad (D_{12} \text{ lift}),$$

$$48 = 2f \quad (\text{BT814 packet ABI}), \quad 96 = 4f \quad (\text{half-flag boundary}).$$

The doubled omnitruncated boundary gives the verified full packet flags, $2 \cdot 96 = 192$. After a cover lift this becomes

$$48k^3 = \text{BT832 lifted packet capacity}, \quad 192k^3 = |W_k|,$$

the W_k order already recorded in the tomatope cover architecture. Thus the Wythoff operations are not decorative variants; they are the compiler's packet expansion states:

$$\text{carrier} \rightarrow \text{edge-axis} \rightarrow \text{mirror lift} \rightarrow \text{packet ABI} \rightarrow \text{flag boundary}.$$

Witness: `bt838_tomotope_wythoff_runtime_ladder.py`.

9 The fuel: matter = magic

Theorem 9.1 (The triality). *In the photon's two-qubit reading, exactly the 4 basis rays are stabilizer states and all 36 ω -rays are magic (no stabilizer state has support 3). Consequently three classifications coincide exactly, on rays [4, 36] and on contexts {4:1, 1:12, 0:27}: the holonet vacuum split, the Schmidt (self-entanglement) strata, and the magic strata. The matter shell is the machine's non-classical resource sector. Witness: `bt822`.*

Theorem 9.2 (Three grades and the cube inside). *Stabilizer fidelity splits the 36 magic rays into grades*

$$\underbrace{8}_{2^q} \text{ deep } \left(F = \frac{2+\sqrt{3}}{6} \right), \quad \underbrace{24}_f \text{ mid } \left(F = \frac{5+2\sqrt{3}}{12} \right), \quad \underbrace{4}_\mu \text{ shallow } \left(F = \frac{3}{4} = \frac{q}{\mu} \right),$$

the shallow grade sitting exactly at the Werner decoherence threshold. The 8 deep rays are precisely the point set of a hyperbolic polar pair (the 24-cell vacuum axis), and the 27 fully-magic contexts stratify by grade signature as $1+8+12+6$ — the cell census of the cube. Witnesses: `bt822`, `bt823`.

Theorem 9.3 (Exact contextuality budget). *The maximum number of contexts a classical $\{0,1\}$ valuation can satisfy exactly-once is exactly $36 = (q!)^2$ (complete branch-and-bound over all 2^{40} markings): classicality saturates at the spread count, the deficit is exactly $\mu = 4$, and the contextual fraction is exactly $1/10 = 1/\Phi_4$. A perfect valuation would be an ovoid, and $W(3,3)$ has none: the true independence number is $\alpha = 7 = \Phi_6$, attained only by the beacon heptads of §12, against the Lovász bound $\vartheta = 10$ — the $\vartheta - \alpha$ gap $3 = q$ is the combinatorial face of Kochen–Specker. Witnesses: `bt818`, `bt823`.*

Theorem 9.4 (No distillation factory: the fuel is the code). *In every standard fault-tolerant architecture the dominant cost is magic-state distillation — a separate factory producing 10^3 – 10^6 raw magic states per logical non-Clifford gate, feeding a Clifford error-correction layer. Here that cost is not paid twice: matter = magic. The 36 magic rays $= (q!)^2$, graded $8+24+4 = \{2^q, f, \mu\}$, are the matter shell — the same D_4 /GKP code register (the $[[240, 81, 4, 3]]_3$ Steinberg module of §10) that protects the logical data. The fuel is therefore structural, and it is genuine magic: the qutrit Strange state $|S\rangle = (|1\rangle - |2\rangle)/\sqrt{2}$ has mana $\log(5/3)$, the single-qutrit maximum (computed from the discrete Wigner function; stabilizer states have mana 0). Its standing density is the contextual fraction $1/\Phi_4 = 1/10$, replenished by the same cycle that maintains the code; the Kochen–Specker deficit $\vartheta - \alpha = 10 - 7 = q = 3$ is the magic per round; and because $W(3,3)$ has no ovoid the magic can never be gauged away. So the dominant cost of fault tolerance is reorganised into the matter=magic identity. And the fuel gauge is a cosmological constant: the $\Phi_4 = 10$ that meters*

the magic density is the same $\Phi_4 = q^2 + 1$ that factors the de Sitter selection cubic (§C.1, Face 1) and the demonstrator readout (Face 6). The machine’s computational fuel and the world’s vacuum selection are one substrate constant. Witness: `analysis/w33_magic_economy.py`.

The resource accounting. Matter = magic has a measurable payoff. Define the non-Clifford premium $P = \text{cost}(\text{non-Clifford})/\text{cost}(\text{Clifford})$. In the surface-code baseline a non-Clifford gate needs a distilled magic state, so $P \sim O(d^3)$ (the distillation factory, $\sim 10^3\text{--}10^4$ qubit-rounds at useful code distances d , occupying 30–90% of the device — the dominant cost of fault tolerance). On this substrate $P = 1$: a non-Clifford gate is a magic injection from the matter shell the code already maintains, costing one error-correction cycle, exactly like a Clifford gate. The distillation factory and its $O(d^3)$ premium are absent; the saving scales as d^3 , and the non-Clifford gate rate equals the Clifford rate (magic per round = $\vartheta - \alpha = q = 3$). Non-Clifford gates are as cheap as Clifford gates because the fuel *is* the code. (Witness: `analysis/w33_magic_resource_accounting.py`.)

The clock renews the magic. What keeps the structural magic from precessing into a Clifford frame is the clock’s aperiodicity. The Boerdijk–Coxeter twist $\theta = \arccos(-2/3)$ — the $q = 3$ simplex angle, $-2/3 = -(q-1)/q$ — has θ/π irrational (Niven, since $\cos \theta = -2/3 \notin \{0, \pm\frac{1}{2}, \pm 1\}$), so the stroboscopic phases $\{n\theta \bmod 2\pi\}$ never repeat, are Weyl-equidistributed (star discrepancy $\rightarrow 0$), obey the Steinhaus three-gap law, and never land exactly on a stabilizer angle $2\pi k/q$. A drive is Clifford-trackable only if it closes a finite cycle (θ a rational multiple of 2π); the irrational q -simplex angle never closes, so each tick injects a fresh cubic (non-Clifford) rotation — the magic is continuously replenished. This is the *temporal* partner of the spatial no-ovoid result: the same $q = 3$ forbids gauging the magic away in space (no ovoid, $\alpha = 7 = \Phi_6 < \vartheta = 10 = \Phi_4$) and in time (no recurrence). The time quasicrystal does not merely clock the machine; it fuels it. (Witness: `analysis/w33_clock_magic_renewal.py`.)

And the code is holographic. The matter shell is not just a code but a *holographic* one. From any pole the 40 rays split $1 + 12 + 27 = \text{self} + \text{gauge boundary} + \text{matter bulk}$, and every bulk ray is screened by exactly $\mu = 4$ boundary rays — the strongly-regular μ -parameter, equal to the code distance $d = 4$. So the bulk matter is reconstructed from the boundary gauge *redundantly*: erase up to $d-1 = 3$ of a bulk ray’s boundary neighbours and it is still recoverable. The causal diameter is 2 (every bulk ray one hop from the boundary, two from the pole), the RT-like depth; and the logical register H_1 is the irreducible Steinberg module ($\dim 81 = q^4$), so there is no local logical operator — logical information lives only nonlocally on the boundary. Redundant bulk-from-boundary recovery, redundancy = distance, an RT-like causal depth, and no local recovery are the holographic signature: the “holo” in holonet is earned. (Witness: `analysis/w33_holographic_code.py`.)

Theorem 9.5 (The architecture is a self-fueling holographic memory). *The three properties above are one object. The matter shell is simultaneously the holographic code (it stores the 81 logical qutrits of the Steinberg module, bulk recovered from the boundary at redundancy $\mu = 4 = d$), the magic fuel (its $36 = (q!)^2$ rays are the non-Clifford resource, so the non-Clifford premium is $P = 1$ — no distillation factory), and the system the irrational Boerdijk–Coxeter clock keeps non-Clifford ($\theta = \arccos(-2/3)$, θ/π irrational). So the machine needs no distillation factory (the fuel is the code), no separate fuel store (the magic is the error-correction redundancy), and never recurs (the irrational clock refreshes the magic): a holographic quantum memory whose error-correction is its magic supply, kept alive by an irrational clock — all controlled by $q = 3$ ($\mu = 4 = d$, $1/\Phi_4 = 1/10$, $\cos \theta = -(q-1)/q$). Moreover the entanglement is geometric: a discrete Ryu–Takayanagi law holds, the entropy of a boundary region being the minimal graph cut, with*

a single ray's RT area the degree $k = 12$, the minimal RT surface (one bulk qutrit) the code distance $d = 4$, the gauge|matter cut $108 = \mu \cdot 27$, and an area law (sub-volume). (Witnesses: `analysis/w33_self_fueling_memory.py`, `analysis/w33_holographic_rt.py`.)

Theorem 9.6 (The memory is the de Sitter spacetime — and its honest price). *The loop closes: the holographic memory's bulk is the de Sitter spacetime it runs in. The collinearity graph $W(3, 3)$ has positive Ollivier–Ricci curvature — computed by optimal transport, $\kappa = 1 - W_1(m_x, m_y) = 1/6 = 2/k$ — the discrete signature of de Sitter / a positive cosmological constant, and it satisfies the Gibbons–Hawking/Gauss–Bonnet closure $E\kappa = 240 \cdot \frac{1}{6} = 40 = v$ (Face 1). So the same positively-curved $W(3, 3)$ is the holographic memory, the RT horizons are the de Sitter horizons, the Boerdijk–Coxeter clock is de Sitter time, and inflation is the expansion: computation and spacetime are one geometry, the machine computing itself because its memory is the de Sitter universe it inhabits. (Witness: `analysis/w33_memory_is_desitter.py`.)*

The honest price. Self-fueling is not a free lunch but a space-for-rate trade. Removing the distillation factory saves $\sim d^3$ qubits (30–90% of a standard device), but the magic is replenished only at the contextual-fraction rate $1/\Phi_4 = 1/10$ per cycle, so the non-Clifford throughput is capped: a circuit whose non-Clifford fraction $f \leq 1/\Phi_4 = 1/10$ runs with no factory and no slowdown (a pure win), while $f > 1/10$ is throttled by $\Phi_4 f = 10f$ (or needs $10f$ times more cores). The remaining costs are routine matter-shell error correction (paid anyway), the clock renewal bandwidth, and the $q = 3$ fresh-magic-per-round deficit. The advantage is real and bounded: the d^3 distillation space is traded for a non-Clifford rate cap at $1/\Phi_4$. (Witness: `analysis/w33_self_fueling_cost.py`.)

10 Memory, protection, and the immune system

Theorem 10.1 (The code register is the Steinberg module). *The companion paper's CSS code $[[240, 81, 4, 3]]_3$ stores its 81 logical qutrits in H_1 of the $W(3, 3)$ 2-complex (40 vertices, 240 edges, 160 in-line triangles). Complete character computation over all 25920 group elements gives*

$$\langle \chi_{H_1}, \chi_{H_1} \rangle = 1.$$

Therefore $H_1 \otimes \mathbb{C}$ is irreducible of dimension 81, hence equal to the unique 81-dimensional irreducible — **the Steinberg module** — with $\chi_{H_1}(g) = \# \text{fixflags} - \# \text{fixpoints} - \# \text{fixlines} + 1$ (the Solomon–Tits alternating sum) verified for every g . The QECC logical space and the Schur-protected register below are one object: any equivariant logical error is a scalar, so only symmetry-breaking noise can touch the code, and the sentinel/clock layers are exactly the symmetry monitors. The homological dictionary then completes: H_0 is trivial (the vacuum pole), H_1 is Steinberg (the register), and H_2 (dim 40, one tetrahedron-boundary 2-cycle per line) is the sign-twisted line module — a fixing symmetry acts on a line's 2-cycle by the parity of its induced S_4 action (verified for all 25920 elements; not the plain permutation module): the top homology is the oriented timetable carrier, feeding the chirality ledger. Witnesses: `bt861_code_register_is_steinberg.py`, `bt862_h2_is_the_line_module.py`.

Corollary 10.2 (Three generations from Steinberg vanishing). $\chi_{\text{St}}(g) = 0$ iff $3 \mid \text{ord}(g)$ (verified for all 25920 elements). Hence any order-3 symmetry of the substrate splits the matter register into eigenspaces of exactly $27 + 27 + 27$ — the three-generation decomposition is forced for every choice of triality element, not selected — and every order-9 element (5760 of them) refines it into nine 9-dimensional sub-generations. In the code's own characteristic the parameters survive: $\text{rank}_3 \partial_0 = 39$, $\text{rank}_3 \partial_1 = 120$, $\dim H_1(\mathbb{F}_3) = 81$, so $[[240, 81, 4, 3]]_3$ is exact over $\mathbb{F}_3$. The number of fermion generations is the defining-characteristic degeneracy of the protected register. Witness: `bt863_generations_from_steinberg_vanishing.py`.

Theorem 10.3 (The dual-torsor Steinberg compiler). *Let $H_{27} = 3_+^{1+2}$ be the canonical Heisenberg O_3 of the point parabolic and let $A_{27} = \mathbb{F}_3^3$ be the canonical translation O_3 of the Bell-line parabolic. Their 27-point and 27-line shells are regular torsors (BT858). On the protected register,*

$$\text{St} \downarrow_{H_{27}} = 3 \mathbb{C}[H_{27}], \quad \text{St} \downarrow_{A_{27}} = 3 \mathbb{C}[A_{27}],$$

because both restricted characters are exactly $\{81^1, 0^{26}\}$. The statement survives in the code's own field and is constructive:

$$H_1(W(3, 3); \mathbb{F}_3) \downarrow_{H_{27}} \cong \mathbb{F}_3[H_{27}]^{\oplus 3}, \quad H_1(W(3, 3); \mathbb{F}_3) \downarrow_{A_{27}} \cong \mathbb{F}_3[A_{27}]^{\oplus 3}.$$

For each torsor the verifier finds three cycle seeds whose 27 group translates add $27 + 27 + 27$ independent classes modulo the 120-dimensional boundary space. Thus the same logical memory has a flat program basis and a curved state basis: A_{27} is the commuting three-trit address compiler, while H_{27} is the noncommuting qutrit-phase compiler.

The center $Z(H_{27}) = \langle z \rangle \cong C_3$ selects a canonical triality. Over $\mathbb{C}$ its central-character blocks are

$$H_1 = H_{1,1} \oplus H_{1,\omega} \oplus H_{1,\omega^2}, \quad \dim H_{1,j} = 27.$$

The trivial block contains the nine linear characters with multiplicity 3; the two nontrivial blocks contain the conjugate 3-dimensional Schrödinger representations with multiplicity 9. Over $\mathbb{F}_3$, where $x^3 - 1 = (x - 1)^3$, the phases coalesce into a canonical unipotent flag. With $N = z - I$,

$$(\text{rank } N, \text{rank } N^2, \text{rank } N^3) = (54, 27, 0),$$

so

$$0 \subset \ker N \subset \ker N^2 \subset H_1, \quad \dim = (0, 27, 54, 81),$$

has three successive quotients of dimension 27. Equivalently, z acts by 27 Jordan blocks of length 3. Complex generations are phase sectors; native qutrit generations are the three layers of one depth-3 memory stack.

The freeness is canonical, but the displayed orbit bases are not: H_{27} and A_{27} are nonisomorphic, and a basis-to-basis matrix still requires a chosen point, line, shell origin, and cycle seeds. No canonical intertwiner is claimed. Witness: `bt865_dual_torsor_steinberg_compiler.py`.




Figure 7: The dual-torsor compiler. Flat program translations and noncommutative state translations give two free rank-3 module coordinates on one Steinberg-protected register. The Heisenberg center reads as three phase sectors over $\mathbb{C}$ and as a three-step unipotent memory flag over $\mathbb{F}_3$.

Theorem 10.4 (Steinberg memory). *The Levi-graph cycle space and the chart-overlap 81-sector are both irreducible and equal to the Steinberg representation of $\mathrm{PSp}(4, 3)$ — the Solomon–Tits homology of the building, whose canonical basis is the 81 apartments through any chamber. Schur’s lemma then forces the uniqueness of the chart–cycle bridge. Restricted to a Sylow 3-subgroup the module is free of rank one: the protected memory is one free generator over the substrate’s ternary symmetry. Witnesses: `bt742`, `bt744`.*

Corollary 10.5 (The oriented timetable spectrum). *Let $P = 3^3:S_4$ be the index-40 line parabolic. The ordinary line permutation module and the oriented top-homology module are the two inductions of its two linear characters:*

$$\mathrm{Ind}_P^{\mathrm{PSp}(4,3)}(1) = 1 \oplus 15 \oplus 24, \quad H_2 = \mathrm{Ind}_P^{\mathrm{PSp}(4,3)}(\mathrm{sign}_{S_4}) = 5_\omega \oplus 5_{\omega^2} \oplus 30.$$

The two 5-dimensional characters are complex conjugates over $\mathbb{Q}(\zeta_3)$; the 30 is rational. Under the outer extension $W(E_6) = U_4(2):2$, the conjugate pair is the restriction of one irreducible 10-dimensional representation, whereas the 30 has two inequivalent extensions. Thus the oriented timetable header reads as a Weyl-visible chiral channel 10 plus a payload $30^\pm$ whose outer parity is not fixed by the projective module alone.

The completed homology spectrum is

$$H_0 = 1, \quad H_1 = \mathrm{St}_{81}, \quad H_2 = 5_\omega \oplus 5_{\omega^2} \oplus 30,$$

and resolves the Euler charge representation by representation:

$$1 - 81 + (5 + 5 + 30) = -40 = -v.$$

BT867 resolves the tempting cache comparison negatively: neither 162-cache branch is one of the two 5-dimensional sectors. The branches are isomorphic copies with identical character overlaps; the outer involution acts on their multiplicity label. The possible identification of the two 30-extensions with BT857’s local gauge bit remains open. Witness: `analysis/bt866_h2_oriented_irreducible_decomposition.py`.

Theorem 10.6 (The cache/transport extension boundary). *Let $H = 3^3:S_4$ be the order-648 Bell-line parabolic. The two 162-slot BT859 cache branches have conjugate cyclic stabilizers and are therefore isomorphic transitive H -sets,*

$$\mathcal{C}_L \cong H/C_4 \cong \mathcal{C}_R.$$

Their permutation characters are equal. In particular, their scalar products with the restrictions of the three BT866 characters $(5_\omega, 5_{\omega^2}, 30)$ are both $(1, 1, 8)$: each cache sees both conjugate 5-sectors equally. “Left” and “right” are copy labels, not internal spectral chiralities.

The parabolic contains one conjugacy class of 81 cyclic C_4 subgroups, and

$$N_H(C_4) = D_8, \quad H/C_4 \longrightarrow H/D_8, \quad 162 \longrightarrow 81$$

is the canonical free two-sheet cache cover. Both copies cover this same stabilizer base. On their union GAP computes

$$C_{\mathrm{Sym}(324)}(H) \cong D_8, \quad 324 = 81 \times 4 = 81 \times 2_{\mathrm{deck}} \times 2_{\mathrm{cache}},$$

with exactly 81 commutant orbits of size four. Thus the four-state fiber carries the symmetry of a square, not merely a numerical pair of bits.

This cache carrier is not the older ternary 162-sector. On one cache fiber the deck swap and its two projectors are

$$D = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad P_{\pm} = \frac{I \pm D}{2}, \quad \mathbb{F}_3^2 = \text{im } P_+ \oplus \text{im } P_-.$$

Since 2 is invertible in $\mathbb{F}_3$, the cache splits; indeed $(D - I)^2 = D - I$. By contrast, packet transport uses

$$N = \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}, \quad N^2 = 0, \quad (I + N)^3 = I,$$

and tensors this indecomposable fiber with the 81-register to obtain the non-split sequence $0 \rightarrow 81 \rightarrow 162 \rightarrow 81 \rightarrow 0$. No change of basis can identify the idempotent cache difference with the nilpotent transport shift. The architecture therefore separates a reversible address/control plane from a memory-bearing ternary temporal data plane by extension class, not by convention. Witness: `analysis/bt867_cache_split_transport_nonsplit_boundary.py`.




Figure 8: The local cache cell and the split/non-split boundary. The D_8 square fiber chooses a reversible address over one of 81 routes; the independent ternary nilpotent updates state while retaining temporal memory. Equal 162-dimensional counts do not imply equivalent modules.

The error layer is native: the $[[240, 81, 4, 3]]_3$ CSS code on the edge qutrits has all four parameters substrate-primitive; the dual Császár/Szilassi cross-check supplies parity-free error detection; and the web's unique non-Ramanujan eigenspace — dimension $15 = g$ — is the spectral sentinel where drift appears first.

11 Time: three clocks and a quasicrystal

Theorem 11.1 (Clock trichotomy). *Of the program's three natural clocks, exactly one is internal: $\mathbb{Z}_{12}$ (the rectangle clock, $12 \mid 51840$) runs the selector and duo machinery, while $\mathbb{Z}_7$ (Császár) and $\mathbb{Z}_{13}$ (Singer) do not divide the group order and act as external references — both with decimal period $6 = q!$, so the digit expansion of $1/7$ is literally the internal clock's shadow modulo its duo bit. Witnesses: `bt774`, `bt803`, `bt807`, `bt819`.*

Theorem 11.2 (The Boerdijk–Coxeter time quasicrystal). *Driving the loop by the BC twist $\theta = \arccos(-2/3)$ per round trip yields a stroboscopic orbit that provably never repeats (Niven), with the Steinhaus three-distance signature at all substrate step counts and — the substrate's signature moment — exactly two gap lengths at $n = 30 = h(E_8)$, the BC ring length, where closure happens*

in S^3 (the 600-cell) rather than the phase circle (deficit 0.0158). This is the photonic sibling of the Fibonacci-drive dynamical topological phase. Witness: `bt820`.

Corollary 11.3 (The machine’s clock is the universe’s clock). *The same BC clock is the de Sitter beat of inflation (Face 7). Its twist $\cos \theta = -(q-1)/q$ is the discrete de Sitter angle, its ring length is the beat $30 = h(E_8)$ (the 600-cell’s 20 rings of 30), and the inflationary cosmology reads off that beat: the e -fold number is $N = 2(v - \Phi_4) = 2h(E_8) = 60$, exactly twice the beat, and the CMB spectral tilt is*

$$1 - n_s = \frac{2}{N} = \frac{1}{30} = \frac{1}{h(E_8)} = \frac{1}{(\text{clock beat})},$$

so the deviation of the primordial spectrum from scale invariance is the inverse of the clock’s period. The two structures even match in kind: the clock is a time quasicrystal (nearly periodic, never exact) just as inflation is quasi-de Sitter (nearly exponential, never exact) — both broken de Sitter by $q = 3$, the clock’s aperiodicity being the cosmology’s tilt. The architecture’s time and the universe’s expansion are one quasicrystal. Witness: `analysis/w33_clock_cosmology.py`.

Corollary 11.4 (A falsifiable clock fingerprint in the CMB). *If the clock couples to the inflaton, its quasiperiodicity imprints a quasi-periodic modulation on the primordial spectrum: peaks at log-period $2\pi/\theta \approx 2.73$ in $\ln k$ (with $\theta = \arccos(-2/3)$) whose spacings obey the Steinhilber three-gap law (two gap lengths near the BC ring $n = 30$) — the time-quasicrystal fingerprint, distinct from smooth slow roll (no peaks) and from single-frequency resonant inflation (a uniform comb, one gap). The log-frequency and the three-gap shape are fixed by $q = 3$; only the amplitude (the coupling) is free. A CMB feature search (Planck, LiteBIRD, CMB-S4) can detect the three-gap comb, bound the coupling, or — finding a clean single-frequency comb — disfavour the clock. (Witness: `analysis/w33_cmb_clock_signature.py`.)*

Theorem 11.5 (The arrow of time). *Without its two-trit feedforward record, the photon’s self-configuration loop is the completely depolarizing channel with unique fixed point I/q — exactly the Bell marginal; with the record it is unitary. Decoherence is self-forgetting, priced at $\log_3 q^2 = 2$ trits per cycle. Witness: `bt823`.*

12 Synchronization: beacons and schedules

The maximum classically coherent constellations are the *beacon heptads*: seven states with all pairwise overlaps exactly $1/3 = 1/q$ — a complete K_7 interference mesh, the Császár toroidal node realized in light. There are exactly 2880 heptads in a single transitive family; each has a $\mathbb{Z}_3 \times \mathbb{Z}_3$ symmetry fixing a unique *master beacon* and cycling two triads ($7 = 1 + 3 + 3$), and single-beacon swaps connect the family into its own exchange network: the reference frames of the network form a network. The measurement timetable is rigid: there exist exactly 36 complete schedules (partitions of the 40 states into 10 contexts) — an exhaustive result that simultaneously proves every spread of $W(3,3)$ is regular — with each context appearing in exactly 9 schedules. *Witnesses: `bt817`, `bt819`.*

Theorem 12.1 (Schedule-overlap law). *Two distinct schedules share exactly 1 or exactly $4 = \mu$ contexts: each schedule has 15 near partners (overlap 4) and 20 far partners (overlap 1), the two relation classes of the double-six diagonal $[1, 15, 20]$, with the exact counting identity $1 \cdot 20 + 4 \cdot 15 = 80 = 10 \times (9 - 1)$. Operationally: a cheap timetable switch preserves 4 of 10 calibrated contexts and a far switch only 1, so a controller hopping among near partners re-calibrates 6 contexts per hop — the timetable library is itself a metric space whose geometry is the double-six relation. Witness: `bt835_schedule_overlap_theorem.py`.*

The Grünbaum–Coxeter cells. Each schedule also has an internal shape. The full schedule stabilizer S_6 is 2-homogeneous on the 10 contexts, but its icosahedral core A_5 — the 600-cell level of the platonic ladder, and *exactly* the automorphism group of the hemi-dodecahedron $\{5, 3\}_5$ — splits the $\binom{10}{2} = 45$ context pairs as $45 = 15 + 30$, and the 15-orbit graph on the 10 contexts is the **Petersen graph**: the edge skeleton of the hemi-dodecahedron, the cell of Coxeter’s exceptional 57-cell. Dually, the hidden 6-set of the S_6 action realizes the hemi-icosahedron $\{3, 5\}_5$ — the cell of Grünbaum’s 11-cell — with incidence count $(6, 15, 10)$: six hidden points, fifteen duads, ten triple-splits = the contexts themselves. The exceptional polytopes’ groups do not embed in $\text{Sp}(4, \mathbb{F}_3)$, but they are substrate arithmetic ($|\text{PSL}(2, 11)| = 660 = k \cdot N_{\text{eff}}$, $|\text{PSL}(2, 19)| = 3420 = k \cdot g \cdot 19$, Heawood genus $H(19) = 20$ = the BC ring count, and $11 = k - 1$ is the Ihara prime of the zeta critical circle $|u| = 1/\sqrt{11}$): the full 11-cell and 57-cell live elsewhere, yet their *cells* are carried by every measurement timetable of this machine, glued by the substrate’s own group instead of $\text{PSL}(2, p)$. *Witness: bt836_gc_hemicells_in_spreads.py.*

Theorem 12.2 (The library is a classical geometry). *The timetable library carries its own rank-3 geometry, and the hemicell structure fibers over it with perfect uniformity: (i) every skew line pair lies in exactly 3 schedules, so the (schedule $\supset$ pair) flag count is $36 \times 45 = 1620$ — the apartment count of the Tits building — and 540×3 ; (ii) the near-partner graph on the 36 schedules is strongly regular $\text{SRG}(36, 15, 6, 6)$, the other classical rank-3 graph of $U_4(2) \cong \text{PSp}(4, \mathbb{F}_3)$, so the control plane (timetables) and the data plane ($W(3, 3) = \text{SRG}(40, 12, 2, 4)$) are the two rank-3 geometries of one group; (iii) each schedule stabilizer contains exactly 6 icosahedral A_5 cores ($216 = 6^3$ in all, each marking its unique schedule via its $[10, 30]$ line orbits) giving 6 distinct Petersen splits; each internal pair is a Petersen edge under exactly 2 of the 6 cores, hence every skew pair of $W(3, 3)$ has exactly $6 = 3 \times 2$ hemi-dodecahedral homes ($3240 = 540 \times 6$ Petersen flags). *Witness: bt837_schedule_library_geometry.py.**

Theorem 12.3 (Two compasses, twelve Petersens, $4 \cdot K_{10}$). $\text{PSp}(4, \mathbb{F}_3)$ has exactly two conjugacy classes of A_5 , both with normalizer S_5 and 216 conjugates, unfused by the outer automorphism, and both fiber over the same 36 schedules: every schedule stabilizer contains $12 = 6 + 6$ icosahedral cores — 6 duad cores (line signature $[10, 30]$, the Theorem-12.2(iii) cores) and 6 pentad cores (signature $[5, 5, 10, 20]$: transitive on the schedule, plus two distinguished pentads of 5 pairwise-disjoint lines outside it). The two 216-sets are non-isomorphic G -sets (rank 10 suborbits $[1, 5, 10, 10, 20, 20, 20, 30, 40, 60]$ vs $[1, 5, 10, 10, 20, 20, 30, 30, 30, 60]$); the Petersen flags form the coset geometry $\text{PSp}(4, \mathbb{F}_3)/D_4$ (3240 flags, dihedral stabilizer). There are 432 distinct pentads in two chiral orbits of 216 (stabilizer order 120); every core pairs one LEFT with one RIGHT pentad, and each pentad serves exactly one core. Pentads are maximal partial spreads: no schedule contains one. The two pentads of one core interlock as $K_{5,5}$ minus a perfect matching — mutual disjointness is forbidden by the overlap law (it would create a 0-overlap schedule) — and the deleted matching consists of 5 skew pairs, i.e. five hypercube charts; over the 216 cores these matchings cover the 540-chart atlas **exactly twice** ($1080 = 540 \times 2$): the routing fabric is readable off the pentad compasses, five charts per needle. Finally, all twelve cores split the 45 internal pairs as Petersen 15-orbits, and the twelve Petersen edge-sets cover every pair exactly $\mu = 4$ times:

$$4 \cdot K_{10} = 12 \text{ Petersens.}$$

Schwenk’s classical obstruction (no edge-partition of K_{10} into 3 Petersen graphs) dissolves at the substrate multiplicity μ . The numerical coincidence $3240 = \#\{\text{triangles of } Q\}$ is refuted at the G -set level (orbits $[360, 2880]$ vs transitive). *Witnesses: bt843_icosahedral_compass.py, bt844_double_five_and_six_petersens.py, bt845_pentad_exchange_and_chart_double_cover.py.*

Theorem 12.4 (The amalgamation ladder). *The 11-cell and 57-cell are universal abstract polytopes: the unique ones with hemi-icosahedral facets and hemi-dodecahedral vertex figures (group $L_2(11)$, no proper quotients) and vice versa ($L_2(19)$). Their incidence is icosahedral-subgroup geometry: each of $L_2(11)$, $L_2(19)$, $U_4(2)$ has exactly two conjugacy classes of A_5 , and in all three the distinguished incidence is intersection in D_{10} with valency 6 — the 11-cell’s vertex–facet relation, the 57-cell’s facet adjacency (the Perkel graph), and the substrate’s compass incidence ($216 + 216$, bipartite 6-regular, $216 \times 6 = 1296 = 6^4$ incident pairs = 36 schedules $\times$ 36 core pairs). The closures of one amalgam $A_5 *_{D_{10}} A_5$ (GAP):*

$$\langle A, B \rangle = \begin{cases} L_2(9) \cong A_6 \text{ (360)} & \text{in } U_4(2): \text{ inside exactly one schedule stabilizer,} \\ L_2(11) \text{ (660)} & \text{the 11-cell,} \\ L_2(19) \text{ (3420)} & \text{the 57-cell.} \end{cases}$$

One local icosahedral amalgam, a $\text{PSL}(2)$ -ladder of completions over $9 = q^2$, $11 = k - 1$ (the Ihara prime), 19 (Heawood, $H(19) = 20$). Universality plus Lagrange ($11, 19 \nmid 25920$) shows the substrate carries the complete local data of both exceptional polytopes while their rank-4 amalgamation is arithmetically forbidden inside $\text{Sp}(4, \mathbb{F}_3)$: the GC polytopes are the substrate’s sibling completions, not subobjects. One step up the universal ladder ($\{\{5, 3\}, \{3, 5\}_5\}$, 10006920 facets) the completion group is $J_1 \times L_2(19)$ — the first sporadic Janko group, whose involution centralizer $2 \times A_5$ is the compass datum. Finally, the compass Petersen 15-orbit carries exactly 12 pentagons splitting under the compass into two 6-orbits, each covering every edge exactly twice: two chiral fully-faced hemi-dodecahedra per needle, not merely skeletons. The dark sector closes the calculus: its 190 line pairs split $[10, 30, 30, 30, 30, 60]$ under the core into a perfect matching of 10 dark charts, two chiral $5 \times K_4$ partitions into skew tetrads, the two dark dodecahedra, and a 6-regular meeting frame; the dark matching is schedule-shadow pairing (each dark line meets exactly 4 schedule lines; matched lines share their shadow), and in K_5 terms the ten shadows are exactly the ten “edge + opposite triangle” configurations — a canonical bijection dark chart $\leftrightarrow K_5$ edge $\leftrightarrow$ schedule line. The timetable is stored three ways (lit-chart transversals, K_5 edges, dark shadows): built-in erasure redundancy. Each chiral tetrad’s four distinguished edges form a K_5 star, the five tetrads of each partition marking the five lit charts exactly once: K_5 vertices are stored three ways (lit charts, left/right tetrads), edges two ways (schedule lines, dark charts) — every A_5 -orbit structure of the pentad core except the dodecahedra and the meeting frame is a K_5 element — and those two fold in as well: each chiral dark dodecahedron is the antipodal double cover of the Petersen (Kneser) graph on the 10 dark charts, with deck transformation the shadow matching (antipode = shadow partner, verified), and the meeting frame is the doubled Johnson graph $T(5)$. The K_5 -relation census is constant on every orbit (matching = same edge, tetrads and meeting = adjacent, dodecahedra = disjoint): the dark sector is the canonical 2:1 cover of the K_5 edge-world, with the Petersen graph appearing three ways in one needle (schedule skeleton, dark Kneser quotient, doubled chiral covers).

Theorem 12.5 (The dark charts are the mirror bus). *The (core, dark chart) slot space — 216 pentad cores $\times$ 10 dark charts = 2160 slots, with every skew pair occurring as a dark chart in exactly 4 cores ($2160 = 540 \times 4$, the repair-atlas factorization) — is **G-set isomorphic to the 2160-slot D_{12} mirror bus** of the middleware: the slot stabilizer has order 12 with the D_{12} profile $\{1:1, 2:7, 3:2, 6:2\}$ and is conjugate in $\text{PSP}(4, \mathbb{F}_3)$ to the antipode-slot stabilizer (direct search over all 25920 elements). The multiplicity-4 agreement is canonical: a dark chart’s transversal tetrad is its schedule shadow, so the 4 host cores’ schedules all contain the pair’s 4 transversals. Hence the compass layer generates the middleware — timetables from its lit side (reconstruction), the transport bus from its dark side — and the mirror bus acquires a hardware identification: one slot = one nee-*

dle's dark chart. Witnesses: `bt856_dark_mirror_bus.py`; further: `bt853_dark_orbit_zoo.py`, `bt854_dark_chart_transversal_duality.py`, `bt855_dark_sector_k5_complete.py`.

Witnesses: `bt848_universal_amalgam_ladder.py`, `.tmp/gap_bt848*.g`; cf. Hartley's classification (17 universal rank-4 locally projective polytopes, 441 quotients) and Hartley–Leemans on $\{5, 3, 5\}$.

The barren rung and the maniplex workaround. Relaxing polytopes to *maniplexes* (flag graphs without the intersection property; the tomatope middleware of this machine is a uniform 4-polytope whose monodromy $18432 = 96 \times 192$ fails that property, giving it *infinitely many* distinct minimal regular covers — the canonical rank-4 cover pathology) does not lift the obstruction: exhaustive search shows $U_4(2)$ admits *no* rank-4 regular maniplex with a 5 in its type, and the amalgam closure $A_6 = L_2(9)$ admits no regular maniplex at any rank — it is the *barren rung* of the $\text{PSL}(2)$ ladder, whose populated rungs are exactly the 11-cell ($q=11$), the 57-cell ($q=19$), and Leemans–Toledo's degenerate single-facet $\{5, 5, 5\}$ family (our control search rediscovers all of them, including both exceptional polytopes). This is why the Grünbaum–Coxeter content of the substrate is *delocalized* over the 36 schedules — Petersen splits, chiral pentads, chart covers, dark dodecahedra — rather than crystallized in one flag object, and why the machine's runtime uses the tomatope: it is the flag structure that survives the obstruction. The tomatope's symmetry type is now exact: $|\text{Aut}| = 96$ (independent centralizer computation on the 192 flags), two flag orbits $[96, 96]$, **class** $2_{\{0,1,2\}}$ — only the cell-color crosses orbits, and the orbits are precisely “flag in a tetrahedron” vs “flag in a hemioctahedron”: the middleware *alternates spherical and projective cells*, the exact situation the classical “locally X ” taxonomy cannot name, and the setting of Mochán's 2-orbit polytopality theory. Moreover the tomatope is a **Cayley maniplex over** $V_4 = \mathbb{Z}_2 \times \mathbb{Z}_2$ (vertex action = full S_4 with a regular Klein subgroup; edge and face actions faithful and transitive): the middleware's vertex layer is a free $\mathbb{F}_2^2$ -torsor, so register XOR-addressing extends from the chart atlas into the runtime with no added structure. Reading Hartley's full classification (seventeen universal locally projective rank-4 polytopes) against the substrate adds three facts: **$\text{Aut}(\text{tomatope})$ is itself a universal polytope group** — isomorphic to $\Gamma(\{\{4, 3\}_3, \{3, 4\}_3\})$, the quotient-free doubly-projective $\{4, 3, 4\}$ universal built from hemicubes and hemicrosses (the tomatope's own projective cell type), with the case-10 group = $C_2 \times \text{Aut}$ of order $192 =$ the flag count; the mixed $\{3, 5, 3\}$ amalgam (icosahedral facets, hemi-dodecahedral vertex figures) *does not exist* — the construction closes into the 11-cell — while the dodecahedral mix is the J_1 giant; and the chart-compass universal $\{\{4, 3\}, \{3, 5\}_5\} = 2^{\mathcal{I}}$ (cube facets = Q_3 charts, hemi-icosahedral vertex figures = compass cells, vertices = $\mathbb{F}_2^6$ on the hidden 6-set) has order $3840 \nmid 25920$: blocked like the rest of the ladder. *Witnesses: `bt849_maniplex_barren_rung.py`, `bt850_tomatope_two_orbit_class.py`, `bt851_tomatope_cayley_test.py`, `bt852_seventeen_universals.py`, `.tmp/gap_bt849*.g`, `.tmp/gap_bt852*.g`.*

Reconstruction, the chart K_5 , and the dark dodecahedra. The pentad core's full signature $[5, 5, 10, 20]$ (lines), $[20, 20]$ (points) is rigid: both chiral pentads cover the *same* 20-point orbit (lit/dark halves of the substrate), the five matching charts' common-transversal bundles overlap to exactly the 10 schedule lines — **a pentad pair generates its timetable** — and the line $\rightarrow$ chart-pair map is a bijection onto the $\binom{5}{2} = 10$ edges of K_5 : the schedule *is* a K_5 on its own charts (4 transversals per chart, 2 charts per line). The 20 dark lines complete the $\{5, 3\}$ family: their 190 pairs split $[10, 30, 30, 30, 30, 60]$ under the core, and exactly two of the 30-orbits are **dodecahedron skeletons** (3-regular, girth 5, diameter 5) — a chiral pair of genuine dodecahedra

beneath every compass needle, one level under the hemi-dodecahedral Petersen splits of the schedule itself. *Witnesses: bt846_pentad_core_anatomy.py, bt847_dark_dodecahedron.py.*

Theorem 12.6 (BT839 GC operation Euler/flag audit). *Treating the full Grünbaum–Coxeter note as a data table gives a second invariant. The base tables for the 11-cell, 57-cell, tomospace partial a, and tomospace partial b are all Euler-neutral:*

$$V - E + F - C = 0.$$

But every Wythoff operation table has a fixed family charge:

$$\chi_{\text{op}}(11\text{-cell}) = -11, \quad \chi_{\text{op}}(57\text{-cell}) = -57, \quad \chi_{\text{op}}(\text{tomotope}) = -4.$$

The omnitruncated rows are full-flag carriers:

$$660 = |\text{PSL}(2, 11)| = 11 \cdot A_5 = kN_{\text{eff}},$$

$$3420 = |\text{PSL}(2, 19)| = 57 \cdot A_5 = k \cdot g \cdot 19,$$

and the tomospace has 96 half-flags, whose double is the verified 192-flag BT814 packet. The 57-cell also locks to the W33 timetable library: BT837 gives 3240 Petersen homes, and

$$3240 + k \cdot g = 3240 + 180 = 3420.$$

Thus one $k \cdot g$ sentinel sheet completes the W33 Petersen-home carrier to the 57-cell flag count.

The regular-polychoron alignment supported by the cell/symmetry data is

$$11\text{-cell} \rightarrow 600\text{-cell}, \quad \text{tomotope} \rightarrow 24\text{-cell}, \quad 57\text{-cell} \rightarrow 120\text{-cell}.$$

The $57 \rightarrow 120$ link is direct at the dodecahedral cell level; the $11 \rightarrow 600$ link is weaker but real, through hemi-icosahedral local structure and the quotient $2I \rightarrow A_5$ rather than through a direct tetrahedral-cell match. The swapped 11/57 orientation is recorded as a lower-scoring dual/vertex-figure hook, not discarded. Witness: bt839_gc_operation_euler_flag_audit.py.

Theorem 12.7 (BT840–BT842 carrier closure). *The three unresolved carriers in the Grünbaum–Coxeter paragraph are now explicit objects.*

(i) The 180 sentinel sheet. The verified Clifford L/R boundary has 36 cells arranged as a 6×6 rook grid. Its zero-overlap relation is not a numerical afterthought: it is the union of the 6 row fibres and 6 column fibres, each fibre carrying the duads of a six-set. Hence

$$12 \cdot \binom{6}{2} = 12 \cdot 15 = k \cdot g = 180.$$

This is exactly the sheet missing from BT837:

$$3240 + 180 = 3420 = |\text{PSL}(2, 19)|.$$

Equivalently, the timetable library gives 216 Petersen cores ($216 \cdot 15 = 3240$), and the rook boundary supplies the missing 12 six-set fibres ($12 \cdot 15 = 180$).

(ii) The 660 carrier. At an apex cell (r, c) of the same 6×6 grid, remove the incident row and column. The apex plus the five nonincident rows and five nonincident columns gives

$$11 = 1 + 5 + 5.$$

Crossing these eleven labels with the verified 60-element Clifford A_5 selector gives

$$11 \cdot A_5 = 11 \cdot 60 = 660 = kN_{\text{eff}}.$$

This is a carrier theorem only: it does not assert an internal $\text{PSL}(2, 11)$ action on $W(3, 3)$.

(iii) The 96 tomotope half-flags. *In the D_4 -root model of the 24-cell, every edge lies in a unique central hexagon plane. The two endpoint axes determine the third, missing axis in that plane, so each edge maps to a Reye incidence (missing axis, hexagon plane). Each of the 48 Reye incidences has exactly two edge lifts:*

$$96 = 2 \cdot 48.$$

Thus the tomotope omnitruncated half-flag count is literally the 24-cell edge lift of the common tomotope/Reye spine.

The hexagon clue supplies the conceptual bridge to the 11-cell. Each central 24-cell hexagon is a six-root cycle C_6 . Completing that cycle to a full K_6 gives 15 duads, exactly the hemi-icosahedral skeleton count of the 11-cell cell. Across the 16 central hexagons,

$$16 \binom{6}{2} = 16 \cdot 15 = 240,$$

the $W33$ edge count and the E_8 root count. The tested dot profile of these 240 duad slots is

$$96_{\langle x,y \rangle=1} + 96_{\langle x,y \rangle=-1} + 48_{\langle x,y \rangle=-2},$$

namely live 24-cell edges, mirror pairs, and repeated opposite-axis Reye incidences. Witnesses: BT840, BT841, and BT842; executable scripts are in `analysis/`.

13 Build sheet and falsifiable witnesses

Components (all standard quantum optics): one heralded single-photon source; one polarizing beam splitter; one symmetric 3-port coupler (tritter); a $0/\tau/2\tau$ delay ladder; one bin-synchronous electro-optic modulator; one polarization rotator set to $\arccos(-2/3)$ in the recirculation loop; single-photon detectors.

Witnesses (kill criteria, all parameter-free):

witness	predicted value	substrate form
trace–Choi visibility of F_3	$1/3$	$1/q$
trace–Choi visibility of X, Z	0	—
Kochen–Specker classical budget	$36/40$	$(q!)^2/v$
beacon-mesh pair visibility	$1/3$ (all 21 pairs)	$1/q$
Werner separability threshold	$3/4$	q/μ
BC-drive gap census at $n=30$	2 gap lengths	$h(E_8)$ ring
Hashimoto transport phases (one tick)	$\arctan \sqrt{10}, \pi - \arctan(\sqrt{7}/2)$	$\Phi_4, \Phi_6; u ^2 = 11$
key-agreement rate	$13/40$	Φ_3/v
D_{12} mirror-slot stabilizer	order 12	projective bus
runtime factorization	$24 \cdot 45 \cdot 48$	$ \mathrm{Sp}(4, \mathbb{F}_3) $
one recursive route digit	≤ 8 reversible hops	$3 + 5$
level- n leaves	40^n	v^n
level- n $W(3, 3)$ instances	$(40^n - 1)/39$	geometric tower
durable commit clock	$4(7^9 - 1)$ ticks	tomotope/Csaszar
compiled stress route	$47 < 48$ moves	BT828 level six
cover-lifted packet capacity	$48k^3$ slots	BT832 storage gauge
Bell-compass prefix router	$648 + 324 + 162 + 162$	code 0, 10, 110, 111
sentinel-aware stress energy	$4 \rightarrow 2$	BT833 reroute
full route sync criterion	$2n \mid 7^n - 1$	BT834 guard band
tomotope Wythoff runtime ladder	4, 12, 24, 48, 96	BT838 packet growth
GC operation Euler charges	$-11, -57, -4$	BT839 table audit
sentinel projector trace	15	g sector
first full-route desync	$n = 5$ remainder 24	f guard band
tomotope cover ABI	48 blocks / 192 flags	cover-invariant

A measured deviation in any row falsifies the corresponding layer.

14 Implications, corrections, and the ethos

This section states what the architecture means — for physics, for computing, for networking, and beyond the laboratory. Every claim below is either (i) a theorem with an executable witness in this paper’s ledger, or (ii) an implication of such a theorem, flagged as such. Nothing here floats free of the verification chain.

14.1 Physics

The vacuum is a gauge choice among five. The decomposition $40 = 1 + 12 + 27$ (self, gauge, matter) that organizes the entire Standard-Model derivation is *one of five* maximal decompositions, indexed by the Schläfli inventory of the cubic surface: 27 registers, 36 schedules, $40 + 40$ building parabolics, 45 polar pairs. The five vacua are connected by an explicit 5×5 double-coset transition matrix. Physically: what looks like “the” vacuum is a parabolic gauge choice, and the theory knows its own alternative phases and the exact combinatorial cost of moving between them. No continuum field theory has this property; here it is finite, exact, and enumerated.

Matter is magic, with three octane grades. The machine’s non-classical resource (magic, in the resource-theoretic sense) coincides *exactly* with its matter sector: the 36 entangled Witting rays are the matter shell, and the 27 matter contexts stratify $[1, 8, 12, 6]$ — the cube’s cell census. The three magic grades — deep $(2 + \sqrt{3})/6$ on $8 = 2^q$ rays, mid $(5 + 2\sqrt{3})/12$ on $24 = f$, shallow $3/4 = q/\mu$ (the Werner threshold) on $4 = \mu$ — mean that “what matter is” and “what fuels

quantum advantage” are one census. A physicist reads this as a structural identification of the fermionic sector with the non-stabilizer sector; an engineer reads the same integers as a fuel gauge. Both readings are forced by the same double cosets.

The Standard-Model spine: one transvection. The preceding paragraphs assemble into a single statement, verified end-to-end in one pass (`bt886_standard_model_spine.py`): *from the pair $(W(3,3), R)$ — the symplectic quadrangle and one long-root transvection — the discrete Standard Model follows with zero free parameters.* The transvection R fixes the 13-point gauge perp-plane and acts freely (nine orbits) on the 27 matter shell; its centralizer $C(R)$ (order 648) is the gauge group, with module $\mathbf{1} \oplus \mathbf{3} \oplus \mathbf{8} = U(1) \times SU(2) \times SU(3)$; its centre $Z(C(R)) = \langle R \rangle = \mathbb{Z}_3$ is the three generations; its normalizer gives the flavor group S_3 and charge-conjugation; the W/Q duality gives parity; the matter shell grades $9+9+9$ with the $\mathbb{Z}_3$ Yukawa rule; and the connection is flat ($\mathbb{Z}_3^2$) along collinear edges, $2T$ -curved on the matter graph Q with a quaternionic field strength. The gauge group, the generations, the flavor symmetry, the Yukawa texture, chirality, parity, and the gauge dynamics are not separate inputs — they are the centralizer, centre, normalizer, duality, grading, and commutator structure of one element of $\text{Sp}(4, \mathbb{F}_3)$, replicated over the 40 points. The gauge group even carries its color/electroweak split: $C(R) = 3^{1+2} : SL(2,3) = 27:24$, with the normal Heisenberg radical (order $q^q = 27$, the color part) fixing the electroweak $\mathbf{1} \oplus \mathbf{3}$ pointwise (the $W/Z/\gamma$ are color-blind, constant on the four lines through p_0) and acting on the gluon octet $\mathbf{8}$, and the Levi $SL(2,3)$ (order $f = 24$, the electroweak part) carrying the $\mathbf{1} \oplus \mathbf{3}$ — color $\times$ electroweak as a parabolic, with substrate orders q^q and f . And the color radical is *the same* Heisenberg group $3^{1+2} = O_3$ that acts simply transitively on the 27 matter shell (the matter torsor): **color is the matter-shell Heisenberg group**, so matter carries color precisely because the 27-shell is a torsor under the color group — color rotations and matter-shell motions are one action of $q^q = 27$. Under the full gauge group the matter shell is the homogeneous space $27 = C(R)/SL(2,3)$ (gauge group over the electroweak Levi): $C(R)$ is transitive with point-stabiliser the electroweak $SL(2,3)$, rank 6 (suborbits $[1, 1, 1, 8, 8, 8]$), module $\mathbb{C}[27] = \mathbf{1} \oplus \mathbf{6} \oplus \mathbf{8} \oplus \mathbf{12}$, and the three electroweak-fixed points form a color-singlet (lepton-like) 3-set while the three color-8 suborbits carry the colored (quark-like) content ($27 = 3 + 24$). One generation’s fermion multiplets are the $C(R)$ -module structure of the 27.

Three generations, matter-blind and gauge-visible. Because the matter register is the Steinberg module (Corollary, §10) and the Steinberg character vanishes on every 3-singular element, *any* order-3 symmetry of the substrate splits matter into exactly $27 + 27 + 27$: the three generations carry identical matter/Steinberg structure — they are indistinguishable to the protected register, which is why generations share gauge charges. Yet the four order-3 conjugacy classes *are* distinguished in the gauge (40-point) sector: the transvections grade it $22 + 9 + 9$ and fix a whole $\text{PG}(2,3)$ plane of $13 = \Phi_3$ points, while the regular classes grade it $16 + 12 + 12$ and fix only $\mu = 4$ points. The *physical* texture triality (Pillar 68’s Yukawa map R) is identified exactly: it is the long-root transvection (a 40-class element = the centre of the point-parabolic’s Heisenberg radical), fixing the 13-point gauge perp-plane $p_0^\perp$ pointwise and acting on the 27-point matter shell with 9 free orbits of 3 — the canonical “fix the gauge, stir the matter” generator. Its eigengrading of $\mathbb{C}[27]$ is $9 \oplus 9 \oplus 9$, and $\mathbb{Z}_3$ Clebsch–Gordan then *forces* the Pillar-68 Yukawa selection rule $T[a, b, v] = 0$ unless the three grades sum to 0 mod 3 (exactly 9 of 27 grade-triples survive); the matter coupling lives on the 36 within-shell tritangent triples (Pillar 69’s Heisenberg-twisted triads, $45 = 9 + 36$). And the gauge side closes it: the centralizer $C(R) = \text{Stab}(p_0)$ of order 648 (the gauge group is what commutes with generation-cycling) acts on the $12 = k$ gauge neighbours with permutation rank 3,

decomposing the gauge module as

$$\mathbb{C}[12] = \mathbf{1} \oplus \mathbf{3} \oplus \mathbf{8} = U(1) \oplus SU(2) \oplus SU(3),$$

the Standard-Model gauge group adjoint (the $\mathbf{1}+\mathbf{3}$ from the A_4 action on the four lines through p_0 , the $\mathbf{8}$ the within-line gluon octet). So R fixes the gauge sector ($12 = 8+3+1$, generation-blind) and grades the matter sector ($27 = 9+9+9$, three generations): the gauge group and the generations are the two eigen-data of one long-root transvection. The gauge sector even carries a parity statement: $C(R)$ acts on the four lines through p_0 as A_4 inside PSp , but as the full S_4 in $\mathrm{PGSp}(4, \mathbb{F}_3) = W(E_6)$ — the odd (parity-violating) 4-line permutations are *exactly* the anti-symplectic W/Q -duality coset, so gauge parity is the duality/chirality $\mathbb{Z}_2$ and lives only in the full $W(E_6)$. The flavor sector even has its charge conjugation: $N_{W(E_6)}(\langle R \rangle)/C(R) = \mathbb{Z}_2$, so there is an element C with $CRC^{-1} = R^{-1}$ that fixes generation grade-0 and swaps grade-1 $\leftrightarrow$ grade-2 — the discrete shadow of CP, conjugating the generation $\mathbb{Z}_3$ exactly as the temporal Bell qutrit's CPT conjugates $\omega \mapsto \bar{\omega}$. The full Standard-Model discrete flavor/parity data — gauge group $\mathbf{1} \oplus \mathbf{3} \oplus \mathbf{8}$, three generations $\mathbb{Z}_3$, generation C , matter chirality, gauge parity — is the normalizer geometry of one long-root transvection. The grading R and its conjugation C together generate the **flavor group** S_3 ($\langle R, C \rangle$, order 6, non-abelian — the minimal non-abelian flavor symmetry of BSM model-building), under which the matter shell decomposes as $\mathbb{C}[27] = 6 \cdot \mathbf{1} \oplus 3 \cdot \mathbf{1}' \oplus 9 \cdot \mathbf{2}$, the standard doublet appearing $q^2 = 9$ times. And the flavor–gauge relationship is exact: R is central in the gauge group $C(R)$, with $Z(C(R)) = \langle R \rangle \cong \mathbb{Z}_3$ — **the three generations are the centre of the gauge group**, acting trivially on the adjoint exactly as the $\mathbb{Z}_3$ centre of colour $SU(3)$ acts on the gluon octet (so the $SU(3)$ centre *is* the generation $\mathbb{Z}_3$). Globally there are 40 such local gauge groups — one per point, a single conjugacy class (homogeneous spacetime) — and the 40 generation-centred long-root $\mathbb{Z}_3$'s generate all of $\mathrm{PSp}(4, \mathbb{F}_3)$: **the 40 points of $W(3, 3)$ are the space of local gauge groups**, each $SU(3) \times SU(2) \times U(1)$ with its generation centre, and the global symmetry is the closure of the local generation symmetries. The connection between adjacent local gauge groups is exact: for collinear points the generation centres commute ($\langle R_p, R_{p'} \rangle = \mathbb{Z}_3 \times \mathbb{Z}_3$), so the gauge connection is **flat along the 240 edges**; for non-collinear points $\langle R_p, R_{p'} \rangle = SL(2, 3) = 2T$ (the binary tetrahedral / 24-cell group), so it is **curved across the 27-per-point matter shell**. Gauge curvature lives exactly on the matter graph Q — flat = collinearity, curved = non-collinearity = matter, with holonomy the 24-cell group. The curvature 2-form $F(p, q) = [R_p, R_q]$ makes this quantitative: it vanishes on collinear (causal) directions and is an *order-4 quaternion unit* of $2T$ (one of $\pm i, \pm j, \pm k$, with $F^2 = -I$ the centre) on every matter pair — an **su(2)-valued field strength supported exactly on the matter graph Q** . Flat causality, quaternionic curvature on matter. The integrated flux (Wilson loop $R_a R_b R_c$) confirms it: all 160 collinear triangles carry order-3 abelian flux (the flat $\mathbb{Z}_3$ line grading), while the 3240 matter triangles of Q carry non-abelian flux of orders $\{2, 4, 6, 12\}$ (counts 180, 180, 1440, 1440) up to $12 = k$ — gauge energy lives on the matter graph. The matter register (Steinberg) couples to this flux only through its non-3-singular part: by Steinberg vanishing, $\chi_{\mathrm{St}}(W) = 0$ on every order-3, 6, 12 loop (3040 triangles, matter-invisible) and is nonzero only on the order-2 loops ($\chi_{\mathrm{St}} = +q^2 = 9$, the polar-pair chirality involution of BT869) and order-4 loops ($\chi_{\mathrm{St}} = -q = -3$); the total coupling $\sum \chi_{\mathrm{St}}(W) = 180 \cdot 9 - 180 \cdot 3 = 1080 = 540 \cdot 2$ is exactly the chart double-cover count — gauge dynamics and the routing fabric meet at 1080. This is the Standard-Model pattern as a representation- theoretic theorem: generations are identical in gauge charge (matter-blind) and differ only through gauge-sector-visible couplings (the Yukawa/mass structure), here the eigenvalue degeneracy of an order-3 symmetry rather than a modeling input. *Witnesses: `bt863_generations_from_steinberg_vanishing.py`, `bt864_triality_census.py`*. Order-6 symmetries see *both* internal gradings at once: an order-6 element factors as g^2 (generation, order 3) times g^3 (chirality, order 2 — the sign-twisted carrier

of H_2), and since χ_{St} vanishes on every 3-singular power, its $\mathbb{Z}_6 = \mathbb{Z}_3 \times \mathbb{Z}_2$ eigengrading factors cleanly: each chirality half $(81 \pm \chi(g^3))/2$ carries 3 equal generations. For the dominant involution type ($\chi(g^3) = q^2$) the chirality split is **45 + 36** — the substrate’s own tritangent and double-six counts — so the left/right matter split is tied to Schläfli geometry; the other type gives 39 + 42. Generation $\times$ chirality = $3 \times 2 = 6 = q!$. The dominant chirality involution is identified exactly: $\text{PSp}(4, \mathbb{F}_3)$ has just two involution classes, and the 45-class (eight fixed points carrying the cube Q_3 , centralizer $2T \times 2T$ of order 576) is the *central involution of a hyperbolic polar pair* $L \leftrightarrow L^\perp$ — so the left/right split of matter, 45 + 36, is sized by the substrate’s own tritangent-plane and double-six counts, and the matter chirality $\mathbb{Z}_2$ is the 24-cell polar-pair swap. *Witnesses: bt868_joint_generation_chirality_grading.py, bt869_involution_chirality_classes.py.*

Time has three hands and an irrational escapement. The internal $\mathbb{Z}_{12}$ rectangle clock, the external $\mathbb{Z}_7$ (Császár) and $\mathbb{Z}_{13}$ (Singer) references, and the Boerdijk–Coxeter drive $\arccos(-2/3)$ together realize a discrete time quasicrystal: a drive that never repeats yet has exactly bounded gap structure (two gap lengths at $n = 30 = h(E_8)$). The arrow of time appears as the unique self-loop fixed point I/q with its two-trit forgetting cost: irreversibility is priced, not postulated. Implication: any laboratory system realizing this architecture is also a clock whose long-term stability is protected by a number-theoretic irrationality rather than by feedback.

The exceptional polytopes are the theory’s siblings, not its parts. The Grünbaum–Coxeter arc (Theorems 12.2–12.4) shows the substrate carries the complete local geometry of the 11-cell and 57-cell — hemi-icosahedra and hemi-dodecahedra in every schedule, with exact incidence — while universality plus Lagrange forbid their global crystallization inside $\text{Sp}(4, \mathbb{F}_3)$. The same local amalgam $A_5 *_{D_{10}} A_5$ closes as $L_2(9)$ (the substrate’s schedule core), $L_2(11)$ (the 11-cell), $L_2(19)$ (the 57-cell), with $J_1 \times L_2(19)$ — a sporadic group — atop the universal ladder. The physical moral is delocalization: the substrate *cannot* host a “GC condensate,” so the icosahedral order is necessarily spread over the 36 schedules as compass structure. This is the same mathematical mechanism by which frustrated phases in condensed matter refuse local order parameters — realized here exactly, with the obstruction an integer $(11, 19 \nmid 25920)$.

Spectral ghosts. The primes that refuse to divide the group still govern its analysis: $11 = k - 1$ is the Ihara prime of the zeta critical circle $|u| = 1/\sqrt{11}$, 19 enters through the Heawood genus $H(19) = 20 =$ the BC ring count, and the chart-web eigenvalue fields are $\mathbb{Q}(\sqrt{\Theta})$ and $\mathbb{Q}(\sqrt{\Phi_{12}})$ ($\Theta = 10$, $\Phi_{12} = 73$, both parameter-closure constants). A spectroscopist’s summary: the machine’s transport spectrum speaks the tier-ladder dialect of the physics paper, and the forbidden primes appear as poles of the zeta function rather than as symmetries.

Gravity counts the matter register. The discrete gravitational bridge is the Matrix-Tree action $S = \frac{M_P^2}{2} \ln \tau(G)$, τ the spanning-tree count, and for the substrate it is exact: the Laplacian spectrum $\{0, 10^{24}, 16^{15}\}$ gives

$$\tau(W(3, 3)) = \frac{10^{24} \cdot 16^{15}}{40} = 2^{81} \cdot 5^{23},$$

where the exponent of 2 is $81 = q^4 =$ the matter-sector (Steinberg) dimension, the base $5 = F_5$ is the tier-ladder prime ($r = 27/80$), and the exponent $23 = f - 1$. The complement matter graph $Q = \text{SRG}(40, 27, 18, 18)$ gives $\tau(Q) = 2^{66} \cdot 3^{39} \cdot 5^{23}$ with 3-exponent $39 = q\Phi_3 =$ the gauge-sector dimension. So the discrete-gravity partition function carries the Standard-Model sector sizes as its

prime-exponent charges — matter (q^4) in the power of 2, gauge ($q\Phi_3$) in the power of 3, cosmological (F_5) shared — and the entropy of the spanning-tree ensemble *counts the registers*. This is not a separate fact from the transport spectrum: the Ihara zeta $\zeta^{-1}(u) = (1-u^2)^{200}(1-u)(1-11u)(1-2u+11u^2)^{24}(1+4u+11u^2)^{15}$ (verified directly by the 480×480 Hashimoto operator) has its non-trivial zeros on the graph-RH circle $|u| = 1/\sqrt{11}$ (the transport sector, $p_{\text{th}} = k-1 = 11$) and its Perron zero at $u = 1$, where the Bass matrix $I - A + 11I = 12I - A$ is the Laplacian, so the Matrix-Tree leading coefficient $10^{24}16^{15} = 2^{84}5^{24} = v \cdot \tau$ recovers exactly the gravity count. Transport and gravity are the off-critical and critical behavior of one zeta: the prime 11 is the transport critical norm, the matter dimension $q^4 = 81$ is the gravity entropy. The substrate is in fact a *pair* of zetas: the complement matter graph $Q = \text{SRG}(40, 27, 18, 18)$ has its own Ihara prime $26 = 2\Phi_3$ (Hashimoto Perron, all complex eigenvalues on $|u| = 1/\sqrt{26}$) and its Perron-point gravity $\tau(Q) = 2^{66}3^{39}5^{23}$ puts the *gauge* dimension $39 = q\Phi_3$ where $W(3, 3)$ puts the matter dimension — complementation (collinearity $\leftrightarrow$ non-collinearity) swaps the matter and gauge entropies, with the cosmological charge 5^{23} shared. The non-backtracking walk census confirms the dynamics: $N_3 = \mu|E| = 960$, $N_5 = |E|q^q n_{\text{even}} = 181440$, $N_n \sim 11^n$. *Witnesses: bt870_spanning_tree_gravity.py, bt872_ihara_zeta_spanning_bridge.py, bt873_dual_zeta_walk_census.py.*

Falsifiability is retail, not wholesale. Each physics reading above is pinned to a tabletop witness with an exact rational value: trace-Choi visibility $1/3 = 1/q$ exactly (not approximately); Kochen–Specker budget $36/40 = (q!)^2/v$ exactly, with deficit μ ; contextual fraction $1/10 = 1/\Phi_4$; Werner threshold $3/4 = q/\mu$; two gap lengths — not three — at drive step 30. A single wrong integer kills the corresponding layer. This is a stronger falsifiability posture than parameter-fitting physics can offer, because there are no parameters to refit.

14.2 Computing

What the machine is. One photon carries a 4-dimensional path–polarization register (states) and a 9-dimensional past–future time-bin register (operators), realizing the same $W(3, 3)$ geometry twice; the operator–operand duality makes program and data literally the same incidence object. Universality is a two-ingredient theorem: three passive optical elements (tritter, phase plate, EOM) generate the full 51840-element Clifford symplectic closure — *exactly*, not approximately — and the matter shell supplies the magic states that promote Clifford to universal. There is no cryostat, no vacuum chamber, no trapped ion: the exotic resource is geometry.

What it costs. The build sheet is a heralded single-photon source, one polarizing beam splitter, one symmetric tritter, a three-rail delay ladder, one electro-optic modulator, one recirculation loop with a fixed irrational rotation, and single-photon detectors. Every component is catalog optics. The machine’s benchmarks are integer-valued: a laboratory can certify the Clifford layer by counting to 51840 and certify the fuel by hitting $36/40$. Self-testing is native because the specification is arithmetic.

The software stack exists and is finite. Programs are braid words (with $\sigma^5 = Z$ exact); the instruction set is the BT828 compiler ABI; the operating system is the timetable library (36 schedules; overlap law 1-or-4 with switching cost 6); synchronization is the beacon-heptad mesh (2880 frames, themselves forming an exchange network); memory protection is the Steinberg sector ($81 = q^4$, a single irreducible — protected by Schur’s lemma, the strongest no-leakage statement representation theory can make). The state/program compiler makes that register directly addressable: $\mathbb{F}_3^3$ supplies three commuting copies of the 27-slot program shell, H_{27} supplies three noncommuting copies

of the state shell, and its center supplies the native generation flag $27 \subset 54 \subset 81$. The oriented timetable header above that memory splits as $5_\omega + 5_{\omega^2} + 30$: the full Weyl runtime fuses the handed pair into a 10-channel and leaves one 30-channel with an outer parity choice. Cache lookup is a distinct split layer: two copies of H/C_4 share an 81-route H/D_8 base, and their 324 slots form 81 four-state D_8 address fibers. The non-split ternary transport operator then updates the selected state; address choice and temporal memory cannot alias because their minimal polynomials differ. The watchdog is the BC loop. Each layer is a theorem; the stack compiles because double cosets compose.

Redundancy is built in, not bolted on. The newest results sharpen the fault story: each pentad core stores its timetable *three independent ways* (lit-chart transversals, the K_5 edge labeling, dark-sector shadows), every skew pair has six Petersen homes and two matching covers, and the dark sector carries its own chart matching and two chiral tetrad partitions (Theorem 12.4 block). Erasure of any single description leaves the schedule reconstructible — redundancy appears as a counting identity, not as an engineering allowance. The middleware’s two packet phases (tetrahedral/hemioctahedral, the tomatope’s $2_{\{0,1,2\}}$ class) confine phase changes to cell hops, which is where the commit boundaries already sit, and the runtime’s vertex layer is a free $\mathbb{F}_2^2$ -torsor (the Cayley structure), so address arithmetic is XOR end to end.

What it is not. No claim is made here of asymptotic speedups over error-corrected gate-model machines, of fault-tolerance thresholds, or of scalability beyond the fractal substitution rule (40^n leaves, diameter $8n$, commit clock $T(n) = 4(7^n - 1)$) — those are stated as exact combinatorial scaling laws, and their physical realization at depth $n > 2$ is an open engineering problem with its first hard boundary already computed (the $n = 5$ desynchronization with remainder $24 = f$). The honest claim is narrower and stranger: a complete, finite, self-verifying computational architecture exists inside one photon’s worth of quantum optics, with zero adjustable parameters.

14.3 Networking

The network is the computer, twice over. The thesis is not rhetorical: the atlas’s transport groupoid and the register’s gate groupoid are the same groupoid, so moving a packet and applying a gate are one operation. The new library-geometry theorem (Theorem 12.2) doubles the statement at the control plane: the 36 timetables form $\text{SRG}(36, 15, 6, 6)$, the *other* rank-3 geometry of the same group, so the control plane and the data plane ($\text{SRG}(40, 12, 2, 4)$) are the two classical geometries of one symmetry. Configuration management inherits strongly-regular guarantees: any two timetables, near or far, have exactly six common near-neighbors — worst-case reconfiguration paths are bounded by the geometry itself.

Routing, addressing, switching. Local routing is e-cube XOR on Q_3 charts (540 of them; the atlas is glued along the 1620 apartments of the Tits building, diameter 5). Addresses are Gray-coded $\mathbb{F}_2^3$ words; the antipode map is bit complement. Timetable switching obeys the overlap law: cheap hops preserve $4 = \mu$ of 10 calibrated contexts, the switch graph is the $U_4(2)$ rank-3 geometry, and every skew pair lies in exactly 3 schedules, 6 Petersen homes, and 2 pentad-matching covers — the router’s spare-path inventory is an exact census, not a provisioning estimate.

Synchronization without a master. The beacon heptads (2880 complete K_7 interference meshes, all pairwise overlaps $1/3$) are reference frames that form their own exchange network

— the network’s clock distribution is itself a network, with no distinguished master node. External discipline comes from the $\mathbb{Z}_7/\mathbb{Z}_{13}$ clocks (both decimal period 6) and the BC quasicrystal drive; internal phase order from C_{12} . A distributed-systems reader will recognize the structure: consensus by geometry, where the “leader election” problem dissolves because the frame family is a single transitive orbit.

Security posture (implication, flagged). Three structural facts have security readings, stated here as directions rather than theorems. (i) Contextuality is tamper-evidence: the KS budget 36/40 is exact, so any channel intervention that classicalizes measurement statistics moves an integer and is detectable in principle. (ii) Gate teleportation stores a unitary in the photon’s self-entanglement and applies it by Bell projection against its own past — a store-and-forward primitive for quantum repeaters that never exposes the gate classically. (iii) The timetable redundancy (three independent reconstructions) is an erasure-code posture against denial: losing a description class does not lose the schedule. Turning these into security proofs requires adversary models this paper does not supply.

Scaling. The fractal holonet substitutes a full 40-node substrate for each leaf: 40^n leaves, $(40^n - 1)/39$ substrate instances, reversible diameter $\leq 8n = 8 \log_{40} N$, durable commit clock $T(n) = 4(7^n - 1)$, with the first guard band an arithmetic necessity at $n = 5$. Because every child inherits the same finite protocol, scaling adds no new control logic — the protocol stack is depth-invariant. This is the architectural property that datacenter fabrics approximate statistically and this geometry has exactly.

14.4 The physics-to-architecture dictionary

The dictionary is the deepest implication: one set of integers reads simultaneously as physics and as engineering, because both readings are projections of the same incidence theorems.

object	physics reading	architecture reading
$1 + 12 + 27$ split	Bell line + meeting + disjoint	pole / I-O ports / fuel+store
240 edges	interaction qutrits	bandwidth = memory
$81 = q^4$ Steinberg	protected sector (Schur)	logical register
dual O_3 restrictions	Heisenberg state phases	state/program compiler
$H_2 = 5 + \bar{5} + 30$	oriented context chirality	timetable header $10 + 30^\pm$
$324 = 81 \times 4$, commutant D_8	split cache multiplicity	four-state address fiber
$0 \rightarrow 81 \rightarrow 162 \rightarrow 81 \rightarrow 0$	ternary temporal extension	non-split state update
36 spreads	measurement bases	timetable OS
SRG(36, 15, 6, 6)	basis-overlap law	control-plane fabric
five vacua	phases of the theory	boot modes
magic strata $8+24+4$	fermion shell grading	fuel octanes
540 skew pairs	two-fold null systems	hypercube charts
1620 apartments	Weyl frames	inter-chart trunks
2880 heptads	equiangular frames	sync beacons
BC drive $\arccos(-2/3)$	quasicrystal time	watchdog clock
compasses (216+216)	icosahedral order, delocalized	needle registers
pentads (chiral 216+216)	maximal partial spreads	F_5 chart caches
dark sector zoo	hidden-half structure	erasure store
tomotope $(2_{\{0,1,2\}}, \text{Cayley}/V_4)$	sphere/projective alternation	packet middleware
11-cell, 57-cell, J_1	universal completions (blocked)	the covers the hardware refuses

The first row carries the deepest identification, drawn from the companion paper *Universal Quantum Computation Through a Single Photon*: the split $1 + 12 + 27$ is not merely the local SRG anatomy of a point — it is the **Bell-line shell** of the photon’s own self-entanglement seed. The temporal Bell qutrit $|\Omega\rangle$ stabilizes one line of $W(3, 3)$; the 40 lines then decompose as 1 (the Bell line itself) + 12 (lines meeting it — the gauge sector) + 27 (lines disjoint from it — the matter sector), with substrate reading $1 + k + q^q$. The vacuum decomposition that organizes the Standard-Model derivation is literally the photon’s view of its own now-context: *self = the seed of self-entanglement*. The Bell-line stabilizer has order $51840/40 = 1296 = \mu^2 q^{q+1}$ — the same $1296 = 6^4$ that counts the compass D_{10} -incidence pairs of Theorem 12.4 (36 schedules $\times$ 36 core pairs), but Theorem 2.3 proves that this is not a regular-action identification. Instead, the projective Bell stabilizer resolves the incidence carrier as $648 + 324 + 162 + 162$, the exact mirror/schedule/left-cache/right-cache prefix decoder, while the point parabolic supplies two persistent 648 address sheets. The two cache words 110 and 111 are now resolved more finely by Theorem 10.6: they name two copies of the same H/C_4 cache, not two inequivalent internal chiral modules. Both double-cover one canonical H/D_8 base of 81 routes, and the complete cache quarter is an exact D_8 -over-81 square bundle. This reversible address plane is provably not the non-split $81 \rightarrow 162 \rightarrow 81$ transport plane, despite their common 162 count. The companion also supplies the master equation at Hilbert-space level ($q^2 = q + q!$: the 9-dimensional past $\otimes$ future space splits into 3 diagonal “now” histories and $6 = q!$ context-changing ones), and a ledger identity: temporal gate teleportation consumes exactly 2 classical trits of memory — the same 2-trit cost that prices the arrow of time in the self-loop fixed point of §11. Storing a state for one’s future self and forgetting one’s past cost the same two trits. The 27 itself now has a dual decoding: the 27 non-collinear *points* form a torsor under the genuine Heisenberg group (the normal 3_+^{1+2} of the point parabolic, acting simply transitively), whose invariant relations reconstruct the whole classical 27-zoo — the 9-triangle fibration, $GQ(2, 4) = SRG(27, 10, 1, 5)$, and the Schläfli graph $SRG(27, 16, 10, 8)$ of the E_6 27-lines world — while the 27 disjoint *lines* of the Bell shell form a torsor under elementary $\mathbb{F}_3^3$ (the line parabolic’s O_3), a flat 3-trit register cube carrying *no* invariant strongly regular graph. Curvature for states, flatness for programs: the matter sector is a Heisenberg manifold on the state side and a ternary address space on the operational side. The register reading is exact: assigning each Bell-shell context its 3-trit address, the four pair relations $[4, 4, 6, 12]$ are the difference-vector classes, and the geometry is a function of register arithmetic alone — meeting contexts ($\pm$ -paired 4-classes) share exactly 1 transversal onto the Bell line, near-skew (12) exactly 2, far-skew (6) exactly 0: base-3 register incidence for matter contexts, twin to the charts’ base-2 Gray arithmetic. *Witnesses: bt858_heisenberg_shell_torsors.py, bt860_bell_shell_register_arithmetic.py.*

Three rows deserve emphasis. The *Steinberg row* is the strongest protection statement available: the register is one irreducible representation, so *any* equivariant leakage operator is zero by Schur’s lemma — protection by symmetry, not by redundancy. The *compass rows* say that what physics sees as delocalized icosahedral order, engineering sees as a navigation instrument: 12 needles per schedule, each carrying Petersen addressing, five-chart caches, and a dark-sector mirror. The *universal-polytope row* states the program’s sharpest structural lesson: the hardware keeps the local physics of the exceptional objects and refuses their global rigidity; the refusal (an integer non-divisibility) is precisely why the architecture is distributed.

14.5 Beyond the laboratory

A different contract between theory and experiment. Every layer of this machine is specified by integers and verified by scripts that re-derive those integers from first principles in seconds. The cultural implication is a publishing model in which papers *run*: a claim is not a sentence but a

reproducible computation, and a correction is a commit. This program has already operated under that contract (see the ethos below) — twice it corrected its own published claims at the source, and the corrected values were more natural than the originals. If the broader physics literature adopted executable claims at even a small scale, entire classes of irreproducibility would become type errors.

Quantum hardware without a building. Architectures that need millikelvin cryostats centralize by economics: only institutions can own them. A universal architecture whose parts list is catalog optics — source, splitters, delay fiber, modulator, detectors — decentralizes by the same economics. The realistic near-term role is not supremacy-class computation but *certified* small-scale quantum information: teaching laboratories that can verify 51840 and 36/40 on a bench; metrology nodes whose clock is protected by an irrational rotation; network testbeds where gates and routes are the same hardware. A machine whose benchmarks are exact integers is also the natural pedagogical machine: a student can check the whole specification.

Communication and computation stop being separate utilities. In this architecture there is no boundary at which computation hands off to networking — the gate groupoid and the transport groupoid coincide, and the fractal substitution makes the computer/network identity scale-invariant. The long-run implication, if any physical system realizes this pattern at depth, is infrastructural: bandwidth and compute become one provisioned quantity, scheduling becomes group theory, and the spectrum-allocation problem for the network’s measurement contexts is solved by a finite library (36 timetables with a known switching metric) rather than by auction or contention.

Society-scale stakes of a zero-parameter physics. Separately from the machine, the substrate program’s physics claims — 54+ observables, 14 falsifiable predictions, no adjustable parameters — carry a specific societal meaning: they are maximally exposed. A theory with no dials cannot retreat; the 2027 neutrino-hierarchy determination, the proton-decay window, the CMB measurements will each either hit the predicted integers or end the corresponding pillar. Public trust in fundamental science is built by exactly this posture — predictions that can die — and the machine half of the program extends the posture to hardware: the device specification is its own audit.

What could go wrong, said plainly. The architecture is exact; its physical realization is not yet demonstrated beyond the component level. The fractal scaling laws are combinatorial theorems, not engineering schedules. The security readings are directions, not proofs. And the identification of the substrate with fundamental physics stands or falls with the prediction table, not with the elegance of the mathematics. This paragraph exists because the program’s ethos (below) requires the failure modes to be listed next to the claims.

14.6 The ethos

Several pre-existing corpus claims failed under exact computation and were corrected at their sources during this program: the independence number of $W(3, 3)$ is 7, not 10 (no ovoid exists; the “perfect graph” block is withdrawn); the Kochen–Specker optimum is 36/40, not 34/40; the pentad-sharing count of BT844 was corrected by the chiral-orbit census of BT845 within one working session; and the tomatope’s framing was corrected from “non-polytopal maniplex” to “uniform polytope with non-polytopal minimal covers” when the literature was read precisely; and the “ $\mathbb{Z}_3$

Berry phase as a discrete second Chern class” was corrected (BT871) — the uniform Berry 2-cochain is a coboundary over $\mathbb{F}_3$, so the physical $\mathbb{Z}_3$ is the order-3 action on the Steinberg register, not a curvature class. In every case the corrected value is *more* substrate-natural than the claim it replaced ($\alpha = \Phi_6$; budget = $(q!)^2/v$; $432 = 216 + 216$ chiral; infinitely many covers as the actual pathology; the $\mathbb{Z}_3$ as a group action rather than a cocycle) — the pattern a real structure shows when probed honestly. The program treats corrections as first-class results: they are committed, logged, and cited like theorems, because a corpus that cannot heal is a corpus that cannot be trusted.

A Verification ledger

Every theorem above is backed by an executable script in **analysis/** with results in **data/**; the chain BT739–BT839 was developed and pushed incrementally with exact arithmetic ($\text{GF}(p)$ elimination, cyclotomic integers, complete branch-and-bound, GAP witnesses for group-theoretic facts).

layer	primary witnesses
duality / carriers	bt817, bt820, bt821
atlas / building / web	bt773, bt777, bt744, bt779
mirror middleware / toмотope runtime	bt814, bt815, bt826
braids / registers	bt740, bt741
vacua / geography / transitions	bt808–bt813, bt816
fuel / contextuality	bt818, bt822, bt823
memory / Steinberg	bt742, bt744
clocks / quasicrystal / arrow	bt774, bt819, bt820, bt823
sync / schedules / overlap law	bt817, bt819, bt835
Grünbaum–Coxeter hemicells / library geometry	bt836, bt837, bt839
two compasses / $4 \cdot K_{10} = 12$ Petersens / chart double cover	bt843–bt845
pentad anatomy / schedule reconstruction / dark dodecahedra	bt846, bt847
amalgamation ladder / universality / J_1	bt848
maniplex barren rung / class $2_{\{0,1,2\}}$ / Cayley over V_4	bt849–bt851
seventeen universals / $\text{Aut}(\text{tomotope})$ universal	bt852
dark orbit zoo / shadow bijection / K_5 -complete dark sector	bt853–bt855
dark charts = mirror bus (G -set iso)	bt856
Heisenberg/ $\mathbb{F}_3^3$ shell torsors, Schläfli decoding	bt858
Bell-shell register arithmetic ($[4, 4, 6, 12]$ named)	bt860
code register = Steinberg ($[[240, 81, 4, 3]]_3$)	bt861
H_2 = sign-twisted lines (homological dictionary)	bt862
three generations = Steinberg vanishing	bt863
triality census (matter-blind, gauge-visible)	bt864
joint generation $\times$ chirality grading	bt868
chirality = polar-pair involution ($45 + 36$)	bt869
spanning-tree gravity $\tau = 2^{81}5^{23}$	bt870
Berry-phase correction ($\mathbb{Z}_3$ is the Steinberg action)	bt871
Ihara zeta = spanning-tree gravity bridge	bt872
two-zeta duality / walk census	bt873
texture triality = long-root transvection	bt874
Yukawa rule = $\mathbb{Z}_3$ grade conservation	bt875
gauge group $1 \oplus 3 \oplus 8$ from $C(R)$	bt876
gauge parity = duality ($A_4 \rightarrow S_4$)	bt877
generation charge-conjugation ($N/C = \mathbb{Z}_2$)	bt878
flavor group S_3 ($\mathbb{C}[27] = 6 \cdot 1 + 3 \cdot 1' + 9 \cdot 2$)	bt879
generations = centre of the gauge group	bt880
spacetime = space of 40 local gauge groups	bt881
gauge connection flat on edges, $2T$ -curved on Q	bt882
curvature 2-form (quaternionic, on Q)	bt883
gauge flux (Wilson loops, $\mathbb{Z}_3$ /lines, $\leq 12/Q$)	bt884
YM–matter coupling $\sum \chi_{\text{St}} = 1080$	bt885
SM spine (one transvection, master theorem)	bt886
gauge group = color : electroweak (27:24)	bt887
color = the matter-shell Heisenberg group	bt888
fermion content $27 = C(R)/SL(2, 3)$ (lepton/quark)	bt889
dual-torsor Steinberg compiler / native generation flag	bt865
oriented H_2 irreducible spectrum / outer lift	bt866
cache D_8 -over-81 fibration / split–non-split boundary	bt867
chirality bookkeeping / Bell–compass parabolic router	bt857, bt859
universality	bt825
fractal computer/network	bt827
runtime compiler / sentinel / commit / cover boundary	bt828–bt834, bt838
corpus corrections	bt818, bt823, bt824

B Architecture Completeness and the Three Physical Residuals

The holonet machine is now specified end to end: the 40-point substrate register and its $W(3, 3)$ commuting-projector Hamiltonian; the $[[240, 81, 4]]_3$ Steinberg code as the native error-correcting register; the local gate set and transversal/holonomy operations; the gauge connection (flat along the 240 edges, curved on the matter graph with binary-tetrahedral $2T$ holonomy) as the physical wiring of multi-qubit interaction; and the discrete dynamics, in which the physical on-shell states are exactly the harmonic forms $\ker D$ (vacuum $\oplus$ Steinberg matter $\oplus$ line module $= 1 \oplus 81 \oplus 40$) — the least-action ground space of the finite spectral action $\text{Tr } f(D^2)$, whose value reproduces the spanning-tree gravity invariant. *As a universal computer the architecture is complete:* universality, the code, the gate set, and the runtime compiler are established (cf. the substrate ledger above).

What remains is not computational but *physical*, and the three residuals of the master manuscript’s completeness boundary (`w33_paper.tex`) are now each resolved or reduced:

- **(R1) the gauge-lattice embedding — resolved.** The machine’s symmetry register carries the canonical symplectic $E_8/2E_8$ datum mod 2, and that datum admits a *unique* $\text{PSp}(4, 3)$ -invariant quadratic refinement, of plus type; hence $\text{PSp}(4, 3) \subset O_8^+(2) = \text{Aut}(E_8/2E_8)$, and the canonical positive-definite lift is E_8 (up to a central $\{\pm 1\}$). The gauge lattice is fixed.
- **(R2) the symmetry / matter register — datum supplied, bridge built.** The quintic moonshine traces $\text{Tr}(2A|V_5) = 74428120$, $\text{Tr}(2B|V_5) = 184024$ are computed; and the matter register is realized as a finite spectral triple (a Hilbert–Schmidt bimodule of dimension 81, first-order condition verified, gauge content $\mathbf{1} \oplus \mathbf{3} \oplus \mathbf{8}$). The residual is now the explicit physical fermion-multiplet assignment, a representation- selection task.
- **(R3) the emergent-spacetime limit — theorem-governed.** On a shape-regular (edge-wise) refinement tower the full gravity-plus-matter spectral action converges term by term to the continuum Einstein–Hilbert + Standard-Model action — by Cheeger–Müller–Schrader (curvature), higher-order Regge (higher-derivative), classical lattice gauge, and the exact finite moments $\{440, 1920, 16320\}$; the spectral action is moreover continuous for Latrémolière’s spectral propinquity. The emergent-spacetime limit is thus the *application* of established convergence theorems, not an open analytic obstruction.

None of the three is a gap in the machine’s *operation* — the holonet runs, corrects, and computes universally as specified — and none is now an open problem in principle: a fixed gauge lattice (E_8), a matter register awaiting only its physical labelling, and an emergent spacetime with a theorem-governed continuum limit. This is the precise sense in which the architecture is complete: a finite, exact, universal machine whose physics closes onto the Standard Model and general relativity through results that are now either proved or reduced to the application of known theorems.

Reading guide for what follows. The remaining subsections develop the architecture in four movements, each a chain of witnessed theorems. (A) *The logical layer is substrate-forced:* the machine is the substrate’s symplectic continuum (the oscillator representation), its fault-tolerant code is the lattice tower $A_2 \subset D_4 \subset E_8$, its universal gate set is the degree-2 symplectic plus degree-3 cubic invariants, the gate group equals the code’s logical Clifford group $\text{Sp}(4, 3)$, and the code lattices sit on the $E_8 \rightarrow \text{Leech} \rightarrow \text{Monster}$ moonshine ladder. (B) *The physical carrier is one massless self-protected photon:* the demonstrator versus the fault-tolerant machine, gates as geometric ($2T$) holonomies, self-error-correction from the substrate-fixed quasicrystal drive (topological, $|C| = q - 1$), why the primitive must be a single massless photon (zero proper time, Wigner

helicity λ), and the carrier as its own geon inside a fractal of nested shells. (C) *The physics and the computer are one object*: the quantum power is the gauge curvature on the matter graph, that discrete curvature is a finite-group lattice gauge theory whose continuum limit is the spectral action’s gauge-plus-gravity field equations (with an exact mass gap f), the whole network is a holographic code — black-hole entropy is the code distance, and the substrate carries a discrete AdS/CFT — and the runtime closes it: the clock is a harmonic oscillator whose frequency $\sqrt{\lambda}$ is the top mass, and the supercycle is the gauge group $\mathrm{Sp}(4, 3)$, so running the machine is the physics. (D) *It is testable now*: a parameter-free near-term falsification protocol on the single-photon demonstrator. Throughout, one substrate number recurs — $q - 1 = \lambda = 2$ sets the drive angle, the topological protection, and the photon’s transverse helicity count — and no constant is fitted.

B.1 The architecture is the substrate’s symplectic continuum

There is a structural reason the machine and the spacetime are different objects built from the same $W(3, 3)$: the substrate has *two* continuum limits. The arithmetic tower $W(3, q)$ (the symplectic quadrangle over $\mathbb{F}_q$, $q = 3^n$) is spectrally rigid — exactly three eigenvalues at every q , no Weyl law — so it never becomes a smooth spacetime; the metric continuum needs the external $K3$ seed (the physics side, `w33_paper.tex`, Prop. on the two continua). The *intrinsic* continuum of $W(3, 3)$ is instead *symplectic*: its Heisenberg matter shell 3^{1+2} carrying the $\mathrm{Sp}(4, 3)$ action is a Weil-representation datum, whose $q \rightarrow \infty$ limit is the metaplectic (oscillator) representation of $\mathrm{Sp}(4, \mathbb{R})$ on $L^2(\mathbb{R}^2)$ — precisely the continuous-variable, Fock/oscillator computation this paper realizes. So the holonet *is* the symplectic continuum of the substrate, exactly as the spacetime is its metric continuum: the universal computer and the universe are the two faces of one finite geometry’s continuum limit. (Witness: `analysis/w33_two_continua_symplectic_metric.py`.)

B.2 The fault-tolerant layer is the substrate’s lattice tower [Face 5]

Because the holonet’s logical layer is a continuous-variable (CV) oscillator computation, its canonical fault-tolerant encoding is the Gottesman–Kitaev–Preskill (GKP) code, which embeds a qudit in an oscillator. A *multimode* GKP code *is* a symplectic lattice in phase space: stabilisers are displacements by lattice vectors, and error correction is closest-point decoding in the symplectic dual lattice. The code quality is set entirely by the lattice, and the established optima are exceptional lattices: the *hexagonal* lattice A_2 for one mode, the densest 4-dimensional lattice D_4 for two modes (strictly better than any product of one-mode codes), and E_8 for four modes [23, 24].

This is not a free design choice for the holonet: the substrate *fixes* the lattice at every scale. The machine is two qutrits, hence two oscillator modes ($L^2(\mathbb{R}^2)$, phase space $\mathbb{R}^4$), so its optimal GKP code lattice is D_4 — and D_4 is exactly the substrate’s matter shell, with $|W(D_4)| = 192$ the tomosphere/flag symmetry and D_4 triality the three-generation structure. The single-mode optimum A_2 is the $q = 3$ (hexagonal, $SU(3)$) lattice of a single qutrit, and two modes stacked, $D_4 \oplus D_4$, sit primitively inside E_8 — the substrate’s mod-2 homology lift (R1), which is also the $K3$ gauge lattice. Thus the tower

$$A_2 \subset D_4 \subset E_8$$

is *simultaneously* the optimal 1-, 2-, 4-mode GKP code lattices (the architecture’s quantum error-correction layer) and $W(3, 3)$ ’s physical lattice tower (single-qutrit $SU(3)$, matter shell, gauge E_8). The error correction of the computer and the gauge structure of the universe are one and the same lattice tower; the QEC code is not chosen, it is the substrate. (Witness: `analysis/w33_gkp_lattice_architecture.py`, verifying the kissing numbers 6, 24, 240 and the embedding $D_4 \oplus D_4 \subset E_8$.)

This also fixes the lattice part of the fault-tolerance threshold. The substrate lattices are *isodual* (symplectically self-dual — the balanced-GKP condition, $\hat{q}$ and $\hat{p}$ protected equally) and densest in their dimensions, so they maximise the nominal coding gain $\gamma = d_{\min}^2 / \det^{1/n}$ among symplectic lattices of that rank:

$$\gamma_{\text{dB}}(A_2) = 0.6, \quad \gamma_{\text{dB}}(D_4) = 1.5, \quad \gamma_{\text{dB}}(E_8) = 3.0$$

(and 6.0 for the Leech lattice). So the holonet’s two-mode code D_4 already buys ~ 1.5 dB, and the four-mode code $E_8 \sim 3$ dB, of squeezing-margin over the trivial square code, for free. The absolute threshold still needs the noise model, the finite-squeezing GKP-state quality and the syndrome-extraction protocol — which we do not fix here — but the lattice contribution, the part that depends only on the code, is the substrate’s and is optimal. (Witness: `analysis/w33_gkp_coding_gain.py`.)

B.3 The universal gate set is degree two plus degree three

The code fixes the fault-tolerant layer; the gate set completes the architecture, and it too is substrate-forced. The continuous-variable universality theorem [25] states that the Gaussian unitaries — those generated by Hamiltonians at most *quadratic* (degree 2) in the mode operators: beamsplitters, phase shifters, squeezers, i.e. the CV Clifford group $\text{Sp}(2n, \mathbb{R})$ — are *not* universal, but adjoining any single generator of degree ≥ 3 (canonically the cubic phase gate $e^{i\gamma q^3}$) generates the full unitary group on the oscillators. A universal CV machine thus needs exactly

$$\underbrace{\text{degree 2}}_{\text{Gaussian / Clifford}} + \underbrace{\text{degree 3}}_{\text{cubic, non-Gaussian}}.$$

The substrate supplies precisely these, as its two lowest matter invariants and no others below degree three: the *degree-2* symplectic form on $\mathbb{F}_3^4$ that $\text{Sp}(4, 3)$ preserves — whose oscillator-rep limit is the Gaussian group $\text{Sp}(4, \mathbb{R})$ that the photon’s tritter, phase plate and modulator already realise (§6.5) — and the *degree-3* E_6 Cartan cubic on the matter $27 = 3 \otimes 3 \otimes 3$, the ε -cubic preserved by $\mathfrak{sl}(3, \mathbb{F}_3)^3$ (`w33_paper.tex`, the 27-cubic and its 45 triads). So the paper’s earlier statement that “the matter shell supplies the magic” is, in the CV picture, exact: the magic is the degree-3 E_6 cubic, the minimal non-Gaussian element, matching the Lloyd–Braunstein criterion on the nose. The substrate’s invariant degrees $\{2, 3\}$ are the universal CV generating set {Gaussian, cubic}. Together with the lattice tower $A_2 \subset D_4 \subset E_8$ this closes the architecture: $W(3, 3)$ fixes both the fault-tolerant code and the universal gate set, leaving no free choice in the logical layer. (Witness: `analysis/w33_cv_universality_cubic.py`.)

Finally, the code and the gate set are not two independent facts: they are welded by a single group. The logical Clifford group of a GKP code on n qudits of dimension d is the symplectic group $\text{Sp}(2n, \mathbb{Z}/d)$ over the logical phase space. For the holonet ($n = 2$, $d = 3$) this is

$$\text{Sp}(4, \mathbb{Z}/3) = \text{Sp}(4, 3), \quad |\text{Sp}(4, 3)| = 51840,$$

which is *exactly* $\text{Aut}(W(3, 3))$, the 2-qudit Clifford group mod Pauli, and the Clifford group the photon’s own tritter, phase plate and modulator already generate (§6.5, “symplectic closure = 51840”). So the substrate’s symmetry group, the machine’s realised gate group, and the GKP code’s logical gate group are *one group*: the gates act on the code by construction. The code-preserving (transversal) Gaussian gates are the lattice automorphisms $\text{Aut}(D_4) = W(F_4)$ of order $192 \cdot 6 = 1152$, whose S_3 factor is D_4 triality — the three-generation structure — and the remaining $[\text{Sp}(4, 3) : \text{Aut}(D_4)] = 45$ logical Cliffords are teleported using the cubic resource. The logical layer of the holonet is thus a single coherent object, every part of it fixed by $W(3, 3)$. (Witness: `analysis/w33_gkp_clifford_coherence.py`.)

B.4 The code lattices are the moonshine ladder [Face 5]

The same lattices carry one last structure that ties the machine to the deepest arithmetic of the substrate. The holonet’s modes are free bosons, and a GKP code is a free-boson theory compactified on its stabiliser lattice; an even lattice Λ of rank r is in turn the lattice vertex operator algebra (VOA) V_Λ of central charge $c = r$ (Frenkel–Kac–Segal), while the stabiliser code itself fixes the corresponding (Narain) code CFT [26, 27]. The holonet’s GKP code lattices are the substrate tower A_2, D_4, E_8 of ranks 2, 4, 8, hence the lattice VOAs of central charge

$$c(A_2) = 2, \quad c(D_4) = 4, \quad c(E_8) = 8.$$

The top rung, the E_8 VOA at $c = 8$, is exactly the base of the moonshine ladder that $W(3, 3)$ already realises on the physics side,

$$V_{E_8} (c = 8) \longrightarrow V_{\Lambda_{24}} (c = 24) \longrightarrow V^\natural (c = 24, \text{Aut} = \mathbb{M}),$$

with $V^\natural$ the Frenkel–Lepowsky–Meurman Monster module (24 free bosons on the Leech torus, $\mathbb{Z}/2$ -orbifolded), and $\Lambda_{24} \supset E_8^3$ the gauge lattice thrice. So the *same* E_8 is the holonet’s optimal 4-mode GKP code, the substrate’s homology/gauge lattice, and the $c = 8$ chiral VOA that begins moonshine; and the top central charge $c = 24 = f = \chi(K3)$ is the $W(3, 3)$ matter count. The computer’s error-correcting code and the Monster module are one vertex-algebra tower — the architecture and the moonshine residual (R2) are the bottom and top of a single ladder of substrate lattices. (Witness: `analysis/w33_gkp_voa_bridge.py`.)

B.5 What is physically needed: demonstrator versus fault-tolerant machine [Face 6]

It is worth being exact about what we are building, because the design is now complete but the *physical* realisation has two genuinely different layers. The single self-entangled photon of the build sheet — polarisation + time-bin qutrits, a tritter (F_3), a delay ladder and an EOM, read out by single-photon detectors — realises the *ideal, un-encoded* logical layer: the $\text{Sp}(4, 3)$ gate set, the routing network, and every parameter-free witness. It is buildable with today’s quantum optics and it demonstrates the *logic*. But a single photon has fixed photon number; it cannot carry a GKP grid state (a many-photon quadrature lattice). So the photonic demonstrator is not, and cannot be, fault-tolerant.

Fault tolerance is the standard two-layer continuous-variable stack — and here both layers are substrate-fixed, with nothing left to choose:

$$\underbrace{240 \text{ squeezed modes}}_{\text{physical}} \rightarrow \underbrace{120 D_4 \text{ GKP pairs}}_{\substack{\text{inner: analog} \rightarrow \text{digital} \\ \text{gain } 1.5 \text{ dB}}} \rightarrow 240 \text{ GKP qutrits} \rightarrow \underbrace{[[240, 81, 4]]_3}_{\text{outer: Steinberg}} \rightarrow 81 \text{ logical qutrits.}$$

The inner code is the GKP code on the matter-shell lattice D_4 (converting continuous displacement noise into discrete qutrit errors, with the 1.5 dB coding gain that eases the squeezing budget); the outer code is the $[[240, 81, 4]]_3$ Steinberg code on the substrate’s 240 edges. What remains is purely physical, and it is the *universal* CV fault-tolerance bottleneck, not a $W(3, 3)$ -specific one:

1. squeezed light at threshold (~ 10 dB; protocol-dependent, 7–20 dB in the literature), lowered by the D_4/E_8 coding gain (1.5/3.0 dB);
2. GKP qutrit *state generation* — measurement-based breeding from squeezed/cat states with photon-number-resolving detection, or matter-assisted preparation: the hard step;

3. the cubic non-Gaussian resource — the degree-3 E_6 “matter-shell magic” — via a $\chi^{(3)}$ medium or magic-state injection;
4. a programmable beamsplitter/phase/squeeze network implementing the $\text{Sp}(4, 3)$ Clifford (mature integrated photonics);
5. homodyne detection for GKP and Steinberg syndrome readout.

This is the precise, honest statement of what the holonet is and what it still needs: a concatenated $\text{GKP}(D_4) \circ \text{Steinberg}[[240, 81, 4]]_3$ continuous-variable qutrit computer in which $W(3, 3)$ fixes *both* code layers, the gate set ($\text{Sp}(4, 3)$ Gaussian + E_6 cubic) and the network — leaving *zero design freedom above the hardware*. The only open problem is the same one the entire CV fault-tolerance field faces, GKP-state generation at threshold squeezing, and the substrate’s optimal lattices ease rather than remove it. (Witness: `analysis/w33_holonet_physical_stack.py`.)

B.6 The gates are geometric: holonomic Clifford from the gauge connection

There is a third, gate-level layer of robustness that the substrate supplies for free, and it explains the name. The substrate gauge connection is flat along the 240 edges and curved on the matter graph, with binary-tetrahedral $2T$ holonomy. That holonomy group is not arbitrary:

$$2T \cong \text{SL}(2, 3) = \text{Sp}(2, \mathbb{Z}/3), \quad |2T| = 24,$$

the centre being $\{\pm I\}$ and the quotient $2T/\{\pm I\} = A_4$; and $\text{Sp}(2, \mathbb{Z}/3)$ is precisely the single-qutrit Clifford group modulo Pauli. So the gauge connection’s holonomy *is* the qutrit Clifford group, and the gates may be realised as *non-abelian holonomies* (Wilczek–Zee geometric phases) traced in the connection [4, 5]. Holonomic gates depend only on the geometric path, not on pulse timing or shape, so they are intrinsically robust to control error: this is fault tolerance at the *gate* level, complementary to the GKP code (analog level) and the Steinberg/color outer code (logical level). The full two-qutrit Clifford $\text{Sp}(4, 3)$ (order 51840) is generated by the matter-graph holonomy on each qutrit together with the entangling network.

The deeper point is the identity of the connection. The *same* connection whose *curvature* furnishes the substrate’s gauge fields furnishes, through its *holonomy*, the computer’s gates: the gauge interactions of the physics and the logic gates of the machine are one and the same non-abelian holonomy. This is the literal content of “holo-net”. (Witness: `analysis/w33_holonomic_gates.py`, constructing $\text{SL}(2, 3)$ and verifying order 24, symplecticity, centre $\{\pm I\}$, quotient A_4 , and the order spectrum $\{1, 1, 8, 6, 8\}$ of $2T$.)

B.7 The quantum power is the gauge curvature

Read together with the physics tier (§14: the gauge connection is flat on the 240 collinear edges and $2T$ -curved on the matter graph Q , with curvature $F(p, q) = [R_p, R_q]$ a quaternion unit), the holonomic identification sharpens into a single statement: *the gates are the gauge generators, and the entangling power is the curvature*. The 40 long-root transvections R_p — one per point, each the generation- $\mathbb{Z}_3$ source of the Standard-Model spine — are exactly the elementary gates (Wilson loops of the substrate connection), and they generate the whole gate group:

$$\langle R_p : p \in W(3, 3) \rangle = \text{Sp}(4, 3), \quad |\cdot| = 51840.$$

Each point has $12 = k$ *flat* (collinear, causal) neighbours, where $[R_p, R_q] = I$ and $\langle R_p, R_q \rangle = \mathbb{Z}_3 \times \mathbb{Z}_3$ is abelian and classically simulable, and $27 = q^q$ *curved* (matter-graph Q , non-collinear) neighbours,

where the curvature $F = [R_p, R_q]$ is an order-4 quaternion unit and $\langle R_p, R_q \rangle = 2T = \text{SL}(2, 3)$ is non-abelian. So the non-Clifford/entangling resource sits *exactly* on the curved directions: a flat pair gives only the abelian $\mathbb{Z}_3 \times \mathbb{Z}_3$ (classically simulable), and the per-gate non-abelian content requires curvature. Quantitatively, a triangle $p \rightarrow q \rightarrow r \rightarrow p$ carries a Wilson loop $W = R_p R_q R_r$ whose order is the discrete curvature flux: the 160 collinear triangles all carry the uniform abelian flux of order 3, while the 3240 matter triangles of Q carry non-abelian flux of orders $\{2, 4, 6, 12\}$ (counts 180, 180, 1440, 1440, up to $12 = k$), and the discrete Yang–Mills action $S = \sum_{\Delta} (1 - \frac{1}{\text{ord } W})$ concentrates $\approx 96\%$ on the matter graph. In the physics reading the same curvature on Q is the $su(2)$ field strength and, through its $\mathbb{Z}_3$ Yukawa grading, the fermion-mass texture — so on the matter graph,

$$\text{curvature} = \text{entanglement} = \text{magic} = \text{mass},$$

all one object, while flatness is causality. (*Honest scope: this is a field-strength statement. Both the flat and the curved Wilson-loop sets generate all of $\text{Sp}(4, 3)$ — the order-3 flat loops are genuine gates — so the curved/flat split is the per-plaquette curvature and the location of the Yang–Mills action, not a global classical-vs-quantum partition of the group.*) The machine’s quantum advantage is not added to the substrate; it lives in the curvature of the Standard-Model gauge field. (Witnesses: `analysis/w33_gauge_curvature_is_computation.py` (40 order-3 transvections, 12 flat / 27 curved per point, curvature order 4 on Q , $\langle R_p \rangle = \text{Sp}(4, 3)$); `analysis/w33_yang_mills_computation.py` (Wilson-loop census, the 96% matter-graph action).)

B.8 The fuel is contextuality: the register is a Wigner phase space and magic is its negativity

There is a second, resource-theoretic answer to “where does the quantum power come from,” and it is exact on the substrate. Howard, Wallman, Veitch and Emerson proved that for odd-prime-dimensional qudits *contextuality* — equivalently, negativity of the Gross discrete Wigner function — is the necessary resource that lifts Clifford computation to universal quantum computation by magic-state injection [10]. The qutrit $d = q = 3$ is the smallest, cleanest case, and the substrate realizes the whole picture geometrically. The discrete phase space of two qutrits is $\mathbb{Z}_3^4$ with $3^4 = 81 = q^4$ points — *exactly* the logical register $H_1 = 81$ of the $[[240, 81, 4]]_3$ code (irreducible under $\text{PSp}(4, 3)$); the Wigner function lives on the same 81 points the machine computes on. The 40 vertices of $W(3, 3)$ are the $(81 - 1)/2 = 40$ nonzero rays of that phase space — the two-qutrit Pauli operators up to phase — so the substrate *is* the two-qutrit Pauli geometry. The Clifford group acts as $\text{Sp}(4, 3) = \text{Aut}(W(3, 3))$, permuting phase points and preserving Wigner-positivity (the classically simulable sector, Gottesman–Knill). And the non-Clifford states the holonet injects — the Hesse/cubic/ T -type magic of the universal gate set — are exactly the Wigner-*negative* states: a stabilizer state such as $|0\rangle$ or $|+\rangle$ has $W \geq 0$, while the qutrit “strange” state $(|1\rangle - |2\rangle)/\sqrt{2}$ has $\min W = -\frac{1}{3}$ and $\text{mana} \ln \frac{5}{3} > 0$. So the machine is quantum precisely where the substrate’s Wigner function goes negative; the magic it spends is the contextuality of the geometry, and the corpus already measures the geometry’s contextual fraction as $4/40 = 1/\Phi_4$. Curvature (the gauge picture) and Wigner-negativity (the resource picture) are two readings of the same fuel. (Witness: `analysis/w33_contextuality_is_the_fuel.py`.)

B.9 From the matter-graph lattice to continuum Yang–Mills: the loop closes

The discrete curvature of the previous subsection is not a metaphor for a gauge theory — it *is* one, and its continuum limit is the gauge sector of the physics. Each matter edge carries the curvature $F = [R_p, R_q]$, an order-4 element of $2T = \text{SL}(2, 3) \subset \text{SU}(2)$; realised as the binary tetrahedral

group of 24 unit quaternions, these are exactly the quaternion units $\pm i, \pm j, \pm k$ (the six order-4 elements, $SU(2)$ -trace 0). So the matter graph Q carries a genuine *finite-group* ($2T$) lattice gauge theory with an $su(2)$ -valued field strength — the computer’s gate fabric, read as a Wilson lattice. The standard Wilson plaquette action $\frac{\beta}{2}\text{Tr}\{1 - P\}$ assigns each plaquette the curvature energy $1 - \frac{1}{2}\text{Re Tr } P$, maximal ($= 1$) on a curved ($\text{Tr} = 0$) plaquette and zero on a flat one, and Wilson’s theorem gives its continuum limit $\frac{\beta}{2}\text{Tr}\{1 - P\} \rightarrow \frac{a^4}{2g^2}F_{\mu\nu}^b F_{\mu\nu}^b$. On the shape-regular edgewise tower (the same refinement that carries the finite spectral action to the continuum) the matter-graph action converges to the continuum $su(2)$ Yang–Mills integral $\frac{1}{4g^2}\int\text{Tr } F^2$, which is precisely the gauge-kinetic a_4 term of the $W(3,3)$ spectral action. Together with the a_2 Einstein–Hilbert term — whose metric variation gave the Einstein field equations (`w33_paper.tex`) — and the a_0 cosmological term, varying $a_0 + a_2 + a_4$ yields the coupled Einstein + $su(2)$ Yang–Mills + Higgs field equations. The loop therefore closes: *the computer’s gate curvature (a $2T$ lattice gauge theory on Q) and the physics gauge-plus-gravity field equations (the continuum spectral action) are the lattice and continuum descriptions of one $W(3,3)$ connection* — discrete gates below, smooth fields above, the edgewise tower between. (*Honest scope: the discrete $2T$ structure and Wilson density are computed; the continuum limit is the established Wilson and edgewise spectral-action results, not re-derived numerically; the local field strength is $su(2)$ per edge while the global holonomy group is $\text{Sp}(4,3)$.*) (Witness: `analysis/w33_lattice_to_continuum_ym.py`.)

And the discrete theory is already *gapped*, exactly: a gauge theory on a finite graph has the graph Laplacian as its kinetic operator, so its mass gap is the algebraic connectivity of the matter graph. Since $Q = \text{SRG}(40, 27, 18, 18)$ is an expander, that gap is large and substrate-fixed — the adjacency spectrum is $\{27^1, 3^{15}, (-3)^{24}\}$, so the Laplacian $L_Q = 27I - A_Q$ has spectrum $\{0^1, 24^{15}, 30^{24}\}$ and

$$\text{mass gap} = 24 = f,$$

the smallest nonzero eigenvalue, at multiplicity $g = 15$ (the gauge-sector dimension), with the top eigenvalue 30 at multiplicity $f = 24$ (the matter dimension). The Laplacian multiplicities are the Standard-Model sector sizes and the gap is the matter count; confinement is structural, forced by the expander geometry of Q , with no continuum limit needed to certify that the gap is nonzero. (*This is the gap of the exact finite theory — the substrate’s shadow of the Yang–Mills mass gap, not a continuum proof.*) (Witness: `analysis/w33_mass_gap_from_matter_graph.py`.)

B.10 Why “holo-net”: the network is a holographic code

The name is earned, not decorative: the architecture is a holographic quantum error-correcting code in three exact senses, each read straight off the substrate spectrum, and each tying a physics fact to a machine component. *(i) Black-hole entropy is the code distance.* A horizon is a stabilizer edge-cut, and the Bekenstein–Hawking law $S_{\text{BH}} = A/4$ has its universal $\frac{1}{4}$ equal to the code’s distance bound $1/d_Z$: the $W(3,3)$ CSS code has $d_Z = 4 = \mu$, so $S_{\text{BH}} = A/d_Z = A/\mu$. The black-hole quarter is the QEC distance. *(ii) Discrete AdS/CFT.* A holographic duality maps a negatively curved bulk to a conformal boundary. The collinearity graph has adjacency spectrum $\{12, 2^{24}, (-4)^{15}\}$; its single *negative* (hyperbolic) eigenvalue -4 has multiplicity $g = 15 = \dim \text{SO}(4, 2)$, the 4-dimensional conformal group. The 15 negative-curvature bulk modes match the 15 conformal-boundary generators exactly — discrete AdS/CFT with no continuum strings. (Dually, the matter graph’s Laplacian $\{0, 24^{15}, 30^{24}\}$ puts the same g/f multiplicities on the gauge mass-gap spectrum.) *(iii) The fractal is the holographic principle.* The nested-shell fractal of §7 is boundary-encodes-bulk made literal: each outer shell’s $[[240, 81, 4]]_3$ code holographically contains the logical information of the inner layers, exactly the concatenated holographic code structure (the Pastawski–Yoshida–Harlow–Preskill

picture), with the substrate fixing the dictionary. So the holonet is a holographic code whose entropy law, conformal boundary, and bulk–boundary map are all substrate arithmetic — which is the literal content of “holo-net”. (Witness: `analysis/w33_holographic_architecture.py`.)

B.11 The boundary central charge is $f = 24$: the Monster CFT is the holonet’s holographic dual

The holographic boundary even has a definite central charge, and it is the substrate’s master invariant. The code lattices climb the moonshine ladder established above — A_2 at $c = 2$, D_4 at $c = 4$, E_8 at $c = 8$, then Leech and the Monster module $V^\natural$ at $c = 24$ — and the top rung fixes the boundary:

$$c = 24 = f = \chi(K3) = \text{rank}(\Lambda_{24}) = 2k = q^q - q,$$

the matter count, the $K3$ Euler number, and the Leech rank, all one number. At $c = 24$ the boundary is the *extremal* holomorphic CFT, whose partition function is $Z = j(\tau) - 744 = q^{-1} + 196884q + \dots$ with graded dimensions $V_0 = 1$, $V_1 = 0$, $V_2 = 196884 = 1 + 196883$ (the McKay vacuum-plus-smallest-irrep identity), and the Monster is its unique symmetry. By Witten’s analysis of pure three-dimensional gravity this extremal $c = 24$ CFT *is* the dual of pure AdS_3 gravity [8], so its graded dimensions count the AdS_3 black-hole microstates — and those dimensions are exactly the quintic moonshine data of the substrate (e.g. $\dim V_5 = 333\,202\,640\,600$ from the master manuscript). Finally, Almheiri–Dong–Harlow showed that AdS/CFT *is* a quantum error-correcting code [9]; the holonet is that code, with bulk $W(3,3)$ and boundary the Monster CFT. So one substrate number, $c = f = 24$, ties the holographic boundary, the moonshine module, the matter sector, and pure 3D gravity into a single object: the holonet is the quantum error-correcting code that AdS/CFT always was, and its boundary is monstrous moonshine. (Witness: `analysis/w33_holographic_central_charge.py`.)

Two consequences sharpen this. First, the entropy is a microstate count: the graded dimensions $\dim V_n$ are the boundary state numbers, so the black-hole entropy is $S(n) = \ln \dim V_n$, and it follows the Cardy law $S \sim 2\pi\sqrt{cn/6} = 4\pi\sqrt{n}$ at $c = 24$ — the Bekenstein–Hawking area law of the BTZ black hole. The leading area term captures the entropy to $\sim 5\%$ (the slow Cardy approach) and the Petersson–Rademacher-corrected form to $\sim 2\text{--}3\%$ by $n \sim 5$, so the moonshine dimensions are a Strominger–Vafa-style microstate count of the substrate’s own horizon, with $S = A/\mu$. Second — and this is why the holographic story is not a separate layer — the holographic *redundancy is the error correction*. A holographic code over-encodes each bulk qudit on many boundary regions, and that over-encoding *is* the code distance; on the holonet the identical redundancy appears three ways at once: the fractal nested shells (each outer $[[240, 81, 4]]_3$ shell re-encoding the inner layers), the timetable triple-storage (§12), and the holographic boundary. So the network’s holographic/fractal redundancy and the computer’s fault tolerance are one mechanism — the precise sense in which the computer is the network. (Witness: `analysis/w33_microstate_count.py`.)

The central charge even explains the generation number. The matter sector’s affine symmetry is the WZW model $\mathfrak{sp}(4)$ at level $k = 12$ — the level is the graph degree — with $c = 10 \cdot 12 / (12 + 3) = 8 = \text{rank}(E_8)$, the same c as $(E_8)_1$. So the matter rung carries one E_8 of central charge, and the boundary’s $c = 24 = f$ is exactly

$$24 = q \cdot 8 = q \cdot \text{rank}(E_8),$$

three generations of that matter E_8 rung; the jump from the matter rung to the full boundary *is* the generation number $q = 3$. The lattice picture makes it concrete: among the $24 = f$ Niemeier (even unimodular rank-24) lattices — the boundary’s possible realizations, their count

equal to the central charge itself — the Leech realization is the extremal Monster CFT, while E_8^3 is literally three copies of the matter E_8 , the three-generation realization. “Why three generations” and “why $c = 24$ ” are one fact: the boundary is q copies of the matter E_8 . (Witness: `analysis/w33_wzw_generation_ladder.py`.)

And that single matter E_8 rung already contains the grand unification and the cosmological constant. The genuine conformal embeddings of $c = 8$ algebras into $(E_8)_1$ are exactly the substrate’s physics: $(E_6)_1 \times (A_2)_1$ is the grand-unified $E_8 \rightarrow E_6 \times \text{SU}(3)_{\text{color}}$ ($248 = 78 + 162 + 8$, the 27 matter and the 8 gluons), while $(D_8)_1 = \text{SO}(16)_1$ is the *cosmological-constant* sector — $\dim \text{SO}(16) = 120 = vq$, with $E_8 = \text{SO}(16) \oplus \mathbf{128}$. (The affine $\mathfrak{sp}(4)_{12}$ shares $c = 8$ but is a distinct CFT, not a conformal subalgebra of E_8 — the central charge matches, the embedding does not.) That $\text{SO}(16)$ sector is where the holography meets thermodynamics: the substrate is discrete de Sitter ($\kappa = 2/k = 1/6 > 0$, with $\sum_{\text{edges}} \kappa = v$ by Gauss–Bonnet), its horizon radiates at the Gibbons–Hawking temperature $T_H = k/2\pi$ — which is exactly the temperature of the modular-flow clock — and its entropy is A/μ , so *time, temperature and the cosmological constant are one thermal fact*, the arrow of time being the de Sitter expansion. On the matter/topological side the same unification recurs: the double $D(2T)$ has 6 abelian and $36 = 4 \cdot 9$ non-abelian anyons (total quantum dimension $f = 24$), but $2T$ is solvable, so braiding alone is non-universal — the magic that lifts it to universality is precisely the Wigner negativity of §B.8. Gauge curvature, contextuality, and topological magic are three names for one resource. (Witnesses: `analysis/w33_e8_conformal_embeddings.py`, `analysis/w33_thermal_cosmology.py`, `analysis/w33_topological_magic.py`.)

Three further consequences close the holographic picture, and they share a single constant $\mu = 4$. (a) *The dark sector is the $\mathbf{128}$* . The visible half of E_8 is $\text{SO}(16)$ ($\dim = 120 = vq$, the gauge/cosmological-constant sector); the remainder is the spinor $\mathbf{128} = 2^{\Phi_6}$, and under $\text{SO}(10) \times \text{SU}(4)$ it is $(\mathbf{16}, \mathbf{4}) \oplus (\overline{\mathbf{16}}, \overline{\mathbf{4}})$ — Standard-Model families (the $\mathbf{16}$) replicated under a *hidden* $\text{SU}(4) = \text{SO}(6)$ with $\mu = 4$ dark colors, so $E_8 = \text{SO}(16) \oplus \mathbf{128}$ is a visible/dark split and the holographic boundary predicts dark matter. (b) *Ryu–Takayanagi is exact and discrete*: a boundary region A has entanglement entropy $S(A) = \delta(A)/\mu$, its minimal edge-cut over the Bekenstein factor (a point gives q , a GQ line gives q^2), with a symmetric Page curve; by the JLMS equality of bulk and boundary modular Hamiltonians the machine reconstructs its own bulk, the entanglement-wedge threshold being the code distance $d_Z = \mu$. (c) *$q = 3$ is selected thermodynamically*: the expanding de Sitter horizon (temperature $T_H = k/2\pi$, curvature $\kappa = 2/k$, entropy A/μ) closes its first-law/Gauss–Bonnet condition $2(q-1)(q^2+1) = (1+q)(1+q^2)$ *only* at $q = 3$ — the thermal clock is self-consistent only there. The unifying thread: the same $\mu = 4$ is the Bekenstein factor, the QEC distance, the RT $1/\mu$ and its reconstruction threshold, the de Sitter entropy factor, *and* the number of dark colors — one holographic quantum. (Witnesses: `analysis/w33_dark_sector_128.py`, `analysis/w33_rt_bulk_reconstruction.py`, `analysis/w33_desitter_q3_selection.py`.)

Pushing the dark sector further collapses it into the bulk itself. (a) The hidden $\text{SU}(4)$ that carries the $\mathbf{128}$ is a confining gauge theory: with $N_c = \mu = 4$ dark colors and the $\mathbf{16}$ -spinor’s $N_f = 16$ flavors its one-loop coefficient is $\beta_0 = (44 - 32)/3 = 4 = \mu > 0$ (versus the visible $\beta_0 = \Phi_6 = 7$), so it is asymptotically free and *confines* — dark matter is a colorless *dark hadron* at the dark confinement scale, a falsifiable alternative to a fundamental WIMP (the relic mass needs UV matching and is not derived). (b) That same $\text{SU}(4) = \text{SO}(6)$ has dimension $4^2 - 1 = 15 = g$, and its non-compact real form is $\text{SO}(4, 2)$, the 4D conformal group — the discrete-AdS/CFT bulk isometry (the 15 negative-curvature modes). So *one* dimension-15 = g symmetry is at once the dark color, the Pati–Salam unifier (lepton number as the fourth color), and the holographic bulk, with its fundamental $\mathbf{4} = \mu$ being both the dark-color count and the spacetime dimension: the dark colors are the spacetime directions. (c) And the information story closes inside that bulk: the discrete RT entropy $S(A) = \delta(A)/\mu$ of a growing radiation region rises, peaks at the balanced cut, and

falls, the *island* (the complementary matter region) taking over at the Page point $|A| = v/2 = 20$ — the $[[40, 20, 11]]$ code — so evaporation is manifestly unitary and information is recovered. Dark matter, the holographic bulk, grand-unified color, the dimension of spacetime, and the Page curve are facets of the one $SU(4) = SO(6) = 15 = g$. (Witnesses: `analysis/w33_dark_matter_mass.py`, `analysis/w33_su4_is_spacetime.py`, `analysis/w33_page_curve_island.py`.)

Three last threads make the dark/bulk $SU(4)$ quantitative, dynamical, and temporal. (a) *A dark-matter scale from E_8* . Matching the dark $SU(4)$ and visible $SU(3)$ couplings at the unification scale gives $\Lambda_{\text{dark}} = M_{\text{GUT}}(\Lambda_{\text{QCD}}/M_{\text{GUT}})^{\beta_0^{\text{vis}}/\beta_0^{\text{dark}}}$; exponentially sensitive to β_0^{dark} , it lands the dark hadron at *tens of GeV* when $\beta_0^{\text{dark}} = 8 = k - \mu$ (the gluon count, $N_f = \Phi_4 = 10$), matching the corpus's ~ 22.8 GeV dark-matter value — a *confining* alternative to the 308 GeV WIMP ($g v_{EW}/k$) the corpus also records. (b) *The expanding holography is dS/CFT* . Because the substrate is de Sitter, the bulk Laplacian modes $\{0^1, 10^{24}, 16^{15}\}$ read as vacuum + matter ($f = 24$, the boundary central charge) + the $15 = g$ de Sitter isometries (the adjoint of $SO(4, 2) = SU(4)$, the same dark/bulk group), with the Monster $c = 24$ CFT at future infinity — the expanding counterpart of the corpus's discrete AdS/CFT. (c) *Time is dark braiding*. The clock $2T$ sits at the bottom of the chain $2T < SU(2) < SU(4)$, so one tick is a topological twist of the $D(2T)$ dark anyons — the modular T -matrix, of order $\text{lcm}(4, 6) = 12 = k$, the clock period. So dark matter, the holographic bulk, the spacetime, and the arrow of time are nested in one $SU(4) = SO(6) = 15 = g$: the dark colors are the spacetime directions, and time is the dark anyons braiding. (Witnesses: `analysis/w33_dark_lambda_gut.py`, `analysis/w33_dscft_modes.py`, `analysis/w33_clock_is_dark_braiding.py`.)

The same expanding boundary is also the origin of the cosmic microwave background. A de Sitter substrate *is* an inflating universe, and dS/CFT organizes its perturbations: of the bulk modes, exactly one is light — the complementary-series vacuum mode ($m^2 = 0 < \frac{9}{4}$), the inflaton — while the matter ($m^2 = 10$) and isometry ($m^2 = 16$) modes are heavy principal-series modes that decay. So there is a unique scalar to drive inflation, and its spectrum is near-scale-invariant precisely because of the de Sitter conformal symmetry of the boundary. The Starobinsky e -fold count is $N = 2(v - \Phi_4) = 2(40 - 10) = 60$, giving

$$n_s = 1 - \frac{2}{N} = \frac{29}{30} = 0.9667, \quad r = \frac{k}{N^2} = \frac{1}{300} = 0.00333,$$

matching Planck ($n_s = 0.9649 \pm 0.0042$) with the tensor amplitude carrying the degree $k = 12$. And the primordial two-point function is a boundary CFT correlator on the *same* Monster $c = 24$ future-infinity boundary that carries the dark/bulk $SU(4)$ — so the CMB fluctuations are correlations in the moonshine module. Inflation, dark matter, and the arrow of time share one expanding dS/CFT boundary. (Witness: `analysis/w33_inflation_dscft.py`.)

These pieces assemble into one early-universe history, with the **128**'s right-handed neutrino as the linchpin. *Inflation* ends and reheats, producing the heavy N_R — the $SO(10)$ -singlet inside each family's **16** in the dark spinor $\mathbf{128} = (\mathbf{16}, \mathbf{4}) \oplus (\overline{\mathbf{16}}, \overline{\mathbf{4}})$. *Seesaw*: N_R has a Majorana mass at the dark/GUT scale $M_R \sim 10^{15}$ GeV $\sim M_{\text{GUT}}$, so the type-I seesaw $m_\nu = m_D^2/M_R \sim 0.05$ eV gives the light neutrino masses (corpus $\sum m_\nu = 58$ meV, cyclotomic PMNS). *Cogenesis*: the *same* N_R decays out of equilibrium with CP violation, seeding both the baryon asymmetry by leptogenesis ($\log_{10} \eta_B = -|E|/(v - k - \lambda) = -240/26 = -9.23$, observed -9.21) and an equal dark hidden- $SU(4)$ asymmetry — so dark matter is asymmetric and $\Omega_{\text{DM}} \sim \Omega_b$ is automatic, with $\Omega_{\text{DM}}/\Omega_b = 82/15$. *CMB*: the single light inflaton mode makes inflation single-field, so the non-Gaussianity is the consistency value $f_{\text{NL}}^{\text{local}} = \frac{5}{12}(1 - n_s) = \frac{5}{6N} = \frac{1}{72}$, with the bispectrum *shape* the Monster $c = 24$ boundary three-point function. So one substrate runs the whole sequence — inflation, the moonshine CMB, the neutrino masses, the matter-antimatter asymmetry, and the dark matter — off the dS/CFT boundary and the **128** spinor, joined at the right-handed

neutrino. (Witnesses: `analysis/w33_cmb_nongaussianity.py`, `analysis/w33_cogenesis.py`, `analysis/w33_neutrino_seesaw_128.py`.)

This whole history hangs on a substrate *scale ladder* and ends in falsifiable numbers. The energy scales descend as closed forms — $M_{\text{Pl}} > M_{\text{GUT}} = v_{EW} 10^{2\Phi_6} > M_R \sim 10^{15} > T_{\text{RH}} > v_{EW} = |E| + 2q > \Lambda_{\text{dark}} > \Lambda_{\text{QCD}}$ — and the thermal history walks down them (inflation, reheating, the N_R lepto/cogenesis, electroweak, dark and QCD confinement). The ladder corrects a loose expectation: $M_R \gg T_{\text{RH}}$, so *leptogenesis is non-thermal* (N_R from inflaton decay), and T_{RH} is a different rung from Λ_{dark} . At the top, the primordial suite is one moonshine sheet from $N = 2(v - \Phi_4) = 60$: $n_s = \frac{29}{30}$, $r = \frac{1}{300}$, $f_{\text{NL}} = \frac{1}{72}$, and the new running $dn_s/d \ln k = -2/N^2 = -\frac{1}{1800}$. At the seesaw rung, the cyclotomic PMNS angles ($\sin^2 \theta_{12} = \frac{4}{13}$, $\sin^2 \theta_{23} = \frac{7}{13}$, $\sin^2 \theta_{13} = \frac{2}{91}$) with the $\mathbb{Z}_3$ Majorana phases ($120^\circ, 240^\circ$) give the neutrinoless-double-beta mass $m_{\beta\beta} \approx 2.3$ meV and the Dirac phase $\delta_{CP} = 14\pi/13 \approx 194^\circ$ — the same phase that drives the cogenesis. So the substrate’s cosmology is a falsifiable program: tiny f_{NL} and running for CMB-S4/LiteBIRD, and $m_{\beta\beta} \approx 2.3$ meV for nEXO/LEGEND. (Witnesses: `analysis/w33_cmb_moonshine_suite.py`, `analysis/w33_scale_ladder.py`, `analysis/w33_neutrinoless_betabeta.py`.)

B.12 The falsifiability ledger, and the demonstrator as an experiment on the theory

A theory of everything is worth only the experiments that can refute it, so the program collects into one ledger: $r = \frac{1}{300}$ and $f_{\text{NL}} = \frac{1}{72}$ and the running $-\frac{1}{1800}$ (CMB-S4/LiteBIRD), $m_{\beta\beta} \approx 2.3$ meV and $\sum m_\nu = 58$ meV and $\delta_{CP} \approx 194^\circ$ (nEXO/LEGEND, DESI, DUNE), a confining tens-of-GeV dark hadron (LZ/XENONnT) with a dark-confinement gravitational-wave background peaking in the LISA band ($f \sim 10^{-4}$ Hz, a fourth GW source beyond the primordial/electroweak/cosmic-string ones), and the collider/proton-decay handles — each killable by a named experiment in 2028–2035.

The sharpest point is that the loop closes: the tabletop single-photon demonstrator *is* one of these experiments, and it can run *now*. It reads the substrate primitives directly as quantized or protected optics. The flagship is the topological frequency-conversion pump: the photon’s three modes form a spin-1, and a two-tone drive pumps energy quanta between the tones at the rate $(C/2\pi)\hbar\omega_1\omega_2$ with C the Floquet-band Chern number; for spin-1 the bands carry $\{+2, 0, -2\}$, so the pump quantum is the *topologically protected* integer $|C| = 2S = \lambda = q - 1 = 2$. A Kochen–Specker test on the 40 Pauli rays gives contextual fraction $4/40 = 1/\Phi_4 = \frac{1}{10}$; the BC clock beats at $\arccos(-\frac{2}{3}) = 131.8^\circ$; the oscillator gate spectrum is $q \pm \sqrt{\lambda} = 3 \pm \sqrt{2}$. So the demonstrator measures λ (twice, once protected), q , and Φ_4 on a bench: if the pump quantum is not 2, or the contextual fraction is not $1/10$, the $W(3,3)$ substrate is falsified at the tabletop. The machine we are building is an experiment on the theory of everything. (Witnesses: `analysis/w33_falsifiability_ledger.py`, `analysis/w33_dark_confinement_gw.py`, `analysis/w33_demonstrator_measures_substrate.py`.)

Three results sharpen this to a theorem, a protocol, and a single picture. *The pump is protected*. A $(2S+1)$ -level qudit driven by two incommensurate tones has extreme Floquet-band Chern number $2S$ (verified for $S = \frac{1}{2}, 1, \frac{3}{2}$ with $|C|_{\text{max}} = 1, 2, 3$), so the qutrit’s pump quantum is $2S = q - 1 = \lambda = 2$ for *any* drive realizing the winding, protected by a finite Floquet gap $\Delta_{\text{min}} \sim \mathcal{O}(1)$ that fixes the spec (drive rate and coherence/temperature below the gap). *The contextuality is a protocol*. The 40 Pauli rays carry the 40 self-dual $\text{GQ}(3,3)$ lines as measurement contexts (four commuting Paulis each, four contexts per ray); a state-independent Kochen–Specker test over them has contextual fraction $4/40 = 1/\Phi_4 = \frac{1}{10}$, a second tabletop falsifier. *And the whole theory is one graph*. Every claim in this paper — the Standard Model, the exceptional/moonshine tower with Monster $c = 24$, the holographic boundary and code, the dark $\text{SU}(4)/\mathbf{128}/N_R$ sector, the cosmology ledger, the runtime, and the bench falsifiers — forms one connected acyclic

dependency graph whose unique root is the selection $q! = 2q$: every node is reachable from $q = 3$. The theory of everything is the unfolding of a single integer, and it ends on numbers a bench and a telescope can check. (Witnesses: `analysis/w33_pump_protection_theorem.py`, `analysis/w33_kochen_specker_protocol.py`, `analysis/w33_grand_dependency_map.py`; graph: `docs/w33_grand_dependency_map.dot`.)

B.13 The geometric origin of the clock: the triangulation ladder $3 \rightarrow 4 \rightarrow 7$

Where does the clock come from, geometrically? From the same place as the substrate integers: closing simplices. The master quadratic factors as $(x - \lambda)(x - \mu) = (x - 2)(x - 4)$ from the seed $q! = 2q$, and the two simplices that build everything are the *triangle* and the *tetrahedron*. Three points close a loop into a 2-simplex; closing that third edge fixes one orientation — it “saves” a single $q = 3$ trit (the Klee–Irwin cycle-clock / trit-saving picture), so the minimal closed cell carries exactly one trit. Four points ($\mu = q + 1$) close the 3-simplex, the tetrahedron — self-dual (K_4 vertices and K_4 faces, both maximally adjacent), genus 0, and the building block of the Boerdijk–Coxeter helix whose stack twist $\arccos(-\frac{q-1}{q}) = \arccos(-\frac{2}{3})$ is the time-quasicrystal clock.

The torus is born when these two combine. The Ringel–Youngs genus of the complete polyhedral graph is $g(K_n) = \lceil (n-3)(n-4)/12 \rceil$, with denominator $12 = k$. At $n = \Phi_6 = 7$ the numerator is

$$(n-3)(n-4) = \mu \cdot q = 4 \cdot 3 = 12 = k, \quad \text{so} \quad g = \frac{\mu q}{k} = \frac{k}{k} = 1,$$

the torus — and the “3” and “4” in that numerator are precisely the triangle ($n-4 = q$) and the tetrahedron ($n-3 = \mu$). The genus is 1 exactly because triangle $\times$ tetrahedron equals the degree. (K_4 , the tetrahedron itself, gives $(1)(0)/12 = 0$, the sphere.) The two genus-1 realizations are the only “complete” toroidal polyhedra: the Szilassi polyhedron, with 7 mutually edge-adjacent faces (the tetrahedron and Szilassi are the *only* polyhedra in which every pair of faces touches), and its dual the Császár polyhedron, with 7 mutually adjacent vertices (K_7). They share $21 = \binom{7}{2}$ edges — the Fano flags — and their incidence is the Heawood graph ($14 = 7+7$ vertices, 21 edges), which is exactly the Fano clock whose Laplacian middle shell gives the oscillator frequency $\omega = \sqrt{\lambda} = \sqrt{2}$ of §B.14. So the two toroidal polyhedra *are* the clock oscillator, and the tetrahedral BC helix is its genus-0 companion: one triangulation ladder $3 \rightarrow 4 \rightarrow 7$, the triangle and the tetrahedron, builds the clock at both genera. (Witness: `analysis/w33_genus_ladder_clock.py`.)

And the ladder takes us all the way home: the same simplices generate the surfaces, the code, and the clock at once. *Surfaces (climbing to genus 3)*. Above the sphere (tetrahedron, $g = 0$) and the torus (Fano/Heawood, $g = 1$) the next maximally-symmetric rung is the Klein quartic, the lowest-genus Hurwitz surface ($g = 3$): a regular map of 24 heptagons / 56 triangles / 84 edges with automorphism group $\text{PSL}(2, 7)$, $|\text{PSL}(2, 7)| = 84(g-1) = \lambda k \Phi_6 = \Phi_6 f = 168$. Its numbers are substrate integers ($24 = f$, $56 = v+k+\mu = \dim E_6^{\text{fund}}$ analogue, $84 = k \Phi_6$), and that $\text{PSL}(2, 7)$ is exactly the clock/readout layer of §B.14 — so the heptagon $\Phi_6 = 7$ threads the whole ladder, the Fano heptad becoming the genus-3 heptagonal surface. *Code (trit-saving is encoding)*. Closing a triangle saves one $q = 3$ trit: on the $W(3, 3)$ chain complex ($40, 240, 160$, $\partial^2 = 0$) the 160 triangles are the weight-3 CSS Z -checks (rank 120) and the vertices the X -checks (rank 39), so the trits no triangle can save are $H_1 = (240-39)-120 = 81 = q^4$, the logical register of the $[[240, 81, 4]]_3$ code — closing a loop and encoding a qutrit are one act. *Clock (the helix is a quasicrystal)*. The tetrahelix twist $\arccos(-\frac{2}{3})$ is Niven-irrational, so the Boerdijk–Coxeter helix never closes in flat 3D — the aperiodic, two-incommensurate-frequency time-crystal clock — yet closes at exactly 30 tetrahedra in S^3 (the 600-cell = 20×30), and $30 = h(E_8)$ is the middle factor of the supercycle $51840 = 24 \cdot 30 \cdot 72 = |\text{Sp}(4, 3)|$. So from the single integer $q = 3$ — three points closing a triangle, saving a trit — unfold the surfaces (genus 0, 1, 3),

the error-correcting code ($H_1 = 81$), and the quasicrystal clock (supercycle $|\text{Sp}(4, 3)|$): geometry, code, and time are one triangulation. (Witnesses: `analysis/w33_klein_quartic_genus3.py`, `analysis/w33_trit_saving_code.py`, `analysis/w33_bc_helix_quasicrystal.py`.)

B.14 The clock is a mass and the supercycle is the gauge group: running the machine is the physics

The movement closes on the runtime itself. The holonet’s clock is not an external metronome — it is a *topological harmonic oscillator* living on the Heawood graph (the incidence graph of the Fano plane $\text{PG}(2, 2)$, the same 7 points and 7 lines that the time-bin guard shell uses). Its Laplacian spectrum is $\{0, (q - \sqrt{\lambda})^6, (q + \sqrt{\lambda})^6, 2q\} = \{0, (3 - \sqrt{2})^6, (3 + \sqrt{2})^6, 6\}$, and on the 12-dimensional middle shell

$$(L_H - qI)^2 = \lambda I, \quad \omega = \sqrt{\lambda} = \sqrt{2}, \quad E_{\pm} = q \pm \sqrt{\lambda} = 3 \pm \sqrt{2},$$

two six-mode $\text{Cl}(1, 1)$ branches — an exact discrete oscillator whose frequency is $\sqrt{\lambda}$. That frequency is not abstract: it is a mass. The heaviest fermion sits at

$$m_{\text{top}} = \frac{v_{EW}}{\sqrt{\lambda}} = \frac{246}{\sqrt{2}} = 173.95 \text{ GeV},$$

so the machine’s clock *rate* is the top-quark mass scale. And the runtime stack $8 \rightarrow 48 \rightarrow 24 \rightarrow 72 \rightarrow 2160 \rightarrow 51840$ tops out at the full Clifford supercycle

$$51840 = 720 \cdot 72 = 24 \cdot \underbrace{30}_{h(E_8)} \cdot \underbrace{72}_{q^2 \cdot 8} = |\text{Sp}(4, 3)|,$$

with the 72-tick frame $= q^2 \times$ the eight-tick word and the 2160 mirror bus $= h(E_8) \times$ frame. The supercycle is the order of the gauge group. So executing the machine for one supercycle is one complete traversal of $\text{Sp}(4, 3)$, ticked by the oscillator whose frequency is the top mass: the holonet does not *simulate* the physics on a clock, its clock is the mass-setting oscillator and one run of its instruction stream is the gauge dynamics. Running the machine and enacting the physics are the same act. (Witness: `analysis/w33_machine_clock_is_mass.py`.)

Yet the clock is a *separate* module, and its arithmetic says why. The clock/readout layer is the Fano plane $\text{PG}(2, 2)$ with symmetry $\text{PSL}(2, 7) = \text{PSL}(3, 2)$ of order $168 = 2^3 \cdot 3 \cdot 7$ (the same layer that carries the $168 = 7 \cdot f$ detector-bin weld), while the gauge/state layer is $\text{Sp}(4, 3)$ of order $51840 = 2^7 \cdot 3^4 \cdot 5$. The clock prime $7 = \Phi_6(3)$ — the external cyclotomic prime that also fixes β_0^{QCD} and the atmospheric angle — *does not divide* the gauge order, so $\text{PSL}(2, 7)$ is no subgroup of $\text{Sp}(4, 3)$: the clock’s Fano symmetry can never be a gauge operation, which is the group-order reason the Heawood clock is a separate homology module rather than a $W(3, 3)$ subgraph (the girth obstruction). The two layers are coprime except for

$$\text{gcd}(51840, 168) = 24 = f,$$

and both realize that shared scale as an order- f subgroup — the single-qutrit Clifford $2T = \text{SL}(2, 3)$ inside $\text{Sp}(4, 3)$, and the Fano point-stabilizer $S_4 (= 168/7)$ inside $\text{PSL}(2, 7)$. So the clock layer $= 7 \times f$ and the gauge layer $= (h(E_8) \cdot 72) \times f$ are welded only at the matter count $f = 24$, which is exactly where the clock–matter spectral resonance lands. The clock and the readout are one Fano layer, coprime to the computation and bonded to it at the 24-cell core. (Witness: `analysis/w33_clock_gauge_two_layer.py`.)

B.15 Four further faces of the machine: time, topological order, scrambling, and periodic orbits

Four short, exact results round out the physical reading of the machine. (i) *Time is the modular flow, and only the clock can source it.* By the Connes–Rovelli thermal-time hypothesis, physical time is the modular (Tomita–Takesaki) automorphism flow of the state, generated by $K = -\ln \rho$. On the maximally mixed gauge/matter sector $\rho = I/81$ one has $K \propto I$, so the modular flow is trivial: the H_1 register is *timeless*. On the Fano/Heawood clock the Gibbs state gives $K = \beta L_H + (\ln Z)I$, so the modular flow is generated by the clock Laplacian and ticks at frequency $\sqrt{\lambda} = \sqrt{2}$ — time emerges from the state, and only the clock layer sources it (KMS verified). (ii) *The matter sector is a non-abelian topological order.* The $2T$ gauge theory on Q is the quantum double $D(2T)$: 7 flux sectors (the conjugacy classes of $2T$), 42 anyons = Catalan $C_5 = 2q\Phi_6$, and total quantum dimension $D = |2T| = 24 = f$. Topological protection is then concrete — logical information is anyon fusion/braiding, with total quantum dimension the matter count. (iii) *The machine is a near-optimal scrambler/decoder.* Because $W(3, 3)$ is Ramanujan, its random walk mixes in the minimum time (total-variation mixing time = 2 = diameter) and operator growth runs at the maximal non-backtracking rate $\sqrt{k-1} = \sqrt{11}$; as a holographic code it is therefore a Hayden–Preskill decoder, with a bulk perturbation recoverable from $O(\log)$ boundary modes essentially at once. (iv) *The spectrum is a periodic-orbit trace formula.* The closed non-backtracking geodesics (particle worldlines) are counted by $N_m = \text{Tr}(B^m)$ on the Hashimoto operator ($N_3 = 960 = 6 \cdot 160$); the Ihara zeta factorizes with all non-trivial zeros on the critical circle $|u| = 1/\sqrt{11}$ — an exact graph Riemann Hypothesis, equivalent to the Ramanujan property, with topological entropy $\ln(k-1) = \ln 11$. (Witnesses: `analysis/w33_thermal_time_clock.py`, `analysis/w33_anyons_from_2T.py`, `analysis/w33_scrambling_decoder.py`, `analysis/w33_periodic_orbit_spectrum.py`.)

These four are not separate curiosities; cross-read against the corpus they are one thermal object, unified by the Hawking temperature $T_H = k/2\pi$. The modular flow of (i) is a KMS state at exactly that T_H , so time, temperature and the clock coincide; the scrambling time of (iii) is $\lambda = 2$ with Page entropy $S_{\max} = E = 240$, the same fast-scrambler the corpus records. The topological order of (ii) is the *non-abelian* Drinfeld double $D(2T)$ ($42 = C_5 = v + \lambda$ anyons, total quantum dimension f), complementary to the *abelian* double $D((\mathbb{Z}/3)^2)$ of dimension $q^\mu = 81$ that is the logical register itself — one double per sector. And the zeta of (iv), counting non-backtracking geodesics $\text{Tr}(B^m)$, certifies the same Ramanujan / graph Riemann Hypothesis as the corpus’s point-counting $\text{Tr}(A^n)$, from the complementary walk operator. So the new thermal-time/modular layer is what binds the pre-existing thermodynamic, topological, and arithmetic pillars into a single structure. (Witness: `analysis/w33_thermal_synthesis.py`.)

B.16 A near-term falsification test: the demonstrator as an experiment on the substrate

The single-photon demonstrator of the build sheet is buildable now and needs no GKP states or squeezing, yet the holonomic-gate and quasicrystal-clock structure make it a *discriminating, parameter-free* test of the substrate hypothesis itself. Three signatures, all measurable on that machine:

1. **Geometric-gate timing-independence (the discriminator).** If the gates are Wilczek–Zee holonomies (the substrate claim), the gate unitary depends only on the *path* traced in control space, not on the schedule. Reparametrising the loop — any speed, any pulse shape with the same path — leaves the gate, read out through the trace–Choi visibility

$V(U) = |\text{Tr } U|/3$, *invariant*; a generic dynamical gate would change. A measured timing-dependence falsifies the holonomic claim.

2. **Holonomy-group fingerprint.** Composing the geometric gates must close into $2T = \text{SL}(2, 3)$: order 24, element-order spectrum $\{1, 1, 8, 6, 8\}$, and a single-qutrit Clifford visibility spectrum $V \in \{0, \frac{1}{3}, \frac{1}{\sqrt{3}}, 1\}$ (with $V(F_3) = \frac{1}{3}$ and $V(X) = V(Z) = 0$). Gate-set tomography reads this out; any other group falsifies the gauge-connection-holonomy identification.
3. **Provably aperiodic clock.** The Boerdijk–Coxeter drive rotates by $\theta = \arccos(-\frac{2}{3}) \approx 131.81^\circ$ per pass. By Niven’s theorem [18] the only rational $\cos(r\pi)$ with r rational are $\{0, \pm\frac{1}{2}, \pm 1\}$, and $-\frac{2}{3}$ is none of them, so θ/π is irrational and the drive *never recurs* — a provably aperiodic time-quasicrystal clock. A measured recurrence at any finite pass falsifies it.

Each signature is fixed by the geometry with no free parameter, and each is within reach of the demonstrator’s existing optics. Together they turn the buildable machine into an experiment *on the substrate*: not a test of whether the device computes, but of whether the gauge connection, the holonomy group, and the quasicrystal drive are the ones $W(3, 3)$ predicts. (Witness: `analysis/w33_near_term_falsification_test.py`.)

B.17 BT1402 runtime-frontier handoff: readout, port, and certificate

The post-BT1377 runtime frontier makes the physical story sharper. The deterministic packet machine is a protected Clifford machine first: the 8-tick optical word, Q6/tomotope packet edge, S_3 phase synchronization, central C_3 Steinberg scheduler, and 51840-window $\text{Sp}(4, 3)$ supercycle are all finite Clifford data. Universal computation enters at the non-Clifford boundary, not by pretending that this Clifford kernel is already universal by itself.

The near-term single-photon demonstrator now has a concrete readout contract. Its Bell-qutrit signatures are

$$V(I) = 1, \quad V(F_3) = \frac{1}{3}, \quad V(X) = V(Z) = 0,$$

and route-controlled Clifford transport deliberately hides route coherence if the Bell legs are discarded: the reduced route register has uniform probabilities, purity $1/3$, and zero ℓ_1 coherence. Coherence is restored only by the quantum eraser on the Bell branch,

$$\frac{|\Omega\rangle + |Z\Omega\rangle + |X\Omega\rangle}{\sqrt{3}},$$

which succeeds with probability $1/3$, restores route ℓ_1 coherence 2, and routes the interferometer to a single output port. The current parametric noise envelope keeps this falsifiable: the baseline readout has $\gamma = 0.92887$ and port-0 probability 0.95258, while the conservative case keeps $\gamma = 0.79198$ and port-0 probability 0.86132.

The non-Clifford resource has likewise become an ABI rather than a slogan. The Hesse-SIC/T token has nine outcomes, is consumed at a Clifford packet boundary, and feeds a Clifford correction plus one T -frame parity bit into the next 8-tick word. Over one 51840-tick Clifford window there are 720 microframes. In the current queue envelope the baseline heavy case still has positive slack (about 55.75 successful tokens per window), and the conservative medium case still has positive slack (about 89.51 tokens per window). This is an engineering envelope, not a claim that the SIC optics or magic-state yield are solved.

Finally, the correction frontier is now stated in certificate language. The best current S_3 gauge witness has score 210 and is locally strict through the checked radius, but the MaxSAT frontier remains witness-only: the example optimality certificate exercises the verifier path and is explicitly synthetic. A project-level global optimum requires an imported solver upper-bound proof matching the witness score. This is the honest architecture boundary after BT1401: the single-photon experiment tests the Bell/route physics, the Clifford machine runs the protected finite kernel, the Hesse-SIC/T port supplies the non-Clifford ABI, and the remaining global optimality claim is a certificate import problem.

Witnesses: `tools/bt1394_reduced_qutrit_demonstrator.py`;
`tools/bt1396_qutrit_quantum_erasure_readout.py`; `tools/bt1385_hesse_sic_t_port_abi.py`;
`tools/bt1391_hesse_sic_t_queue_model.py`; `tools/bt1395_s3_maxsat_bound_pathway.py`;
`tools/bt1400_qutrit_erasure_noise_sensitivity.py`.

B.18 BT1403 eraser-lift port: the Hesse boundary is the readout boundary

The Hesse port is not a second computer attached to the holonet after the fact. It is the quantum-eraser measurement boundary lifted by one qutrit phase label. BT1396 gives the first half of the boundary: the route register is incoherent if the Bell legs are traced out, and coherence returns only when the Bell branch is erased by the projector

$$\frac{|\Omega\rangle + |Z\Omega\rangle + |X\Omega\rangle}{\sqrt{3}}.$$

This is a three-branch measurement, with branch labels

$$\Omega, \quad Z\Omega, \quad X\Omega.$$

BT1385 gives the non-Clifford half: the Hesse-SIC/T token has nine outcomes. The ABI statement is therefore the exact lift

$$9 = 3 \times 3, \quad h = 3r + p,$$

where $r \in \{0, 1, 2\}$ selects the route/Bell branch and $p \in \{0, 1, 2\}$ is the added phase label. Equivalently, the nine Hesse outcomes are exactly 3 route branches times 3 phase labels. The feed-forward record maps the outcome to a two-trit Clifford correction $X^r Z^p$ plus the one T -frame parity bit needed by the non-Clifford injection.

This is the physical meaning of the non-Clifford port in the single-photon architecture. The same interferometric readout that proves the route qutrit is coherent also supplies the boundary at which a non-stabilizer qutrit resource can be consumed. The Clifford runtime remains an 8-tick packet machine; the Hesse measurement writes a side record, restores the packet ABI, and hands the next word back to the protected scheduler. Thus the architecture is not “Clifford machine plus mysterious magic box” but

$$\text{Bell-branch eraser} \longrightarrow \text{Hesse } 3 \times 3 \text{ lift} \longrightarrow \text{next Clifford packet word.}$$

The boundary remains honest: this identifies the ABI and alphabet of the port, not the optical SIC implementation or magic-state factory yield. (Witness: `tools/bt1403_hesse_port_eraser_lift.py`.)

B.19 BT1404 holonet scope: one Hesse grid is one microframe

BT1404 makes the eraser-lift port visible as a packet oscilloscope. The timing identity is

$$9 \text{ Hesse outcomes} \times 8 \text{ packet ticks} = 72 \text{ ticks},$$

exactly one BT1385/BT1391 microframe. In plain terms: 9 Hesse outcomes times 8 packet ticks fill the whole scope. Thus the entire Hesse outcome alphabet fits in one microframe, with each cell $h = 3r + p$ occupying one 8-tick return word.

The scope word is:

0	erase Bell branch
1	record route trit r
2	record phase trit p
3	apply X^r frame correction
4	apply Z^p frame correction
5	store T -frame bit $(r + p \bmod 2)$
6	restore Clifford ABI
7	hand off next packet word.

This is the first machine-surface form of the non-Clifford boundary: not just a nine-outcome label, but a one-microframe live trace compatible with the 720 microframes in a 51840-tick Clifford window. The generated static scope is `docs/bt1404_holonet_scope.html`. It remains an ABI display, not a claim of physical SIC-optics hardware.

B.20 BT1405 continuous Q6 path router

BT1374 turned packet digits into executable Q6 edge addresses, but deliberately left one boundary open: an address list is not yet a continuous walk. BT1405 closes that boundary at the tomatope ABI level. For every compiled BT828 packet program, orient the packet edges to minimize the total Hamming connector length, insert shortest Q6 connector walks between consecutive packet edges, and map every traversed edge back through the BT1371 table.

The six-digit stress route keeps the six BT1374 packet edges

$$175, 133, 56, 37, 142, 77$$

but now they are waypoints on a single hypercube walk:

$$6 \text{ packet edges plus } 10 \text{ connector edges} = 16 \text{ Q6 steps.}$$

In words: 6 packet edges plus 10 connector edges give the continuous carrier trace. Since the tomatope body budget is 48 ticks, the continuous trace leaves

$$48 - 16 = 32$$

slack ticks inside the same body. Thus the route is no longer merely a set of six distinct Q6 addresses; it is a contiguous Q6 path from 110111 to 010011, with every traversed edge carrying a tomatope flag and block. This is still an ABI route certificate, not a waveguide layout or detector-loss model. (Witness: `tools/bt1405_continuous_q6_path_router.py`.)

B.21 BT1406 tomatope-body edge pulse scheduler

BT1406 uses the slack exposed by BT1405 as timing structure. Each Q6 edge in the continuous route is not a bare edge label; it expands into the ternary microcycle

0		LOAD_FLAG: load tomatope flag, block, and transversal
1		FLIP_Q6_AXIS: perform the one-bit Q6 traversal
2		LATCH_VERTEX: latch the target vertex and ABI record.

Therefore the stress route obeys

$$16 \text{ Q6 edge traversals} \times 3 \text{ pulse phases} = 48 \text{ body ticks.}$$

In words: 16 Q6 edge traversals times 3 pulse phases equals 48 body ticks. The old 32-tick slack becomes the timing shell: the stress route has no idle body tick after expansion. Its packet-edge load ticks are

$$0, 6, 12, 21, 27, 45,$$

so the final packet edge occupies ticks 45, 46, 47: load tomatope flag 180 in block 45, flip 010111 $\rightarrow$ 010011, and latch the target vertex. This is a timing ABI, not a calibrated optical pulse-width, detector-jitter, or loss model. (Witness: `tools/bt1406_tomatope_body_edge_pulse_scheduler.py`.)

B.22 BT1407 microframe transaction composer

BT1407 closes the whole oscillator frame. BT1406 fills the first 48 tomatope-body ticks. The remaining BT1300 local-lift epilogue is

$$24 = 3 \times 8,$$

exactly three Hesse return words. In words: 48 Q6 body pulse ticks plus 3 Hesse return words equals 72 ticks.

For the six-digit stress route, the final target digit is 4. Hence $4 \bmod 3 = 1$, so the selected Hesse route branch is $r = 1$. The epilogue is the phase fanout along that branch:

frame ticks	h	(r, p)	correction
48–55	3	(1, 0)	$X^1 Z^0$
56–63	4	(1, 1)	$X^1 Z^1$
64–71	5	(1, 2)	$X^1 Z^2$.

Thus ticks 0–47 execute the Q6/tomatope carrier body, while ticks 48–71 execute the Hesse return words that erase the branch, record route and phase, apply the Pauli-frame correction, store the T -frame bit, restore the Clifford ABI, and hand off the next packet word. This is a full-frame ABI transaction, not a SIC-optics implementation. The witness is the BT1407 generator and JSON certificate.

B.23 BT1408 Witting contextual communication bridge

BT1408 imports the Witting-polytope communication scheme into the same runtime ABI. The external protocol uses 40 Witting rays and 40 orthogonal tetrads. Around any selected ray the other side of the delayed-query round splits as

$$1 \text{ same ray} + 12 \text{ compatible orthogonal rays} + 27 \text{ incompatible rays} = 40.$$

In words: 1 same ray plus 12 compatible orthogonal rays plus 27 incompatible rays make the whole Witting communication desk. Hence the accepted key-agreement rate is

$$\frac{1 + 12}{40} = \frac{13}{40},$$

and the rejected sector is 27/40, the same q^q matter shell that the holonet already uses as contextual fuel.

The four slots of an accepted Witting tetrad are not a new bus. They are the four local residues in the packet ABI:

$$\text{tomotope_flag} = 4 \text{tomotope_block} + (\text{mirror_slot} \bmod 4).$$

Thus an accepted communication round has the handoff

$$\text{Witting query pair} \rightarrow \text{common tetrad} \rightarrow \text{mirror slot} \rightarrow \text{Q6 body} \rightarrow \text{Hesse epilogue}.$$

The contextuality audit remains the corrected BT823 ceiling 36/40, with deficit 4 and contextual fraction 1/10; the older 34/40 value from the communication paper is a useful historical warning, not the local theorem. BT1408 is an ABI bridge, not a security proof or a ququart hardware certificate.

B.24 BT1409 Witting duplex admission scheduler

BT1409 separates the two acceptance clocks inside the Witting communication bridge. The state-query clock is the BT1408 shell:

$$\frac{13}{40} = \frac{1 \text{ same ray} + 12 \text{ compatible rays}}{40}.$$

The basis-query clock is smaller. A selected Witting ray lives in exactly four tetrads, hence

$$\frac{4}{40} = \frac{1}{10}$$

is the witness aperture. In words: 13/40 is communication throughput, while 1/10 is contextual witness aperture. The complementary basis shadow is

$$\frac{36}{40},$$

which matches the corrected BT823 noncontextual ceiling. Thus the old “36/40” number is not a communication success rate; it is the basis retry shadow and the classical boundary.

For one selected ray, serving every compatible state as a BT1407 transaction uses

$$13 \times 72 = 936$$

frame ticks. Auditing only the four witness apertures uses

$$4 \times 72 = 288$$

frame ticks. Rejected choices need not consume a full frame; they are classified as retry shadow or matter-shell context unless an implementation chooses to log them. This is a scheduler certificate, not a cryptographic security proof.

B.25 BT1410 Witting delayed-query frame compiler

BT1410 makes the delayed-query step operational. The protocol in [28] first lets the two sides query Witting rays, then delays the final basis measurement until the queried states have been compared. The holonet reading is a compiler:

$$(\text{Alice ray}, \text{Bob ray}) \mapsto \text{common Witting tetrad} \mapsto \text{BT1407 frame}.$$

If the two rays share no tetrad, the pair is a retry-shadow event and no full frame is opened. If they are distinct compatible rays, they share exactly one tetrad, so the selected basis is forced. If they are the same ray, they share four tetrads; a two-bit aperture selector chooses which witness basis to audit.

There are two tables, and confusing them loses the physics. The logical ordered pair table has

$$40 \cdot 40 = 1600$$

rows, split as

$$40 \text{ same-ray pairs} + 480 \text{ compatible distinct pairs} + 1080 \text{ retry-shadow pairs}.$$

Thus the accepted logical table has 520 rows and rate $520/1600 = 13/40$. The physical frame table is basis-local:

$$40 \text{ tetrads} \times 4 \text{ Alice query slots} \times 4 \text{ Bob query slots} = 640.$$

Equivalently, every Witting tetrad is a 4×4 frame tile with four diagonal witness apertures and twelve off-diagonal data handshakes. Across all forty tetrads this is

$$160 \text{ diagonal witness records} + 480 \text{ off-diagonal data records}.$$

The extra $640 - 520 = 120$ basis-local records are not extra throughput. They are the same-ray contextual aperture: each of the 40 rays has four basis contexts, so the physical compiler carries three extra context choices per ray above the single logical same-ray identity.

Once the common basis is selected, the later four-outcome measurement slot is just the packet slot:

$$\text{tomotope_flag} = 4 \text{tomotope_block} + (\text{mirror_slot} \bmod 4).$$

Therefore the Witting desk is now a packet admission ROM. Off-diagonal entries open communication frames; diagonal entries open contextual audit frames; both land in the same BT1407 72-tick body-plus-Hesse-epilogue transaction. This matches the feed-forward/programmable-loop physical idiom of time-bin photonic qudit processors [29], but BT1410 itself remains a finite compiler certificate rather than a loss model, detector calibration, or cryptographic security proof.

B.26 BT1411 Witting basis analyzer unitaries

BT1411 supplies the physical analyzer that BT1410 deliberately left abstract. For every Witting tetrad

$$B = (r_0, r_1, r_2, r_3)$$

define U_B by taking row j to be the conjugate ray $r_j^\dagger$. Then

$$U_B r_j = e_j,$$

so detector slot j is exactly the packet residue `mirror_slot` mod 4. This is the standard linear-optics measurement move: an analyzer unitary rotates the chosen orthonormal basis to the detector basis, after which fixed detectors read the slots. General finite-dimensional unitaries are optically synthesizable by beam-splitter/phase-shifter meshes [30, 31]; BT1411 shows the Witting case is much more structured than a generic 4×4 mesh.

The forty analyzers have the exact hardware split

$$1 + 12 + 27.$$

There is one computational direct-rail analyzer, twelve one-direct-rail plus complement-tritter analyzers, and twenty-seven four-three-rail contextual analyzers. Equivalently, the nonzero-entry histogram is

$$4^1, \quad 10^{12}, \quad 12^{27},$$

where the exponent records how many analyzers have that many nonzero matrix entries. No Witting analyzer uses all 16 entries of a generic dense four-mode unitary. The entry alphabet is only

$$0, \quad 1, \quad \frac{\pm 1}{\sqrt{3}}, \quad \frac{\pm \omega}{\sqrt{3}}, \quad \frac{\pm \omega^2}{\sqrt{3}}, \quad \omega^3 = 1.$$

Thus the physical reading is sharper than “program a universal interferometer.” The Witting communication desk is a sparse analyzer ROM:

$$\begin{aligned} \text{delayed-query tetrad} &\longmapsto \text{four-mode Witting analyzer} \\ &\longmapsto \text{detector slot} \longmapsto \text{BT1374 packet residue.} \end{aligned}$$

The twelve middle analyzers literally have one bypass rail and a three-rail tritter-like complement; the twenty-seven contextual analyzers are balanced three-rail rows, each dark on one rail. This is a physics certificate for the single-photon ququart analyzer bank, not a chip calibration, loss budget, or beam-splitter-angle synthesis.

B.27 BT1412 toroidal Q_4 oscillator boundary

BT1412 returns to the 4×4 toroidal square because that square is not a decorative drawing of the local packet router. With ordinary boundary conditions the knight graph has only 24 edges and mixed degree; with toroidal wraparound it is exactly Q_4 . The existing closed toroidal knight tour is therefore a 16-tick Gray Hamilton clock on the four-cube: each packet tick flips one Q_4 bit, and the board parity $(r + c) \bmod 2$ is the same as the Q_4 word parity. Hence parity alternates on every clock edge.

The new splice is that the two-dimensional cells of this same local router have the right oscillator boundary:

$$\#F_2(Q_4) = \binom{4}{2} 2^2 = 24 \text{ } Q_4 \text{ square faces} = q(q-1)(q+1).$$

BT802 already proved that the tetrahedral oscillator is allowed exactly at the mod-12 marks

$$\{0, 3, 4, 7\}$$

and that the neighborly genus level $h = q = 3$ is forbidden by the discriminant 145. Lifting those four marks onto the 24-plaquette Q_4 shell gives the exact aperture

$$\frac{4}{24} = \frac{1}{6},$$

which is the local “one-sixth” selector window in the clock face, not a fitted probability. Removing the forbidden genus level then lands on the toroidal dual edge invariant:

$$24 - 3 = 21.$$

That 21 is not numerology. It is the edge channel preserved by the dual toroidal polyhedra:

$$\text{Császár} : (V, E, F) = (7, 21, 14), \quad \text{Szilassi} : (V, E, F) = (14, 21, 7).$$

Császár is the maximal vertex-adjacency boundary, Szilassi is the maximal face-adjacency boundary, and geometric duality swaps V and F while keeping $E = 21$ intact. In this exact local reading the tetrahedral oscillator sits between the two toroidal boundaries: the Q_4 plaquette shell supplies the 24-slot face clock, the forbidden $q = 3$ genus gap removes the non-realizable handle, and the surviving 21-channel edge bus is precisely the dual boundary shared by the two toroidal polyhedra.

The error-code boundary is also now sharper. A Hamilton cycle in a labeled hypercube is a cyclic Gray code, but a snake-in-the-box code is an induced path or cycle constraint [14, 15]. The full BT1412 Q_4 Gray clock is not claimed to be such a snake: it uses 16 cycle edges while Q_4 has 32 edges, so there are 16 non-cycle chords. The protected projection is instead the every-other-tick even layer

$$0, 6, 3, 5, 9, 15, 10, 12,$$

an eight-word parity-even code whose cyclic steps have Hamming distance 2 and whose pairwise minimum distance is 2. Thus the physical story is: toroidal Q_4 Gray clock for packet motion, every-other parity projection for single-bit-error detection, and the 1/6-aperture oscillator boundary tying the 24-plaquette shell to the 21-edge Császár/Szilassi channel. BT1412 is an ABI certificate, not a waveguide layout, a calibrated oscillator, or a new snake-code optimum. (Witness: `tools/bt1412_toroidal_q4_oscillator_boundary.py`.)

B.28 BT1413 Q_4 plaquette-tomotope compiler

BT1413 makes the 24-plaquette boundary executable. BT1363 had already shown that Q_4 face-edge incidence modulo the antipodal translation is the tomotope/Reye medial layer, with 48 middle blocks. BT1371 had already shown that the 192 tomotope flags are a bijective Q_6 edge address table. The BT1413 compiler is just the composed map:

$$24 \text{ } Q_4 \text{ plaquettes} \cdot 4 \text{ edge incidences}/2 = 48 \text{ middle blocks},$$

$$48 \text{ middle blocks} \cdot 4 \text{ flag residues} = 192 \text{ tomotope}/Q_6 \text{ flags}.$$

The generated certificate checks that each middle block has exactly two Q_4 lifts, that the three ternary sheets carry 64 flags each, that each sheet hits all 16 tomotope face labels exactly once at block level, and that every flag is a one-bit Q_6 edge address. The result is not an optical layout. It is the local address compiler that lets the toroidal Q_4 shell talk directly to the tomotope/ Q_6 packet ABI.

B.29 BT1414 Csaszar-Szilassi dual port

BT1414 uses the active/ground split that BT1317 already found in the tomotope packet:

$$192 = 168 + 24.$$

The active part is now a concrete dual analyzer port:

$$21 \text{ shared toroidal edge channels} \cdot 2 \text{ orientations} \cdot 4 \text{ residues} = 168.$$

In Császár mode the seven analyzers are vertex-neighborhood checks: each vertex of K_7 sees six incident edge channels, and each channel is seen by two vertex analyzers. In Szilassi mode the same 21 channels are read as the dual complete face-adjacency checks: seven face analyzers, again six channels each, again two analyzers per channel. The current BT1318 metric-axis records fix Császár vertex 6 and Szilassi face 4, so the crossed calibration channel is the shared edge $\{4, 6\}$.

The remaining 24 flags are not discarded. They are the BT1412/BT1413 Q4 plaquette guard band. Thus the local physical port has a sharp shape:

$$21 \cdot 2 \cdot 4 \text{ active dual-boundary slots} + 24 \text{ Q4 plaquette guards.}$$

B.30 BT1415 even-projection Steinberg/CSS syndrome ledger

BT1415 finishes the front end by making the every-other Q4 projection into a check ledger. The allowed clock states are the eight even words

$$0, 6, 3, 5, 9, 15, 10, 12.$$

They form a binary distance-2 parity layer: every allowed word has parity zero, and any one-bit Q4 clock fault flips the syndrome bit. This check is binary front-end logic; the memory register is still the existing $\mathbb{F}_3$ Steinberg/CSS object.

BT1375 supplies the memory clock. The central C_3 action on the Steinberg register has 27 three-cycles and nilpotent rank profile

$$54, 27, 0.$$

Therefore the exact syndrome ledger is

$$27 \text{ Steinberg central cycles} \cdot 8 \text{ even Q4 states} = 216,$$

and the BT1414 guard band completes the existing W33 CSS edge ledger:

$$216 + 24 = 240.$$

This is the first complete local front end in the paper:

$$\begin{aligned} \text{Q4 plaquettes} &\rightarrow \text{tomotope/Q6 flags} \\ &\rightarrow \text{Császár/Szilassi port} \\ &\rightarrow \text{even-parity syndrome} \\ &\rightarrow \text{Steinberg/CSS edge ledger.} \end{aligned}$$

I also checked the external Golden Quartic/Moebius-ball electron line of work. Otto's 2022 paper connects a golden quartic, icosahedral coefficients, chirality, and a proposed Moebius-ball electron geometry [36]; his 2025 preprint asks whether AI can compare such extended electron models [37]. A 2025 review frames this family as Zitterbewegung/field-dynamics modelling rather than established particle physics [38]. The holonet use is deliberately conservative: closed-loop topology and chirality are good design warnings, but the repo's exact quartic frontier remains the two independent D_4 quartic atoms in the local minimal-magic audit. No electron mass, spin, or waveguide claim depends on the Moebius-ball model. (Witnesses: `tools/bt1413_q4_plaquette_tomotope_face_compiler.py`, `tools/bt1414_csaszar_szilassi_dual_physical_port.py`, and `tools/bt1415_even_projection_steinberg_syndrome_layer.py`.)

B.31 BT1416 sparse CSS intertwiner

BT1416 turns the BT1415 ledger from a count into matrices. It rebuilds the canonical W33 edge-chain CSS carrier:

$$H_X = d_1 : C_1 \rightarrow C_0, \quad H_Z = d_2^T : C_1 \rightarrow C_2,$$

over $\mathbb{F}_3$. The generated sparse certificate has shapes

$$H_X \in \mathbb{F}_3^{40 \times 240}, \quad H_Z \in \mathbb{F}_3^{160 \times 240},$$

and verifies

$$\text{rank}(H_X) = 39, \quad \text{rank}(H_Z) = 120, \quad H_X H_Z^T = 0.$$

Thus the carrier is still the exact edge code

$$[[240, 240 - 39 - 120, 3]]_3 = [[240, 81, 3]]_3.$$

The new map is the typed 240×240 ledger interwiner:

$$\Pi(e_i) = e_i, \quad 0 \leq i < 240.$$

Rows $0, \dots, 215$ are the $\mathbb{F}_2$ Q4 parity front end, with check vector $(1, 1, 1, 1)$, rank 1, and the eight even Q4 words as a distance-2 kernel. Rows $216, \dots, 239$ are the guard tail addressing the same $\mathbb{F}_3$ edge carrier. This is deliberately not a field coercion: the binary clock syndrome and ternary CSS stabilizer live in different fields, and the interwiner is the finite scheduling boundary between them. (Witness: `tools/bt1416_css_sparse_intertwiner_matrices.py`.)

B.32 BT1417 linear-optical dual-port primitive synthesis

BT1417 converts the BT1414 port into an objectwise single-photon optical primitive list. The active dual port is not a vague optical metaphor; it has the exact inventory

$$\begin{aligned} &21 \text{ edge-channel couplers,} \quad 42 \text{ oriented phase latches,} \\ &168 \text{ active detector bins,} \quad 24 \text{ guard apertures.} \end{aligned}$$

The 21 couplers are the shared K_7 edge channels. The 42 latches choose the two directions on each channel. The four residue bins behind each oriented latch are the tomatope flag residue. The 24 Q4 guard apertures are separate from the active mesh.

The analyzer matrices are 7×21 incidence matrices. In Császár mode a row is a vertex star; in Szilassi mode a row is the dual face star. In both modes the Gram law is

$$AA^T = 5I_7 + J_7,$$

so every analyzer sees six channels and any two analyzers overlap in exactly one channel. This is the chip-level object language the previous ABI was missing: coupler, direction latch, residue demultiplexer, detector bin, and guard aperture. It is still not a calibrated waveguide mask, loss budget, or detector-efficiency model. (Witness: `tools/bt1417_linear_optical_dual_port_primitives.py`.)

B.33 BT1418 finite D4-quartic magic injection

BT1418 closes the non-Clifford aperture without importing the continuum Moebius-ball electron hypothesis. The repo-local minimal-magic audit already found the exact signed frontier: two independent irreducible D_4 quartic atoms, with no shared quadratic subfield and splitting field

$$D_4 \times D_4.$$

The BT1415/BT1416 guard band fits those atoms exactly:

$$2 \text{ quartic atoms} \cdot 4 \text{ algebraic branches} \cdot 3 \text{ qutrit phases} = 24 \text{ guard apertures.}$$

Then the internal D_4 orientation torsor expands each aperture to the tomatope flag bus:

$$24 \text{ guard apertures} \cdot 8 \text{ } D_4 \text{ orientations} = 192 \text{ oriented tomatope tokens.}$$

This is the exact finite counterpart of the golden-quartic topology that is safe to use in the paper. For one quartic atom, the Steinberg lift gives

$$4 \text{ branches} \cdot 27 \text{ central cycles} \cdot 8 \text{ } D_4 \text{ orientations} = 864,$$

which is exactly the repo's Golden D_4 /Weyl shell. With both independent atoms the shell is 1728, but the two atoms stay independent; BT1418 does not collapse them into one continuum object. The physical reading is that the 24 guard apertures are the non-Clifford injection rail, while the 216 parity rows remain the cheap binary front-end syndrome and the $\mathbb{F}_3$ CSS carrier keeps its commutation proof. (Witness: `tools/bt1418_d4_quartic_magic_injection_frontier.py`.)

B.34 BT1697 typed packet ABI

Reading the paper back as an architecture exposes the missing promoted object. The header introduced in §4

(payload, Pauli frame, chart word, apartment hop, mirror slot, clock phase)

is not only a description of fields. It is a fixed-width transaction ABI whose verified projections are all the packet layers above. The master frame identity is

$$72 = 48 + 24 = 16 \cdot 3 + 3 \cdot 8.$$

The first 48 ticks are the continuous Q6/tomatope body:

$$16 \text{ Q6 edge traversals} \times 3 \text{ pulse phases} = 48,$$

and the last 24 ticks are the selected Hesse branch:

$$3 \text{ Hesse return words} \times 8 \text{ ticks} = 24.$$

The same transaction header receives the Witting communication desk as an admission ROM,

$$1600 = 40^2 = 520 + 1080, \quad 520/1600 = 13/40,$$

while the basis-local physical table is

$$640 = 40 \cdot 4 \cdot 4 = 480 + 160.$$

The extra $160 - 40 = 120$ physical rows are the same-ray contextual aperture, not extra throughput. At the physical front end the dual toroidal port is

$$192 = 168 + 24 = 21 \cdot 2 \cdot 4 + 24,$$

and at the memory front end the CSS edge ledger is

$$240 = 216 + 24 = 27 \cdot 8 + 24.$$

Finally the non-Clifford guard rail is typed by the finite quartic atoms,

$$24 = 2 \ D_4 \text{ atoms} \cdot 4 \text{ branches} \cdot 3 \text{ qutrit phases}, \quad 192 = 24 \cdot 8.$$

Thus 24 is not one untyped object renamed four times. It is the common guard boundary at which Hesse epilogue, Q4 guard flags, CSS guard-tail rows, and D_4 -quartic magic apertures meet while retaining distinct source certificates and field semantics. Globalizing the local 48-block packet then recovers the mirror runtime:

$$2160 = 45 \cdot 48, \quad 51840 = 24 \cdot 2160 = 24 \cdot 45 \cdot 48.$$

So the packet header is the finite operating-system object of the holonet: compute, route, admit, check, inject, and commit are typed projections of one transaction. BT1697 is an ABI theorem, not a calibrated optical layout. (Witness: `analysis/bt1697_holonet_typed_packet_abi.py`.)

B.35 BT1698–BT1700 execution stack

The ABI can now be executed, lowered, and recursively scaled. BT1698 reads the existing microframe rows as a deterministic state machine. The body is not a bag of 48 ticks; it is exactly 16 chained Q6/tomotope edges, each with the three-phase instruction word

$$\text{LOAD_FLAG} \rightarrow \text{FLIP_Q6_AXIS} \rightarrow \text{LATCH_VERTEX}.$$

Every committed target vertex is the next edge source. The epilogue is exactly three Hesse return words, each with the eight-tick Clifford word

$$\text{ERASE, ROUTE, PHASE, X-CORR, Z-CORR, T-BIT, RESTORE, NEXT}.$$

Thus the packet is a finite transition system:

$$72 = (16 \cdot 3) + (3 \cdot 8).$$

BT1699 then lowers the same ticks onto the physical single-photon frame already used by the Witting/Fano automaton:

ticks	hardware role
0...7	source switch
8...31	program/delay
32...47	analyzer or OAM body
48...63	detector or Hesse handoff
64...71	dark-reference closeout.

Each tick has a symbolic loss hook, and the last eight ticks carry the dark reference. The guard boundary is welded objectwise:

$$\text{port guard flag } (168 + i) = \text{CSS edge row } (216 + i) = D_4 \text{ magic aperture } i, \quad 0 \leq i < 24.$$

This proves that the repeated 24 is a typed interface between physical port, memory syndrome, and non-Clifford resource rail.

BT1700 is the fractal compiler law. Substituting a $W(3,3)$ packet fabric at each site gives, at depth n ,

$$\begin{aligned} N(n) &= 40^n, & T(n) &= 72 \cdot 40^n, \\ B(n) &= 48 \cdot 40^n, & G(n) &= 24 \cdot 40^n, & B(n) &= 2G(n), \end{aligned}$$

and the route/scheduler clocks

$$R(n) = 8n, \quad S(n) = 2160 \cdot 40^n, \quad C(n) = 51840 \cdot 40^n.$$

The local operating system never changes: every scale keeps the same packet ABI, the same 48/24 body/guard split, and the same typed guard handoff. This is the computer/network architecture in one sentence: a single-photon packet operating system recursively tiled by $W(3,3)$. (Witnesses: `analysis/bt1698_holonet_packet_state_machine.py`, `analysis/bt1699_holonet_abi_to_hardware_lowering.py`, and `analysis/bt1700_recursive_holonet_packet_compiler.py`.)

B.36 BT1701–BT1703 runtime safety layer

The execution stack has three immediate runtime consequences. First, BT1701 renders the packet as an actual trace artifact at `docs/bt1701_holonet_packet_trace_visualizer.html`. The page is generated from the verifier data, not drawn by hand: every tick in the 72-tick packet is colored by the physical stage

$$8 + 24 + 16 + 16 + 8,$$

and the same page records the 48/24 split, the 24-row guard weld, and the recursive compiler law.

Second, BT1702 proves the scheduler boundary. The collision-free recursive key is

$$\text{extended slot} = \text{packet address} \cdot 2160 + \text{polar sheet} \cdot 48 + \text{body slot}, \quad \text{phase} = \text{body slot} \bmod 3.$$

With this key, recursive packets do not collide. If one intentionally reuses a single finite 2160-slot physical bus, the packet address must become a time slice:

$$(\text{packet time slice}, \text{local mirror slot}, \text{phase}).$$

This is the correct engineering boundary. The recursive compiler gives a collision-free addressing fabric; it does not turn one finite bus into simultaneous infinite parallel hardware.

Third, BT1703 classifies symbolic faults at the ABI boundary:

fault	finite exit
loss	retry frame or local reprogram retry
dark click	local dark-reference termination
parity fault	CSS syndrome handoff.

The verifier injects 72 loss hooks, 8 dark-reference clicks, and 24 guard-weld parity faults. Every symbolic fault has one finite exit. This is still not a stochastic noise model; it is the routing table that a calibrated noise model must plug into. (Witnesses: `analysis/bt1701_holonet_packet_trace_visualizer.py`, `analysis/bt1702_holonet_scheduler_collision_audit.py`, and `analysis/bt1703_holonet_fault_propagation_simulator.py`.)

B.37 BT1704–BT1706 replay and retry economics

The next layer turns runtime safety into an operational model. BT1704 replays the 72-tick event log twice from the same initial state. The two executions produce identical tick logs and identical final registers. The final cursor is

HESSE_WORD:5:DONE,

with final corrections X^1 , Z^2 , and time-frame bit 1. This is the determinism test the packet ABI needed: the trace is not merely ordered; it is replayable.

BT1705 simulates one shared 2160-slot mirror bus under FIFO time division. One packet consumes one 2160-slot service slice. Three profiles are tested:

profile	packets	max queue
depth-1 burst	40	40
depth-2 wavefront	1600	1561
depth-2 sequential	1600	1.

All profiles are collision-free, every packet is served once, and Jain fairness is exactly 1. The result is deliberately engineering-facing: the finite bus is fair and schedulable, but bursty recursive traffic pays finite queueing latency.

BT1706 attaches symbolic rates to the BT1703 fault exits. For the nominal profile

$$p_{\text{loss}} = 10^{-5}, \quad p_{\text{dark}} = 10^{-6}, \quad p_{\text{parity}} = 5 \cdot 10^{-6},$$

the expected retry/reprogram load per packet is

$$64 p_{\text{loss}} = \frac{2}{3125},$$

and the expected CSS handoff load is

$$24 p_{\text{parity}} = \frac{3}{25000}.$$

The verifier also includes an engineering 10%-guard profile and the exact guard-budget limit profile. These are not experimental thresholds; they are the exact ABI economics once a bench supplies rates. (Witnesses: `analysis/bt1704_holonet_packet_replay_runner.py`, `analysis/bt1705_holonet_shared_bus_time_division_simulator.py`, and `analysis/bt1706_holonet_retry_economics.py`.)

C The exceptional skeleton and its physics

The architecture above shows that a single photon on $W(3, 3)$ is a universal computer, network, and clock. This section turns from the *machine* to the *world*: it shows that the same $q = 3$ substrate carries an exact tower of exceptional structures — and a layer of falsifiable physics.

Roadmap. The argument is one ascending ladder, each rung an executable witness.

1. *Why $q = 3$* (§C.15): three independent selections — geometric ($q! = 2q$), quantum-resource (minimal magic dimension), and holographic ($c = 24$) — all force $q = 3$.

2. *The qubit/qutrit crossover and the genus tower*: the multi-qubit contextuality ladder meets the $\{3, 7\}$ Hurwitz tower at the genus-7 Macbeath surface; the $\{3, 8\}$ octagon tower is its qubit twin; and the substrate selects the vertex figure $n - 6 \mid k = 12$.
3. *The Eisenstein body*: the Witting polytope (240 E8 roots, 2160 bus edges, 40 $W(3, 3)$ rays) is the substrate over $\mathbb{Z}[\omega]$, welded to E8 two ways (Eisenstein $\omega = C^{10}$ and the 120 icosians).
4. *The exceptional rungs* $E_6 \rightarrow E_7 \rightarrow E_8$: the Hessian polytope (27 lines), the Klein quartic (28 bitangents), and the icosian 600-cell; their Weyl groups telescope $51840 \rightarrow \times 56 \rightarrow \times 240$.
5. *Above E8*: the complex Leech (Eisenstein, $12 = k$), the Suzuki chain from the three-qubit $G_2(2)$, and the Monster at $c = 24 = f$.
6. *The physics* (§C.12): the Standard Model from trinification, gauge unification, the neutrino sector, and the falsifiable scorecard — with one honest tension and its resolution.

The whole climb is drawn in `docs/w33_exceptional_tower.svg`, and its honest epistemic status (framework / derived / measurable) is kept by `analysis/w33_exceptional_tower_ledger.py`: most rungs are exact theorems wearing substrate labels, a smaller set is computed from $q = 3$, and only the measurable layer below falsifies.

Results at a glance. The substrate’s falsifiable predictions, all fixed by $q = 3$ arithmetic with no free parameters ($N = 2(v - \Phi_4) = 60$ e-folds):

Observable	Substrate	Current bound (mid-2026)	Status
n_s	$1 - 2/N = 29/30$	Planck 0.9649 ± 0.0042	consistent ($\sim 0.4\sigma$)
r	$k/N^2 = 1/300$	< 0.036	consistent; LiteBIRD
f_{NL}	$5/(6N) = 1/72$	-0.9 ± 5.1	consistent
$dn_s/d \ln k$	$-2/N^2 = -1/1800$	-0.0045 ± 0.0067	consistent
$\sum m_\nu$	≈ 0.06 eV (NO)	DESI < 0.072 eV	consistent (resolved)
$m_{\beta\beta}$	2–4 meV	< 36 –156 meV	consistent
Ω_{DM}/Ω_b	82/15	≈ 5.36	$\sim 2\%$
$\sin^2 \theta_W$	$3/8 \rightarrow 0.231$	0.23122	consistent (run)
$\tau_p(p \rightarrow e^+ \pi^0)$	$\sim 4.6 \times 10^{35}$ yr	$> 2.4 \times 10^{34}$ (SK)	testable (Hyper-K)
contextual fraction	$1/\Phi_4 = 1/10$	demonstrator	testable
pump Chern C	$2S = \lambda = 2$	demonstrator	testable

Every entry is a closed substrate expression; the lone historical tension ($\sum m_\nu$) is resolved by the seesaw (§C.12), and three rows are sharp near-term falsification handles.

C.1 The one object and its seven faces

Before the details, the organizing claim — and it frames the entire paper, not just this section. Everything here — the selection of $q = 3$, the substrate constants, the gauge group, the neutrino masses, the fault-tolerant code, the benchtop experiment, and inflationary cosmology — is *one structure* seen from seven sides. That structure is the $q = 3$ Eisenstein object: the Witting polytope $3\{3\}3\{3\}3\{3\}3$ over $\mathbb{Z}[\omega]$ (240 vertices = E8 roots), its symmetry the Shephard–Todd group #32 (order $155520 = 3|\mathrm{Sp}(4, 3)| = 3|\mathrm{Aut}(W(3, 3))|$), equivalently the degree-two cyclotomic skeleton $\{\Phi_3, \Phi_4, \Phi_6\}$ whose indices $\{3, 4, 6\}$ are the crystallographic complex periods. Its seven faces (Fig. 9):

1. **Selection** (cosmology/geometry, §C.15). $\Phi_4(q) = q^2 + 1$ factors the de Sitter closure cubic $(q - 3)\Phi_4(q)$ and the GQ point count $(q + 1)\Phi_4(q) = 40$; the periods $\{3, 4, 6\}$ are the only crystallographic complex periods, and 3 is the only prime among them.
2. **Constants** (arithmetic). Two number-sets fall out of the one object: the cyclotomic *values* $\{\Phi_3, \Phi_4, \Phi_6\}(3) = \{13, 10, 7\}$ (the register, the $\text{Sp}(4)$ θ , the Fano 7; sum 30), and the Witting *degrees* $\{12, 18, 24, 30\} = \{k, h(E_7), c=f, h(E_8)\}$ (product 155520). The top degree 30 = $h(E_8)$ is the cyclotomic sum $13+10+7$.
3. **Gauge** (the Standard Model, §C.12). E_6 enters as the Hessian polytope (27, a sub-rung of the Witting tower $8 \rightarrow 27 \rightarrow 240$), and the gauge dimensions *are* Witting degrees: $\dim \text{SU}(3)^3 = 24 = c=f$ and $\dim \text{SM} = 8+3+1 = 12 = k$; the two non-Coxeter degrees $\{12, 24\}$ are exactly the SM and trification dimensions, with $\sin^2 \theta_W = 3/8$ and three generations from the S_3 triality.
4. **Neutrino** (particle masses, §C.12). The Majorana hierarchy ratio is $\Phi_3/q^2 = 13/9$ and the PMNS deformation is $1/\Phi_3$ — the same $\Phi_3 = 13$.
5. **Code** (fault tolerance, §C). The GKP tower $A_2 < D_4 < E_8$ is the Eisenstein tower: A_2 the $q = 3$ hexagonal (one-qutrit) lattice, D_4 the matter shell, E_8 the Witting polytope. The boundary charge $c = 24 = f$ is the Witting degree-24.
6. **Demonstrator** (experiment). The benchtop meters the skeleton: the contextual fraction on the $40 = (q+1)\Phi_4$ $W(3, 3)$ rays is $1/\Phi_4 = 1/10$, and the Thouless-pump Chern number on the $A_2 < D_4 < E_8$ ladder is $\lambda = 2$.
7. **Cosmology** (inflation, §C.12). The e-fold number is $N = 2(v - \Phi_4) = 2h(E_8) = 60$ — twice the E_8 Coxeter number (the top Witting degree), built from the GQ count $v = 40$ and $\Phi_4 = 10$. The CMB observables are then closed forms in N and the Witting degree $k = 12$: $n_s = 1 - 2/N = 29/30$, $r = k/N^2 = 1/300$, $f_{NL} = 5/(6N) = 1/72$, running $-2/N^2 = -1/1800$.

So the $q = 3$ substrate is not a list of coincidences spanning cosmology, gauge physics, particle masses, computation, and metrology; it is one Eisenstein object presenting seven faces. (Witnesses: `analysis/w33_eisenstein_grand_synthesis.py`, `analysis/w33_gauge_sixth_face.py`, `analysis/w33_cosmology_seventh_face.py`.) The faces are not independent because they share the same integers: the load-bearing ones — $k = 12$, $\Phi_4 = 10$, $\Phi_3 = 13$, $c=f = 24$, $v = 40$ — each serve three or more faces at once, and the whole bookkeeping is one periodic table over $q = 3$ (Table 1; witness `analysis/w33_substrate_periodic_table.py`). The subsections that follow detail each face and the connections between them — including a benchtop bridge that turns the neutrino mass ratio into a single-photon measurement (`analysis/w33_neutrino_demonstrator_bridge.py`), which a contextuality simulation on the 40 $W(3, 3)$ rays realises explicitly, returning $\Phi_4 = 10$ and $13/9$ from the incidence geometry (`analysis/w33_contextuality_simulation.py`).

The edge, and how it closes. A synthesis is worth its boundary. Every geometric integer of the seven faces is a cyclotomic value $\Phi_d(3)$, a GQ count, a Witting degree, or a simple combination (e.g. $|\text{PSL}(2, 7)| = 168 = \Phi_6 \cdot f$, $\dim E_8 = 248 = 240 + 8$). The two quantities that seemed to escape — the Monster minimal dimension 196883 and the fine-structure constant ($1/\alpha \approx 137$), which live at the $c = 24 = f$ moonshine ceiling — turn out to close too, at the integer level. The key is that the *Leech kissing number factors entirely into substrate constants*,

$$196560 = 6 \mu q^2 \Phi_3 \Phi_4 \Phi_6 = 6 \cdot 4 \cdot 9 \cdot 13 \cdot 10 \cdot 7,$$

so the complex (Eisenstein) Leech — the $\mathbb{Z}[\omega]$, rank-12= k , $6 \cdot \text{Suz}$ lattice atop the substrate's tower — is built from $\{\mu, q, \Phi_3, \Phi_4, \Phi_6\}$. Then the lift parameter is forced, $\tau = \mu q^2 \Phi_6 = 252 = 32760/(\Phi_3 \Phi_4)$; the Monster dimension is $196883 = 196560 + \mu q^4 - 1$ (with $\mu q^4 = h(E_7)^2 = 324$ the classical j/Leech gap); and the fine-structure *integer* is a cyclotomic combination, $1/\alpha = 137 = \Phi_3 \Phi_4 + \Phi_6 = 2^{\Phi_6} + q^2$. So no free *integer* parameter remains: even the famous 137 is $\{\Phi_3, \Phi_4, \Phi_6\}$ arithmetic, and the moonshine ceiling is the substrate constants. The one genuinely non-integer input is the 0.036 of the measured 137.036 — the QED running, dynamics rather than arithmetic. *That* is the honest residue: the substrate fixes the integer, the renormalization flow supplies the rest. (The third value $\Phi_6(3) = 7$ also surfaces in the toroidal thread: the Császár polyhedron is K_7 on the torus and the dual Szilassi has 7 faces, with C_2 symmetry of order $2 = \lambda = q - 1$.) Witnesses: `analysis/w33_alpha_closure.py`, `analysis/w33_monster_leech_second_layer.py`, `analysis/w33_eisenstein_stress_test.py`. And the single most decisive near-term test — benchtop, calibration-free, three faces — is the contextual fraction $1/\Phi_4 = 1/10$, now a concrete single-photon protocol (`analysis/w33_contextuality_protocol.py`).

Table 1: The periodic table of the $q = 3$ Eisenstein object (excerpt): the load-bearing integers, their closed expression, and the faces they serve (1 Selection, 2 Constants, 3 Gauge, 4 Neutrino, 5 Code, 6 Demonstrator, 7 Cosmology).

Value	Expression	Class	Faces
2	$q - 1 = \Phi_1$	cyclotomic	6 (pump Chern λ)
10	$q^2 + 1 = \Phi_4$	cyclotomic value	1, 2, 6 (θ , $1/\Phi_4$)
12	$q(q + 1) = k$	Witting degree	3, 5, 7 (dim SM, D_4 , $r=k/N^2$)
13	$q^2 + q + 1 = \Phi_3$	cyclotomic value	2, 4 (register, Φ_3/q^2)
24	$q^3 - q = c=f$	Witting degree	2, 3, 5 (dim $\text{SU}(3)^3$, charge)
30	$h(E_8) = \Phi_3 + \Phi_4 + \Phi_6$	Witting degree	2, 7 ($N/2$)
40	$(q + 1)\Phi_4 = v$	GQ count	1, 6, 7 (rays, $N=2(v-\Phi_4)$)
60	$2(v - \Phi_4) = 2h(E_8)$	derived	7 (e-folds N)
240	E8 roots = Witting	derived	1, 5 (gauge lattice)
155520	$3 \text{Sp}(4, 3) $	the object	1–7 (ST#32 order)

The closure. Taken to its end, the arithmetic of the substrate is a single forced integer unfolding into the world. From $q = 3$ alone — itself the unique common solution of the selection census — the entire integer content follows in closed form: $\lambda = q - 1$, $\mu = q + 1$; the cyclotomic skeleton $\{\Phi_3, \Phi_4, \Phi_6\} = \{13, 10, 7\}$; the GQ counts $k = 12$, $v = 40$; the exceptional $c=f = 24, 27$, $h(E_7) = 18$, $h(E_8) = 30$; the Witting degrees $\{12, 18, 24, 30\}$ (product 155520) and 240 E8 roots; the e-fold number $N = 60$; and the moonshine ceiling — Leech kissing $196560 = 6\mu q^2 \Phi_3 \Phi_4 \Phi_6$, $\tau = \mu q^2 \Phi_6 = 252$, Monster 196883, and the fine-structure integer $1/\alpha = 137 = \Phi_3 \Phi_4 + \Phi_6$. There is *no free integer parameter*. What remains is not arithmetic but dynamics, and we name it plainly: the QED running (the 0.036 of 137.036), the absolute mass scales (the substrate fixes ratios and textures, not the overall eV/GeV), and the exact B – L VEV direction (the neutrino 13/9 versus the generic cubic-form ~ 1.25). The substrate fixes the arithmetic; physics supplies the flow and the scales. And one benchtop measurement — the contextual fraction $1/\Phi_4 = 1/10$ — decides it. (Witness: `analysis/w33_master_closure.py`.)




Figure 9: The $q = 3$ substrate is one Eisenstein object — the Witting polytope, its symmetry ST#32, the cyclotomic skeleton $\{\Phi_3, \Phi_4, \Phi_6\}$ — seen from seven sides. The selection of $q = 3$, the substrate constants, the Standard-Model gauge group, the neutrino mass hierarchy, the fault-tolerant code lattice, the benchtop experiment, and inflationary cosmology are faces of the same structure, sharing the same invariants ($240 = E_8$ roots, $\{\Phi_3, \Phi_4, \Phi_6\}$, the Witting degrees $\{12, 18, 24, 30\}$, $155520 = 3|\text{Aut}(W(3, 3))|$).

C.2 BT1707–BT1709 contextuality and Hesse crossover

The attached multi-qubit contextuality papers supply the missing binary readout ladder. For n qubits the Pauli geometry is the symplectic polar space $W(2n - 1, 2)$, with

$$|P_n| = 4^n - 1, \quad |L_n| = \frac{(4^n - 1)(4^{n-1} - 1)}{3}.$$

Thus the first six layers have

n	$ P_n $	$ L_n $	contextual carrier
1	3	0	single Pauli triple
2	15	15	doily / Mermin–Peres square
3	63	315	split Cayley hexagon
4	255	5355	DW(5,2), Heawood/Coxeter cores
5	1023	86955	PG(4,2) point–hyperplane core
6	4095	1396395	two hexagons and a $K_{7,7}$ carrier

The new architectural identity is not only the count 63. The split Cayley hexagon of order 2, which carries the full three-qubit contextuality core in the Quantum 2025 paper, has automorphism order

$$12096 = 168 \cdot 72 = 7 \cdot 1728.$$

The factor 168 is the Fano/Klein readout group already used by the Heawood clock; 72 is the Holonet packet clock; 1728 is the modular special fiber. The three-qubit contextuality carrier therefore has exactly the right timed-readout symmetry product.

The same paper exposes a local configuration of type $(24_2, 16_3)$. Its incidence count is

$$24 \cdot 2 = 16 \cdot 3 = 48.$$

That is the already-verified tomotope/Holonet interface:

$$48 = 12 \text{ tomotope edge axes against } 16 \text{ faces} = 16 \text{ Q6 body edges} \cdot 3 \text{ pulse phases}.$$

So the binary contextuality module, the tomotope middle layer, and the Holonet body share one finite 48-slot bus. This is an incidence bridge, not yet a claimed graph isomorphism.

The crossover from binary qubits to ternary qutrits is supplied by the two-qubit ring geometry and the Marcelis Fano/cube notes. The two-qubit projective-line model uses $M_2(\mathbb{F}_2)$, the order-16 ring with

$$16 = 6 \text{ units} + 10 \text{ zero divisors},$$

and its Pauli doily admits the three exact decompositions

$$15 = 9 + 6 = 10 + 5 = 8 + 7.$$

In the cube/Fano route, deleting two hyperovals leaves a 9-point $\text{AG}(2, 3)$ grid; adding the line at infinity gives $\text{PG}(2, 3)$ with 13 points. This lands exactly on the Holonet ABI:

$$\text{AG}(2, 3) = \mathbb{F}_3^2 = 9 = \text{Hesse outcome field} = \text{Pauli feed-forward frame},$$

and

$$\text{PG}(2, 3) = 9 + 4 = 13.$$

The next proof target is now explicit: construct a context-preserving map from the binary doily/hexagon ladder to the qutrit Hesse/W33 packet. (Witnesses: `analysis/bt1707_qubit_contextuality_ladder.py`, `analysis/bt1708_hexagon_tomotope_contextual_bus.py`, and `analysis/bt1709_binary_to_hesse_qutrit_crossover.py`.)

C.3 The $\{3, 7\}$ Hurwitz tower and the genus-7 crossover

The genus $2 \leq g \leq 14$ polyhedral-embedding survey of Bokowski and H. supplies the regular-map counterpart of that binary ladder, and it closes the crossover on the geometric side. Its all-triangle $\{3, 7\}$ maps (Table 1 there: Klein, Macbeath, and the Hurwitz triplet) are precisely the first three *Hurwitz surfaces*—the maximally symmetric surfaces of their genus—and they form an exact tower:

g	(V, E, F)	rotation Aut	q
3	(24, 84, 56)	$\text{PSL}(2, 7) = 168$	7
7	(72, 252, 168)	$\text{PSL}(2, 8) = 504$	8
14	(156, 546, 364)	$\text{PSL}(2, 13) = 1092$	13

For any $\{3, 7\}$ map $V = 12(g - 1)$, $E = 42(g - 1)$, $F = 28(g - 1)$, and the orientation-preserving automorphism group has order $84(g - 1) = 2E$ (the dart count), so $g = 1 + |\text{Aut}|/84$.

The middle rung—the genus-7 Macbeath surface, $\text{PSL}(2, 8)$ —is the exact 3-qubit $\leftrightarrow$ qutrit crossover the contextuality papers ask for. Its combinatorics are pure substrate:

$$V = 72 = \text{the packet/frame clock}, \quad E = 252 = \tau, \quad F = 168 = \lambda k \Phi_6 = \text{PSL}(2, 7),$$

i.e. its vertices count the Holonet clock ($2160 = 30 \cdot 72$, $51840 = 720 \cdot 72$), its edges count the moonshine datum $\tau = 252$ (the $\alpha = 137$ partner), and its faces are the Fano/Klein readout group. Hence

$$F \cdot V = 168 \cdot 72 = 12096 = |\text{G}_2(2)| = |\text{Aut}(\text{split Cayley hexagon})| = \text{PSL}(2, 8) \cdot f = 504 \cdot 24,$$

so the genus-7 surface's (faces) $\times$ (vertices) is exactly the automorphism order of the three-qubit contextuality core of the previous subsection; its derived group $\text{G}_2(2)' = \text{PSU}(3, 3)$ has order $6048 = \tau \cdot f$.

The three Hurwitz moduli are the crossover made arithmetic:

$$\{7, 8, 13\} = \{\Phi_6, 2^3, \Phi_3\}.$$

The flanks $\Phi_6 = 7$ and $\Phi_3 = 13$ are the qutrit cyclotomic primitives (the heptagon/torus clock layer, and the Eisenstein norm prime $\text{PG}(2, 3) = 13$ of §BT1709), while the middle modulus $2^3 = 8 = \text{GF}(8)$ is precisely the *three-qubit dimension*. So the binary three-qubit rung sits between the two ternary flanks, and the dimension-14 exceptional group G_2 ($\dim \text{G}_2 = 14 = 2\Phi_6 =$ the top genus) governs both ends—the genus-14 Hurwitz triplet and, via $\text{G}_2(2)$, the genus-7 three-qubit hexagon. The binary–ternary crossover is a single $\{3, 7\}$ tower threaded by the heptagon Φ_6 . (Witness: `analysis/w33_hurwitz_tower_qubit_crossover.py`.)

C.4 Two Schläfli families, a vertex-figure selection, and the Witting body

The same survey contains the *qubit* twin of that qutrit tower. Its all-triangle $\{3, 8\}$ maps (Dyck $g=3$, Fricke–Klein $g=5$, R8.1/R8.2 $g=8$) satisfy $V = 6(g - 1)$, $E = 24(g - 1)$, $F = 16(g - 1)$, and the vertex figure is an octagon: $n = 8 = 2^3$ is the *three-qubit Hilbert dimension* $\text{GF}(8)$, exactly as the heptagon $n = 7 = \Phi_6$ is the qutrit vertex figure. The genus-8 rung R8.1 has $42 = 2q\Phi_6$ vertices (the $D(2T)$ anyon count) and $168 = \lambda k \Phi_6 = \text{PSL}(2, 7)$ edges, with structure group $\text{PSL}(3, 2) = \text{GL}(3, 2)$ —the Fano/three-qubit group. So the two Schläfli families are the qubit/qutrit pair, $\{3, 8\} (2^3)$ and $\{3, 7\} (\Phi_6)$, hinged on $\text{PSL}(2, 7) = 168$. (Witness: `analysis/w33_qubit_schlaefli_tower_3_8.py`.)

The substrate selects the vertex figure. A reflexible $\{3, n\}$ map of genus g has $E = 6n(g - 1)/(n - 6)$ and $|\text{Aut}|_{\text{full}} = 4E$, so E is integral only when $(n - 6) \mid 6n(g - 1)$. Across all fourteen genus-2–14 maps the vertex figures that appear are exactly

$$n \in \{7, 8, 9, 10, 12\} = 6 + \{1, 2, 3, 4, 6\} = 6 + \{d : d \mid k = 12\} = \{\Phi_6, 2^3, q^2, \Phi_4, k\},$$

and the unique gap in $7 \leq n \leq 12$ is $n = 11$, the one value with $n - 6 = 5 \nmid k = 12$. Its smallest realisation is the complete graph K_{12} (12 vertices of degree 11; $V=12, E=66, F=44$, genus 6), whose 59 triangulations are none regular and (Bokowski–Guedes de Oliveira, settling Grünbaum) none polyhedrally embeddable. Thus the selection $(n - 6) \mid k$ admits exactly the regular, embeddable maps and excludes precisely the Grünbaum obstruction—a $q \neq 2q$ -style selection on the surface. (Witness: `analysis/w33_genus_vertex_figure_selection.py`.)

The functor through the Fano heptad. The binary-to-qutrit context map that §BT1709 and BT1710–1712 leave open has a concrete structure group: the Fano heptad $\text{PSL}(2, 7) = \text{GL}(3, 2)$ of order 168. By the ATLAS it is maximal in $\text{U}_3(3) = \text{G}_2(2)'$, hence inside Aut of the split Cayley hexagon (the three-qubit core $W(5, 2)$: 63 points, 315 contexts); it is the genus-3 Klein rotation group; it is the genus-7 Macbeath face count; and it acts on the three-qubit label space $\mathbb{F}_2^3$. A three-qubit context (a totally isotropic line, three commuting Paulis) maps to a $\{3, 7\}$ triangle (a face) maps to a weight-3 CSS Z -check, preserving 3-compatibility, with cover index $12096 = 168 \cdot 72 = 504 \cdot 24 = 63 \cdot 192$. (Witness: `analysis/w33_macbeath_hexagon_functor.py`.)

The Witting body. All of this lives in one complex polytope. The Witting polytope $3\{3\}3\{3\}3\{3\}3$ in $\mathbb{C}^4$ over the Eisenstein integers $\mathbb{Z}[\omega]$ (every edge a $3\{3\}$ triangle—the “3” is q) has 240 vertices = the E8 roots = E , 2160 three-edges = the holonet mirror bus = $h(\text{E}_8) \cdot \text{frame} = 30 \cdot 72$, and $240 = 40 \cdot 6$ vertices falling into 40 hexagonal diameters—the 40 Witting rays = the totally isotropic 1-spaces of $W(3, 3) = \text{GQ}(3, 3)$ ($v = 40$). Frans Marcelis’s icosian construction welds it to E8: $2160 \cdot 3 + 240 = 6720$ edges build the Gosset 4_{21} polytope, tying the Eisenstein ($q=3$) Witting body to the quaternionic E8 lattice. Its symmetry group is exactly q times the substrate gauge group,

$$|\text{Aut}(\text{Witting})| = 155520 = 3 \times 2 \cdot \text{U}_4(2) = \mathbb{Z}_q \times \text{Sp}(4, 3) = q \cdot |W(3, 3) \text{ Aut}|,$$

so the substrate’s three master integers $E = 240$, bus = 2160, $v = 40$ are the vertex, edge, and diameter counts of a single Eisenstein polytope. (Witness: `analysis/w33_witting_polytope_substrate.py`.)

C.5 The register atlas, the explicit E8 weld, and the Hesse engine

Three threads close the surface picture. *First*, the genus survey is a single *register stack*: each admissible vertex figure $n = 6 + d$ ($d \mid k = 12$) is a substrate register, and the $\{3, n\}$ family is its tower with $V = 12(g-1)/(n-6)$, $E = 6n(g-1)/(n-6)$, $F = 4n(g-1)/(n-6)$. The five registers are

$$\begin{aligned} n = 7 &= \Phi_6 \text{ (qutrit)}, & 8 &= 2^3 \text{ (qubit)}, & 9 &= q^2 \text{ (single-qutrit phase space)}, \\ 10 &= \Phi_4 \text{ (dim Sp(4), contextual denominator)}, & 12 &= k \text{ (the degree)}, \end{aligned}$$

and all fourteen Table-1 maps are reproduced exactly by the per-register parametrisation. (Witness: `analysis/w33_register_atlas_3n.py`.)

Second, the identification “Witting body = E8” is now explicit. The Coxeter element C of $W(\text{E}_8)$ has order $h = 30$, and because the exponents $\{1, 7, 11, 13, 17, 19, 23, 29\}$ are coprime to 3, the element $\omega = C^{10}$ has order 3 with eigenvalues ω, ω^2 only — a *fixed-point-free* isometry, the multiplication-by- ω that turns E8 into a rank-4 Eisenstein lattice. Explicit computation on the 240 roots gives $240/3 = 80$ ω -triangles and (under C^5 , order 6) exactly $240/6 = 40$ *hexagons*; each hexagon spans one complex line, so the 40 hexagons are the 40 Witting rays = the 40 isotropic 1-spaces of $W(3, 3)$. Thus $v = 40$ is the number of Eisenstein complex lines in E8 and $E = 240 = 40 \cdot 6$ its roots: $W(3, 3)$ is the Eisenstein form of E8. (Witness: `analysis/w33_e8_eisenstein_witting_weld.py`.)

Third, the qutrit registers carry the contextuality engine. The Hesse configuration — the affine plane $\text{AG}(2, 3)$, a $(9_4, 12_3)$ configuration of 9 points and 12 lines in four parallel classes — is the single-qutrit discrete phase space: its 9 points are the register $n = 9 = q^2$, its 12 lines are the $n = 12 = k$ contexts, and its four parallel classes are the $4 = (q^2 - 1)/(q - 1)$ mutually unbiased bases. A single qutrit is non-contextual; the two-qutrit substrate $\mathbb{F}_3^4$ ($\text{Sp}(4, 3) = W(3, 3)$) with its 40 isotropic rays carries the state-independent contextuality, with contextual fraction $1/\Phi_4 = 1/10$, so the vertex figures 9, 10, 12 are exactly the qutrit phase space, its contextual denominator, and its context count. (Witness: `analysis/w33_hesse_mermin_contextuality.py`.)

C.6 The Eisenstein tower $\text{qutrit} \rightarrow \text{E}_6 \rightarrow \text{E}_8$, the Golay gap, and the quadrangle ladder

The Witting body extends downward into a complete Eisenstein ladder. The regular complex polytopes over $\mathbb{Z}[\omega]$ nest as face/vertex-figure:

polytope	dim	vertices	symmetry
$3\{3\}3$ (Möbius–Kantor)	2	8	$24 = 2T = f$
$3\{3\}3\{3\}3$ (Hessian)	3	27	$648 = 27f$
$3\{3\}3\{3\}3\{3\}3$ (Witting)	4	240	$155520 = q \text{Sp}(4, 3) $

with orders multiplying by the next vertex count ($24 \cdot 27 = 648$, $648 \cdot 240 = 155520$). The middle rung’s 27 vertices are the 27 lines on a cubic surface = the minuscule E_6 rep, and its automorphism group is $W(\text{E}_6)$ with $|W(\text{E}_6)| = 51840 = |\text{Sp}(4, 3)|$. Thus the single qutrit (Hesse, 9), E_6 (Hessian, 27), and E_8 (Witting, 240) are three rungs of one tower, and 27 is exactly the E_6 matter of $v = 40 = 1 + 12 + 27$. (Witness: `analysis/w33_hessian_polytope_e6.py`.)

The vertex-figure gap $n = 11 = K_{12}$ is not empty algebraically. Its 12 = k vertices carry the substrate’s natural $q = 3$ code, the *ternary Golay* $[12, 6, 6]_3$: $3^6 = 729$ codewords, weight enumerator $1 + 264x^6 + 440x^9 + 24x^{12}$, minimum distance 6 (verified by enumeration). Its 132 weight-6 hexads are the Steiner system $S(5, 6, 12)$ with Mathieu group M_{12} , and doubling $k = 12 \rightarrow f = 24$ gives the binary Golay $[24, 12, 8]$, M_{24} , $S(5, 8, 24)$ — the $c = 24$ Monster charge. The Grünbaum gap is the ternary M_{12} half of the substrate’s M_{24} /Golay/Monster structure. (Witness: `analysis/w33_ternary_golay_m12_grunbaum.py`.)

Finally, $W(3, 3) = \text{GQ}(3, 3)$ sits in a quadrangle ladder 27–40–45. $\text{GQ}(2, 4)$ has 27 points = the 27 lines on a cubic (= E_6 ; intersection graph $\text{SRG}(27, 10, 1, 5)$, collinearity the complement Schläfli graph $\text{SRG}(27, 16, 10, 8)$), with $\text{Aut} = W(\text{E}_6) = 51840 = |\text{Sp}(4, 3)|$ — the group of the $D = 5$ E_6 black-hole/qubit entropy. The substrate $\text{GQ}(3, 3)$ is $\text{SRG}(40, 12, 2, 4)$, and $\text{GQ}(4, 2)$ (the dual, 45 points) is the other flank; the four-qubit $W(7, 2)$ carries the same 27/45 split through its 135 $\text{DW}(5, 2)$. (Witness: `analysis/w33_generalized_quadrangle_ladder.py`.)

C.7 A guardrail on the 27, and the Klein quartic as the E_7 rung

Two honest refinements close this arc. *First, a guardrail.* It is tempting to read $v = 40 = 1 + 12 + 27$ as realising the $\text{E}_6/27$ -lines Schläfli graph inside $W(3, 3)$. It does *not*: fixing a point of $\text{GQ}(3, 3)$, the 12 collinear points induce $4K_3$ and the 27 non-collinear points induce an 8-regular graph that is *not* strongly regular — in particular not the 16-regular Schläfli graph $\text{SRG}(27, 16, 10, 8)$. So $40 = 1 + 12 + 27$ is an exact count (the matter cone $28 = 1 + 27$), and 27 matches the E_6 representation dimension and the $\text{GQ}(2, 4)$ point count, but it is *not* a subgraph embedding. The genuine bridge is group-theoretic: the substrate’s projective gauge group and the E_6 group are the same simple group $\text{U}_4(2) = \text{PSp}(4, 3) = \text{PSU}(4, 2)$, of order 25920, with $\text{Sp}(4, 3) = 2.\text{U}_4(2)$, $W(\text{E}_6) = \text{U}_4(2):2$, and $\text{Aut}(\text{GQ}(2, 4)) = \text{U}_4(2):2$ all order 51840. (Witness: `analysis/w33_27_not_schlaefli_group_bridge.py`.)

Second, the E_7 rung. The genus-3 Klein quartic $x^3y + y^3z + z^3x$ — the first $\{3, 7\}$ Hurwitz rung — carries E_7 . Its automorphism group $\text{PSL}(2, 7)$ acts transitively on three distinguished point-sets that are all substrate integers:

$$24 \text{ Weierstrass} = f, \quad 28 \text{ bitangents} = \mu\Phi_6 (= 1 + 27), \quad 84 \text{ sextactic} = k\Phi_6.$$

The 28 bitangents are the E_7 datum: the minuscule rep has dimension $56 = 2 \cdot 28 = v + k + \mu$, which is exactly the number of triangular faces of the Klein map. So the exceptional trinity

27 lines (cubic) $\rightarrow \text{E}_6$, 28 bitangents (quartic) $\rightarrow \text{E}_7$, 120 tritangents $\rightarrow \text{E}_8 (= 120 \text{ icosians}, 240 = 2 \cdot 120)$,

closes the substrate’s exceptional ladder, with Weyl orders $|W(E_6)| = 51840 = |\mathrm{Sp}(4, 3)|$, $|W(E_7)| = 56 \cdot 51840$, and $|W(E_8)| = 240 \cdot |W(E_7)| = 696729600$ — the multipliers being 56 ($= 2 \times$ bitangents) and 240 ($= |E| = E_8$ roots = Witting vertices). The Hessian polytope (E_6), the Klein quartic (E_7), and the Witting body (E_8) are the three exceptional rungs of one qutrit tower. (Witness: `analysis/w33_klein_quartic_e6_e7_trinity.py`.)

C.8 The three exceptional rungs made explicit

Each rung of $E_6 \rightarrow E_7 \rightarrow E_8$ can now be built and verified. *E_6 as trinification.* The guardrail showed the 27 does not sit in $W(3, 3)$; it sits in $\mathrm{GQ}(2, 4)$ — the Schläfli graph $\mathrm{SRG}(27, 16, 10, 8)$, group $W(E_6) = \mathrm{U}_4(2):2$ — and its substrate content is trinification,

$$27 = (3, \bar{3}, 1) + (1, 3, \bar{3}) + (\bar{3}, 1, 3) = 9 + 9 + 9,$$

three bifundamental nonets under $\mathrm{SU}(3)^3 \rtimes S_3$. Each nonet is $\mathbb{F}_3 \times \mathbb{F}_3 = \mathrm{AG}(2, 3)$, a Hesse single-qutrit phase space, and the S_3 triality permuting them is the three generations. So the 27-lines/ E_6 is three Hesse-9s welded by triality. (Witness: `analysis/w33_e6_trinification_schlaefli.py`.)

E_7 as a symplectic space. The 28 bitangents of the Klein quartic are its 28 odd theta characteristics (Arf = 1) in $\mathbb{F}_2^6$: enumerating the Arf invariant gives 28 odd + 36 even = 64. The group $\mathrm{Sp}(6, 2) = W(E_7)/\{\pm 1\}$ (order 1451520, half of $|W(E_7)| = 2903040$) permutes the 28 transitively, and the minuscule $56 = 2 \cdot 28 = v + k + \mu$ — the Freudenthal symplectic double — is the Klein map’s 56 faces. (Witness: `analysis/w33_e7_theta_bitangents.py`.)

E_8 as the icosians. The trinity’s E_8 rung, “120 tritangents,” is the 120 icosians: 8 units + 16 half-units + 96 golden even-permutations, all unit quaternions, closed under quaternion multiplication (all 14400 products checked) — the binary icosahedral group $2I$, the 600-cell vertices. The icosian construction gives $240 = 2 \cdot 120$ E_8 roots = the Witting vertices, so alongside the Eisenstein weld ($\omega = C^{10} \rightarrow 40$ rays) the quaternionic weld gives $E_8 = 2I$ -doubled = the Witting body, with the 600-cell the Boerdijk–Coxeter clock. One body, two arithmetics — Eisenstein ($q = 3$) and icosian (golden quaternion). (Witness: `analysis/w33_icosian_e8_witting.py`.)

C.9 The magic square, the complex Leech, and the icosian clock

These exceptional rungs sit inside three larger frames. *The magic square.* E_6, E_7, E_8 are the octonion row of the Freudenthal–Tits magic square (dims 78, 133, 248), and the substrate’s two welded arithmetics are its $\mathbb{C}$ and $\mathbb{H}$ columns: Eisenstein ($q = 3$) gives $E_6 = \mathrm{magic}(\mathbb{O}, \mathbb{C})$ (the Hessian/ trinification 27), the icosian (quaternion) gives $E_7 = \mathrm{magic}(\mathbb{O}, \mathbb{H})$, and $E_8 = \mathrm{magic}(\mathbb{O}, \mathbb{O})$. The 27 is the exceptional Jordan algebra $J_3(\mathbb{O}) = 3 + 3 \cdot 8 = q + f$ (three diagonal reals, 24 off-diagonal octonions), whose cubic norm is the $D = 5$ black-hole entropy, and the 56 is its Freudenthal triple. (Witness: `analysis/w33_magic_square_substrate.py`.)

The complex Leech. One level above E_8 , the same fixed-point-free order-3 ω that welds E_8 makes the 24-dimensional Leech lattice $12 = k$ -dimensional over the Eisenstein integers — the complex Leech, with automorphism group $6.\mathrm{Suz} \leq \mathrm{Co}_0 = 2.\mathrm{Co}_1$. The Suzuki chain $\mathrm{G}_2(2) \rightarrow \mathrm{J}_2 \rightarrow \mathrm{G}_2(4) \rightarrow \mathrm{Suz}$ (on 36, 100, 416, 1782 points) starts at $\mathrm{G}_2(2) = \mathrm{U}_3(3):2$, $|\mathrm{G}_2(2)| = 12096 = 168 \cdot 72$ — the split Cayley hexagon, the three-qubit contextuality core — and climbs to Suz , the complex Leech group. So the substrate’s $q = 3$ Eisenstein arithmetic threads E_8 (8d) $\rightarrow$ complex Leech ($12 = k$ complex $= 24 = f$ real) $\rightarrow \mathrm{Co}_0 \rightarrow \mathrm{Monster}$ ($c = 24 = f$), with the three-qubit hexagon as its first rung. (Witness: `analysis/w33_complex_leech_suzuki_chain.py`.)

The icosian clock. Finally the runtime closes on the 600-cell: its 120 vertices are the gate set $2I$, its 600 tetrahedral cells are 20 rings of 30 (the Boerdijk–Coxeter beat, $30 = h(E_8)$), its $720 = 6!$

edges are the supercycle width, and its $240 = 2 \cdot 120$ E8 roots/Witting vertices are the mirror-bus addresses — so $2160 = 30 \cdot 72$ and $51840 = 24 \cdot 30 \cdot 72 = 720 \cdot 72 = |\mathrm{Sp}(4, 3)|$ are the polytope’s own counts. The machine’s clock *is* the icosian 600-cell. (Witness: `analysis/w33_icosian_600cell_runtime.py`.)

C.10 The ceiling, the Suzuki tower, and the G2 thread

The climb tops out, telescopes, and is held together by one group. *The ceiling.* The complex-Leech / Co0 tower ends at the Monster $M = \mathrm{Aut}(V^\natural)$, the $c = 24$ holomorphic CFT, whose graded dimension is $J = j - 744 = q^{-1} + 196884q + \dots$ with the moonshine head $196884 = 1 + 196883$ and $21493760 = 1 + 196883 + 21296876$. Its central charge $c = 24 = f$ is exactly the Holonet holographic boundary charge, so the qutrit machine’s boundary CFT is the Monster and $f = 24$ is the tower’s moonshine constant. (Witness: `analysis/w33_monster_moonshine_ceiling.py`.)

The Suzuki tower. The chain $G_2(2) \rightarrow J_2 \rightarrow G_2(4) \rightarrow \mathrm{Suz}$ is the automorphism tower of four nested strongly regular graphs, $\mathrm{SRG}(36, 14, 4, 6) \prec \mathrm{SRG}(100, 36, 14, 12) \prec \mathrm{SRG}(416, 100, 36, 20) \prec \mathrm{SRG}(1782, 416, 100, 96)$, in which each valency is the previous vertex count and each λ the previous valency, so the parameters telescope $14 < 36 < 100 < 416 < 1782$. The base $\mathrm{SRG}(36, 14, 4, 6)$ anchors the substrate: its 36 vertices are the *even* theta characteristics of the Klein quartic (the E_7 companion of the 28 bitangents), and its valency $14 = \dim G_2$. (Witness: `analysis/w33_suzuki_tower_srg.py`.)

The G2 thread. $G_2 = \mathrm{Aut}(\mathbb{O})$ recurs at every level, keyed by the heptad $\Phi_6 = 7$: its 7-dim rep is the Fano plane (the imaginary octonions); $\dim G_2 = 14 = 2\Phi_6$ is the genus of the $\{3, 7\}$ Hurwitz triplet; $G_2(2) = 12096$ is the split Cayley hexagon (the three-qubit core, $\mathrm{SRG}(36, 14)$ valency $= \dim G_2$); $G_2(4)$ is the Suzuki rung; and the octonion column gives $E_8 = \mathrm{magic}(\mathbb{O}, \mathbb{O})$. So one group, keyed by $7 = \Phi_6$, runs $\mathrm{Fano} \rightarrow \text{genus-14} \rightarrow \text{three-qubit} \rightarrow \mathrm{Suzuki} \rightarrow E_8$. (Witness: `analysis/w33_g2_thread.py`.)

C.11 The ledger, the Standard Model, and the tower as one figure

Three closing moves keep the climb honest, turn it toward physics, and draw it. *The ledger.* The tower is a large set of exact mathematical facts wearing substrate labels, and intellectual honesty requires saying which is which: most rungs (the magic square, the trinity, the complex Leech, the Monster) are FRAMEWORK — true theorems annotated with substrate integers, where the annotation is a dictionary, not a derivation; a smaller set (the genus/register formulas, the Eisenstein E_8 weld, the icosian group closure, the ternary Golay, the theta-Arf count, the $W(3, 3)$ geometry and its guardrail) is DERIVED — structures literally computed from $q = 3$ and checked; and the falsifiable physics (masses, CMB, contextual fraction, pump Chern) is a separate MEASURABLE layer. *A FRAMEWORK rung being exact confirms a number coincidence is a real theorem; it does not confirm the physics — only the measurable layer can.* (Witness: `analysis/w33_exceptional_tower_ledger.py`.)

The Standard Model. The exceptional E_6 sits at the head of the trinification chain $E_6 \rightarrow \mathrm{SU}(3)_C \times \mathrm{SU}(3)_L \times \mathrm{SU}(3)_R \rtimes S_3 \rightarrow \mathrm{SU}(3) \times \mathrm{SU}(2) \times \mathrm{U}(1)$, with substrate-clean dimensions: the trinification adjoint is $3 \cdot 8 = 24 = f$ and the Standard-Model gauge dimension is $8 + 3 + 1 = 12 = k$. The 27 is one generation ($9 + 9 + 9 = \text{quarks/leptons/Higgs}$) and the three generations are the S_3 triality of the three $\mathrm{SU}(3)$ factors — not inserted, but the same triality that permutes the three Hesse nonets. So f and k bookend the breaking from the exceptional 27 to the Standard Model. (Witness: `analysis/w33_standard_model_from_trinification.py`.)

The figure. The whole session is one ascending ladder — $q = 3 \rightarrow \text{genus/ register tower} \rightarrow$

Witting body $\rightarrow E_6/E_7/E_8 \rightarrow$ complex Leech $\rightarrow Co_0 \rightarrow$ Monster ($c = 24$), with a downward branch to the Standard Model — run the full height by three threads (the Eisenstein order-3 ω , $G_2 = \text{Aut}(\mathbb{O})$, and the simple group $U_4(2) = \text{PSp}(4, 3)$). It is drawn in `docs/w33_exceptional_tower.dot` and rendered to `docs/w33_exceptional_tower.svg`. (Witness: `analysis/w33_exceptional_tower_synthesis.py`.)

C.12 The measurable layer: a scorecard and a unification [Faces 2, 3, 7]

Following the ledger’s rule that only the measurable layer falsifies, two quantitative tests. *The scorecard.* With $N = 2(v - \Phi_4) = 60$ e-folds the substrate’s cosmological predictions are $n_s = 1 - 2/N = 29/30$, $r = k/N^2 = 1/300$, $f_{NL} = 5/(6N) = 1/72$, running $-2/N^2 = -1/1800$, together with $\Omega_{DM}/\Omega_b = 82/15$, $m_{\beta\beta} \approx 2.3 \text{ meV}$, and $\sum m_\nu \approx 0.101 \text{ eV}$. Against mid-2026 data these are mostly CONSISTENT (n_s at $\sim 0.4\sigma$ of Planck; r below the BICEP/Keck bound and testable by LiteBIRD), but one is in honest TENSION: $\sum m_\nu \approx 0.101 \text{ eV}$ satisfies Planck ($< 0.12 \text{ eV}$) yet exceeds the DESI 2024 bound ($< 0.072 \text{ eV}$) — a genuine falsification handle, not a confirmation. (Witness: `analysis/w33_measurable_scorecard_2026.py`.)

The unification. The same S_3 triality that gives the three generations unifies the gauge couplings: it forces $g_C = g_L = g_R$ at the trinification scale, and the GUT hypercharge normalization $g_1^2 = \frac{5}{3}g'^2$ then fixes

$$\sin^2 \theta_W = \frac{g'^2}{g^2 + g'^2} = \frac{3/5}{1 + 3/5} = \frac{3}{8}$$

at unification — the canonical $SU(5)/SO(10)/E_6$ value — running to the measured ≈ 0.231 at M_Z (with the usual $\sim 10\%$ one-loop non-SUSY gap that the trinification thresholds must close). So the three generations and coupling unification are one triality. (Witness: `analysis/w33_trinification_unification.py`.)

C.13 Three live tests, scrutinised

The neutrino tension, both ways. The scorecard’s one red flag deserves honest scrutiny. Externally, $\sum m_\nu \approx 0.10 \text{ eV}$ is in $\sim 1.5\text{--}2\sigma$ tension with the tight DESI-2024 ΛCDM bound ($< 0.072 \text{ eV}$) but consistent with the looser $w_0 w_a \text{CDM}$ bound ($< \sim 0.16 \text{ eV}$) that DESI’s own data prefer — so whether the theory is falsified here depends on the dark-energy model. Internally, a naive geometric mass cascade ($r = 0.5225$) predicts $\Delta m_{21}^2/\Delta m_{31}^2 = 0.21$ against the observed 0.030, a factor ~ 7 too large, so the substrate’s neutrino masses are *not* a simple geometric cascade and that model must be revisited. Both flags are recorded, not smoothed. (Witness: `analysis/w33_neutrino_desi_scrutiny.py`.)

Unification, honestly. Minimal one-loop SM running does not unify — the couplings meet pairwise at $\sim 10^{13}, 10^{14}, 10^{17} \text{ GeV}$, a ~ 4 -order spread. The two-step chain $E_6 \rightarrow SU(3)^3 \rightarrow \text{SM}$ inserts one intermediate scale, giving two free scales that fit the two low-energy constraints ($\sin^2 \theta_W$ and α_s); so trinification unifies, with $\sin^2 \theta_W = 3/8$ the exact boundary, at the price of predicting M_I and M_{GUT} — the honest status of every non-SUSY GUT. (Witness: `analysis/w33_trinification_two_step_unification.py`.)

The demonstrator as a meter. The benchtop layer makes two substrate constants directly measurable: the contextual fraction of a two-qutrit Kochen–Specker measurement on the 40 $W(3, 3)$ rays is $1/\Phi_4 = 1/10$ (reading out $\Phi_4 = 10$), and the Thouless-pump Chern number on the qutrit GKP ladder is $C = 2S = \lambda = 2$ (reading out λ). Both are integer/unit-fraction — robust, calibration-free — so a measured $1/10$ and 2 confirm the substrate’s contextuality and topology,

and any deviation falsifies them on a table. (Witness: `analysis/w33_demonstrator_substrate_constants.py`.)

C.14 Fixing the neutrino, predicting the proton, and the balance sheet

The neutrino, fixed. The DESI tension was a model artifact. The geometric cascade mismatched the Δm^2 hierarchy (0.21 vs 0.030) and forced $\sum m_\nu \approx 0.10$ eV; the substrate’s type-I seesaw with a hierarchical Dirac Yukawa instead *squares* the hierarchy, driving the lightest mass small. For strong normal ordering ($m_1 \lesssim 0.009$ eV) the spectrum is fixed by oscillation data and $\sum m_\nu \approx 0.06$ eV — below the DESI bound, so the tension is *resolved* — with $m_{\beta\beta} \approx 2\text{--}4$ meV. The corrected prediction is $\sum m_\nu \approx 0.06$ eV (strong NO), not 0.10. (Witness: `analysis/w33_neutrino_seesaw_texture.py`.)

The proton. The two-step fit gives $M_{\text{GUT}} \sim 10^{16}$ GeV, so $\tau_p \sim (1/\alpha_{\text{GUT}}^2) M_{\text{GUT}}^4/m_p^5 \approx 4.6 \times 10^{35}$ yr via $p \rightarrow e^+\pi^0$ — above the Super-Kamiokande bound (2.4×10^{34} yr, which already requires $M_{\text{GUT}} \gtrsim 10^{15.7}$ GeV) and within Hyper-Kamiokande’s reach. A sharp falsifiable GUT-scale prediction. (Witness: `analysis/w33_proton_lifetime_gut_scale.py`.)

The balance sheet. After building, auditing, applying, testing, and scrutinising the tower: the cosmological and electroweak predictions are CONSISTENT; the one DESI tension is FIXED; proton decay, the contextual fraction 1/10, and the pump Chern 2 are TESTABLE near-term handles; the $q = 3 \rightarrow$ Monster tower is exact FRAMEWORK (a magnificent dictionary, not a confirmation of physics); and the open problems are the precise neutrino texture, the GUT scales, and — above all — the bridge from the framework to why these numbers are the physical world, which is the program’s defining task. (Witness: `analysis/w33_state_of_the_theory.py`.)

C.15 Why $q = 3$, and one triality for the neutrino [Faces 1, 4]

The bridge: $q = 3$ is triply forced. The defining task — why is the $q = 3$ substrate the physical world? — has a sharp answer in the form of a convergence. Three independent first-principles selections each return $q = 3$: (i) *geometric*, the master equation $q! = 2q$ is equivalent to $(q-1)! = 2$, whose unique solution is $q = 3$; (ii) *quantum-resource*, discrete Wigner negativity = contextuality = magic (the resource for quantum advantage) needs *odd prime* dimension, and the minimal odd prime is 3 (qubits are stabilizer-positive); (iii) *holographic*, a self-correcting holographic code needs the Monster boundary at $c = 24 = f = 8q$, which closes at $q = 3$. Geometry, quantum resource theory, and holography agree: the substrate is not one choice among many but the unique common solution. (Witness: `analysis/w33_q3_triple_selection.py`.)

One triality, the whole neutrino sector. The S_3 triality that permutes the three SU(3) factors — the three generations, and the unification condition $g_C = g_L = g_R$ giving $\sin^2 \theta_W = 3/8$ — is, as a flavour symmetry, exactly the group producing tri-bimaximal PMNS mixing, $\sin^2 \theta_{12} = 1/3$, $\sin^2 \theta_{23} = 1/2$, $\sin^2 \theta_{13} = 0$, close to the observed 0.304, ~ 0.5 , 0.022 (the small θ_{13} the Φ_3 -scale deformation). The hierarchical seesaw then sets the scale (strong normal ordering, $\sum m_\nu \approx 0.06$ eV). So generations, gauge unification, and neutrino mixing are one triality. (Witness: `analysis/w33_neutrino_tbm_s3.py`.)

The proton signature. The two-step fit gives $M_I \sim 10^{13.5}$ GeV and $M_{\text{GUT}} \sim 10^{16}$ GeV; non-SUSY gauge-mediated decay makes $p \rightarrow e^+\pi^0$ the dominant channel (branching $\sim 30\text{--}60\%$, with $p \rightarrow \bar{\nu}K^+$ suppressed for lack of superpartners), and $\tau_p \approx 4.6 \times 10^{35}$ yr sits squarely in the Hyper-Kamiokande window. (Witness: `analysis/w33_proton_branching_two_scale.py`.)


C.16 Pass 1 — three closures: the neutrino matrix, the selection census, the threshold number

Reading the next four subsections. They make four passes over three of the five faces of §C.1 — the **neutrino** (Face 3), the **selection** and its constants (Faces 1–2), and the **code** (Face 4) — each pass sharper than the last: closure (Pass 1), invariant (Pass 2), skeleton (Pass 3), and the cubic form / Witting degrees / measured curve (Pass 4). The same three threads run through all four; the fifth face, the demonstrator, ties them together at the end (§C.12 and the bridge witness).

Three sharper statements close gaps left open above.

(1) *The neutrino texture, pinned to one matrix.* The mixing pattern (§C.15) and the mass scale were established separately; an earlier single-parameter geometric cascade gave $\Delta m_{21}^2/\Delta m_{31}^2 = 0.21$, about $7\times$ the observed 0.030 — recorded honestly as a wrong-model flag. That flag is now closed by a *single* matrix. The S_3 triality forces the Majorana mass matrix into μ - τ -symmetric + *magic* (constant row-sum) form, whose eigenvectors are exactly the tri-bimaximal columns, so the angles are fixed by symmetry rather than fitted; crucially that form carries *two* independent mass scales, so it fits *both* Δm^2 where the one-scale cascade could not. Diagonalising the exact TBM matrix numerically returns $\sin^2 \theta = (1/3, 1/2, 0)$ and $\Delta m_{21}^2/\Delta m_{31}^2 = 0.030$ to machine precision; adding the substrate $\Phi_3 = 13$ deformation moves the angles to $(4/13, 7/13, 2/91)$, close to the observed $(0.307, 0.546, 0.0220)$, with the ratio unchanged. (Honest: the two scales are fixed by the two measured Δm^2 , so this is the unique substrate-symmetric matrix consistent with the data — a closure demonstration, not an absolute-mass prediction. Witness: `analysis/w33_neutrino_texture_pinned.py`.)

The ratio as a partial prediction. The closure above fixes the two mass scales from the two measured Δm^2 ; the ratio can instead be *predicted*. Feeding the substrate’s own mild Yukawa hierarchy — up-type singular values $10:5:1$ (Pillar 68), with $Y_\nu = Y_{\text{up}}$ at the GUT scale — into a type-I seesaw with a minimal flat S_3 -symmetric M_R , with *no* neutrino-sector input, the seesaw squares the hierarchy and returns $\Delta m_{21}^2/\Delta m_{31}^2 \approx 0.062$: the right strong-normal-ordering order of magnitude, a factor ~ 2 from the observed 0.030; the down-type ratios $6:2:1$ give 0.012, so the observed value is *bracketed* (against the 0.21 — a factor 7 — of the geometric cascade). The residual factor ~ 2 is named, not removed: the M_R texture and $Y_\nu \neq Y_{\text{up}}$ running. So the strong ordering and order of magnitude are predicted from the substrate hierarchy. (Witness: `analysis/w33_neutrino_seesaw_prediction.py`.)



two independent mass scales $\Rightarrow$ fits *both* Δm^2 : $\Delta m_{21}^2/\Delta m_{31}^2 = 0.030$ (observed), *not* the 0.21 of the one-scale cascade

Figure 10: The neutrino sector closed by one matrix. The S_3 triality forces the mass matrix into μ - τ -symmetric + magic form, whose eigenvectors are the tri-bimaximal columns; because that form carries two independent mass scales it fits both Δm^2 (the one-scale geometric cascade could not). The $\Phi_3 = 13$ deformation moves the angles onto the observed values.

(2) *The $q = 3$ selection census.* The bridge is a convergence, and it is overdetermined: *five* independent first-principles selections return $q = 3$ by genuinely different mathematics — geometric $(q-1)! = 2$ (a forcing equation), minimal odd prime (a minimality fact), holographic $c = 24 = 8q$ (a charge match to the Monster ceiling), thermodynamic de Sitter closure $2(q-1)(q^2+1) = (1+q)(1+q^2) \Leftrightarrow (q-3)(q^2+1) = 0$ (a cubic, a *different* polynomial from the factorial, so not one equation

in disguise), and a new *Eisenstein complex-polytope* selection: the regular complex polytopes with order-3 reflections over $\mathbb{Z}[\omega]$ form the exceptional tower Möbius–Kantor $3\{3\}3$ (8 vertices, $2T$) $\rightarrow$ Hessian $3\{3\}3\{3\}3$ ($27 =$ lines on a cubic $= E_6$) $\rightarrow$ Witting $3\{3\}3\{3\}3\{3\}3$ ($240 = E_8$ roots), with Witting symmetry the Shephard–Todd group #32 of order $155520 = 3 \times |\mathrm{Sp}(4, 3)|$. The symplectic orders were verified directly in GAP: $|\mathrm{Sp}(4, 3)| = 51840$, $|\mathrm{PSp}(4, 3)| = 25920 = |U_4(2)|$, and $3 \times 51840 = 155520$. (Witness: `analysis/w33_q3_selection_census.py`.)

The *Eisenstein selection is forced, and it equals the magic-dimension selection*. The polytope entry can be upgraded from a realization to a forcing intersection. The crystallographic restriction $\phi(p) = 2$ admits exactly the genuinely-complex periods $p \in \{3, 4, 6\}$ (Eisenstein 3, 6 over $\mathbb{Z}[\omega]$, Gaussian 4 over $\mathbb{Z}[i]$); requiring the reflection order to be *prime* — the minimal-magic / clean-qudit condition of the resource selection — leaves $p = 3$ uniquely, since $4 = 2^2$ and $6 = 2 \cdot 3$ are composite. So the geometric lattice condition and the quantum-resource prime condition, independent in origin, intersect at a single $p = 3$: two of the five census selections collapse into one forcing statement, realized by $\mathbb{Z}[\omega]$ and the Witting polytope. (GAP-verified: $\phi^{-1}(2) \cap [3, \infty) = \{3, 4, 6\}$, prime part $\{3\}$. Witness: `analysis/w33_eisenstein_forcing.py`.)




Figure 11: The $q = 3$ selection census: five first-principles selections, each a different kind of mathematics (two forcing equations of distinct degree, a minimality fact, a holographic charge match, and the Eisenstein complex-polytope tower $2T \rightarrow E_6 \rightarrow E_8$), all return $q = 3$. The substrate value is overdetermined.

(3) *The fault-tolerance threshold as a lab number*. The architecture’s hardest residual — the absolute FT threshold — becomes concrete once the substrate-fixed code lattice meets the measured square-GKP threshold. The best reported square-GKP + surface-code threshold is 9.9 dB of squeezing (Noh–Chamberland 2022; 11.2 dB under minimum-weight matching), and 9.5 dB has already been demonstrated in hardware. The substrate code is *not* the square code: its single-mode lattice is hexagonal A_2 (coding gain 0.6 dB), its natural two-mode code is D_4 (1.5 dB), and its four-mode code is E_8 (3.0 dB), so the threshold shifts down to $\sim 9.3/8.4/6.9$ dB respectively — all at or below the 9.5 dB already reached. Universality needs only the minimal Lloyd–Braunstein set, Gaussian (degree-2, the $\mathrm{Sp}(4, 3)$ Weil rep) plus one cubic gate (degree-3, E_6), with cubic-phase magic states the sole distilled resource. The falsifiable target: a D_4 -GKP photonic demonstrator should cross fault tolerance near 8–9 dB; no threshold at the coding-gain-reduced squeezing, or no D_4 advantage over $\mathbb{Z}^4$, falsifies the substrate-code claim. (Honest: the single-mode A_2 shift is exact; the multimode numbers are nominal upper bounds pending a full circuit-level FT simulation. Witness: `analysis/w33_holonet_ft_threshold_budget.py`.)




Figure 12: The fault-tolerance threshold as a lab number. The measured square-GKP + surface-code threshold (9.9 dB, Noh–Chamberland 2022; 11.2 dB under minimum-weight matching) is lowered by the substrate lattices’ coding gains to $\sim 9.3/8.4/6.9$ dB for $A_2/D_4/E_8$ — all at or below the 9.5 dB squeezing already demonstrated in hardware. (Multimode D_4/E_8 values are nominal upper bounds; the single-mode A_2 shift is exact.)

C.17 Pass 2 — sharper: a Majorana invariant, a cyclotomic merge, a code curve

Each of the three closures admits a sharper statement.

The neutrino factor of two closes on a substrate invariant. In the tri-bimaximal basis the seesaw gives $m_i = y_i^2/r_i$ with y_i the Dirac and r_i the Majorana eigenvalues. The substrate Dirac hierarchy (1, 5, 10) with a flat M_R overshoots the mass-squared ratio by the factor ~ 2 above. The Majorana eigenvalue ratio the seesaw needs is $r_2/r_3 = 1.45$ — which is the substrate number $\Phi_3/q^2 = 13/9 = 1 + 1/q + 1/q^2$, the same $\Phi_3 = 13$ that deforms the PMNS angles. Fixing $r_2/r_3 = \Phi_3/q^2$ (not fitted) gives $\Delta m_{21}^2/\Delta m_{31}^2 = 0.030$, $m_2 = 8.6$ meV $= \sqrt{\Delta m_{21}^2}$, and $\sum m_\nu \approx 59$ meV — the observed strong ordering. (Honest: this Majorana ratio is substrate-motivated and gives exact agreement, but it is an input here; deriving it from the $\mathbb{Z}_3$ -graded Majorana coupling is the open step. Witness: `analysis/w33_neutrino_majorana_texture.py`.)

Two of the selections are one cyclotomic structure. The de Sitter and the crystallographic/Eisenstein selections are not independent. The de Sitter closure cubic factors as $(q-3)\Phi_4(q)$ with $\Phi_4(q) = q^2+1$ (complex roots $\pm i = \zeta_4$, the Gaussian period); the crystallographic periods $\{3, 4, 6\}$ are exactly the three degree-two cyclotomics $\{\Phi_3, \Phi_4, \Phi_6\}$; and Φ_4 already factors the GQ(q, q) point count $v = (q+1)\Phi_4(q)$ ($= 40$ at $q = 3$). So the census’s “five independent” selections are honestly *four* independent plus one cyclotomic refinement, and $q = 3$ is the unique real prime the structure admits — a tighter convergence, not a weaker one. (Witness: `analysis/w33_desitter_crystallographic_unify.py`.)

The D_4 code as a measured curve. The threshold number has an explicit code and curve behind it. The D_4 stabilizer lattice $\{x \in \mathbb{Z}^4 : \sum x_i \text{ even}\}$ (generators $\times \sqrt{\pi}$) has exact invariants $d_{\min}^2 = 2$, $\det = 4$, kissing 24, coding gain 1.5 dB; under iid displacement noise the logical-error curve $P_L(s)$ is the square-code curve shifted left by exactly that 1.5 dB (Fig. 13), so D_4 reaches any target error at up to 1.5 dB less squeezing and the surface-GKP threshold 9.9 dB becomes ~ 8.4 dB — below the 9.5 dB already demonstrated. The relative D_4 -versus-square advantage is the falsifiable handle. (Honest: the lattice quantities are exact and the 1.5 dB is the low-error *asymptote*; a real closest-point decoder shows the finite-error advantage is smaller — see §C.18. Witness: `analysis/w33_d4_gkp_error_curve.py`.)




Figure 13: The D_4 -GKP logical-error curve as a spec. Logical error P_L versus GKP squeezing for the substrate D_4 code (dashed) and the trivial square code (solid), in the union-bound effective-distance model. The D_4 curve is the square curve shifted left by its *asymptotic* 1.5 dB coding gain. This idealised constant shift is corrected by a real decoder in §C.18 (Fig. 14): the actual advantage is error-rate dependent and smaller at realistic error rates.

C.18 Pass 3 — deeper: a cyclotomic skeleton, a projective Majorana ratio, a decoder correction

Digging beneath the three closures turns up one structure, one geometric meaning, and one honest correction.

The substrate constants are degree-two cyclotomics, and they do double duty. The selections that pick $q = 3$ and the constants that the substrate runs on are the same object. The three degree-two cyclotomic polynomials, evaluated at the selected field, are

$$\Phi_3(3) = 13, \quad \Phi_4(3) = 10, \quad \Phi_6(3) = 7, \quad 13 + 10 + 7 = 30 = h(E_8),$$

i.e. the substrate’s core invariants (register $\Phi_3 = 13$, the $\text{Sp}(4)$ $\theta = 10$, the Fano/Hurwitz 7), summing to the E_8 Coxeter number. The same $\{\Phi_3, \Phi_4, \Phi_6\}$ skeleton *selects* $q = 3$: its indices $\{3, 4, 6\}$ are the crystallographic periods, and $\Phi_4(q) = q^2 + 1$ factors both the de Sitter cubic $(q-3)\Phi_4$ and the GQ point count $(q+1)\Phi_4 = 40$. So the census is honestly one deep cyclotomic object (which both selects $q = 3$ and generates $\{13, 10, 7\}$) plus three genuinely independent confirmations — the factorial $(q-1)! = 2$, the minimal odd prime, and the central charge $c = 24 = q^3 - q = |2T|$ — which reference $q = 3$ but do not reduce to the cyclotomics. (Witness: `analysis/w33_cyclotomic_skeleton_census.py`.)

The Majorana ratio is projective over affine. The neutrino number $\Phi_3/q^2 = 13/9$ has a geometric meaning. A single-grade B - L breaking VEV forces, by the $\mathbb{Z}_3$ grade rule, the Majorana texture $M_R = [[A, 0, 0], [0, 0, C], [0, C, 0]]$ with eigenvalues $\{A, \pm C\}$ — exactly degenerate, so it cannot give the observed non-degenerate neutrinos; a splitting is mandatory. The splitting that works is $r_2/r_3 = |\text{PG}(2, 3)|/|\text{AG}(2, 3)| = 13/9$: the democratic (trivial-rep) right-handed neutrino couples to all $\Phi_3 = 13$ projective points, the others to the $q^2 = 9$ affine ones. Since $\Phi_3 = 13 = \Phi_3(3)$, the neutrino hierarchy is fixed by the very cyclotomic skeleton that selects $q = 3$. (Honest: the degeneracy

is exact and $13/9$ now has a geometric meaning, but identifying M_R 's eigenvalues with these counts from the cubic form is still open. Witness: `analysis/w33_majorana_grade_derivation.py`.)

A real decoder corrects the threshold advantage. Replacing the union-bound shift with an actual Conway–Sloane closest-point decoder over sampled Gaussian noise (both codes at equal density) shows the D_4 -versus-square advantage is *error-rate dependent*: it grows from ~ 0.6 dB at logical error 3×10^{-2} to ~ 0.9 dB at 10^{-3} , tracking the union bound *with* kissing numbers (D_4 's 24 versus the square's 8 suppresses the finite-error gain) and reaching the nominal 1.5 dB only asymptotically. So the constant-1.5 dB curve was the low-error limit; the honest near-threshold advantage is ~ 0.7 – 1.0 dB, with the full coding gain realised deep below threshold. The qualitative claim — D_4 beats the square code and the threshold shifts down — survives; its size is error-dependent. (Witness: `analysis/w33_d4_decoder_montecarlo.py`.)

C.19 Pass 4 — deepest: the cubic form, the Witting degrees, and the measured curve

Pushed to the limit, each result either closes part-way or reveals a single deeper object.

The cubic form proves the degeneracy and bounds the rest. Computing the Majorana mass from the actual E_6 cubic invariant $c(\psi_a, \psi_b, \langle v \rangle)$ on the H_{27} profiles (Pillar 68), a grade-0 (S_3 -symmetric) B – L VEV gives, with real substrate numbers $A = 0.0017$, $C = 0.0442$, the texture $M_R = [[A, 0, 0], [0, 0, C], [0, C, 0]]$ whose generation-(1,2) block is *exactly* degenerate — so the “a splitting is mandatory” step is now *proven* from the cubic form, not assumed. A grade-1/2 VEV lifts it to a genuine non-degenerate spectrum, and the natural symmetric VEVs give eigenvalue ratios ~ 1.25 — near, but not equal to, $\Phi_3/q^2 = 1.444$; the exact $13/9$ needs a particular B – L direction. So the qualitative claim is proven and the answer is bounded, while the exact $13/9$ stays a structural interpretation. (Witness: `analysis/w33_majorana_cubic_form.py`.)

The substrate constants are the Witting degrees. The cyclotomic skeleton has a companion: the Witting complex reflection group (Shephard–Todd #32, the $q = 3$ Eisenstein polytope's symmetry, $|G| = 155520 = 3|\text{Sp}(4, 3)|$) has the four invariant degrees $\{12, 18, 24, 30\}$, whose product is $|G|$, and these are

$$12 = q(q+1) = k, \quad 18 = h(E_7), \quad 24 = q^3 - q = |2T| = f = c, \quad 30 = h(E_8) = \Phi_3(3) + \Phi_4(3) + \Phi_6(3).$$

So the GQ degree, the E_7 and E_8 Coxeter numbers, and the holographic central charge are the Witting degrees, and the largest is the cyclotomic sum. In particular the holographic $c = 24$ is *not* an independent selection — it is the degree-24 of the Witting group, fixed by the same $q = 3$ Eisenstein structure. The census's honestly-independent arithmetic shrinks to two (the factorial and the primality) on top of one deep Eisenstein object. (Witness: `analysis/w33_witting_degrees_unify.py`.)

The advantage is a curve, measured to 10^{-6} . Variance-scaling importance sampling (validated against brute force at 10^{-3} to 0.0%) pushes the real Conway–Sloane decoder deep below threshold. The D_4 -versus-square advantage grows $0.48 \rightarrow 0.74 \rightarrow 0.92 \rightarrow 1.03 \rightarrow 1.11 \rightarrow 1.17$ dB as the logical error falls $10^{-1} \rightarrow 10^{-6}$ (Fig. 14), tracking the union bound and reaching only ~ 1.2 dB even at 10^{-6} — the full 1.5 dB is realised only at far lower error. So the honest engineering spec is a curve: D_4 buys ~ 0.7 dB near threshold, ~ 1.2 dB at 10^{-6} , the full coding gain only asymptotically. (Witness: `analysis/w33_d4_advantage_curve.py`.)

C.20 Self-error-correction and why the primitive is one

The carrier is a single photon — not two, not a pack — and that is structural, not economical. Two ingredients explain why the universe's primitive is one.




Figure 14: The measured D_4 -versus-square advantage as a function of target logical error, from a real Conway–Sloane closest-point decoder (importance-sampled below 10^{-3} , validated against brute force). The advantage is error-rate dependent — ~ 0.7 dB near threshold (10^{-2}), ~ 1.2 dB at 10^{-6} — approaching the nominal 1.5 dB coding gain (dashed) only asymptotically, because D_4 ’s larger kissing number suppresses the finite-error gain. This corrects the constant-shift idealisation of Fig. 13.

Self-entanglement instead of inter-particle entanglement. A photon carries several internal registers (polarization, path, time-bin, frequency, transverse mode), and entanglement among them needs no second particle. Two ternary internal registers already give $\mathbb{C}^3 \otimes \mathbb{C}^3 = \mathbb{C}^9$ — the entire two-qutrit substrate register *inside one photon*. The geometry is realised within a single indivisible quantum, which is exactly the operator–operand self-duality of §2.

Self-error-correction, and no-cloning as a feature. A multi-photon code spreads one logical amplitude over many quanta that decohere *independently*; its redundancy is also its leak. A single photon cannot be cloned — and here that is the asset, not the obstacle: the logical amplitude is one indivisible excitation, with no independent copies to dephase against each other. Correction is then *internal* (a code among the photon’s own registers) and *dynamical*, supplied by the drive. The Boerdijk–Coxeter loop (§11) is not a periodic clock but a genuine *time quasicrystal*: its twist $\theta = \arccos(-\frac{2}{3})$ is substrate-fixed — at $q = 3$, $-\frac{2}{3} = -(q-1)/q = 2 \cdot (-1/q)$, the pairwise vertex dot of the regular q -simplex (the tetrahedron = the $q=3$ simplex) — and the stroboscopic orbit $\{n\theta \bmod 2\pi\}$ obeys the Steinhaus three-gap law at every N (exactly two gaps at $n = 30 = h(E_8)$), is Weyl-equidistributed, and never recurs (Niven). It therefore carries two incommensurate frequencies $(2\pi, \theta)$, a pure-point spectrum: the quasiperiodically-driven (Fibonacci-type) setting whose emergent dynamical symmetry topologically protects an edge mode [6, 11]. The photon self-protects its logical qutrit through its own substrate-fixed drive — no copies, no external apparatus. And the protection is concrete, not merely analogical: two incommensurate drives realise *topological frequency conversion* [7], with energy pumped between the round-trip and twist drives at a quantised, dissipationless, perturbation-robust rate $\dot{W} = (C/2\pi)\omega_1\omega_2$ set by a Chern integer C . We compute it. Because the BC twist is a rotation (the holonomy group $2T \subset \text{SU}(2)$), the qutrit carrier transforms as the spin-1 representation, and the spin-1 two-tone pump has band Chern numbers

$$(C_-, C_0, C_+) = (+2, 0, -2)$$

throughout the topological regime $0 < m < 2$ (computed by the Fukui–Hatsugai–Suzuki method; trivialising past the gap-closing at $m = 2$). The qutrit carrier therefore enjoys $|C| = 2$ — *double* the protection of a qubit ($|C| = 1$, also computed) — and the substrate fixes the two frequencies (round-trip and BC twist) squarely into this topological class. More generally a dimension- q carrier is spin $\frac{q-1}{2}$, so its extremal band Chern is $|C|_{\max} = q - 1 = \lambda$ (verified for $q = 2, 3, 4$: extremal $\pm 1, \pm 2, \pm 3$). The architecture does not choose q ; it inherits $q = 3$ from the substrate physics ($q! = 2q$, KO-dimension $2q = 6$) and reaps the enhanced protection as a consequence. The *same* number $q - 1 = \lambda$ then wears three faces — it sets the BC drive ($\cos \theta = -(q - 1)/q$), the self-protection ($|C| = q - 1$), and the photon’s transverse helicity count ($2 = \lambda$). (*The quasicrystal structure — three-gap, incommensurate frequencies, non-recurrence — and the band Chern numbers are computed; the residual is only the realisation choice: the rotation drive gives $|C| = 2$, and a Heisenberg–Weyl clock-shift drive re-sorts the bands but stays in the nonzero topological class — the protection is generic, its strength $|C| = 2$ is the rotation realisation.*) (Witnesses: `analysis/w33_self_correction_quasicrystal.py`, `analysis/w33_topological_pump_test.py`, `analysis/w33_qutrit_topological_pump.py`.)

C.21 Why the primitive is one *massless* photon

One more layer answers *why a photon* and not some other quantum. Wheeler’s one-electron universe (relayed to Feynman in 1940) imagined every electron as a single worldline weaving forward and backward through time. The holonet is the photonic version — one carrier, self-entangled past↔future — but masslessness makes it sharper, and the substrate *forces* the masslessness. Three facts lock together.

Null worldline, zero proper time. A massless carrier moves on $ds^2 = 0$, so its proper time $\tau = t\sqrt{1 - \beta^2} \rightarrow 0$: its entire worldline is one proper-time instant, and its past and future are proper-time-*simultaneous*. That is exactly what lets a single photon entangle its past and future registers — a relation at one proper-time point. Wheeler’s massive electron must literally weave; the massless photon is all-at-once. Self-reference wants masslessness.

Gauge invariance forces it. The photon is the gauge boson of the substrate connection of §6 — the one whose holonomy is the Clifford gates ($2T = \text{SL}(2, 3)$) and whose curvature is the gauge fields. The substrate gauge symmetry is unbroken (the holonomy is the full Clifford group), and gauge invariance forces an unbroken gauge boson to be massless, removing the longitudinal mode (Wigner’s classification).

The helicity count is the substrate’s. A massless vector has Wigner little group $\text{ISO}(2)$ and exactly *two* transverse helicities ± 1 ; a massive vector has $\text{SO}(3)$ and *three*. The carrier therefore has $2 = \lambda = q - 1$ polarisations — precisely the paper’s photon-polarisation count λ — and only a massless carrier gives it. (Its clock is then the optical frequency: $\tau = 0$ but the phase $\phi = \omega t$ still accrues, so the time-bin/frequency registers — the very degrees of freedom that self-entangle — are the photon’s internal clock. A testable corollary: in vacuum the self-entanglement carries no proper-time decoherence, i.e. its visibility is worldline-length independent up to lab-frame loss and dispersion.)

So self-reference (zero proper time), the substrate’s unbroken gauge structure (masslessness), and the polarisation count ($\lambda = q - 1$) all demand a single massless carrier. The universe’s primitive is one photon. (Witness: `analysis/w33_why_one_massless_photon.py`.)

C.22 The single carrier, the geon, and the fractal of nested shells

Three pictures of the same machine fall out, each tested. *(i) Minimal activation.* If one carrier’s worldline braids through the substrate’s topology to compute, how few controls suffice? The gate group is $\text{Sp}(4, 3)$ (order 51840), and it is 2-generated: a search over the natural symplectic “switches” (one-mode Fourier/phase, two-mode CZ/SUM) finds 13 generating *pairs* — e.g. $\langle F_2, F_1 \cdot CZ \rangle$ reaches all 51840. So flipping *two well-chosen switches* in the right temporal pattern (a binary word, scheduled by the Boerdijk–Coxeter quasicrystal) reaches every gate; the naive pair $\langle F_1, CZ \rangle$ does *not* (it generates only a 108-element subgroup), so the pair must be a generating one. Moreover the patterns are *short*: the Cayley diameter of $\text{Sp}(4, 3)$ under such a pair is only 14, so every one of the 51840 gates is a word of at most 14 switch-flips (mean ≈ 10), and the diameter is the same across the generating pairs tested — the worst-case compilation length is substrate-fixed, with the hardware entering through *which* operations are the two switches, not through the length. The switches are two elementary braids of the carrier’s worldline, robust because braiding depends only on the topology — the holonomic robustness again, and the photoelectric effect is the hinge that lets the same word run on a massless photon (all-at-once) or a massive electron braided through a machine’s wiring. As literal braids the minimal hardware is tiny: a single qutrit needs three junctions ($B_3 \twoheadrightarrow \text{PSL}(2, \mathbb{Z}) \rightarrow \text{SL}(2, 3) = 2T$, the same $2T$ as the gauge-connection holonomy), and four symplectic-transvection junctions already generate the full $\text{Sp}(4, 3)$. Their interaction graph (an edge where two junctions braid, $\langle v_i, v_j \rangle = \pm 1$; a non-edge where they commute) is a *path* P_4 , so the minimal hardware is four junctions *in a line* — consecutive braid, non-consecutive commute. To run the machine on a given fabric is then to place that path: any wiring containing a simple four-junction path hosts it, which is every grid, ring, mesh, hypercube and almost every random fabric. Topology does genuinely matter in the pathological case — a pure star has no four-path, and the symplectic geometry forbids reshaping the interaction graph into a three-leaf star (mutually commuting junctions are isotropic, at most a Lagrangian) — but every real fabric qualifies, so “runs on any classical machine” becomes a concrete graph search for a four-junction path.

(ii) The carrier as its own self-contained shell. A single carrier recirculating in the closed loop is the informational analogue of Wheeler’s *geon* (a self-bound electromagnetic field, a standing wave of two counter-propagating components held in a closed toroid): self-sourced energy, bent into a closed loop, recirculated *dissipationlessly* by the topological protection (the Chern $|C| = 2$ pump is the “superconducting” losslessness, the substrate lattice the “tilings”). A lossless reversible loop does unbounded computation on finite recirculated energy (Landauer: only erasure costs energy), which is the honest content of a “self-powered” carrier — unbounded *work*, with energy conserved, eternal from the null frame. (A single optical photon is ~ 56 orders from a literal *gravitational geon*; that forms at the Planck scale, where the substrate’s derived Einstein equations govern.)

(iii) The network is a fractal of nested shells. Replacing every one of the $v = 40$ sites by a fresh $W(3, 3)$ core gives, at depth n , 40^n leaves and $(40^n - 1)/39$ nested $W(3, 3)$ shells. Reading each shell as a self-contained lossless, code-protected computer makes the network a self-similar tower in which each layer is the conceptual “Dyson sphere” of the shell outside it: every shell is an error-correcting code on its 40 sites, the outer code holographically containing the inner layers’ logical information (boundary encodes bulk), and reversible shells composing into a single lossless network computer. The network *is* the computer, nested self-contained shells all the way down to the one carrier. (Witnesses: `analysis/w33_two_switch_generation.py`, `analysis/w33_photon_geon_dyson_sphere.py`, `analysis/w33_fractal_nested_dyson_spheres.py`. Honest scope: the “Dyson sphere” is informational/holographic — self-containment, losslessness, boundary-encodes-bulk — not a literal energy-harvesting megastructure, and no step creates energy.)

C.23 Pass 5 — the capstone: machine = world, gravity to a number, and a fittable sky template

Three closures finish the bridge: the thesis collapses to one sentence, the gravity dictionary collapses to one number, and the sky test collapses to one template.

The machine is the world. The whole program is a single chain, each link a separate witness: $q = 3$ is *forced* (the crystallographic restriction caps the admissible Φ_n at $n \in \{3, 4, 6\}$, and primality selects $q = 3$); this gives the degree-2 cyclotomic skeleton $\{\Phi_3, \Phi_4, \Phi_6\} = \{13, 10, 7\}$ and the Witting object; which unfolds into the seven faces with *no free integer parameter* (even $1/\alpha = 137 = \Phi_3\Phi_4 + \Phi_6$ is cyclotomic, and the Leech kissing number $196560 = 6\mu q^2\Phi_3\Phi_4\Phi_6$); the matter shell is simultaneously the $[[240, 81, 4, 3]]_3$ code, the magic fuel (matter = magic, no distillation factory), and the clock-renewed resource — one self-fueling holographic memory; whose bulk has positive Ollivier curvature $\kappa = \Lambda = 1/6$, Newton constant $G = 1/2^q$, de Sitter entropy $S = f = 24$. The theorem: *from the single forced integer $q = 3$, the substrate is the self-fueling holographic memory whose bulk is the de Sitter universe it computes — machine and world are one object — with no free integer parameter, decided by two experiments.* The residue is purely dynamical (the QED running 0.036, the absolute mass scales). (Witness: `analysis/w33_machine_is_world.py`.)

Gravity, closed to a number. Given the holographic dictionary — horizon area $A = k = 12$ (the gauge causal screen of the $1+12+27$ split) and de Sitter entropy $S_{\text{dS}} = c = f = 24$ (the boundary central charge) — the Bekenstein–Hawking relation $S = A/4G$ inverts to a single dimensionless Newton constant

$$G = \frac{A}{4S} = \frac{k}{4f} = \frac{12}{96} = \frac{1}{8} = \frac{1}{2^q},$$

the inverse Hilbert dimension of the D_4 matter shell ($2^q = 8$): gravity couples as one over the states the horizon resolves. With $\Lambda = 2/k = 1/6$ (the computed Ollivier curvature) and the de Sitter radius $\ell = 1/\sqrt{\Lambda} = \sqrt{6}$, the three gravitational numbers are all q -integers and mutually consistent: $S = A/4G = 24 = f$, $\Lambda\ell^2 = 1$, $E\Lambda = 240/6 = 40 = v$ (the Gauss–Bonnet closure), $GS_{\text{dS}} = q$. Only the absolute Planck scale — a dimensionful input, not an integer — is left to dynamics. (Witness: `analysis/w33_gravity_dictionary.py`.)

The sky test, as a template. The clock’s CMB fingerprint (§11) sharpens from a qualitative three-gap comb to an explicit, pipeline-ready modulation of the primordial spectrum,

$$\frac{\delta P}{P}(\ln k) = A \cos(\omega_1 x + \phi_1) [1 + b \cos(\omega_2 x + \phi_2)], \quad x = \ln(k/k_*),$$

with *both* log-frequencies pinned by $q = 3$: $\omega_1 = \theta = \arccos(-\frac{2}{3}) = 2.3005$ (the Boerdijk–Coxeter twist) and $\omega_2 = 2\pi/\text{beat} = 2\pi/30 = 0.2094$ (the $h(E_8)$ ring envelope, the same beat as $1 - n_s = 1/30$). Their ratio $\omega_1/\omega_2 = 15\theta/\pi = 10.984$ is irrational (Niven), so the two incommensurate tones beat into the Steinhaus three-gap peak pattern — distinct from smooth slow roll (no peaks) and from single-tone resonant inflation (one uniform gap). Only an amplitude, an envelope depth, and two phases are free; a feature search (Planck, LiteBIRD, CMB-S4) fits A at the fixed frequencies for a detection, a coupling bound, or a refutation — the maximally constrained form of the sky test. (Witnesses: `analysis/w33_cmb_template.py`, `analysis/w33_gravity_dictionary.py`, `analysis/w33_machine_is_world.py`.) *Honest scope*: a consolidation — each link is a separately-witnessed result with its own scope (the forcings exact; the seven faces structural; the α closure at the integer level; the holographic/de Sitter and area/entropy identifications the substrate’s central geometric claim; the CMB feature conditional on the clock–inflaton coupling). Given those, machine = world, $G = 1/2^q$, and the two-tone template are forced.

C.24 Pass 6 — sharpening the capstone: a fit, a derived frequency, and a running G

Each capstone claim has a soft spot; this pass hardens three of them.

The sky template touches data with an error bar. Product-to-sum turns the modulation into an exact spectral *triplet*: $\delta P/P = A \cos(\omega_1 x + \phi_1) + \frac{Ab}{2} \cos((\omega_1 + \omega_2)x + \dots) + \frac{Ab}{2} \cos((\omega_1 - \omega_2)x + \dots)$ — three lines at $\{\omega_1 - \omega_2, \omega_1, \omega_1 + \omega_2\} = \{2.091, 2.301, 2.510\}$, a carrier split by the ring envelope (the frequency-domain face of the three-gap comb). Fitting the central line to a cosmic-variance-limited mock Planck-TT vector ($\ell = 30\text{--}2000$, $f_{\text{sky}} = 0.7$, $k = \ell/D_A$, transfer $T = 1$) gives a sensitivity $\sigma_A \simeq 1.3 \times 10^{-3}$, i.e. a 95% upper limit $A < 2.6 \times 10^{-3}$ in the optimistic case; an injected 1% modulation is recovered at 8σ , so the pipeline works. Realistic acoustic transfer (projection + Silk damping) weakens this by a factor of a few-to-ten, landing near the few-percent level current feature searches reach — but the substrate’s prediction is now a *constrainable amplitude at fixed, $q = 3$ -locked frequencies*, the first contact with data carrying an error bar. (Witness: `analysis/w33_cmb_template_forecast.py`.)

The slow frequency is derived, not asserted. The template’s $\omega_2 = 2\pi/30$ was an identification; it is now a theorem of the Boerdijk–Coxeter helix. Building the tetrahelix from regular tetrahedra by the recurrence $v_{n+4} = \text{reflect}(v_n, \text{plane}(v_{n+1}, v_{n+2}, v_{n+3}))$, the per-tetrahedron screw twist has cosine *exactly* $-\frac{2}{3}$, so the substrate clock angle $\theta = \arccos(-\frac{2}{3})$ is the BC-helix twist (the same $-(q-1)/q$). In flat space $\theta/2\pi$ is irrational, so the helix never closes — the fast tone $\omega_1 = \theta$, the time quasicrystal; in the curved 600-cell the same helix closes after exactly 30 tetrahedra (600 cells = 20 rings $\times$ 30, and $30 = h(E_8) = \Phi_3 + \Phi_4 + \Phi_6$), the slow tone $\omega_2 = 2\pi/30$. Both tones share one geometric origin, and the ratio $\omega_1/\omega_2 = 30\theta/2\pi = 15\theta/\pi = 10.984$ is forced and irrational (in 30 ticks the fast phase wraps a non-integer 10.984 times). Only the tick $\leftrightarrow$ e-fold normalisation remains an identification. (Witness: `analysis/w33_bc_helix_omega2.py`.)

Newton’s constant runs over the tower. $G = k/4f = 1/2^q$ was a single number; over the GKP shell ladder $A_2 < D_4 < E_8$ it is a curve. Treating each shell as a horizon whose microstate count is its kissing number ($= \# \text{roots} = h \cdot \text{rank}$: 6, 24, 240) and keeping the self-similar gauge-screen area $A = k = 12$, Bekenstein–Hawking gives

$$G_s = \frac{A}{4S_s} = \frac{k}{4 \text{kiss}_s} = \frac{q}{h_s \text{rank}_s} : \quad G(A_2) = \frac{1}{2}, \quad G(D_4) = \frac{1}{8} = \frac{1}{2^q}, \quad G(E_8) = \frac{1}{80},$$

exact on all three rungs. The matter shell D_4 sits at the physical $G = 1/2^q$; gravity weakens up the tower by $\mu = 4$ then $\Phi_4 = 10$. With the Coxeter number as scale, the seed A_2 ($h = 3$) is the UV (strong, Planckian) and the full E_8 gauge shell ($h = 30$) the IR (weak) — gravity strongest in the UV, irrelevant in the IR, as a gravitational coupling should run ($G_s \sim 1/h_s \text{rank}_s$). So the substrate predicts not just the value of G but its scale dependence. (Witness: `analysis/w33_newton_running.py`.) *Honest scope*: the forecast uses a simplified (cosmic-variance-limited, $T = 1$) noise model; the BC twist and 600-cell closure are computed geometry while the e-fold=tick coupling remains an assumption; the G -running is a discrete three-point trend on the established tower, with area= k and entropy=kissing the same holographic dictionary that fixes $G(D_4) = 1/2^q$ and the UV/IR-via- h reading interpretive. Given those, the bound, ω_2 , the ratio, and the running are forced.

C.25 Pass 7 — three edges pressed: a resolution reality, the last identification closed, and an honest hierarchy

Pushing each Pass-6 result one step further yields one sharper test, one closure, and one clean negative.

What the CMB can and cannot measure. The substrate signal is a rigid spectral triplet (carrier at ω_1 , sidebands at $\omega_1 \pm \omega_2$, amplitudes $A : b/2 : b/2$). A matched filter co-adds all three, but the Fisher gain over the carrier alone is only $\sqrt{1 + b^2/2} = 1.086$ (for $b = 0.6$) — the sidebands carry little power, so the CMB’s job is to *bound the amplitude*. The closure itself — the spacing $\omega_2 = 2\pi/30$, the direct test of the 600-cell — requires a log- k window at least one beat wide, $\Delta x \geq 2\pi/\omega_2 = \text{beat} = 30$ e-folds, whereas the CMB spans only $\Delta x \sim 4\text{--}7$; the triplet is unresolved there. Even adding μ -distortions ($k \sim 10^4$) and 21 cm reaches only $\Delta x \sim 18\text{--}20$. So measuring the closure number 30 is a target for a wide-lever joint analysis, not the CMB alone — an honest correction to “the spacing confirms closure.” (Witness: `analysis/w33_cmb_triplet_matched_filter.py`.)

The last identification is a corollary. The CMB template’s one remaining asserted input — that one inflationary e-fold = one Boerdijk–Coxeter tick — follows from the machine=world capstone. de Sitter space has a single timescale $1/H = \ell_{\text{dS}} = \sqrt{6}$ ($\Lambda = 1/6$); since the holographic memory’s bulk *is* that de Sitter spacetime, the substrate clock is the expansion and runs at H , so ticks = $Ht = N$ — exactly one tick per e-fold, because there is one clock. Then everything is forced with no free normalisation: the phase advances $\theta = \arccos(-\frac{2}{3})$ per e-fold (ω_1), the ring closes after 30 e-folds (ω_2), inflation lasts $N = 2\text{beat} = 60 = 2(v - \Phi_4)$ e-folds, and the tilt’s two readings agree, $2/N = 1/\text{beat} = 1/30$, precisely because $N = 2\text{beat}$. (The rate is H , the expansion, not the thermal $H/2\pi$ — the machine=world choice.) Only the coupling amplitude A stays free. (Witness: `analysis/w33_efold_tick.py`.)

The gravity hierarchy: exact at the tower, open at 10^{17} . Extending the Newton ladder to the Leech ceiling, the step factors are all substrate constants: $G(A_2)/G(D_4) = \mu = 4$, $G(D_4)/G(E_8) = \Phi_4 = 10$, $G(E_8)/G(\Lambda_{24}) = q^2\Phi_3\Phi_6 = 819$. So the cumulative spans are cyclotomic and exact: $G(A_2)/G(E_8) = \mu\Phi_4 = 40 = v$ (the gravity-to-gauge hierarchy) and $G(A_2)/G(\Lambda_{24}) = \mu q^2\Phi_3\Phi_4\Phi_6 = 32760 = \text{kiss}(\Lambda_{24})/\text{kiss}(A_2)$, tying back to the α -closure factorisation $196560 = 6\mu q^2\Phi_3\Phi_4\Phi_6$. But 3×10^4 is ~ 13 orders short of the Planck/electroweak 10^{17} : the discrete tower fixes the *geometric* gravity hierarchy ($v = 40$), not the absolute one. The missing orders need the dynamical scale (the named residue) or an exponential (e^N , $N = 60 \Rightarrow 10^{26}$, overshooting; 10^{17} needs $N \sim 39$) — neither lands without tuning. A genuine open edge, reported plainly. (Witness: `analysis/w33_gravity_hierarchy.py`.) *Honest scope:* the triplet gain and resolution thresholds are exact/standard-order; the e-fold=tick closure rests on machine=world plus de Sitter’s single timescale; the hierarchy result is exact arithmetic on the kissing numbers (positive) with a real 10^{17} shortfall (negative) left open.

C.26 Pass 8 — the hierarchy’s exponent, a designed closure test, and one clock for the whole sky

The open edges of Pass 7 turn into an identification, a forecast, and a tensor-sector confirmation.

The missing exponent is a substrate integer. Pass 7 found the Planck/electroweak hierarchy needs ~ 39 e-folds of exponential separation that the power-law tower does not supply. That integer is $q\Phi_3 = 3 \cdot 13 = 39$: writing the hierarchy in the e-fold currency, $\ln(M_{\text{Pl}}/M_{\text{EW}}) = q\Phi_3 = 39$, so $M_{\text{Pl}}/M_{\text{EW}} = e^{39} \approx 8.7 \times 10^{16}$ and the predicted electroweak scale is $M_{\text{Pl}}e^{-39} \approx 141 \text{ GeV}$ — among the W, Z , Higgs and top masses. It closes the e-fold budget exactly: inflation $N = 60 = 2\text{beat} = q(\Phi_3 + \Phi_6)$ splits into the Planck→EW descent $q\Phi_3 = 39$ and the remaining $q\Phi_6 = 21$ e-folds (with $\Phi_3 + \Phi_6 = 20 = 600/30$ BC rings). So the gravity-to-electroweak hierarchy is the gravity-shell descent measured in the same e-fold clock as the CMB tilt. *Honest:* an integer-level match/postdiction ($q\Phi_3 = 39$ brackets the observed $\ln$, 38.4 at v to 39.3 at $\sim 100 \text{ GeV}$, to $\sim 1\text{--}2\%$), not a derivation of why the EW scale sits there. (Witness: `analysis/w33_hierarchy_exponential.py`.)

The closure as a designed experiment. Measuring the number 30 (the triplet spacing ω_2) is a Fisher forecast: $\sigma(\text{beat})/\text{beat} = \text{beat}\sqrt{12}/(2\pi\rho\Delta x)$, with $\rho = A\sqrt{N_{\text{modes}}}/2$. For $A \sim 10^{-2}$, the CMB ($\Delta x \sim 7.6$) and CMB-S4 only bound the amplitude, while a dark-ages 21 cm interferometer ($k \sim 100$, $\Delta x \sim 14$, $N_{\text{modes}} \sim 10^{14}$) formally reaches $\sigma(\text{beat})/\text{beat} \ll 0.1$. But *no* probe spans a full beat ($\Delta x_{\text{max}} \sim 14 < 30$), so the closure is extrapolated from less than one envelope period — an extreme-frontier (lunar-farside/space 21 cm) target, not near-term. (Witness: `analysis/w33_closure_forecast.py`.)

One clock for the whole sky. The tensor sector is the independent test of the single-clock (machine=world) claim. All four inflationary observables are powers of the same beat = 30 with substrate-constant coefficients:

$$1 - n_s = \frac{1}{\text{beat}} = \frac{1}{30}, \quad r = \frac{1}{\Phi_4 \text{ beat}} = \frac{1}{300}, \quad n_t = \frac{-1}{2^q \Phi_4 \text{ beat}} = \frac{-1}{2400}, \quad \frac{dn_s}{d \ln k} = \frac{-1}{2 \text{ beat}^2} = \frac{-1}{1800} = -\frac{2}{N^2}.$$

The single-field consistency $r = -8n_t$ holds *identically* — because $2^q = 8$ — so scalars and tensors share the one de Sitter clock. This makes $r = 1/300 = 0.0033$ (below the current bound $r < 0.036$, a LiteBIRD/CMB-S4 target) and $n_t = -1/2400$ concrete predictions, with the falsification handle: an independently measured $n_t \neq -r/8$ would mean two clocks (multi-field), breaking the single-clock content of machine=world. (Witness: `analysis/w33_tensor_clock.py`.) *Honest scope:* the hierarchy exponent is a postdiction; the closure forecast is order-of-magnitude with a $\Delta x < \text{beat}$ extrapolation caveat; r and the running are prior results and $r = -8n_t$ is a standard identity — new is the one-clock packaging in beat = 30, the $2^q = 8$ making the consistency exact, and $n_t = -1/2400$ as a testable prediction.

C.27 Pass 9 — the breakthrough: the complete primordial spectrum, normalisation included

Three sharpenings and one break into the residue.

The hierarchy exponent, partly derived. Splitting the gravity→electroweak ladder at the corpus’s trinitification scale, the upper rung is a clean integer: $\ln(M_{\text{Pl}}/M_{\text{GUT}}) = \Phi_6 = 7$, i.e. $M_{\text{GUT}} = M_{\text{Pl}}e^{-\Phi_6} \approx 1.1 \times 10^{16} \text{ GeV}$ — deriving the GUT scale from gravity by a substrate integer. The total $\ln(M_{\text{Pl}}/M_{\text{EW}}) = q\Phi_3 = 39$ then makes the desert $q\Phi_3 - \Phi_6 = 32 = 2^q\mu$, and $39 = q^2 + \text{beat}$ — every rung an integer. (Witness: `analysis/w33_hierarchy_derivation.py`.)

A dated falsification line. The prediction $r = 1/300 = 0.0033$ is invisible now (0.26σ) but a 3.3σ detection for LiteBIRD (~ 2035 , $\sigma(r) \sim 0.001$) and $\sim 6.7\sigma$ for CMB-S4; LiteBIRD also separates $1/300$ from the nearest substrate alternatives ($0, 1/100, 1/1800, 1/3600$) at $\sim 3\text{--}7\sigma$. By ~ 2035 , r either sits at $1/300$ or is bounded < 0.0013 , refuting it. (Witness: `analysis/w33_r_falsification.py`.)

The clock is over-determined. Beyond $r = -8n_t$ (relation I, $2^q = 8$), the running obeys a second slow-roll identity with no new input: $1 - n_s = 2/N$ on $N = 2 \text{ beat} = 60$ forces $dn_s/d \ln k = -2/N^2 = -(1 - n_s)^2/2 = -1/1800$ (relation II). Four observables from two inputs (beat, Φ_4) linked by two consistency relations — an over-determined, falsifiable system. (Witness: `analysis/w33_overdetermined_clock.py`.)

The breakthrough: the amplitude is a substrate integer. The one pure number that sets the absolute size of all cosmic structure — the heart of the long-standing “absolute scale” residue — is the scalar amplitude, and it too is an integer in the exponent:

$$\ln(1/A_s) = \Phi_3 + \Phi_6 = 20 \implies A_s = e^{-20} = 2.06 \times 10^{-9},$$

versus Planck’s 2.10×10^{-9} ($\ln(10^{10}A_s)$: substrate 3.026 vs observed 3.044 ± 0.014 , a 1.3σ agreement), with $20 = \Phi_3 + \Phi_6 = N/q = v/2 = 600/\text{beat}$ (the Boerdijk–Coxeter ring count). So the

complete primordial spectrum is fixed by substrate integers $\{q, \text{beat}, \Phi_4, \Phi_3 + \Phi_6\}$ with no continuous parameter:

$$A_s = e^{-20}, \quad 1 - n_s = \frac{1}{30}, \quad r = \frac{1}{300}, \quad n_t = -\frac{1}{2400}, \quad \frac{dn_s}{d \ln k} = -\frac{1}{1800}.$$

And the amplitude pins the inflationary energy scale, $V^{1/4} = (24\pi^2 \epsilon A_s)^{1/4} M_{\text{Pl}} \approx 3.9 \times 10^{16} \text{ GeV} \approx M_{\text{GUT}} = M_{\text{Pl}} e^{-\Phi_6}$, so inflation happens at the GUT scale and the inflaton's energy leaves the residue. (Witness: `analysis/w33_complete_primordial_spectrum.py`.) *Honest scope:* $A_s = e^{-20}$ is an integer-level match (20 reproduces $\ln(1/A_s) = 19.98$ to $< 0.1\%$, 1.3σ on $\ln(10^{10} A_s)$); the deep reason for the exponent — plausibly the de Sitter horizon mode count/entropy or the e-folds-per- q -sector N/q — is conjectural and the next derivation target. But the last major dimensionless cosmological number is now a substrate integer, completing the primordial spectrum and denting the scale residue: a real break, flagged as a match whose dynamical origin is open.

C.28 Pass 10 — the scale residue closed to one unit: amplitude as a law, all masses as integers, two exponents as one

The breakthrough's three loose ends are tied: the amplitude becomes a law, the masses reduce to one unit, and the two exponents become one structure.

The amplitude is a law, not a coincidence. Exactly, $A_s = 1/(\epsilon S_{\text{dS}})$: the scalar amplitude is the inverse of the slow-roll ϵ times the de Sitter horizon entropy ($A_s = H^2/8\pi^2 \epsilon M_{\text{Pl}}^2$, $S_{\text{dS}} = 8\pi^2 (M_{\text{Pl}}/H)^2$), so $A_s \sim 10^{-9}$ is the inverse of a horizon entropy $S_{\text{dS}} \sim 10^{12}$ nats times $\epsilon \sim 1/4800$ — a derivation of the smallness. And the substrate value is the clean law $(A_s)^q = e^{-N}$: the q -th power of the amplitude equals the inverse total expansion e^{-60} ($q = 3$ the spatial dimensions/sectors), i.e. $A_s = e^{-N/q} = e^{-20}$. These fix $H_{\text{inf}} \sim 1.4 \times 10^{13} \text{ GeV}$ and $V^{1/4} \sim 10^{16} \text{ GeV}$ at the GUT scale. (Witness: `analysis/w33_amplitude_entropy.py`.)

All masses reduce to one unit. A dimensionful number cannot come from integers — a unit must be chosen — so the achievement is reducing *every* scale to one unit by substrate-integer exponents: $M_{\text{GUT}} = M_{\text{Pl}} e^{-\Phi_6}$, $M_{\text{EW}} = M_{\text{Pl}} e^{-q\Phi_3}$, $m_p \sim M_{\text{Pl}} e^{-(v+\mu)}$ ($\ln(M_{\text{Pl}}/m_p) = 44.0 \approx v + \mu = 44$). Choosing any one scale as the unit fixes all others by integers, so the theory has *one* dimensionful input (a units choice) and *zero* free dimensionless parameters. Gravity's weakness follows: $\alpha_G = (m_p/M_{\text{Pl}})^2 = e^{-2(v+\mu)} = e^{-88} = 6 \times 10^{-39}$. (Witness: `analysis/w33_scale_reduction.py`.)

The two exponents are one structure. The amplitude exponent (20) and the hierarchy exponent (39) are both partitions of the single inflationary e-fold count $N = 60 = 2 \text{ beat}$: $N = q(N/q) = 3 \cdot 20$ (amplitude) and $N = q\Phi_3 + q\Phi_6 = 39 + 21$ (hierarchy), locked by $39 = N - q\Phi_6 = q \cdot 20 - q\Phi_6$. So one GUT-scale inflaton rolling $N = 60$ e-folds fixes the whole exponential layer at once: its energy scale gives the hierarchy, its e-fold flow gives the normalisation $A_s = e^{-N/q}$, its clock gives the tilts. (Witness: `analysis/w33_exponent_unification.py`.) *Honest scope:* $A_s = 1/(\epsilon S_{\text{dS}})$ and $(A_s)^q = e^{-N}$ are exact/structural with the horizon entropy's value the remaining input; the GUT (7) and EW (39) exponents are clean while the proton $44 = v + \mu$ admits several substrate forms (flagged) and ultimately traces to QCD running; the $\{20, 39\}$ unification is exact partition arithmetic with the explicit $V(\phi)$ the next step. The robust result: the absolute-scale residue is now exactly one units choice, with zero free dimensionless parameters.

C.29 Pass 11 — the potential identified (Starobinsky R^2), the entropy derived, the tower audited

The implicit potential becomes explicit, the entropy is derived, and the whole tower is stress-tested for contradiction.

The inflaton is the Starobinsky scalaron. The substrate is Starobinsky (R^2) inflation with $N = 2 \text{ beat} = 60$: all four observables coincide *exactly* with the model at $N = 60$ — $1 - n_s = 2/N = 1/30$, $r = 12/N^2 = 1/300$, $n_t = -3/2N^2 = -1/2400$, $dn_s/d \ln k = -2/N^2 = -1/1800$ — the cyclotomic forms equalling the Starobinsky values because $\text{beat} = q\Phi_4$ and $N = 2 \text{ beat}$ give $\Phi_4 \text{ beat} = N^2/12 = 300$. The potential is the plateau $V = V_0(1 - e^{-\sqrt{2/3}\phi/M_{\text{Pl}}})^2$ with a *super-Planckian* excursion $\Delta\phi \sim 5.4 M_{\text{Pl}}$ (excluding small-field models) and $V_0^{1/4} \sim 10^{16} \text{ GeV}$. Crucially, Starobinsky inflation is $f(R) = R + R^2/6M^2$ *gravity* — the inflaton is the scalaron of curvature itself — so the substrate’s inflaton is geometric, exactly machine = world. (Witness: `analysis/w33_starobinsky.py`.)

The horizon entropy, derived (with an honest negative). Granting $A_s = e^{-N/q}$ and the Starobinsky $\epsilon = q/4N^2$, the exact law $A_s = 1/(\epsilon S_{\text{dS}})$ derives both $S_{\text{dS}} = (4N^2/q)e^{N/q} \sim 2.3 \times 10^{12}$ and $H_{\text{inf}}/M_{\text{Pl}} = (\pi\sqrt{2q}/N)e^{-N/2q} \sim 5.8 \times 10^{-6}$ ($H_{\text{inf}} \sim 1.4 \times 10^{13} \text{ GeV}$) purely from N, q . But the holographic shortcut — identifying this bulk entropy with the boundary central charge $f = 24$ — *fails* ($S_{\text{dS}}/f \sim 10^{11}$, not a clean integer); the exponent traces instead to the GUT-scale inflation $V^{1/4} \sim M_{\text{Pl}}e^{-\Phi_6}$. (Witness: `analysis/w33_horizon_entropy_derivation.py`.)

The tower is internally consistent. An adversarial audit computes nine over-determined quantities (beat, N , $1 - n_s$, r , n_t , running, ϵ , the Planck/EW exponent, the A_s exponent) by 2–4 independent routes each: all agree, no internal contradiction — the agreements non-trivial (e.g. $r = 1/(\Phi_4 \text{ beat})$ equals $12/N^2$ only because $\Phi_4 \text{ beat} = N^2/12$). The tightest *external* point is $n_s = 0.9667$ versus Planck 0.9649 ± 0.0042 (0.42σ), the sharpest near-term discriminator (CMB-S4 tests it at $\sim 1\sigma$). (Witness: `analysis/w33_consistency_audit.py`.) *Honest scope:* the four-observable equality with Starobinsky-at- $N = 60$ is exact, and $f(R) = R^2$ is the natural action-level reading (motivated, not derived from the substrate action); the entropy reduction is exact with the $e^{N/q}$ origin tied to GUT-scale inflation (holographic $f = 24$ route a clear negative); the audit certifies self-consistency (not truth — the integers could be coincidence), so a single off measurement breaks it.

C.30 Pass 12 — the residue as one number, the scalaron meets the neutrinos, and the tower as one plot

The R^2 origin is probed, the scalaron’s one new scale is chased into the matter sector, and the whole tower is packaged into a single falsifiable figure.

The residue is one number (and the anomaly does not fix it). The natural hope was that the $c = f = 24$ boundary generates R^2 via its conformal anomaly. It does not: the anomaly induces an R^2 coefficient of order $c/2880\pi^2 \sim 10^{-3}$, whereas Starobinsky needs $M_{\text{Pl}}^2/12M^2 \sim 6 \times 10^8$ (the scalaron mass $M \sim 2.8 \times 10^{13} \text{ GeV}$ from A_s) — about twelve orders too weak. So machine = world is not promoted to a derivation through the anomaly. The positive content: the R^2 coefficient, the scalaron mass M , and the A_s exponent are the *same* single number ($A_s = N^2 M^2 / 24\pi^2 M_{\text{Pl}}^2$), so the entire residual freedom of the cosmological tower is one quantity. (Witness: `analysis/w33_r2_anomaly.py`.)

The scalaron meets the neutrinos. That one new scale is $M \sim 2.8 \times 10^{13} \text{ GeV}$, with $\ln(M_{\text{Pl}}/M) \sim 13 \sim \Phi_3$ (full M_{Pl}). It lands in the type-I seesaw window ($M_R \sim 10^{14} - 10^{15} \text{ GeV}$ for the corpus’s $\sum m_\nu \sim 0.1 \text{ eV}$), so the Starobinsky scalaron can reheat by decaying into right-handed neutrinos: inflation feeds the light-neutrino masses (seesaw) and the baryon asymmetry (leptogenesis). One scale ties the end of inflation, neutrino masses, and matter genesis. (Witness: `analysis/w33_scalaron_seesaw.py`.)

The tower as one plot. The substrate is the point $(n_s, r) = (0.9667, 1/300)$ — the $N = 2 \text{ beat} = 60$ point on the Starobinsky line $r = 3(1 - n_s)^2$ (Fig. 15). Today it sits 0.42σ from Planck and

below the r bound; a CMB-S4/LiteBIRD measurement ($\sigma(n_s) \sim 0.002$, $\sigma(r) \sim 0.0005$) detects $r = 1/300$ at 6.7σ , pins the point, and measures $N = 60 \pm 3.6$ (testing $N = 2$ beat against $N = 50$ at 2.8σ). (Witness: `analysis/w33_ns_r_forecast.py`.)




Figure 15: The cosmological tower as one point in the (n_s, r) plane. The substrate predicts the $N = 2$ beat = 60 point (red) on the Starobinsky line $r = 3(1 - n_s)^2$ (blue): $(n_s, r) = (0.9667, 1/300)$. It sits 0.42σ from Planck's n_s (grey band) and below the current r bound. A CMB-S4/LiteBIRD measurement ($\sigma(n_s) \sim 0.002$, $\sigma(r) \sim 0.0005$; red ellipse) detects $r = 1/300$ at 6.7σ , pins the point, and measures $N = 60 \pm 3.6$ – testing $N = 2$ beat against $N = 50$ at 2.8σ .

Honest scope: the anomaly result is a genuine negative (the 12-order gap robust to scheme factors) with the residue-unification exact; the scalaron mass is fixed, but $\ln(M_{\text{Pl}}/M) \sim \Phi_3$ is convention-dependent and $M \sim M_R$ is an order-of-magnitude coincidence (a mechanism, not yet a precise lock); the (n_s, r) point and line-membership are exact, the forecast sensitivities published. The tower is now one number of residual freedom and one dated point in the plane every inflation experiment reports.

C.31 Pass 13 — one particle for inflation and matter, the baryons, and the master ledger

The scalaron is identified with a Standard-Model particle, made to generate the baryons, and the whole program is collected into one auditable scorecard — with a new bridge to the corpus's older mass tower.

The scalaron is the lightest right-handed neutrino. The scalaron mass $M \sim 2.8 \times 10^{13}$ GeV sits at the light end of the substrate's seesaw RHN spectrum: $\ln(M_{\text{Pl}}/M) \sim \Phi_3 = 13$ (the lightest N_1) and $\ln(M_{\text{Pl}}/M_{R3}) \sim q^2 = 9$ (the heaviest $N_3 = v^2/m_{\nu 3} \sim 10^{15}$ GeV) — the cyclotomic skeleton bracketing the seesaw. Identifying the scalaron with N_1 clears Davidson–Ibarra ($M_1 > 10^9$ GeV) by eight orders, so inflation, reheating, the seesaw, and leptogenesis are *one* $\sim 10^{13}$ GeV sector driven by one particle. (Witness: `analysis/w33_scalaron_is_rhn.py`.)

The inflaton makes the baryons. Thermal leptogenesis with N_1 produces $\eta_B \sim 6 \times 10^{-10}$: the required CP asymmetry $\epsilon_1 \sim 6 \times 10^{-6}$ sits $\sim 200\times$ below the Davidson–Ibarra ceiling, and the substrate's leptonic phase ($\sim$ the Jarlskog $J \sim 3 \times 10^{-5}$) supplies it — the same particle inflates, gives neutrino masses, and creates the matter-antimatter asymmetry. (Witness: `analysis/w33_baryon_asymmetry.py`.)

The master ledger and the tier bridge. Eighteen predictions — the Starobinsky $N = 60$ tower (A_s, n_s, r, n_t , running, $f_{NL} = 1/72$) and the particle sector ($1/\alpha = 137$, Jarlskog 3×10^{-5} , $\sin^2 \theta_W = 0.231$, $\sum m_\nu = 0.101$ eV, $\Omega_{\text{DM}}/\Omega_b = 82/15$, $M_{\text{Pl}}/M_{\text{EW}} = e^{39}$, the scalaron, η_B , the contextual fraction $1/10$, the proton) — are tallied: thirteen CONSISTENT, two TEST-SOON (r /LiteBIRD, proton/Hyper-K), one TEST-FUTURE (n_t), one LAB, and exactly one TENSION ($\sum m_\nu$ vs DESI 2024 < 0.072 eV). And a novel bridge: the corpus's mass-tier tower ratio $r = q^q/(l^\mu F_5) = 27/80$ is exactly $r = q^3/(\lambda v)$, built from the SRG(40, 12, 2, 4) invariants of the W(3, 3) collinearity graph (the 2 is λ , the common-neighbour count, with $l^\mu F_5 = \lambda v = 80$) — so the multiplicative tier tower and the additive cyclotomic skeleton are both W(3, 3), each tier $\ln(\lambda v/q^3) = 1.086$ e-folds. (Witness:

analysis/w33_falsification_scorecard.py.) *Honest scope:* the scalaron $= N_1$ is a consistent identification (passes Davidson–Ibarra), η_B an order-of-magnitude estimate (exact value needs the Yukawa texture), the ledger a faithful audit (reporting the DESI tension and lab-only entries), and the tier bridge an exact identity with the integer alignment of the two mass quantizations ($q\Phi_3 = 39 = 35.9$ tiers) left open.

C.32 Pass 14 — the tension resolved, three hierarchies reconciled, and dark matter completed

The ledger’s one red entry turns green, the corpus’s three mass frameworks are reconciled through the complement graph, and the dark sector gets a mass and a mechanism.

The neutrino tension dissolves. The substrate predicts *normal* ordering with the splitting ratio $\Delta m_{31}^2/\Delta m_{21}^2 = 2\Phi_3 + \Phi_6 = 33$ (observed 32.6, 0.8%) and a near-massless lightest state ($m_1 \sim 0$), giving the NH-minimum $\sum m_\nu = \sqrt{\Delta m_{21}^2} + \sqrt{\Delta m_{31}^2} = 8.7 + 49.8 = 58$ meV — comfortably below the DESI 2024 bound (< 72 meV). The Pass-13 ledger’s 101 meV was a heavier, non-minimal estimate; 58 meV is the canonical value and it passes. Falsifiable three ways: $\sum m_\nu < 58$ meV, an inverted ordering, or a JUNO ratio $\neq 33$. (Witness: analysis/w33_neutrino_nh_minimum.py.)

Three hierarchies, one graph. The Planck–electroweak gap is written three ways: canonical base-10 ($\log_{10}(M_{\text{GUT}}/M_{\text{EW}}) = 2\Phi_6 = 14$, $M_{\text{Pl}}/M_{\text{GUT}} = 496 = \dim SO(32)$), cyclotomic e-fold ($\ln(M_{\text{Pl}}/M_{\text{EW}}) = q\Phi_3 = 39$), and the multiplicative tier tower — all giving the same ~ 38 –39 e-fold gap (38.4 vs 39, 1.4%). The tier base is now a complement-graph invariant: $r = q^q/(\lambda v) = \bar{k}/(\lambda v) = 27/80$, where $q^q = 27 = \bar{k} = v - k - 1$ is the degree of the complement SRG(40, 27, 18, 18) (the E_6 fundamental; $\bar{\lambda} = \bar{\mu} = 18 = h(E_7)$) and $\lambda v = 80 = l^\mu F_5$. The cross-base integer alignment stays open (39 e-folds = 35.4 tiers). (Witness: analysis/w33_hierarchy_three_frameworks.py.)

Dark matter: a mass and a mechanism. The dark-matter mass is $m_{\text{DM}} = M_Z/\mu = \Phi_3\Phi_6/\mu = 91/4 = 22.8$ GeV, a Z-funnel WIMP at one quarter of the Z mass; the candidate is the dark fermion of the E_6 **27** shell (the 8-dimensional eigenspace at -1 , with $q^2 = 9$ dark families). Z-portal freeze-out reproduces the substrate density $\Omega_{\text{DM}} = \mu/g = 4/15$, i.e. $\Omega_{\text{DM}}h^2 = 0.12$ (Planck 0.120) and $\Omega_{\text{DM}}/\Omega_b = 82/15$. It is testable by LZ/XENONnT (~ 2028). (Witness: analysis/w33_dark_matter.py.) *Honest scope:* the 58 meV follows from the cyclotomic ratio plus normal ordering and $m_1 \sim 0$; the three-hierarchy reconciliation is a 1.4% numerical agreement with the complement-graph grounding of r the genuine new identity (cross-base alignment open); $\Omega_{\text{DM}} = \mu/g$ and $m_{\text{DM}} = M_Z/\mu$ are substrate formulas with the relic abundance matched at the density-fraction level, and a 22.8 GeV Z-coupled WIMP requires a *suppressed* Z coupling to satisfy both the relic density and the invisible- Z width (standard for Z-funnel models) — a flagged constraint, not yet a full freeze-out calculation.

C.33 Pass 15 — the dark matter stress-tested, the lightest neutrino, and the mass ladder

Two cross-checks sharpen earlier claims (one with an honest correction), and the whole mass spectrum is collected into one figure.

The dark matter, stress-tested. $m_{\text{DM}} = M_Z/\mu = 22.8$ GeV slots into the mass web ($\sim m_{\text{charm}} h(E_7)$; $\sim v + 1$ e-folds below M_{Pl}), collider-invisible between b and top. *Correction:* it is *off* the Z resonance ($2m_{\text{DM}} = M_Z$ needs $m_{\text{DM}} = M_Z/2 = 45.6$ GeV), so the Pass-14 “Z-funnel” is a generic off-resonance Z-portal. The LEP invisible width ($N_\nu = 2.984 \pm 0.008$) caps the DM– Z coupling at $g_{\text{DM}}/g_\nu \lesssim 1/\mu$ (a substrate ratio — a mostly-singlet state), and the relic density

then comes from the $q^2 = 9$ dark families' co-annihilation rather than Z-portal alone. (Witness: `analysis/w33_dark_matter_constraints.py`.)

The lightest neutrino is light, not massless. Does the $\mathbb{Z}_3$ grade rule force $m_1 = 0$? *No*: with a grade-0 B–L VEV the selection rule $a + b \equiv 0$ populates the “0/(1,2)-swap” texture in *both* m_D and M_R , both rank 3, so the seesaw is rank 3 and $m_1 \neq 0$. What is forced is $m_1 = m_2$ (the cyclotomic degeneracy) at the light-pair scale $(y_1/y_3)^2 m_3 \sim 1.2 \text{ meV}$; lifting it (a grade-1/2 VEV, Pillar 68/69) leaves $m_1 \sim 1 \text{ meV}$. So $(m_1, m_2, m_3) \sim (1, 8.7, 49.8) \text{ meV}$, $\sum m_\nu \sim 59 \text{ meV}$ — still below DESI; Pass-14's $m_1 \sim 0$ sharpened to a meV-scale lightest neutrino. (Witness: `analysis/w33_neutrino_lightest_mass.py`.)

The mass ladder. Every scale sits at a substrate-integer depth $\ln(M_{\text{Pl}}/M)$ (Fig. 16): M_{GUT} at $\Phi_6 = 7$, N_3 at $q^2 = 9$, the scalaron N_1 at $\Phi_3 = 13$, M_Z at $q\Phi_3 = 39$, $m_{\text{DM}} = M_Z/\mu$ just below, the proton at $v + \mu = 44$, the neutrinos at the seesaw floor ~ 68 — the gaps themselves cyclotomic ($\Phi_6, 2\Phi_3, \sim f$). (Witness: `analysis/w33_mass_ladder.py`.)




Figure 16: The mass spectrum of the universe as a cyclotomic descent from the Planck scale. Each scale sits at a substrate-integer depth $\ln(M_{\text{Pl}}/M)$: M_{GUT} at $\Phi_6 = 7$, the heaviest right-handed neutrino N_3 at $q^2 = 9$, the inflaton scalaron N_1 at $\Phi_3 = 13$, the electroweak scale at $q\Phi_3 = 39$, the dark matter M_Z/μ just below, the proton at $v + \mu = 44$, and the light neutrinos at the seesaw floor ~ 68 . The gaps between rungs are themselves cyclotomic ($\Phi_6, 2\Phi_3, \sim f$). One axis, one descent, the whole mass spectrum from $q = 3$.

Honest scope: the DM mass-web ratios are exact-cyclotomic, the off-resonance a correction, $g \lesssim 1/\mu$ a standard LEP bound on a substrate ratio, the dark-family relic qualitative; the neutrino result is an honest negative (m_1 not forced to zero) refined to $m_1 \sim 1 \text{ meV}$ (the lifting needs the grade-1/2 VEV magnitude); the ladder is a summary of the Pass 8–14 exponents, each with its own scope, the softest rungs the DM (~ 40.8 , non-integer) and the neutrino floor (~ 68).

C.34 Pass 16 — the lightest neutrino pinned, the $0\nu\beta\beta$ mass, and the descent to dark energy

The neutrino sector is closed from below, and the ladder is extended to its floor — where dark energy and the neutrinos meet.

The lightest neutrino, pinned. Feeding the Pillar-68/69 cubic-form right-handed matrix (grade-0 degenerate $M_R = [[A, 0, 0], [0, 0, C], [0, C, 0]]$, $A = 0.0017$, $C = 0.0442$, lifted by a grade-1 B–L component) through the seesaw with the up-type Dirac texture pins the lightest neutrino at $m_1 \sim 2 \text{ meV}$ (light, not zero) and $\sum m_\nu \sim 56\text{--}58 \text{ meV}$ — and m_1 lands at the same meV scale as $m_{\beta\beta}$ and the dark-energy scale. (Witness: `analysis/w33_neutrino_lightest_pinned.py`.)

The $0\nu\beta\beta$ mass. With the cyclotomic PMNS angles, the $\mathbb{Z}_3$ Majorana phases ($120^\circ, 240^\circ$), and the pinned $m_1 \sim 2 \text{ meV}$, $m_{\beta\beta} = |c_{12}^2 c_{13}^2 m_1 + s_{12}^2 c_{13}^2 m_2 e^{i120^\circ} + s_{13}^2 m_3 e^{i240^\circ}| \sim 1.4 \text{ meV}$ — the nonzero m_1 pushing it down (phase cancellation) from the $m_1=0$ value 2.3 meV into the normal-ordering trough. A far-future nEXO/LEGEND target; $m_{\beta\beta} > 10 \text{ meV}$ would falsify. (Witness: `analysis/w33_betabeta_refined.py`.)

The descent to dark energy. Extending the ladder, the dark-energy scale $\rho_\Lambda^{1/4} \approx 2.24 \text{ meV}$ sits ~ 71 e-folds below M_{Pl} — the same meV floor as the lightest neutrino ($\sim 2 \text{ meV}$) and $m_{\beta\beta}$ ($\sim 1.4 \text{ meV}$), all within a factor ~ 2 (Fig. 17). The vacuum-energy hierarchy $\log_{10}(\rho_\Lambda/M_{\text{Pl}}^4) \sim -123$ carries $vq = 120$ as its leading integer. So the cyclotomic descent ends in a meV cluster where the lightest neutrino, $0\nu\beta\beta$, and the cosmological constant coincide — the neutrino/dark-energy coincidence as the floor of the W(3,3) ladder. (Witness: `analysis/w33_cc_floor.py`.)

Honest scope: the cubic-form M_R is exact but the lift is the open B–L direction, so $m_1 \sim 1\text{--}3 \text{ meV}$ is a range (central ~ 2); $m_{\beta\beta} \sim 1.3\text{--}2.3 \text{ meV}$ tracks it; $vq = 120$ is the *leading* integer of the CC density (a ~ 3 -order residual — the CC problem not fully closed), the ~ 71 e-fold placement exact, and the meV-floor coincidence genuine but within a factor ~ 2 , not a derived equality.

C.35 Pass 17 — the cosmological constant exact, neutrino dark energy, and the final ledger

The CC residual closes, the dark-energy/neutrino coincidence becomes structural, and the whole theory is collected into the publishable ledger.

The cosmological constant is $-vq$, exact. The Pass-16 ~ 3 -order residual was the full-versus-reduced Planck-mass choice. With the *reduced* Planck mass ($M_{\text{Pl}}^2 = 1/8\pi G$, the one in General Relativity), $\log_{10}(\rho_\Lambda/M_{\text{Pl}}^4) = -120.15$, so to 0.1% in the log it is exactly $-vq = -120 = \dim(\text{adj } SO(16)) = 4 \text{ beat}$ — the famous “120 orders of magnitude” *is* the substrate point count vq . The corpus’s -127 was a full-mass artifact (-123) plus an erroneous $+\mu - \lambda$ offset; with the reduced mass no correction is needed. The residual 0.15 is the $O(1)$ (Ω_Λ, H_0) cosmology. (Witness: `analysis/w33_cc_exact.py`.)

Neutrino dark energy. Since $vq = 4 \text{ beat}$, the dark-energy *scale* is $\rho_\Lambda^{1/4} = M_{\text{Pl}} 10^{-\text{beat}} = M_{\text{Pl}} 10^{-30} = 2.44 \text{ meV}$ — exactly $\text{beat} = 30 = h(E_8)$ decades below the Planck scale, the same floor as the lightest neutrino $m_1 \sim 2 \text{ meV}$. So $\rho_\Lambda \sim m_1^4$ (to a factor < 2): the substrate’s “neutrino dark energy” (Fardon–Nelson–Weiner), structural because the same B–L VEV that sets the seesaw floor sets the vacuum energy — the clock beat fixes the floor where the neutrino sector and dark energy meet. (Witness: `analysis/w33_neutrino_dark_energy.py`.)

The final ledger. Twenty-five predictions — the complete primordial spectrum (Starobinsky $N = 60$), the gauge/CP sector, the mass ladder from M_{Pl} to the cosmological constant, the neutrinos (now resolved), the dark sector, and the vacuum ($\text{CC} = -vq$) — tally to twenty CONSISTENT, two TEST-SOON ($r/\text{LiteBIRD}$, m_{DM}/LZ), two TEST-FUTURE, one LAB, with *zero* tensions and




Figure 17: The full cyclotomic descent from the Planck scale to the cosmological constant. Every rung is a substrate-integer number of e-folds below M_{Pl} : M_{GUT} at Φ_6 , the scalaron N_1 at Φ_3 , the electroweak scale at $q\Phi_3$, the proton at $v + \mu$, the heavy neutrino m_{ν_3} at the seesaw floor. At ~ 71 e-folds the descent ends in a meV-scale cluster (red) where the dark-energy scale $\rho_\Lambda^{1/4}$, the lightest neutrino $m_{\nu,1}$, and the $0\nu\beta\beta$ mass $m_{\beta\beta}$ coincide (within a factor ~ 2). The vacuum-energy hierarchy $\log_{10}(\rho_\Lambda/M_{\text{Pl}}^4) \sim -123$ carries $vq = 120$ as its leading integer. One ladder, the Planck scale to dark energy.

zero falsified, and *zero* free dimensionless parameters. The decisive near-term test is $r = 1/300$ at LiteBIRD (~ 2035). (Witness: `analysis/w33_final_scorecard.py`.) *Honest scope*: $-vq = -120$ is an integer-level postdiction (CC measured, vq matches to 0.1% with the reduced mass; the 0.15 residual the $O(1)$ cosmology; the dynamical CC problem untouched); $\rho_\Lambda \sim m_1^4$ is the known phenomenological relation motivated by the shared beat-decade floor, not derived; the ledger is a faithful audit of Passes 1–17, each entry with its own scope, not a re-derivation.

C.36 Pass 18 — a mechanism for the cosmological constant, the joint sky test, and the standalone distillation

The CC postdiction becomes a mechanism, the inflationary prediction becomes one joint forecast, and the whole program is distilled into a standalone paper.

A mechanism for the cosmological constant. The substrate’s Hodge spectrum has an exact boson–fermion balance: the gauge sector ($f = 24$ modes at gap $\Phi_4 = k - r = 10$) and the matter sector ($g = 15$ at gap $\mu^2 = k - s = 16$) carry equal vacuum energy, $f\Phi_4 = g\mu^2 = 240 = |\text{Roots}(E_8)|$, so the leading M_{Pl}^4 vacuum energy *cancels* (structural supersymmetry, $f(k - r) = g(k - s)$). The balance breaks only at the mass-ladder floor $M_{\text{SUSY}} \sim M_{\text{Pl}} 10^{-\text{beat}} = 2.4 \text{ meV}$ (the neutrino/dark-energy scale), so $\rho_\Lambda = M_{\text{SUSY}}^4 = M_{\text{Pl}}^4 10^{-4 \text{ beat}} = M_{\text{Pl}}^4 10^{-vq}$: the 120 orders are 4 beat (TeV-SUSY

would give 10^{-64} , far too large). Holographically $\rho_\Lambda/M_{\text{Pl}}^4 = 1/S_{\text{ds}}$ with $S_{\text{ds}} = 10^{vq}$. (Witness: `analysis/w33_cc_mechanism.py`.)

The joint sky test. The substrate fixes one point $(\ln 10^{10} A_s, n_s, r) = (3.026, 0.9667, 1/300)$ on the Starobinsky line. The joint Planck+CMB-S4+LiteBIRD forecast: today 1.4σ (consistent), with $A_s = e^{-20}$ the tightest single constraint (1.3σ); by ~ 2035 , $r = 1/300$ is a 6.7σ detection pinning the point, and A_s becomes the most exposed ($\sim 1.8\sigma$) — a measurement off any of $\{r = 1/300, \text{the Starobinsky line}, A_s = e^{-20}\}$ falsifies. (Witness: `analysis/w33_inflation_joint_forecast.py`.)

The standalone distillation. The eighteen passes are distilled into a self-contained manuscript, `w33_everything.tex` (“Physics from One Integer”): the $q = 3$ selection, the cyclotomic skeleton, the couplings and mixing, the Starobinsky primordial spectrum, the mass ladder from M_{Pl} to the cosmological constant, the dark/baryon sector, and the falsification ledger (25 predictions, zero free dimensionless parameters, zero falsified) — the publishable capstone, compiling independently of this engineering paper.

Honest scope: the boson–fermion balance $f\Phi_4 = g\mu^2 = 240$ is an exact SRG theorem (the leading cancellation real); the breaking-at-beat-decades is a scale-level mechanism that explains the 120 as 4 beat *given* the meV floor (the neutrino-floor input), reframing the CC problem rather than solving it from nothing; the joint forecast uses published CMB-S4/LiteBIRD sensitivities; the standalone paper is a summary, not new derivations.

C.37 Pass 19 — the floor derived from the 600-cell, and the dated decision

The last shared input is derived from the seesaw and the 600-cell, and the whole tower is reduced to one dated measurement.

The beat-decade floor, derived. The meV floor that the CC mechanism (Pass 18) and the neutrino sector share is not an input: it is the type-I seesaw combination of the substrate’s own exponents. The lightest neutrino is $m_1 = y_1^2 v^2 / M_R$, so

$$\ln(M_{\text{Pl}}/m_1) = \ln(M_{\text{Pl}}/v) + \ln(M_R/v) - 2 \ln y_1 = q\Phi_3 + 2\Phi_3 + \Phi_6 = (q+2)\Phi_3 + \Phi_6 = 72,$$

the electroweak descent ($q\Phi_3$), the right-handed scale ($2\Phi_3$ above v , from $M_R = M_{\text{Pl}}e^{-\Phi_3}$, the scalaron), and a neutrino Dirac Yukawa $y_1 = e^{-\Phi_6/2} \approx 0.03$. This ~ 72 e-folds equals beat decades ($30 \ln 10 = 69$) to $\sim 4\%$ — and beat $= 30 = h(E_8)$ is the Boerdijk–Coxeter helix ring length on the 600-cell ($600 = 20 \times 30$). So the same 600-cell ring that sets the inflationary clock sets the floor; the CC’s breaking scale is closed up to the one Yukawa y_1 . (Witness: `analysis/w33_floor_derivation.py`.)

The dated decision. The whole cosmological tower reduces to one number by ~ 2035 . The substrate predicts $r = 12/N^2 = 1/300$ ($N = 2 \text{ beat} = 60$); the $\sigma(r)$ timeline — Simons Observatory (~ 2027 , $r/\sigma \sim 1$), LiteBIRD (~ 2035 , 3.3σ), CMB-S4 (6.7σ) — gives three verdicts (Fig. 18): (A) $r \approx 1/300$ on the line $\Rightarrow$ the tower stands and beat $= 30$ is in the sky; (B) $r < 10^{-3} \Rightarrow$ the Starobinsky $N = 60$ spectrum is falsified; (C) $r \text{ detected} \neq 1/300 \Rightarrow$ the beat $\rightarrow N$ origin is wrong. (Witness: `analysis/w33_decisive_r_test.py`.) These distillations are collected in the standalone manuscript `w33_everything.tex`, now with a derived/matched/open table and references.

Honest scope: the seesaw is standard and the EW/RH exponents are the substrate’s, so the floor is derived up to the one Yukawa $y_1 = e^{-\Phi_6/2}$ (a charm-scale neutrino Dirac coupling, motivated by Φ_6 not independently derived), and $72 \approx \text{beat decades}$ to $\sim 4\%$ (the base-10-vs- e factor the residual); the decision tree uses published $\sigma(r)$ forecasts (foregrounds/delensing may move the dates), with $r = 1/300$ and the Starobinsky line the exact predictions.




Figure 18: The theory lives or dies on r by ~ 2035 . The substrate predicts $r = 12/N^2 = 1/300$ ($N = 2 \text{ beat} = 60$, Starobinsky). A LiteBIRD/CMB-S4 measurement ($\sigma(r) \sim 10^{-3}$, a $3\text{--}7\sigma$ detection of $r = 1/300$) yields three clean verdicts: (A) $r \approx 1/300$ on the line $r = 3(1 - n_s)^2$ confirms the tower; (B) $r < 10^{-3}$ falsifies the Starobinsky $N = 60$ spectrum; (C) r detected away from $1/300$ breaks the $\text{beat}=30 \rightarrow N=60$ origin.

C.38 Pass 20 — the one residual input, the floor base settled, and the look-elsewhere bound

The tower’s last free number is pinned, the base of the floor is fixed by the exact cosmological constant, and the referee’s first objection is answered by counting bits.

The one residual input. The seesaw floor (Pass 19) reduced the neutrino/dark-energy/CC scale to $(q+2)\Phi_3 + \Phi_6 = 72$ e-folds up to a single input, the lightest neutrino’s Dirac Yukawa y_1 . The $\mathbb{Z}_3$ texture (Pillar 68) is a *theorem* that fixes only the Dirac *ratios* (the grade rule, up-type $\sim 9:2:1$), never the overall Higgs-VEV scale — so y_1 is genuinely the one normalisation the geometry leaves free. The value the floor wants, $y_1 = e^{-\Phi_6/2} = 0.030$ ($m_{D1} = y_1 v = 7.4$ GeV, charm-scale), is the substrate’s own: a *half* of the GUT exponent $\Phi_6 = q^2 - q + 1 = 7$ ($M_{\text{GUT}} = M_{\text{Pl}} e^{-\Phi_6}$), the coupling at the geometric mean of the M_{Pl} and M_{GUT} scales. And the floor is only *logarithmically* sensitive: a factor-3 change in y_1 moves it by $\sim \Phi_6 \approx 7$ e-folds (m_1 by ~ 10), so the ~ 2 meV scale needs y_1 only within a factor ~ 3 of $e^{-\Phi_6/2}$. So the whole tower reduces to *one* residual number, a charm-scale neutrino Dirac Yukawa equal to $e^{-(\text{half the GUT exponent})}$. (Witness: `analysis/w33_neutrino_dirac_yukawa.py`.)

The floor base, settled by the CC. Pass 19 left the base (e-folds vs decades) open, noting $72 \approx \text{beat decades}$ to $\sim 4\%$. The cosmological constant removes the ambiguity: it is *exact in base 10*, $\log_{10}(\rho_\Lambda/M_{\text{Pl}}^4) = -vq = -120$ to 0.1% (Pass 17, reduced M_{Pl}), and $vq = 4 \text{ beat}$, so the dark-energy scale is $\rho_\Lambda^{1/4} = M_{\text{Pl}} 10^{-\text{beat}} = 2.44$ meV — *exactly* $\text{beat} = 30$ decades below M_{Pl} . The seesaw’s 72 e-folds is the same depth in natural log ($72/\ln 10 = 31.3$ decades $\approx \text{beat}$); the 4% is the *e-vs-10* base mismatch, not a second number. Decades is primary *because* the 0.1% -exact statement is the CC and it is base-10 ($vq = v \cdot q$ is an integer point count), so the 600-cell BC ring counts decades. (Witness: `analysis/w33_floor_base.py`.)

The look-elsewhere bound. A numerologist can always fit *one* number; the honest question is ~ 25 observables against one fixed integer alphabet. Counting bits (MDL): the clean measured matches (m_p/m_e to 10^{-4} , the CC to 10^{-3} , $\sin^2 \theta_W$ and M_Z to 2×10^{-3}) carry ~ 100 bits of precision evidence; and the *same* integers recur — $\text{beat} = 30$ sets n_s, r, n_t , running and the CC; Φ_3, Φ_6 set $\sin^2 \theta_W, M_Z, \theta_{23}, \Delta m_{31}^2/\Delta m_{21}^2, A_s$ — so the trials factor is paid *once per integer* (~ 12), not per observable (~ 25). Granting (generously) each integer a free choice from a pool of 50, the

vocabulary costs ~ 62 bits, against ~ 100 bits of evidence: net $\sim +40$ bits, odds $\sim 10^{12}$ against chance. The *only* null that washes the signal out is a fresh free integer per observable — exactly what a fixed, $q = 3$ -forced alphabet forbids. (Witness: `analysis/w33_look_elsewhere.py`.)

Honest scope: y_1 is an input (motivated by Φ_6 as a half-GUT-exponent, not derived), needed only within a factor ~ 3 ; the floor-base tie-break is principled (the CC is the 0.1%-exact base-10 statement) with the residual the e -vs-10 factor and y_1 ; the look-elsewhere bound uses order-of-magnitude bit counts with deliberately generous priors (not a rigorous frequentist p). The robust claim is the *over-determination*: one ~ 12 -integer $q = 3$ alphabet predicts ~ 25 numbers carrying ~ 100 bits of precision.

C.39 Pass 21 — the computed look-elsewhere, the pinned Yukawa, and the tower on one axis

Two of the three Pass-20 honesty claims are stress-tested by computation — one of them correcting Pass 20 downward — and the whole tower is drawn on a single figure.

The look-elsewhere pool, computed (and Pass 20 tempered). Pass 20 granted the numerologist a “pool of 50” by hand and got net $\sim +40$ bits (10^{12}). Enumerating the *actual* reachable set of the 13 $q = 3$ -forced integers under a fixed simple grammar (a/b , $a \cdot b$, $a \pm b$, $1/a$, a^2 , $\sqrt{a}$, e^{-a}) gives ~ 14 values per e-fold in the coupling band (depth 1). At that density the *loose* few-percent matches ($1-n_s$, $\sin^2 \theta_{23}$, A_s , $\Omega_{\text{DM}}/\Omega_b$, Δm^2) have $E_i \sim 0.5$ –1 each — individually unremarkable, 5/5 matching is only $P \sim 0.6$. The statistical weight is entirely in the *six high-precision* matches (m_p/m_e , $1/\alpha$, m_H , the CC, $\sin^2 \theta_W$, M_Z ; windows 10^{-4} – 2×10^{-3}), whose summed expected accidentals $E \approx 0.30$ make all six matching a Poisson tail $P \approx 8 \times 10^{-7}$. So the over-determination is *real but carried by precision, not count*: this honestly tempers Pass 20’s 10^{12} to a computed $\sim 10^{-6}$, $\sim 4.8\sigma$. (Witness: `analysis/w33_alphabet_pool.py`.)

The pinned Yukawa. Pass 20 left $y_1 = e^{-\Phi_6/2} = 0.030$ as the tower’s one free input. A second, data-anchored route pins it: fix the Dirac texture ratio 9:2:1, pin the heaviest eigenvalue by matching the atmospheric mass $m_{\nu 3} = \sqrt{\Delta m_{31}^2}$ to the seesaw $y_3^2 v^2/M_R$ with $M_R = M_{\text{Pl}} e^{-\Phi_3}$ (the scalaron), and read $y_1 = y_3/9 \approx 0.017$, $m_1 = m_{\nu 3}/81 \approx 0.6$ meV. This agrees with the a priori half-GUT-exponent 0.030 to a factor ~ 1.8 — two independent routes on the same coupling, so y_1 is not independently free. Honest limitation: a single M_R over-hierarchises (81:4:1 vs the observed $\sqrt{33} \approx 5.7$), so the spectrum needs the Majorana M_R texture and the residual freedom moves there — a *partial* closure. (Witness: `analysis/w33_yukawa_closure.py`.)

The tower on one axis. Figure 19 puts every measured observable’s fractional agreement on a log axis (mass/coupling postdictions at 0.01–0.3%, inflation/neutrino at 1–5% within their errors) above the dated kill-shot timeline (LZ ~ 2028 , SO ~ 2027 , JUNO, DESI/CMB-S4, nEXO/LEGEND, LiteBIRD ~ 2035) — the over-determination and the falsifiability in one view, now the cover figure of `w33_everything.tex`. (Witness: `analysis/w33_tower_falsification_figure.py`.)

Honest scope: the computed look-elsewhere depends on the grammar (depth-2 raises the density and softens the count) — the grammar-stable core is that the six sub-0.3% matches are not a simple-grammar accident; y_1 is pinned only to a factor ~ 2 and the spectrum still rests on the (here-underived) Majorana texture; the figure mixes integer/LO postdictions with measurement-limited points on a fractional-deviation axis (stated in its caption).

C.40 Pass 22 — the depth-stable significance, an honest neutrino-mixing failure, and dated discovery years

Three stress-tests, including one that fails: the look-elsewhere is made grammar-robust, the neutrino-mixing residual is shown *not* to close at the minimal level, and the falsification schedule is given exact years.

The depth-stable significance. Pass 21’s look-elsewhere grew with grammar depth. Ranking each match by its *shortest* description (cost = integer leaves) and counting at fixed cost removes that: the alphabet reaches only $N(\leq 2) = 224$ values at cost 2, a count that does not change when the grammar deepens, so a cost-2 match’s bound $N(\leq 2) \cdot 2\delta$ is depth-invariant. The strongest matches are cost-2 ($\sin^2 \theta_W = q/\Phi_3$, M_Z , the CC = vq , m_H), and substrate observables hit cost ≤ 2 far more than random targets of equal precision (70% vs 38%). *But the tempering continues:* only four matches have a depth-stable bound below 1, so no single numerical match is decisive — the coincidences are corroborative, the theory’s weight is in the derivations. (Witness: `analysis/w33_md1_shortest.py`.)

An honest neutrino-mixing failure. Pass 21 framed the residual as “the single lift direction.” Running the full seesaw $m_\nu = m_D m_R^{-1} m_D^T$ with the up-type Dirac block and the cubic-form M_R ($A=0.0017$, $C=0.0442$) lifted by one grade-1 component, and extracting the PMNS angles, shows this is *too small*: across the whole lift scan, in both the graded and diagonal Dirac textures, *no single lift* reproduces all four of $(\Delta m_{31}^2/\Delta m_{21}^2, \theta_{12}, \theta_{13}, \theta_{23})$ — where θ_{13} nears its observed 8.6° the ratio explodes past 200 and $\theta_{23} \rightarrow 0$. The masses sit at the meV floor but the mixing does not come out. So the residual is the *full* neutrino Yukawa/Majorana texture (which the Pillar-65 optimisation reaches at PMNS error 0.006), not one number — an honest negative result that resizes the residual. (Witness: `analysis/w33_seesaw_pmns_joint.py`.)

Dated discovery years. Interpolating the published $\sigma(t)$ forecasts and solving $\text{pred}/\sigma(t) = 3$ (and 5): $\Sigma m_\nu = 58$ meV reaches $3\sigma \sim 2028$ and $5\sigma \sim 2031$ (DESI+CMB-S4), and $r = 1/300$ reaches $3\sigma \sim 2032$ and $5\sigma \sim 2034$ (SO→CMB-S4→LiteBIRD); JUNO pins $\Delta m_{31}^2/\Delta m_{21}^2 = 33$ by ~ 2030 and LZ covers $m_{DM} = 22.8$ GeV by ~ 2028 (Fig. 20). The tower is largely decided by the mid-2030s. (Witness: `analysis/w33_killshot_dashboard.py`.)

Honest scope: the depth-stable bounds are $O(1)$ even for the best matches (corroborative, not a proof); the neutrino-mixing test is a genuine *failure* of the minimal cubic-form M_R (the masses come out, the angles do not); the discovery years use published forecasts that may slip, and the Σm_ν “detection” assumes the NH-minimum value is true (current data could instead yield an exclusion, which would falsify).

C.41 Pass 23 — the predicted mixing, the derivation ledger, and the separation from SO(10)

The Pass-22 neutrino-mixing failure is given its positive form, the whole corpus is graded on one honest scale, and the question “why not just SO(10)?” is answered.

The mixing, predicted with zero parameters. Pass 22 showed the minimal cubic-form M_R fails the PMNS angles; the angles’ proper form is the substrate’s three zero-parameter integer formulas, $\sin^2 \theta_{12} = \mu/\Phi_3 = 4/13$, $\sin^2 \theta_{13} = \lambda/(\Phi_6\Phi_3) = 2/91$, $\sin^2 \theta_{23} = \Phi_6/\Phi_3 = 7/13$, each matching the global fit to under 1σ (the reactor angle to 0.07σ , the sharpest single test), with all denominators from $\Phi_3 = 13$. They satisfy a zero-parameter relation *between* two angles, $\sin^2 \theta_{12} + \sin^2 \theta_{23} = 11/13$ ($11 = k - 1$, Ihara) — a cross-angle prediction a generic model leaves free. So the mixing is not fitted (the Pillar-65 optimisation fits it; these formulas do not): the angles are fixed by $q = 3$. The one honest gap is the leptonic Dirac δ_{CP} — the substrate’s clean CP datum $\sin \delta = 15/17$

is the *quark* phase — so δ_{CP} is the open forward test (DUNE/T2HK ~ 2030 – 2035). (Witness: `analysis/w33_pmns_prediction.py`.)

The derivation ledger. Grading the corpus on one scale — DERIVED (geometry-forced, no free parameter), POSTDICTED (a zero-parameter integer formula equal to a measured value), FITTED (parameters tuned), INPUT (a dimensionful anchor) — the 25 major results split 8 DERIVED, 12 POSTDICTED, 2 FITTED, 3 INPUT. The *structure* is derived (4D via the KO-dimension, three generations via $\text{Sp}(4, 3) = W(E_6)$, the gauge group, the cyclotomic skeleton, Starobinsky inflation, the boson–fermion balance behind the CC); most *numbers* are zero-parameter postdictions; only the full CKM/PMNS matrices are fitted, and only the electroweak scale and one neutrino Yukawa are input. So the honest backbone is derived structure + zero-parameter numbers, with very few fits — the claim the Pass-22 audit pointed to. (Witness: `analysis/w33_derivation_ledger.py`.)

The separation from $SO(10)$. Much of the tower (seesaw, unification, Starobinsky-like inflation, a dark-matter candidate) is shared with generic $SO(10)$, which carries ~ 27 free Yukawa/Higgs parameters; the substrate fixes them by $q = 3$, leaving ~ 1 , so it is a single *point* in the $SO(10)$ family’s parameter space. Every zero-parameter relation it forces is a discriminator: the PMNS relation $11/13$ and $\sin^2 \theta_{13} = 2/91$ (JUNO/DUNE ~ 2030); the inflation *point* $(A_s, n_s, r) = (e^{-20}, 1 - 1/30, 1/300)$ at $N = 60$ where generic Starobinsky has a *line* (LiteBIRD ~ 2035); $m_{\text{DM}} = M_Z/\mu = 22.8$ GeV (LZ ~ 2028); and the CC $= -vq$, which $SO(10)$ does not predict at all. The single cleanest discriminator is the inflation point — three primordial numbers fixed at once. (Witness: `analysis/w33_vs_so10.py`.)

Honest scope: the angle formulas are zero-parameter postdictions (matched, but unfitted), δ_{CP} is genuinely not predicted; the DERIVED/POSTDICTED line is a judgement (structure vs. value); and $SO(10)$ is a family that can fit any single number, so the separation is a rigidity/Occam argument (fewer parameters, sharper relations), not a logical impossibility.

C.42 Pass 24 — machine = world: the code is the matter algebra, and why 3

Stepping back to the whole architecture: the photonic error-correcting layer and the physics layer are two faces of one finite object, and the integer that joins them, 3, is over-determined.

The code is the E_6 root system. The Holonet/QEC track’s finite code (the $[[66, 8, 3; 5]]_3$ subsystem code) lives on the genus-6 triangulation of the complete graph K_{12} : $V = 12$ vertices, $E = 66$ edges, $F = 44$ triangular faces, genus 6 (a triangulation, $F = 2E/3$; Euler $12 - 66 + 44 = -10 = 2 - 2g$). *Every* count is a substrate integer: the 12 vertices are the valency $k = q(q+1)$ (the $\mathbb{Z}_3 \times (\mathbb{Z}_2)^2$ local fibre); the 66 edge-qudits are $\dim SO(12) = |\text{roots}(E_6)| - \text{rank}(E_6)$; the 44 faces are $v + \mu$; the 6 genus holes / parity symbols are $\text{rank}(E_6) = \text{the KO-dimension} = 2q$; the code length $72 = |\text{roots}(E_6)| = \text{the seesaw-floor e-folds } (q+2)\Phi_3 + \Phi_6$; the distance $d = 3 = q$; the parent logical count $k = 13 = \Phi_3$ (the electroweak-mixing integer, $\sin^2 \theta_W = q/\Phi_3$, $M_Z = \Phi_3 \Phi_6$); and the rate $(k-1)/k = 11/12$ carries the same $11 = k - 1$ as the PMNS relation $11/13$. So the universe’s fault-tolerant code *is* the E_6 matter Lie algebra (72 symbols = 72 roots, 6 parity = rank) on the genus-6 surface of K_{12} , the protected logical information is the Standard Model, and the three fermion generations *are* the distance-3 redundancy. The D -series threads it: $SO(10) = 45$ (GUT), $SO(12) = 66$ (code payload), $SO(14) = 91 = M_Z$. (Witness: `analysis/w33_machine_world_bridge.py`.)

Why 3 — over-determined. The same 3 is the color number, the generation number, and the code distance, and it is the *minimum* for three independent reasons: the substrate rule $q! = 2q$ (uniquely $q = 3$, shared value $6 = 2q = \text{KO-dimension}$); CP violation, which by Kobayashi–Maskawa needs ≥ 3 generations (the CKM phase count $(n-1)(n-2)/2$ is first nonzero at $n = 3$); and fault tolerance, which needs distance ≥ 3 to correct one error. Three different *kinds* of constraint — arithmetic, representation-theoretic, coding-theoretic — all select 3, so $d = q = n_{\text{gen}}$ is a structural

Table 2: The machine = world dictionary. Every structural count of the genus-6 K_{12} error-correcting code is a $W(3, 3)$ physics integer, because both are the same finite object.

Code / geometry object	Count	Physics identity
K_{12} vertices (Z -checks)	12	$k = q(q+1)$ valency; $\mathbb{Z}_3 \times (\mathbb{Z}_2)^2$ fibre
edges (payload qudits)	66	$\dim \text{SO}(12) = \text{roots}(E_6) - \text{rank}(E_6)$
triangular faces (X -checks)	44	$v + \mu$ (proton-scale exponent)
genus = parity symbols	6	$\text{rank}(E_6) = \text{KO-dim} = 2q$
code length n	72	$ \text{roots}(E_6) = (q+2)\Phi_3 + \Phi_6$ (seesaw floor)
distance d	3	$q = \text{number of generations}$
parent logical k	13	Φ_3 ($\sin^2 \theta_W = q/\Phi_3$, $M_Z = \Phi_3 \Phi_6$)
surface logical $2g$	12	k (valency)
rate $(k-1)/k$	11/12	$11 = k-1$, as in PMNS $\sin^2 \theta_{12} + \sin^2 \theta_{23} = 11/13$

inevitability, not a choice. (Witness: `analysis/w33_q3_triple_minimality.py`.)

The architecture closes on itself. $W(3, 3) = \text{SRG}(40, 12, 2, 4) = \text{GQ}(3, 3)$ is one object read four ways — combinatorics (the graph), physics (the KO-dim-6 spectral triple $\rightarrow 4\text{D} + \text{SM} + \text{cosmology}$), computation (the genus-6 K_{12} code), and dynamics (the QCA of topological index 27 running universal Rule-110). The QCA runs the code that protects the physics, and the QCA’s conserved index 27 *is* the E_6 matter shell ($40 = 1 + 12 + 27$) it protects: $27 = 27 = 27$. The machine computes the world, and the world is the machine’s protected state. (Witness: `analysis/w33_substrate_four_faces.py`.)

Honest scope: the arithmetic is exact (genus, Euler, E_6 roots/rank, $d = q$, $k = \Phi_3$, $27 = 27 = 27$); “machine = world” is a structural identification — the same finite integers organize both layers, strong evidence they are one object — not a causal derivation that the universe must be a code; the $44 = v + \mu$ entry is the one loose match; and the $q! = 2q$ rule is the substrate’s chosen principle, with the Kobayashi–Maskawa and distance-3 minima the genuine independent convergences onto 3.

C.43 Pass 25 — the protected 27 is one Standard-Model generation, and charge is $1/3$ because $q = 3$

Pass 24 identified the code’s protected matter shell with the E_6 fundamental **27**. This pass shows that **27** carries *exactly* the Standard-Model fermion quantum numbers, and that the most basic facts of the electromagnetic and strong sectors — charge quantised in $1/3$, fractionally-charged quarks, confinement — are one consequence of $q = 3$.

The 27 is one generation. The standard branching $E_6 \rightarrow \text{SO}(10) \rightarrow \text{SU}(5) \rightarrow \text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ gives **27** = **16** + **10** + **1**, and the **16** of $\text{SO}(10)$ is one complete generation (Table 3): $16 = g + 1$ Weyl states, of which $15 = g$ are chiral, plus the right-handed neutrino; the **10** + **1** are vector-like exotics. The multiplicities are the substrate’s Hodge integers: the chiral count is $g = 15$, the gauge adjoint is $f = 24 = \dim \text{SU}(5)$, and the boson–fermion balance $f\Phi_4 = g\mu^2 = 24 \cdot 10 = 15 \cdot 16 = 240 = |\text{roots}(E_8)|$ (the structural supersymmetry behind the cosmological constant) is the $\text{SU}(5)$ adjoint balanced against one generation. Three copies ($\text{Sp}(4, 3)$, the three generations) give $3 \cdot 27 = 81 = q^4 = \dim H_1$, the massless-matter homology. (Witness: `analysis/w33_e6_27_standard_model.py`.)

Charge is $1/3$ because $q = 3$. The defining integer $q = 3$ is the order of the colour centre $\mathbb{Z}_3$ (the SRG 3-grading, $\text{Sp}(4, 3) = \text{three copies of } \mathbb{F}_3^4$, colour $\text{SU}(3)$). Computing $Q = T_3 + Y$ for every fermion (Table 3), *every* charge is a multiple of $1/q = 1/3$, and the fractional part obeys

Table 3: The protected **27** as one Standard-Model generation. The **16** of SO(10) is one family; every electric charge $Q = T_3 + Y$ is a multiple of $1/q = 1/3$, and the fractional part is $-t/3 \bmod 1$ with t the colour triality.

Multiplet	(SU3, SU2) _Y	states	triality t	Q	$Q \bmod 1$
Q (quark doublet)	$(\mathbf{3}, \mathbf{2})_{+1/6}$	6	1	$+\frac{2}{3}, -\frac{1}{3}$	$\frac{2}{3}$
u^c	$(\bar{\mathbf{3}}, \mathbf{1})_{-2/3}$	3	2	$-\frac{2}{3}$	$\frac{1}{3}$
d^c	$(\bar{\mathbf{3}}, \mathbf{1})_{+1/3}$	3	2	$+\frac{1}{3}$	$\frac{1}{3}$
L (lepton doublet)	$(\mathbf{1}, \mathbf{2})_{-1/2}$	2	0	$0, -1$	0
e^c	$(\mathbf{1}, \mathbf{1})_{+1}$	1	0	$+1$	0
ν^c (RH neutrino)	$(\mathbf{1}, \mathbf{1})_0$	1	0	0	0
total = 16 = $g + 1$; chiral = 15 = g			+ 10 + 1 exotics = 27		

the triality law $Q \bmod 1 = -t/3$: colour singlets ($t = 0$ — leptons, hadrons) carry integer charge and are free, while colour triplets ($t \neq 0$ — quarks) carry fractional charge in $\frac{1}{3}\mathbb{Z}$ and are confined, since only triality-0 states are asymptotic. So the $1/3$ charge quantum is the $1/q$ of the $\mathbb{Z}_3$ centre, fractional charge is the signature of non-zero triality, and confinement is the rule that only $\mathbb{Z}_3$ -trivial states are observable — charge quantisation, the fractional quark charge, and confinement are *one* consequence of $q = 3$. (Witness: `analysis/w33_charge_quantization_z3.py`.)

Honest scope: the $E_6 \rightarrow \text{SM}$ branching and the triality-charge relation are standard group theory; the substrate content is that the matter group is E_6 (the **27** = q^a shell, derived earlier) and that the multiplicities $g = 15$, $f = 24$, $q^4 = 81$ and the charge quantum $1/q$ are the $W(3, 3)$ integers, so the Standard-Model content, its multiplicities, and its charge quantisation are substrate-forced rather than inserted. The exotic **10 + 1** is assumed heavy, and the confinement *dynamics* is asserted through the triality-0 selection rule, not derived from QCD.

C.44 Pass 26 — why the Standard Model is consistent, and the weak angle as a **27**-trace

Having placed one generation in the E_6 **27** (Pass 25), two deep facts follow with no extra input: the Standard Model’s anomaly cancellation, and its weak mixing angle.

Anomaly cancellation, inherited. A chiral gauge theory is mathematically consistent only if its gauge and gravitational anomalies cancel — five delicate sum rules over the fermion content. For one generation with the **27**’s hypercharges they all vanish exactly:

$$\text{SU}(3)^3 = 0, \quad \text{SU}(3)^2\text{U}(1) = 0, \quad \text{SU}(2)^2\text{U}(1) = 0, \quad \text{grav}^2\text{U}(1) = 0, \quad \text{U}(1)^3 = 0.$$

In the bare Standard Model these look like arithmetic accidents of the hypercharges; here they are automatic, because the matter lives in the E_6 fundamental **27**, which is an anomaly-free representation (E_6 has no cubic Casimir). The substrate forces the matter group to be E_6 (the **27** = q^a shell), so the cancellation is *inherited*, not checked generation by generation; and the hypercharges that make it work, $Y = (\frac{1}{6}, -\frac{2}{3}, \frac{1}{3}, -\frac{1}{2}, 1)$ — quantised in $\frac{1}{6}$, i.e. $6Y = (1, -4, 2, -3, 6)$ — are the unique anomaly-free assignment, the $q = 3$ -quantised charges of the **27**. (Witness: `analysis/w33_anomaly_cancellation.py`.)

The weak angle is a 27-trace. At a grand-unified scale the Weinberg angle is a pure group-theory trace, $\sin^2 \theta_W = \text{Tr}(T_3^2)/\text{Tr}(Q^2)$, with no dynamics. Evaluated over one generation (the **27**),

$$\sum T_3^2 = 2, \quad \sum Q^2 = \frac{16}{3}, \quad \sin^2 \theta_W = \frac{2}{16/3} = \frac{3}{8} = \frac{q}{q^2 - 1},$$

the exact SU(5) tree value, here a cyclotomic ratio. The measured low-energy value is the substrate's *other* ratio, $\sin^2 \theta_W(M_Z) = q/\Phi_3 = 3/13 = 0.2308$ (vs. 0.2312 measured), and the renormalisation-group running carries the unified $3/8 = 0.375$ down to $3/13$ at the Z mass: the denominator runs from $q^2 - 1 = 8$ at unification to $\Phi_3 = q^2 + q + 1 = 13$ at low energy, a shift of $5 = F_5$. So the weak angle is q over a quadratic in q at both ends of the descent — $q/(q^2 - 1)$ at the top, q/Φ_3 at the bottom, the same numerator $q = 3$. (Witness: `analysis/w33_weinberg_from_27.py`.)

Honest scope: the anomaly cancellation and the SU(5) tree value $\sin^2 \theta_W = 3/8$ are standard group theory; the substrate content is that the matter group is E_6 (an E_6 **27** is anomaly-free, so the consistency is inherited) and that the weak angle equals $q/(q^2 - 1)$ at unification and q/Φ_3 at M_Z , both cyclotomic ratios, with the F_5 denominator shift. The running that connects them is the standard SM/SUSY-GUT RG, not derived here; the measured q/Φ_3 is a separate (very good) postdiction.

C.45 Pass 27 — the three forces meet at $M_{\text{Pl}}e^{-\Phi_6}$ with $\alpha_{\text{GUT}}^{-1} = f$, and the proton decays

Pass 26 fixed the weak angle as a **27**-trace; the deeper test is whether the three running couplings actually converge — and what it predicts for proton decay.

Unification at the ladder rung, with $\alpha_{\text{GUT}}^{-1} = f$. Running the one-loop renormalisation group from M_Z (in SU(5) normalisation, $\alpha_1^{-1} = 59.0$, $\alpha_2^{-1} = 29.6$, $\alpha_3^{-1} = 8.5$), the plain Standard-Model content does *not* unify — its three pairwise crossings are spread over a factor $\sim 10^4$ in scale (10^{13} – 10^{17} GeV, the famous near-miss). But with the substrate's structural-supersymmetric content — the SUSY-like spectrum implied by the exact boson–fermion balance $f\Phi_4 = g\mu^2 = 240$ — the three crossings coincide to a factor 1.2 at

$$M_{\text{GUT}} \simeq 2 \times 10^{16} \text{ GeV} \simeq M_{\text{Pl}}e^{-\Phi_6}, \quad \alpha_{\text{GUT}}^{-1} \simeq 24.3 = f = \dim \text{SU}(5).$$

So the scale at which the forces merge is the substrate's mass-ladder rung $M_{\text{Pl}}e^{-\Phi_6}$ (the grand-unification e-fold, $\Phi_6 = q^2 - q + 1 = 7$; the exponent $\ln(M_{\text{Pl}}/M_{\text{GUT}}) = 6.3$ vs $\Phi_6 = 7$ to $\sim 10\%$), and the *strength* at which they merge is the gauge-mode count $f = 24 = \dim \text{SU}(5)$ — the unified inverse coupling is the dimension of the unified group. And the content doing the unifying is the *same* boson–fermion balance $f\Phi_4 = g\mu^2$ that cancels the leading cosmological constant (Pass 18): one spectrum, two consequences. (Witness: `analysis/w33_gauge_unification.py`.)

The proton decays — a dated kill-shot. Unification makes the proton decay through the same heavy gauge bosons, at a rate fixed by those two numbers: $\tau(p \rightarrow e^+\pi^0) \sim M_{\text{GUT}}^4/(\alpha_{\text{GUT}}^2 m_p^5) \sim 2.6 \times 10^{35} \text{ yr}$ (or $\sim 3 \times 10^{36} \text{ yr}$ at the one-loop scale), i.e. $\sim 10^{35}$ – 10^{36} yr . This is above the Super-Kamiokande bound $\tau > 2.4 \times 10^{34} \text{ yr}$ (not excluded) and squarely within Hyper-Kamiokande's reach ($\sim 10^{35} \text{ yr}$ by the late 2030s–2040): a detection near 10^{35} yr confirms the $M_{\text{Pl}}e^{-\Phi_6}$ scale, a null beyond $\sim 10^{36} \text{ yr}$ falsifies the minimal picture. So grand unification stakes itself on one dated experiment, like the inflationary $r = 1/300$. (Witness: `analysis/w33_proton_decay_test.py`.)

Honest scope: one-loop MSSM-like running unifying near 2×10^{16} with $\alpha_{\text{GUT}}^{-1} \sim 24$ –25 is textbook, and plain SM does not unify; the substrate content is that M_{GUT} is the ladder rung $M_{\text{Pl}}e^{-\Phi_6}$, that $\alpha_{\text{GUT}}^{-1} = f = \dim \text{SU}(5)$, and that the unifying SUSY-like spectrum is the balance $f\Phi_4 = g\mu^2$. The proton-lifetime band $\sim 10^{35}$ – 10^{36} yr carries a 1–2 order-of-magnitude hadronic and threshold uncertainty; the exact M_{GUT} , α_{GUT} and the precise running are not derived in detail here.

C.46 Pass 28 — the Higgs sector, one structural supersymmetry, and the state of the solution

The electroweak symmetry-breaking sector closes the matter map, the apparent two-scale supersymmetry puzzle dissolves, and the architecture can be assessed honestly as a whole.

The Higgs and the hierarchy. The electroweak sector is a cyclotomic descent like the rest: the Higgs mass is $m_H = vq + \mu + 1 = 125$ GeV, the electroweak scale is $v_{EW} = 2(vq + q) = 246$ GeV (so $m_H/v_{EW} \simeq \frac{1}{2}$ and the quartic $\lambda = m_H^2/2v_{EW}^2 = 0.129 \simeq \mu^2/(vq + \mu) = 16/124$), the electroweak scale sits at $q\Phi_3 = 39$ e-folds below the Planck scale ($\ln(M_{Pl}/v_{EW}) = 38.4$ — the hierarchy is the substrate integer $q\Phi_3$), and the Higgs-vacuum metastability scale ($\sim 10^{11}$ GeV) is $h(E_7) = 18$ e-folds below M_{Pl} ($\ln(M_{Pl}/10^{11}) = 18.6$, with $18 = h(E_7) = \bar{\lambda} = \bar{\mu}$ of the complement SRG). (Witness: `analysis/w33_higgs_sector.py`.)

One structural supersymmetry, no superpartners. Gauge unification (Pass 27) needs SUSY-like running — which in ordinary low-energy supersymmetry means $\sim$ TeV superpartners — while the cosmological constant (Pass 18) is the broken balance at the meV floor: a 10^{15} mismatch. The resolution is that the substrate’s supersymmetry is *structural*, the exact Hodge mode balance $f\Phi_4 = g\mu^2 = 240$, *not* a particle supersymmetry: it predicts no light superpartners, it cancels the leading vacuum energy and the Higgs quadratic divergence by mode-counting, and it breaks only at the meV floor. The cosmological constant fixes that breaking at meV ($\rightarrow 10^{-vq}$) and *forbids* TeV ($\rightarrow 10^{-62}$, fifteen orders too large): the absence of TeV superpartners is a prediction, and the two-scale tension dissolves — there is no particle-SUSY scale, only the structural balance and its meV breaking. (Witness: `analysis/w33_two_susy_scales.py`.)

The state of the solution. From the single selected integer $q = 3$ the substrate DERIVES the structure of the Standard Model and cosmology, MATCHES essentially every measured dimensionless number as a zero-parameter integer postdiction, takes as irreducible INPUT one dimensionful scale (the Planck mass) plus one charm-scale neutrino Yukawa, and leaves a short, explicit list of OPEN dynamical questions (Table 4). So the honest verdict: a complete, falsifiable, *parameter-free* framework — zero free dimensionless parameters, one dimensionful input, six dated kill-shots ($r = 1/300$, Σm_ν , m_{DM} , proton decay, the Δm^2 ratio, δ_{CP}) decided by the 2030s–2040s. It is solved as a closed predictive framework with dated tests, *not* as a complete dynamical theory with gravity derived: the continuum Einstein–Hilbert lift is the principal open piece. (Witness: `analysis/w33_state_of_the_solution.py`.)

Table 4: The state of the solution. From $q = 3$: DERIVED structure, MATCHED zero-parameter numbers, irreducible INPUT, and the OPEN dynamical questions.

DERIVED	$q = 3$; 4D (KO=2q); 3 generations (Sp(4,3)); gauge group; matter = 27 ; charge 1/3 & confinement; anomaly cancellation; Starobinsky $N=60$; balance 240; machine=world code
MATCHED	$1/\alpha=137$; $\sin^2 \theta_W=3/13 \rightarrow 3/8$; $\alpha_s=9/76$; $M_Z=91$; $m_H=125$; $v_{EW}=246$; $m_p/m_e=1836$; PMNS 4/13, 2/91, 7/13; $\Sigma m_\nu=58$; Δm^2 ratio 33; CC = $-vq$; $A_s=e^{-20}$, n_s , $r=1/300$; Ω_{DM}/Ω_b ; $m_{DM}=M_Z/\mu$; $M_{GUT}=M_{Pl}e^{-\Phi_6}$; $\alpha_{GUT}^{-1}=f$; hierarchy $q\Phi_3$; metastability $h(E_7)$
INPUT	the Planck mass (one dimensionful scale); the neutrino Dirac Yukawa y_1 (pinned to $\sim 2\times$)
OPEN	continuum Einstein–Hilbert gravity lift; precise intermediate spectrum for exact unification; from-nothing CC cancellation; leptonic δ_{CP} ; full CKM/PMNS matrices (fitted)

Honest scope: the grading is the project’s own judgement (the Pass-23 scale extended to the full architecture); MATCHED means a zero-parameter integer equals a measured number (a postdiction, the value known), not a from-nothing derivation; the Higgs/ v_{EW} entries are integer postdictions; the structural-SUSY hierarchy protection is asserted by mode-counting, not proven loop-by-loop; and continuum gravity is not derived. So “solved” means a closed, falsifiable, parameter-free framework with dated tests — a strong claim, honestly bounded.

C.47 Pass 29 — gravity is the spectral action, and the continuum gap is two theorems

The principal open item of Pass 28 was the continuum gravity lift. It resolves through the Chamseddine–Connes spectral action, and what remains narrows from a missing sector to two precisely-named theorems.

Gravity is the substrate’s spectral action. For the almost-commutative geometry $M^4 \times F$ with $F = W(3, 3)$, the bosonic spectral action $S = \text{Tr } f(D^2/\Lambda^2)$ expands, by the heat-kernel (Seeley–DeWitt) theorem, into local curvature integrals weighted by the finite $W(3, 3)$ Hodge moments. The three leading coefficients *are* the three gravitational terms:

$$a_0(\Lambda^4, M_0) \rightarrow \text{CC}, \quad a_2(\Lambda^2, M_0) \rightarrow \text{Einstein–Hilbert } \left(\frac{1}{16\pi G} \sim \Lambda^2 v\right), \quad a_4(\Lambda^0, M_1, M_2) \rightarrow R^2 \text{ Starobinsky+Weyl}$$

and the same expansion gives the Yang–Mills and Higgs terms — so gravity and the Standard Model are *one* spectral action. The weights are the substrate’s Hodge spectrum $\{0^1, 10^{24}, 16^{15}\}$: $M_0 = \text{Tr } 1 = 40 = v$, $M_1 = \text{Tr } L = 480$, $M_2 = 6240$. Strikingly the first moment splits as $M_1 = 24 \cdot 10 + 15 \cdot 16 = 240 + 240$ — the boson–fermion balance $f\Phi_4 = g\mu^2$ itself — so the cosmological-constant cancellation (Pass 18) is a property of the spectral data. Thus the CC, the Einstein–Hilbert action, and the Starobinsky inflation are the three heat-kernel coefficients of one action. (Witness: `analysis/w33_gravity_spectral_action.py`.)

The continuum gap is two theorems, not a sector. The gravity *dynamics* is therefore derived. What is *not* yet proven narrows to two precisely-located theorems: **(T1)** the a_2 positivity / curved-4D refinement — proving the Einstein–Hilbert coefficient on the curved tower gives a *positive* Newton constant ($\frac{1}{16\pi G} \sim \Lambda^2 v$), with the spectral data pinned and the geometric-route convergence established but the curved positivity open; and **(T2)** the emergence of the spacetime manifold M^4 from the discrete substrate (Connes reconstruction of a smooth manifold from the finite datum) — the substrate supplies $F = W(3, 3)$, while M^4 is currently tensored in, not derived. So Pass 28’s single open item splits: the gravity dynamics moves to DERIVED, and the open frontier becomes (T1)+(T2). (Witness: `analysis/w33_continuum_gap.py`.)

Honest scope: the Chamseddine–Connes spectral action and the $a_0/a_2/a_4 = \text{CC}/\text{EH}/R^2$ identification are established machinery; the substrate content is that $F = W(3, 3)$, so the weights are the specific moments $\{40, 480, 6240\}$ and the $M_1 = 240 + 240$ split is the structural supersymmetry. This pass shows gravity *is* the substrate’s spectral action and locates the frontier precisely; it does *not* prove (T1) or (T2). The continuum gravity lift is thus narrowed from a missing sector to two named theorems — the genuine, honestly-stated edge of the solution.

C.48 Pass 30 — the emergent spacetime is K3, and (T2) reduces to one convergence theorem

The deeper of the two residual theorems, (T2) the manifold emergence, is taken on: the continuum 4-manifold is identified, its topology is shown to be substrate data, and the open question reduces to a single named convergence statement.

The emergent spacetime is K3. The continuum manifold M^4 is not a free choice: it is the K3 surface, and every topological invariant is a $W(3,3)$ integer. The Euler characteristic $\chi(K3) = 24 = f$ (the gauge-mode count, $\dim \text{SU}(5)$); the signature $\sigma(K3) = -16 = -\mu^2$ (the matter gap); the Betti numbers $(1, 0, 22, 0, 1)$ with $b_2^+ = 3 = q$; and the K3 intersection lattice $E_8(-1)^2 \oplus H^3$, whose two E_8 summands carry $2 \times 240 = 480 = M_1 = \text{Tr } L$ — the substrate’s first heat-kernel moment (Pass 29). So the manifold’s topology *is* the substrate’s spectral data. The spectral action’s topological gravity term, the Gauss–Bonnet integral $\frac{1}{32\pi^2} \int \text{GB} = \chi$, is fixed on K3 to $\chi = 24 = f$: the gauge-mode count is the Euler characteristic of spacetime. And the substrate realises K3 concretely as an edgewise triangulation tower (the BT984–1135 program), whose homology is forced to K3’s Betti numbers at every refinement level. (Witness: `analysis/w33_manifold_emergence_k3.py`.)

(T2) *reduces to spectral-propinquity convergence.* Latrémolière’s *spectral propinquity* (Math. Ann. 2023) is a metric on spectral triples for which the Dirac spectrum *and* the spectral action $\text{Tr } f(D^2/\Lambda^2)$ are *continuous* functionals. So the gravity+matter action (Pass 29) is continuous for the propinquity, and (T2) reduces to the single statement: does $[W(3,3) \otimes (\text{edgewise-K3 tower})]$ converge in the propinquity to $K3 \otimes W(3,3)$? If it does, the spectral action — Einstein–Hilbert, the cosmological constant, R^2 and the Standard Model — converges *rigorously* and $M^4 = K3$ emerges from the discrete substrate. The inputs are in place: the limit (K3, topology = substrate integers), the action’s continuity (Latrémolière), the triangulation tower, the symmetry densification $\text{Sp}(4, \mathbb{F}_q) = W(E_6) \rightarrow \text{Sp}(4, \mathbb{R}) \rightarrow \text{AdS}_4 \rightarrow \text{Minkowski}$, and term-by-term geometric convergence. So the *whole theory’s open frontier* is now exactly two named theorems: **(T1)** the a_2 curved positivity (Newton constant $G > 0$) and **(T2’)** the spectral-propinquity convergence of the $W(3,3) \times K3$ tower — everything else derived or matched. (Witness: `analysis/w33_propinquity_reduction.py`.)

Honest scope: “the continuum is K3” is a candidate identification (strong invariant matches — $\chi = f$, $\sigma = -\mu^2$, $b_2^+ = q$, lattice $= 2E_8 \oplus 3H = M_1$ — plus the corpus’s heterotic–K3 dictionary and the triangulation tower), not a proven uniqueness; spectral propinquity and the continuity of the spectral action are Latrémolière’s theorems, but the convergence of *this* tower (T2’) is the residual, not proven here. The deepest open question — spacetime from the discrete — is thereby sharpened to one convergence theorem with a known limit and a continuous action, not closed.

C.49 Pass 31 — the convergence mechanism, demonstrated in 4D, and (T2’) narrowed to one input

Attacking (T2’) directly: the discrete→continuum spectral convergence it relies on is demonstrated explicitly in 4D, and on the K3 tower the necessary conditions are checked, leaving a single analytic input.

The convergence mechanism, demonstrated. (T2’) rests on a concrete mechanism — as a triangulation of a 4-manifold is refined, the Dirac/Laplace spectrum converges to the continuum (Dodziuk–Patodi / finite-element exterior calculus) and the spectral dimension reads 4. On the flat 4-torus $T^4 = \lim Z_n^4$, where every spectrum is explicit, the rescaled graph-Laplacian eigenvalues converge to the continuum $\sum_i k_i^2$ at the Dodziuk–Patodi rate $O(1/n^2)$ (each doubling of n cuts the eigenvalue error by ~ 4 : e.g. the lowest mode $0.950 \rightarrow 0.987 \rightarrow 0.997 \rightarrow 0.999$), and the heat trace gives a spectral dimension $d_s = -2 d \log Z / d \log t = 4.00$ in the well-resolved regime. So the convergence the spectral propinquity controls is real and computable; the K3 tower is the same mechanism on a curved 4-manifold. (Witness: `analysis/w33_spectral_convergence_demo.py`.)

(T2’) *narrowed to the metric input.* Applying this to the K3 triangulation tower, five necessary conditions for propinquity convergence are established: (C1) spectral dimension 4 (each level tri-

angulates the closed 4-manifold K3, the 4D Weyl law); (C2) Betti numbers stable = (1, 0, 22, 0, 1), topologically forced; (C3) $\chi = 24 = f$ stable; (C4) closed pseudomanifold (each codimension-1 face in exactly two top faces); (C5) the combinatorial Hodge eigenvalues converge to the continuum K3 spectrum (Dodziuk–Patodi, the mechanism above). The one remaining analytic input is the quantum-*metric* convergence (M): spectral propinquity is a metric on metric spectral triples, so beyond the spectrum it needs the Lipschitz seminorm — the combinatorial distance from the Dirac commutator $[D, \cdot]$ — to converge to the geodesic distance on K3. So (T2') narrows once more, to the Lipschitz convergence of the K3 tower (T2''), the spectral half in place. The whole theory's frontier is then **(T1)** the a_2 curved positivity and **(T2'')** the K3-tower Lipschitz convergence — everything else (spectrum, dimension, topology, and all of physics) derived or matched. (Witness: `analysis/w33_k3_convergence_checklist.py`.)

Honest scope: T^4 is flat, so the demonstration shows the spectral/dimension convergence cleanly but not K3's curved subtleties; (C1)–(C4) are established facts about the K3 tower, (C5) is the eigenvalue convergence whose *mechanism* the T^4 demo exhibits but whose K3-specific computation (large, curved) is not run; and (M), the Lipschitz / quantum-metric convergence, is genuinely open — the named residual. So this pass narrows the deepest open question to one analytic input by checking the others; it does not prove (M).

C.50 Pass 32 — cracking the frontier: both theorems collapse to one finite condition

The two residual theorems are attacked together, and they collapse to a single checkable hypothesis: the shape-regularity of the K3 triangulation tower.

(T1) *is manifest, and K3 is a vacuum solution.* The Einstein–Hilbert coefficient is $\frac{1}{16\pi G} \sim f_2 \Lambda^2 M_0$, where $f_2 = \int_0^\infty f(u) du$ is the cutoff moment (positive for *any* positive cutoff: 1.00 for e^{-u} , 0.33 for a bump, 0.89 for a Gaussian) and $M_0 = \text{Tr}_F(1) = v = 40 > 0$. So $G > 0$ is manifest from positive spectral data — the curvature R is the *integrand* of a_2 , its sign-fixing *coefficient* is the positive $f_2 M_0$, and the curved-positivity worry dissolves. Moreover K3 is Ricci-flat (Calabi–Yau, holonomy $\text{SU}(2)$): $R = 0$, $\int R = 0$, so K3 is a *vacuum* Einstein solution — a gravitational instanton — with zero background cosmological constant (the a_0 term cancelled by the boson–fermion balance), the observed CC being the meV-floor breaking. (Witness: `analysis/w33_newton_positivity.py`.)

(T2'') *is the taxicab problem, cured by shape-regularity.* A naive cubic grid gives the *taxicab* (L^1) distance, overestimating the Euclidean by $\sqrt{d}$ (a *scale-invariant* ratio — $\sqrt{2}$ in 2D), so refining a cubic mesh never recovers Euclidean geometry: that is why the metric is harder than the spectrum. A shape-regular *simplicial* triangulation cures it — its geodesic metric converges to the Riemannian metric (Gromov–Hausdorff / FEEC), the diagonal edge turning the taxicab $2a$ into the exact Euclidean $\sqrt{2}a$. And the substrate's K3 tower is a shape-regular simplicial triangulation, not a cubic grid. (Witness: `analysis/w33_metric_taxicab.py`.)

The collapse. One property — shape-regularity of the $W(3,3) \times K3$ tower — *simultaneously* delivers the Hodge-eigenvalue convergence (Dodziuk–Patodi, C5) *and* the metric convergence (Gromov–Hausdorff/FEEC, M), hence by the spectral propinquity the spectral-action convergence and the emergence of the Riemannian manifold $M^4 = K3$, with $G > 0$ already manifest. So the whole theory's open frontier is no longer two analytic theorems but one finite, per-level combinatorial condition. The chain runs

$$q=3 \rightarrow W(3,3) \rightarrow \text{finite spectral triple} \rightarrow \text{spectral action} = \text{gravity+SM} \rightarrow M^4=K3 \xrightarrow{[\text{shape-regular}]} \text{Riemannian } M^4$$

its one unproven link being whether the explicit K3 triangulation tower is shape-regular. In this

framework, the deepest questions — where spacetime comes from, why gravity is positive — become a finite check on an explicit simplicial complex plus established analysis. (Witness: `analysis/w33_frontier_collapse.py`.)

Honest scope: this is a *reduction*, not a closure. (T1)’s positivity assumes a standard positive cutoff (the mild fermion-doubling sign bookkeeping aside); (T2’’)’s reduction uses the established FEEC / Gromov–Hausdorff / propinquity convergence for shape-regular towers; and the shape-regularity of the *specific* BT984–1135 K3 tower at every level is a finite but unrun combinatorial check — the single remaining step. The K3 identification keeps its candidate caveat. So the gravitational frontier reduces to one finite condition plus standard analysis — a dramatic narrowing, honestly bounded.

C.51 Pass 33 — the last link is a theorem: the gravity derivation closes to two narrow conditions

The single finite condition is run, and the gravity arc closes: no open analytic theorem remains in the chain from $q = 3$ to Riemannian spacetime.

Edgewise subdivision is shape-regular. The K3 tower is built by edgewise refinement, and edge-wise subdivision is the classical shape-regular scheme: Edelsbrunner–Grayson (2000) proved its sub-simplices fall into a *bounded* number of similarity classes, so the aspect ratio is uniformly bounded across all levels. Concretely, in 2D the edgewise subdivision of a triangle yields four sub-triangles all *similar* to the parent, so the chunkiness $\rho = \text{diam}/r_{\text{in}}$ is exactly constant (3.464) at every level; in 3D the corner sub-tetrahedra stay regular ($\rho = 4.899$ always) and the central-octahedron pieces add only a bounded number of shapes under the canonical scheme (a naive recursive diagonal drifts, which is why the *canonical* Edelsbrunner–Grayson scheme is the shape-regular one). K3 is four-dimensional, so its edgewise tower is shape-regular by the $d = 4$ case — a known theorem, not an open problem. (Witness: `analysis/w33_edgewise_shape_regularity.py`.)

The chain closes to two narrow conditions. So the gravity arc — Pass 28 (open sector) $\rightarrow$ 29 (gravity *is* the spectral action) $\rightarrow$ 30 (manifold = K3) $\rightarrow$ 31 (4D convergence demonstrated) $\rightarrow$ 32 (collapse to one finite condition) $\rightarrow$ 33 (that condition is a theorem) — ends with *no open analytic theorem* in the chain. What remains is two narrow, well-defined conditions, neither an open theorem: **(1)** the candidate identification that the substrate’s continuum is specifically K3 (supported by $\chi = f$, $\sigma = -\mu^2$, $b_2^+ = q$, the lattice $2E_8 \oplus 3H = M_1$, the hyperkähler structure, and the heterotic–K3 dictionary), and **(2)** the confirmation that the BT984–1135 edgewise refinement is the canonical Edelsbrunner–Grayson subdivision. Given both, the chain runs unbroken:

$$q=3 \rightarrow W(3,3) \rightarrow \text{spectral action} = \text{gravity} + \text{SM} \rightarrow M^4 = \text{K3} \xrightarrow{[\text{shape-regular}]} \text{Dodziuk–Patodi} + \text{GH}/\text{FEEC} \rightarrow \text{propinquity}$$

The whole of physics — the Standard Model, cosmology, and gravity as the spectral action whose continuum is Ricci-flat K3 with positive Newton constant — follows from the single integer $q = 3$. (Witness: `analysis/w33_last_link_closed.py`.)

Honest scope: this is the closing *status*, not a claim of complete proof. The two remaining conditions are genuine — the K3 identification is a strong candidate, not a uniqueness proof, and the scheme confirmation is a combinatorial property of the construction, not re-verified here — and several analysis steps cite established theorems (Edelsbrunner–Grayson, Dodziuk–Patodi, FEEC, Gromov–Hausdorff, Latrémolière) applied to the substrate’s tower, with the Newton positivity assuming a standard positive cutoff. “Closed to two narrow conditions” means no *open* analytic theorem remains in the chain, the residuals being a candidate identification and a scheme confirmation — a faithful, honestly-bounded closing status.

C.52 Pass 34 — the system: the substrate as a computer and network architecture

Stripped of physics, the substrate is a complete computer and network. Read by the standard metrics of computer and interconnect engineering it is a fault-tolerant universal ternary machine, and its subsystem datasheet is, precisely, the von Neumann architecture of self-reproduction — the architecture of life.

The interconnect. The wiring is $GQ(3, 3) = SRG(40, 12, 2, 4)$: 40 nodes of radix (degree) 12, diameter 2 (any-to-any in two hops, deterministic via the line incidence), vertex and edge connectivity both 12 (it survives any eleven simultaneous failures, the maximum for degree 12), second eigenvalue $\lambda_2 = 2$ far below the Ramanujan bound $2\sqrt{11} = 6.63$ (a *better-than-Ramanujan* expander, near-optimal bisection), and vertex- and edge-transitive under $Sp(4, 3)$, self-dual as a generalized quadrangle (perfect load balancing, no hot spots), with 240 links. Against a 6×6 torus (radix 4, diameter 6) or a hypercube Q_6 (radix 6, diameter 6), it pays radix 12 to buy diameter 2 and maximal resilience — the low-diameter, high-resilience corner (flattened-butterfly/dragonfly class) with SRG regularity. (Witness: `analysis/w33_interconnect_network.py`.)

The processor. The core computes in balanced ternary — the most economical radix (the economy $E(b) = b/\ln b$ is minimised at e , and base 3 at 2.731 beats binary and base 4 at 2.885; the Setun choice), with the 3-grading $\{-1, 0, +1\}$ the balanced-ternary digit. Its native word is $27 = 3^3$ (a tryte of three qutrits, the E_6 register), and its instruction set is the universal minimum: degree-2 (Gaussian/Clifford, from $Sp(4, 3)$) gates are free and classically simulable, and a single degree-3 (cubic, the E_6 cubic form on the **27**) gate makes them universal (Lloyd–Braunstein). The power source is contextuality — the qutrit register is a Wigner phase space, and $W(3, 3)$ is a contextuality structure (its smallest contextual sub-configuration is the doily $GQ(2, 2)$), supplying the magic quantum advantage requires. (Witness: `analysis/w33_ternary_processor.py`.)

Table 5: The Holonet system datasheet: the substrate as a fault-tolerant universal ternary computer.

Subsystem	Specification	Substrate object
Processor	balanced ternary; word $27=3^3$; ISA deg-2 (Clifford) + deg-3 (magic) = universal	$q=3$ radix; E_6 27 ; $Sp(4, 3)$
Interconnect	40 nodes, radix 12, diameter 2, 11-fault, >Ramanujan expander	$GQ(3, 3)=SRG(40, 12, 2, 4)$
Memory / ECC	$[[66, 8, 3]]_3$ qutrit surface code, distance 3	genus-6 K_{12} face code
Clock	two-gap golden-ratio quasicrystal, self-clocking, beat = 30	Boerdijk–Coxeter / (ϕ)
Universality	Turing-complete (Rule-110 on the BC tape)	UTM tape mapping

The architecture of life. That datasheet (Table 5) is the von Neumann architecture of a self-reproducing automaton, which requires three things: a universal computer to read instructions, a universal constructor to build from a description, and a copyable error-corrected description (heredity). The substrate supplies all three — Rule-110/UTM (universal computation), the degree-2+degree-3 gate set with the network (universal construction), and the $[[66, 8, 3]]_3$ code (the error-corrected genome). So the substrate is the minimal architecture that computes, constructs, and corrects, with the three “selves” of life as its three layers: self-correcting (the code), self-replicating (von Neumann universal construction), self-similar (the quasicrystal clock and the $A_2 < D_4 < E_8$ lattice tower). (Witness: `analysis/w33_holonet_system.py`.)

Honest scope: the interconnect and radix-economy metrics are exact and computed; the gate-set universality (Lloyd–Braunstein) and contextuality-as-fuel (Howard et al.) are standard quantum-computing facts mapped onto the substrate’s $Sp(4, 3)$ /cubic/ $W(3, 3)$ structure; physical gate and wiring realisations are an implementation question. “Architecture of life” means the von Neumann

architecture of self-reproduction — the rigorous computational notion (universal computation + construction + error-corrected heredity) — not literal biochemistry.

C.53 Pass 35 — the datapath: the router, the clock, and the reliability budget

Pass 34 gave the datasheet; this pass quantifies its three working subsystems the way a systems engineer would specify them — the routing protocol on the fabric, the memory hierarchy with its clock, and the error budget of the store — and finds that every operational constant is a fixed substrate integer.

The router. The $GQ(3,3)$ fabric carries a table-free routing protocol. A node’s address is a projective point of $\mathbb{F}_3^4$ — four balanced-ternary digits — and the forwarding decision is a single ALU operation: two nodes are directly linked iff their symplectic inner product $B(x,y)$ vanishes, so the *address is the route* (the parabolic-router duality, no routing table). Diameter 2 makes every delivery one or two hops. The path diversity is exact and strongly-regular: between adjacent nodes the direct link plus $\lambda = 2$ length-2 detours; between non-adjacent nodes $\mu = 4$ internally-disjoint length-2 paths — native *four-way shortest multipath* for load spreading and instant failover — and by Menger’s theorem the vertex connectivity 12 gives twelve internally-disjoint paths in all, so the protocol survives any eleven simultaneous failures with guaranteed delivery. Oblivious minimal routing on a diameter-2 graph is deadlock-free with a single virtual channel, and vertex/edge-transitivity under $Sp(4,3)$ makes load perfectly balanced. So the protocol is a 4-trit address, a one-operation symplectic forwarding test, two-hop delivery, four-way multipath, twelve-way resilience — every constant $(2,4,12)$ an SRG parameter. (Witness: `analysis/w33_router_protocol.py`.)

The clock and memory. The memory is a conventional hierarchy — the $27 = 3^3$ tryte register (three qutrits, the E_6 word) at the top, the $[[66,8,3]]_3$ protected store (8 fault-tolerant logical qutrits in 66 physical) as the working set, and the Boerdijk–Coxeter / Fibonacci tape (the UTM tape) as the long store, addressed quasicrystallinely. The clock is *unconventional*: not a periodic crystal but a quasicrystal oscillator, a two-gap timing built from the golden ratio $\varphi = 1.618$ whose tick sequence is the Fibonacci word (substitution $S \rightarrow L, L \rightarrow LS$), aperiodic (no resonances — spread-spectrum, EMI-quiet) and self-similar at every scale. Its key engineering property is the *three-gap theorem*: the inter-tick intervals of the golden rotation take at most three distinct values at every horizon (verified here to $N = 500$), so the clock’s jitter is bounded to at most three intervals — a deterministic, low-discrepancy schedule ($O(\log N)$, optimal for a 1-D sequence, since φ is the most irrational number and so the slowest to resonate). The supercycle is beat = 30, the 600-cell Boerdijk–Coxeter ring. (Witness: `analysis/w33_memory_clock.py`.)

The reliability budget. The store’s error budget is clean. The code distance is 3, so it corrects one fault per cycle with certainty — every single-qutrit error has a unique syndrome (528 of them) and every weight-2 error is detected. The depolarizing reliability is the standard distance-3 quadratic, $P_L \sim A p^2$ with $A \sim \binom{66}{2} = 2145$, so the pseudo-threshold (the physical rate below which encoding helps, $P_L < p$) is $p_{th} = 1/A \approx 5 \times 10^{-4}$: below it the encoded store beats the bare qutrit (at $p = 10^{-4}$, $P_L \sim 2 \times 10^{-5}$, a $5\times$ gain), above it encoding hurts — the honest crossover. For the erasure / photon-loss channel a photonic build actually faces, surface codes tolerate $\sim 25\text{--}50\%$ erasure, so the $\sim 0.2\%$ -per-component optical loss leaves $\sim 78\text{--}88\%$ survival and the code runs by post-selection. And reliability scales: below threshold the logical error falls super-exponentially as the code is enlarged (the threshold theorem), the enlargement being the GKP concatenation tower $A_2 < D_4 < E_8$ (the moonshine ladder), driving the logical error to zero. With a 1 GHz clock and $p = 10^{-4}$ this already gives a multi-year mean time between logical failures. (Witness: `analysis/w33_reliability_threshold.py`.)

Honest scope: the routing constants ($\lambda = 2, \mu = 4$, connectivity 12) and the three-gap prop-

erty are exact and computed; the routing dictionary (header = address, forward = symplectic test, deadlock-freedom from diameter 2), the three-gap theorem, the golden-rotation $O(\log N)$ discrepancy, and the threshold theorem are standard results applied to the substrate objects. The coefficient $A \sim \binom{66}{2}$ is a generous upper count of weight-2 error pairs — the true value is smaller — so $p_{th} \approx 5 \times 10^{-4}$ is a conservative bound; the MTBF depends on the assumed clock and physical error rate.

C.54 Pass 36 — the silicon questions: the ISA enumerated, the power budget, the floorplan

Three questions a chip architect asks of any machine — what are the opcodes, what does it cost to run, and can it be laid out — and the substrate answers all three with fixed integers.

The ISA, enumerated. Pass 34 called the instruction set degree-2 (free) + degree-3 (universal); this pass *enumerates* it. The degree-2 (Clifford) layer is not an abstract “free” class but a concrete finite opcode group: the symplectic group $\text{Sp}(4, 3)$ acting on the substrate’s $\mathbb{F}_3^4$ phase space (the $W(3, 3)$ points). Building it by closure from its symplectic transvections — one shear $x \mapsto x + aB(x, v)v$ per substrate point (the 40 points of $\text{GQ}(3, 3)$) — gives a group of order *exactly* $51840 = |W(E_6)| = 3^4(3^2 - 1)(3^4 - 1)$: the free part of the ISA *is* the Weyl group of E_6 . The per-qutrit Clifford (mod Paulis) is the smaller $\text{Sp}(2, 3) = \text{SL}(2, 3)$ of order $24 = f$. On top sits a single non-Clifford macro-op, the degree-3 E_6 cubic form $\det(X)$ on the **27**-register (the unique E_6 cubic invariant); by Gottesman–Knill the Clifford group is exactly the classically-simulable opcodes, and by Lloyd–Braunstein the one cubic op lifts the set to universal. So the ISA is fully specified: a **27** = 3^3 operand register, a Clifford opcode group of order $51840 = |W(E_6)|$, and one magic opcode. (Witness: `analysis/w33_isa_encoding.py`.)

The power budget. A machine that claims to be the architecture of life must have a metabolism. From Landauer’s principle: first, balanced ternary costs nothing thermodynamically — the bound is $kT \ln 2$ per bit, a trit erasure costs $kT \ln 3$, but a trit carries $\log_2 3 = 1.585$ bits, so the cost per bit is $kT \ln 3 / \log_2 3 = kT \ln 2$ *exactly*; the radix is base-independent in energy, and base 3’s advantage is purely in digit count, not dissipation. Second, the only irreducible cost is error correction, and it is the metabolism: the degree-2 Clifford datapath is unitary, hence reversible and Landauer-free (Bennett), and the $\llbracket 66, 8, 3 \rrbracket_3$ genome is copied reversibly; the one irreversible step is syndrome extraction, which measures and resets $n - k = 58$ syndrome qutrits per cycle, exporting their entropy at $(n - k)kT \ln 3$ per cycle. That is exactly Schrödinger’s negative entropy / Prigogine’s dissipative structure — a living system stays ordered by exporting entropy, and the QEC cycle *is* that export. The metabolic rate is $\sim 2.6 \times 10^{-19}$ J/cycle at 300 K, a power of $\sim 2.6 \times 10^{-10}$ W per logical block at 1 GHz (and $\sim 9 \times 10^{-15}$ W at a 10 mK dilution-fridge temperature). (Witness: `analysis/w33_power_budget.py`.)

The floorplan. A topology is only as good as it is buildable, and the decisive number is the bisection width. For $\text{GQ}(3, 3)$ it is *exactly* 100: the spectral lower bound $(n/4)(k - \lambda_2) = (40/4)(12 - 2) = 100$ is met by an explicit balanced 20|20 cut of 100 edges (spectral + Kernighan–Lin), so the minimum bisection equals the spectral bound. That is 100 of the 240 links (42%) across any halving — enormous cross-sectional bandwidth, the hallmark of a good expander — but the same number is the cost: by Thompson’s VLSI theorem a 2-D layout needs area $\gtrsim (\text{bisection}/2)^2 \sim 2500$ wire-crossing units, and a graph whose bisection grows with its size is intrinsically *non-planar* (an expander cannot be laid out flat with short wires). Against a 6×6 torus (bisection 12, planar, diameter 6) or a hypercube Q_6 (bisection 32), $\text{GQ}(3, 3)$ ’s 100 buys diameter-2 all-to-all reach at the price of non-planarity. The resolution is the substrate’s own: a high-bisection, non-planar fabric is exactly what a *photonic* realisation is for — light routed in 3-D / free space (or in wavelength) has

no planar crossing penalty — so the optical Holonet is *forced* by the bisection, not a convenience. (Witness: `analysis/w33_noc_floorplan.py`.)

Honest scope: the group orders ($|\mathrm{Sp}(4, 3)| = 51840 = |W(E_6)|$, $|\mathrm{Sp}(2, 3)| = 24$) and the bisection (100, spectral bound met by an explicit cut) are exact and computed; Gottesman–Knill, Lloyd–Braunstein, the base-independence of the per-bit Landauer cost, Bennett reversibility, and Thompson’s area bound / expander non-planarity are standard. The substrate content is the identifications (Clifford group = $\mathrm{Sp}(4, 3) = W(E_6)$, magic = E_6 cubic, irreversible cost localised to the 58 syndrome qutrits, the bisection forcing a photonic medium); Landauer is a thermodynamic lower bound (real devices dissipate orders more), and the numerical power rate assumes the chosen clock and temperature.

C.55 Pass 37 — the boundary questions: complexity class, distributed clock, I/O, and one group

Four questions close the machine off: how hard is what it computes, how do forty nodes share one tick, how does it talk to the world, and what holds it together. The fourth answer subsumes the rest — one group is the whole machine.

The complexity class. The separating invariant is the Gross qutrit Wigner function. The degree-2 Clifford layer keeps it non-negative: all 12 single-qutrit stabilizer states have Wigner ≥ 0 (computed), so by Veitch–Mari–Emerson–Gross the Clifford datapath is efficiently classically simulable — it lives in $\mathbf{P}$ (qudit Gottesman–Knill), with zero mana and no contextuality. The degree-3 cubic magic gate breaks exactly this: its resource (the Strange state) has a negative Wigner entry $-\frac{1}{3}$ and mana $\ln(5/3) = 0.5108 > 0$, and by Howard–Wallman–Veitch–Emerson (2014) Wigner negativity *is* contextuality for odd-prime qudits and is necessary for speed-up. So Clifford alone = $\mathbf{P}$, Clifford + one cubic = $\mathbf{BQP}$ -universal, and the gap is exactly the mana the cubic injects — which $W(3, 3)$ supplies; classically sampling the magic-fuelled output collapses the polynomial hierarchy. (Witness: `analysis/w33_complexity_advantage.py`.)

The distributed clock. A network needs one tick shared by all 40 nodes. Diameter 2 makes clock distribution a two-hop broadcast. The natural distributed clock is neighbour-averaging $W = A/k$, whose second eigenvalue magnitude is exactly $\frac{1}{3}$ (the spectrum $\{12, 2, -4\}$ of A gives $\{1, \frac{1}{6}, -\frac{1}{3}\}$ for W), so the phase spread contracts threefold every round — geometric, expander-fast convergence (confirmed by simulation). Against silent faults the connectivity 12 keeps the clock connected through 11 failures; against malicious (Byzantine) faults the bounds are tighter and exact — $n \geq 3t+1$ (Lamport–Melliar-Smith / Dolev–Halpern–Strong) and connectivity $\geq 2t+1$ (Dolev) give $t = \min(\lfloor 39/3 \rfloor, \lfloor 11/2 \rfloor) = \min(13, 5) = 5$ traitors, the connectivity binding: 11 dead nodes but 5 lying ones, with $\mu = 4$ disjoint distribution routes. (Witness: `analysis/w33_clock_distribution.py`.)

The I/O and boundary. The readout interface is the line geometry: the 40 totally-isotropic lines of $\mathrm{GQ}(3, 3)$ (computed: 4-point cliques, 4 lines per point) are the measurement *contexts* — 40 maximal commuting sets, one readout basis each. The air-gap carrier is a photon’s orbital angular momentum in the sector $\{\ell = -1, 0, +1\}$, exactly the balanced-ternary digit, so one photon is one trit; by the Holevo bound a qutrit conveys at most $\log_2 3 = 1.585$ bits per use and the OAM-trit saturates it. The logical port is the code boundary: $[[66, 8, 3]]_3$ exposes $k = 8$ fault-tolerant logical qutrits ($= 8 \log_2 3 = 12.68$ bits/block, 12.68 Gbit/s at 1 GHz), the $n = 66$ interior protected. (Witness: `analysis/w33_io_boundary.py`.)

One group runs the whole machine. The deepest fact is that the *same* symmetry group — $W(E_6)$, order 51840, realised as the symplectic group on $\mathbb{F}_3^4$ (the order closed in the Pass-36 ISA witness, $|\mathrm{Sp}(4, 3)| = 51840 = |W(E_6)|$) — is simultaneously the processor’s Clifford gate group, the network’s automorphisms (it preserves collinearity and acts transitively on the 40 nodes,

verified), the memory’s code automorphisms (transitive on the 40 line-contexts), and the readout’s contextuality symmetry (those same 40 lines), and it is transitive even on the 160 point-on-line flags — the joint (register, readout-basis) states. So the computer *is* the network *is* the memory: one group, four faces (and, by Solovay–Kitaev, an efficient compiler target too). (Witness: `analysis/w33_one_group_machine.py`.)

Honest scope: the Wigner positivity, the mana $\ln(5/3)$, the $\frac{1}{3}$ contraction, the 40 line-contexts, and the transitivities (40 nodes, 40 lines, 160 flags) are computed; Veitch–Mari simulability, Howard’s negativity $\Leftrightarrow$ contextuality, the Byzantine bounds $n \geq 3t + 1$ / connectivity $\geq 2t + 1$, the Holevo bound, $\text{Aut}(\text{GQ}(3,3)) = W(E_6) = \text{PSp}(4,3):2$, and Solovay–Kitaev are established. The substrate content is the identifications and that one group realises all four roles; BQP is the standard placement of a universal quantum machine, the classical-hardness conditional on PH not collapsing, and bit/skew rates assume the stated clock.

C.56 Pass 38 — the advantage measured, the scheduler, and the falsifiable datasheet

Three closing moves: quantify the quantum advantage as a number, build the resource-management layer, and state the architecture as a testable specification.

The advantage, measured. Pass 37 placed the machine in BQP; this pass turns that into a computed separation via the *robustness of magic* $R(\rho) = \min\{\sum_i |x_i| : \rho = \sum_i x_i \sigma_i, \sigma_i \text{ pure stabilizer}\}$, the cost parameter of the Pashayan–Wallman–Bartlett quasiprobability simulator (cost $\sim R^2$ per resource). By linear programming over the 12 single-qutrit pure stabilizer states, the substrate’s degree-3 resource — the Strange state $(|1\rangle - |2\rangle)/\sqrt{2}$ — has $R = 3$ *exactly* (with an exact 6-term decomposition; the sanity check $R = 1$ for a stabilizer state). Since robustness is submultiplicative, a circuit with t cubic magic gates costs the classical simulator $\sim R^{2t} = 9^t$, exponential in the magic count, while the quantum machine runs in time $\sim t$: every cubic gate multiplies the classical cost by 9, and exact classical sampling collapses PH. So the resource is quantified — robustness 3, mana $\ln(5/3)$, a per-gate classical overhead of 9. (Witness: `analysis/w33_provable_advantage.py`.)

The scheduler. The 240 links partition *exactly* into the 40 lines, each a 4-clique (every edge on a unique line), so the wiring is 40 four-way buses. The collinearity graph is 1-factorable into 12 perfect matchings (computed), giving a conflict-free point-to-point schedule of exactly 12 slots (each node drives one of its 12 ports per slot), perfectly balanced by vertex/edge-transitivity. For readout, a spread of 10 disjoint lines covers all 40 nodes (so 10 contexts measure in parallel), but the 40 contexts do *not* resolve into 4 parallel spreads (proven by backtracking) — so the readout is not fully parallelizable, and that non-resolvability *is* measurement incompatibility, the contextuality that fuels the magic, appearing as a hard scheduling obstruction. (Witness: `analysis/w33_scheduler_os.py`.)

The falsifiable datasheet. The architecture fits on one consistency-checked ledger (the headline constants (40, 12, 2, 4), bisection 100, Holevo $\log_2 3$, optimal radix 3, Byzantine 5, mana $\ln(5/3)$ all re-derived and cross-checked here), and it separates two kinds of claim. Theorems (verifiable, not falsifiable): $|\text{Sp}(4,3)| = 51840 = |W(E_6)|$, diameter 2, connectivity 12, distance 3, bisection 100, Clifford = P / +cubic = BQP, 1-factorability. Predictions (falsifiable — a build could refute them, each with a criterion): radix-12 diameter-2 wiring; the OAM-trit saturating Holevo at 1.585 bit/photon; the $[[66, 8, 3]]_3$ pseudo-threshold near 5×10^{-4} ; the magic robustness 3/mana $\ln(5/3)$; and 5-Byzantine/11-crash tolerance. So the architecture is a complete, internally consistent, and *testable* computer specification. (Witness: `analysis/w33_architecture_capstone.py`.)

Honest scope: the robustness $R = 3$, the bus/1-factor/spread/non-resolvability facts, and the ledger constants are computed here (the heavy $|\text{Sp}(4,3)|$ closure is referenced from the Pass-36

witness); the Pashayan–Wallman–Bartlett R^2 cost and submultiplicativity, and the post-selection PH-collapse hardness, are established. The theorem/prediction split is the honest core: the graph- and group-theoretic facts are theorems; the physical predictions are what a hardware realisation could refute. “Testable” means the refutation criteria are operational, not that a build has been performed.

C.57 Pass 39 — the planetary computer: the universal VM, the everyone-node fractal, and the bench test

Three moves turn the 40-node machine into civilization-scale infrastructure, and a companion paper (`holonet_practical_implications.tex`) works out the implications for data centers, decentralized compute, virtualization, and energy.

The first milestone, cut loose. The headline prediction becomes a costed benchtop experiment. The substrate predicts a contextual fraction $CF = 1/10 = 1/\Phi_4$; distinguishing it from the classical null $CF = 0$ is a one-sided binomial test needing $n \geq k^2 p(1-p)/p^2$ valid coincidences — 81 for 3σ , 225 for 5σ , or ~ 256 –289 raw photons at 78–88% optical survival. Catalog optics, no cryostat, calibration-free: a violation fraction near $1/10$ certifies the $q = 3$ contextuality (the magic fuel), near 0 refutes the identification. (Witness: `analysis/w33_demonstrator_experiment.py`.)

The planetary computer. The fractal law H_n (replace each of the 40 points by H_{n-1}) gives 40^n leaf cores, $(40^n - 1)/39$ instances, and routing diameter exactly $8n = 8 \log_{40} N$ (verified). Seating N needs level $\lceil \log_{40} N \rceil$: humanity (8×10^9) and all devices (3×10^{10}) both fit at level 7 (1.64×10^{11} slots), diameter 56 hops, with $W(E_6)$ transitive on the leaves so no node is privileged — consensus is structural, a node joins by splicing. And it boots on legacy hardware: the Clifford layer is the Gottesman–Knill stabilizer formalism (polynomial), so the whole architecture runs as a classical emulator on any computer; only the quantum advantage costs more, 9^t for t cubic gates, a knob from $t = 0$ (a fully classical node) upward. (Witness: `analysis/w33_fractal_datacenter.py`.)

One number, two worlds. The engineering datasheet and the Standard Model are arithmetic in the same nine integers (forced by $q! = 2q$): $q = 3$ is the radix and the three generations; $k = 12$ the network radix and $\sin^2 \theta_W = 3/8$; $f = 24$ the per-qutrit Clifford order and $1/\alpha = 137 = \Phi_4 \Phi_3 + \Phi_6$; $|W(E_6)| = 51840 = 24 \cdot 2160$; bisection 100; beat $30 = h(E_8)$; $1/10 = 1/\Phi_4$. So one benchtop measurement tests the computer and the cosmos at once. (Witness: `analysis/w33_physics_bridge.py`.)

Honest scope: the shot-count decision statistics, the fractal scaling, the 9^t emulation cost, and the shared-integer arithmetic are computed; the contextual fraction $1/10$, the fractal substitution law (BT827), and the physics identifications are corpus results used as inputs, several at integer level. “Boots on a dinosaur” means the architecture layer is polynomial-time classical, not that legacy hardware delivers quantum advantage.

C.58 Pass 40 — the machine made runnable: the emulator and three frontier applications

This pass ships the architecture as a *runnable* program and pushes three frontier applications from conjecture to verified witnesses; the companion paper `holonet_practical_implications.tex` collects the practical case.

The universal VM, executable. `analysis/holonet_node.py` is a runnable holonet node: it builds the 40-node $GQ(3, 3)$ fabric and routes a packet in two hops (address is route, $\mu = 4$ multipath); runs a qutrit register as a Clifford stabilizer tableau over $\mathbb{F}_3$ that evolves by the symplectic group in *polynomial* time (a Fourier+SUM circuit builds a valid entangled qutrit-Bell stabilizer

state, checked by mutual commutation; gate timing stays flat from 4 to 32 qutrits); exposes the magic dial (classical cost 9^t); and self-reproduces by splicing a $W(3,3)$ child (address +1 digit, diameter +8). That the file executes on a laptop is the demonstration: the architecture of life boots on ordinary hardware, and the network of many such emulators is, as one $W(E_6)$ -symmetric object, the planetary computer.

Mining-free consensus. `analysis/w33_consensus_ledger.py`: $W(E_6)$ vertex-transitivity makes a ledger leaderless (leader election dissolves); agreement is two hops plus a 1/3-per-round averaging contraction (~ 19 rounds for a part in 10^9), Byzantine-safe to 5. Proof-of-work pays $\sim 10^9$ J/transaction (Bitcoin ~ 150 TWh/yr); the holonet’s floor is the $\sim 2.6 \times 10^{-19}$ J syndrome export — ~ 27 orders lower, security from geometry not burned power.

The substrate code is a DNA code. `analysis/w33_dna_storage_code.py`: a homopolymer-free DNA encoding has density $\log_2 3 = 1.585$ bit/nt — the substrate’s exact Holevo/OAM alphabet — and the balanced-ternary digit $\{-1, 0, +1\}$ maps to a homopolymer-free, GC-balanced base alphabet (verified on a 200-trit strand), with $[[66, 8, 3]]_3$ supplying distance-3 correction. A substrate-encoded strand *is* a holonet memory block.

The quantum-internet specification. `analysis/w33_quantum_internet.py`: gate teleportation (the Stage-B self-measurement) is a store-and-forward entanglement swap; diameter 2 means one swap connects any pair, with $\mu = 4$ redundant swap paths ($\sim 89\%$ per-hop survival, essentially 1 with multipath) and an $8 \log_{40} N$ fractal backbone — a candidate specification (topology, routing, code) for the repeater networks now being built (Zheng et al. 2025; Photonic-TELUS; Qunnect-Cisco).

Honest scope: the emulator’s routing, polynomial-time Clifford evolution, validity check, and fractal growth are computed and run; the consensus graph facts, the DNA density and homopolymer-free mapping, and the swap/multipath/backbone counts are computed. The applications themselves are directions, not built systems: the 9^t magic layer is not executed (it is the quantum-advantage regime), the energy comparisons use external estimates and a thermodynamic floor, the DNA match is an encoding hypothesis, and the quantum-internet mapping is a specification. “Boots on a dinosaur” means the architecture layer is poly-time classical, not that legacy hardware yields quantum advantage.

C.59 Pass 41 — proof of life: the machine corrects, teleports, and reproduces, executed

Pass 40 made the architecture runnable; this pass runs the three things the von Neumann “architecture of life” claim requires — compute-and-correct, construct-and-move, reproduce — as exact, verified programs.

It corrects errors. `analysis/holonet_qec_demo.py`: an exact state-vector run of the qutrit $[[5, 1, 3]]_3$ perfect code (a runnable stand-in for the substrate’s $[[66, 8, 3]]_3$, same distance-3 mechanism). The four stabilizers (cyclic shifts of X, Z, Z^{-1}, X^{-1}, I) commute; the codespace is 3-dimensional (one logical qutrit); all 40 single-qutrit Pauli errors give 40 distinct syndromes; and every error is decoded and corrected to fidelity 1. The memory is a running code, not a parameter triple.

It teleports. `analysis/holonet_teleport_demo.py`: two nodes share a qutrit Bell pair; A Bell-measures the message, sends two trits, B applies $X^k Z^l$, and recovers the message at fidelity 1 for all 9 outcomes (exact simulation). The message is destroyed at A (no-cloning $\Rightarrow$ teleport-only VM migration); chained, it is entanglement swapping — the quantum-internet repeater primitive executed.

It reproduces. `analysis/holonet_quine.py`: the program carries its own source as an internal

description (a genome); running it constructs a byte-identical child, which constructs an identical grandchild — a verified self-reproduction fixed point. Interpreter = universal computer, the write = universal constructor, the genome string = copyable description, the $[[66, 8, 3]]_3$ code (just shown correcting errors) = error-corrected heredity, the fractal $W(3, 3)$ splice = variation.

Honest scope: the QEC and teleportation are exact, ideal-syndrome state-vector simulations (no measurement error/loss); $[[5, 1, 3]]_3$ stands in for the larger surface code (same mechanism). The quine fixed point is real and verified; it demonstrates the *logical* architecture of self-reproduction (von Neumann’s three components), not biology. Everything here is classical, exact, and runs today; the quantum advantage remains the priced 9^t dial.

C.60 Pass 42 — the benchtop number derived, and the quantum packet delivered

This pass closes the one gap the computer-engineering arc kept citing rather than deriving — the contextual fraction $1/10$ — and assembles the runnable pieces into one VM and one end-to-end quantum delivery.

The contextual fraction, from first principles. `analysis/w33_contextual_fraction.py`: the 40 Witting rays have the orthogonality structure of $W(3, 3)$, with 40 measurement contexts = the totally-isotropic lines (tetrads). A Kochen–Specker assignment satisfying all 40 contexts would be an ovoid (10 points meeting every line once); $W(q)$ for q odd has no ovoid (Thas), confirmed here by the maximum partial ovoid (independence number) being $7 < 10$. By exact integer programming over the 40 rays, the most contexts any global 0/1 assignment can satisfy is exactly 36, so the contextual fraction is $CF = (40 - 36)/40 = 4/40 = 1/10 = 1/\Phi_4$. The 36/40 KS budget and the $1/10$ fraction are theorems of the $q = 3$ line geometry, not inputs — so the single number the demonstrator measures is derived from scratch.

The VM does all three. `analysis/holonet_node.py` now exposes `correct()` (a full $[[5, 1, 3]]_3$ cycle, fidelity 1), `teleport()` (a qutrit moved to a partner, the two classical trits routed over the fabric, fidelity 1), and `reproduce()` (a fractal $W(3, 3)$ splice) as node methods — one runnable object that computes-and-corrects, constructs-and-moves, and reproduces.

A quantum packet, delivered. `analysis/holonet_quantum_packet.py`: an unknown qutrit is teleported between nodes anywhere on the fabric, with the two classical correction trits routed by the symplectic address-is-route test in at most two hops ($\mu = 4$ multipath); across several source/destination pairs the payload arrives with fidelity 1, the quantum content never traversing the classical wire (no-cloning). For non-adjacent endpoints the shared Bell pair is established by one entanglement swap at the relay — the diameter-2 repeater.

Honest scope: the contextual fraction is the logical/possibilistic one (max contexts admitting a consistent global assignment), computed from the incidence geometry and equal to the corpus $1/10$, here derived; the QEC and teleportation are exact ideal state-vector simulations; the classical routing is exact; the entanglement-swap pair establishment is modelled (the corpus Stage-B primitive). The quantum advantage remains the priced 9^t dial.

C.61 Pass 43 — consensus run against an adversary, the KS inequality, and the holonet CLI

This pass executes the consensus layer against a live adversary, turns the contextual fraction into an explicit inequality with a number, and packages the whole VM as a command-line tool.

Consensus, executed. `analysis/holonet_consensus_demo.py`: on a simulated 40-node fabric, neighbour-averaging contracts the disagreement by $1/3$ per round (measured), five silent nodes do

not stop agreement (connectivity 12 survives 11 crashes), and five *malicious* nodes broadcasting oscillating ± 5 values are outvoted — each honest node runs the trimmed-mean (MSR) rule (discard the five highest and five lowest neighbour values) and the honest nodes still converge. Six malicious nodes break it, so the boundary sits exactly at $t = 5$, the Dolev connectivity bound $2t + 1 \leq 12$. Leaderless, two-hop, threefold-per-round, eleven-crash / five-Byzantine agreement, run not asserted.

The KS inequality, with a number. `analysis/w33_ks_inequality.py`: over the 40 contexts let S count those returning exactly one outcome. A noncontextual model pre-assigns each ray a fixed 0/1 value, and by exact integer programming the maximum is $S \leq 36$ (the best classical strategy lights 11 rays and fails exactly 4 contexts); quantum mechanically every context is an orthonormal basis, so $S = 40$ with certainty (state-independent). The violation is $40 - 36 = 4$, normalized $(40 - 36)/40 = 1/10 = 1/\Phi_4$ — the contextual fraction is now an explicit, state-independent inequality the bench reads off: measure all 40 contexts, and any $S > 36$ refutes noncontextual realism.

The holonet CLI. `analysis/holonet_cli.py` packages the VM as a tool: `holonet route A B, teleport, correct, reproduce, verify` (self-tests the whole stack to PASS/FAIL), and `info` (the datasheet). The universal machine as something anyone can install and run.

Honest scope: the $1/3$ contraction is the measured spectral fact; the crash and Byzantine runs are executed simulations with explicit adversaries (the $t = 5$ Byzantine convergence is empirical for the MSR rule, matching the connectivity bound, not a general robustness proof). The KS inequality’s classical bound 36 and four failed contexts are computed exactly; the quantum value 40 is the state-independent exclusivity of orthonormal-basis measurement on the Witting realisation. The CLI is a wrapper over the classical/exact VM; the quantum advantage remains the priced 9^t dial.

C.62 Pass 44 — the experiment simulated, the magic dial turned, the tool installed

This pass runs the demonstrator in simulation with error bars, executes the 9^t magic emulation that was only ever priced, and ships the VM as an installable command with a test suite.

The demonstrator, simulated. `analysis/holonet_ks_experiment.py`: the Cabello–Severini–Winter witness $\chi = \sum_v P(v \text{ fires in its context})$ obeys $\chi \leq \alpha = 7$ for any noncontextual model ($\alpha =$ the independence number of the $W(3, 3)$ exclusivity graph, computed), while quantum mechanically the 40 projectors sum to $10\mathbb{I}$, so $\chi = 10$ for any state — a gap 10 vs 7 tolerating 30% depolarizing noise. Simulating the heralded-photon run (maximally-mixed input, shot noise, visibility, dark counts), only ~ 21 detected events per ray (~ 840 photons) clear the noncontextual bound at 5σ , with $\hat{\chi} \rightarrow 10$. The first milestone is now a simulated experiment with error bars.

The magic dial, turned. `analysis/w33_magic_dial.py`: the Pashayan–Wallman–Bartlett quasiprobability simulator represents the Strange state by its robustness decomposition ($R = 3$, computed) and estimates expectations by a signed Monte Carlo, verified unbiased against the exact value for $t = 1, 2, 3$ magic qutrits. Each t -gate sample lies in $[-3^t, 3^t]$ (the realized maximum is exactly 3^t), so by Hoeffding the sample budget for fixed precision scales as $R^{2t} = 9^t$ — the priced advantage, executed and measured, not asserted.

The tool, installed. `holonet_cmd.py` plus a `console_scripts` entry in `setup.py` make `holonet` a real command after `pip install -e .`; and `tests/test_holonet_vm.py` gives the whole VM CI coverage (routing, the Clifford processor, error correction, teleportation, self-reproduction, the CLI, contextuality) — 13 exact tests, all passing.

Honest scope: the noncontextual bound $\alpha = 7$ and the magic robustness $R = 3$ are computed; the quantum value $\chi = 10$ is the state-independent Hoffman/Lovász value on the Witting realisa-

tion; the KS experiment models standard heralded-photon shot statistics; the magic-dial estimator is executed and verified unbiased with the 9^t Hoeffding sample bound from the per-sample range. The quantum advantage itself is still the 9^t dial — now run on this machine for small t .

C.63 Pass 45 — the threshold curve measured, the quickstart, and the unified Machine paper

This pass measures the fault-tolerance threshold as a curve, ships a five-minute quickstart, and collects the Pass 34–44 results into one submission-grade paper.

The threshold, measured. `analysis/holonet_threshold_demo.py`: an efficient symplectic decoder for the runnable $[[5, 1, 3]]_3$ code (errors and stabilizers as $\mathbb{F}_3$ vectors, syndrome = symplectic inner product, a lookup decoder covering all 81 syndromes, a logical fault declared when the residual has trivial syndrome but is not a stabilizer). Monte-Carlo over the depolarizing rate gives a *measured* curve $P_L(p) \sim Ap^2$ with $A \approx 8$, crossing break-even $P_L = p$ at a pseudo-threshold $p_{th} = 1/A \approx 0.12$ — the honest crossover, now a point on a plot. Syndrome-measurement noise at the same per-round rate removes the single-round break-even (motivating repeated extraction). The $[[5, 1, 3]]_3$ value (~ 0.12) is the small stand-in; the substrate's $[[66, 8, 3]]_3$ has $A \sim \binom{66}{2} = 2145$, hence $p_{th} \sim 5 \times 10^{-4}$ — same p^2 mechanism, different size.

The quickstart. `HOLONET.md`: install the `holonet` command, run `holonet verify`, and a tour of every runnable witness (route / correct / teleport / reproduce / consensus / threshold / contextuality / magic dial) — zero to a verified, running node in five minutes.

The unified Machine paper. `holonet_machine.tex` collects the computer-engineering arc (Pass 34–44) into one self-contained, submission-grade document with the one-object figure: the one group, the enumerated ISA, the non-planar floorplan, the measured threshold, the I/O boundary, the P-to-BQP complexity placement with the derived 1/10 contextual fraction, the runnable VM, consensus, and the falsifiable-predictions table.

Honest scope: the decoder and the logical-error test are exact; $P_L(p)$ is a Monte-Carlo estimate with the Ap^2 fit and break-even read off (single-round; circuit-level faults are future work); the Machine paper restates computed results and cites the standard theorems. The quantum advantage remains the priced 9^t dial.

C.64 Pass 46 — the minimal machine, the matched efficiency, the circuit-level threshold, and the scorecard

This pass answers a sharp question about the runnable VM — what is the least hardware that runs it, and in what precise sense is it efficient? — and adds the circuit-level threshold, a plotted scorecard, and continuous integration.

The minimal machine. `analysis/w33_minimal_architecture.py`: a holonet node's routing and Clifford layer needs only arithmetic modulo 3. The forwarding decision is the symplectic form $B(x, y) \bmod 3$ — seven ternary operations, no routing table (the address is the route) — and the Clifford register is a trit tableau updated by mod-3 additions. So the entire instruction set is three primitives $\{\text{ADD}_3, \text{SUB}_3, \text{MUL}_3\}$, with no FPU, cache, OS, or bus. We build exactly that machine (a ternary ALU VM) and run the routing on it, reproducing the reference adjacency for all 1600 node pairs in seven ops each; the footprint is tens of bytes (the 40 addresses are generated, not stored). The node fits an 8-bit microcontroller or a 1970s 4-bit CPU and runs *most* efficiently on a balanced-ternary computer (Setun), where the $\{-1, 0, +1\}$ digit is native.

Efficiency by matching. `analysis/w33_vm_speedup.py`: the VM cannot out-compute its substrate, but it is more efficient *per logical operation* than the general-purpose CPU+OS+network

stack it runs on, for data-movement-bound workloads, because it eliminates the abstraction tax. Routing state is 0 versus a conventional $N^2 = 1600$ -entry table; the forwarding op is a constant seven mod-3 operations; a logical operation crosses the bus once (route = gate = memory) versus the von Neumann ~ 5 (fetch, fetch, compute, write back, send); and it needs no cache, MMU, hypervisor, or routing protocol because the geometry supplies addressing, isolation (Schur), and scheduling ($W(E_6)$ transitivity) for free. ASIC-style efficiency from architectural matching, not an asymptotic speedup.

The circuit-level threshold and the scorecard. `analysis/holonet_ft_threshold.py`: with one noisy syndrome round there is no break-even, but measuring each syndrome trit three times and majority-voting restores the threshold ($p_{th} \sim 0.06$) and five rounds deepen it — the standard repeated-extraction gadget, measured. `analysis/holonet_scorecard.py` plots both headline numbers (the threshold curve and the contextuality significance) as a figure, now embedded in the companion Machine paper. A GitHub Actions workflow runs `holonet verify`, the fast witnesses, and the test suite on every push.

Honest scope: the minimal-architecture footprint and the routing/efficiency counts are exact (the von Neumann ~ 5 crossings is the standard architectural model); the circuit-level threshold is Monte-Carlo with the exact decoder under phenomenological noise (a full faulty-gate model is more detailed). “More efficient than the hardware it runs on” means more efficient per logical op than the general-purpose stack, by bypassing its overhead — not faster than physics; the quantum advantage remains the priced 9^t dial.

C.65 Pass 47 — the router on a 4-bit CPU, the ternary encoding tax, and the live scorecard

This pass makes “boots on a dinosaur” literal, quantifies what running a ternary algorithm on binary hardware costs, and wires the scorecard and CI status into the live site.

The router on a 4-bit CPU. `analysis/w33_tritcpu_emulator.py`: we build a minimal Intel-4004-flavoured machine (sixteen 4-bit registers, a small RAM, opcodes $\{\text{LDI, LD, ST, ADD, SUB, MUL, MOD}_3, \text{HLT}\}$), assemble the holonet forwarding test $B(x, y) \bmod 3$ as a 22-instruction program (every value stays in 0..15, so a 4-bit datapath suffices), and execute it with a full trace. It reproduces the reference adjacency for all 1600 node pairs in 22 cycles; on the real 1971 Intel 4004 (~ 740 kHz) that is $\sim 240 \mu\text{s}$ per routing decision — about four thousand holonet routes per second on a fifty-year-old four-bit chip. The classical layer runs on hardware older than its users; only the quantum advantage needs photons.

The ternary encoding tax. `analysis/w33_ternary_energy.py`: a balanced trit carries $\log_2 3 = 1.585$ bits, but a binary machine needs 2 bits per trit (three of four states used), so the bit-traffic tax is $2/\log_2 3 = 1.26$ (26% overhead, 25% wasted states), plus a $\sim 5\%$ digit-count tax from the radix economy. Crucially, Landauer’s bound is base-independent ($kT \ln 2$ per bit = $kT \ln 3$ per trit), so there is *no* fundamental energy win from ternary — the $1.26\times$ is strictly the *encoding* overhead a binary host pays and a native ternary substrate (the Setun digit) would not. That is the exact, honest sense in which the holonet “wants ternary hardware.”

The live scorecard. The scorecard figure and the CI status badge are now embedded in `docs/index.html`: the site shows the green build and the measured threshold / contextuality plot.

Honest scope: the CPU emulator is a faithful 4-bit register machine (program verified on all 1600 pairs; the 4004 timing is the historical order-of-magnitude figure); the encoding-tax numbers are exact arithmetic, and the honesty is that ternary gives no fundamental per-information energy advantage (Landauer is base-independent) — the tax is encoding overhead only. The quantum advantage remains the priced 9^t dial.

C.66 Pass 48 — the assembler, packet energy, and the live router

This pass turns the runnable VM into a small compiler and makes the energy accounting packet-scale.

The assembler. `analysis/w33_holonet_asm.py`: the router

$$B(x, y) = x_0y_1 - x_1y_0 + x_2y_3 - x_3y_2 \pmod{3}$$

is now compiled through a tiny *holonet-asm* layer to two targets. The first is the Pass 47 four-bit machine, still a 22-instruction program with primitive MUL and MOD₃. The second is stricter: a 6502-style eight-bit accumulator target with no multiply and no mod opcode. It synthesizes multiplication and reduction from loads, stores, add, subtract, compare, and branches, then verifies the resulting program on all 1600 ordered pairs of $W(3, 3)$ addresses. The worst case executes 185 primitive accumulator steps. This is not a cycle-accurate MOS 6502 emulator; it is the compiler-level statement that the classical router needs only the ordinary integer-control substrate of an old eight-bit machine.

The packet-scale energy bill. `analysis/w33_packet_energy.py`: Pass 47 gave the per-trit binary encoding tax $2/\log_2 3 = 1.262$. Pass 48 pushes it through the minimal Holonet control packet. A worst-case two-hop route reads two four-trit addresses at each hop (16 trit reads), the Holonet body contributes 16 Q6 edges times 3 pulse phases (48 trits), and the Hesse/Pauli epilogue contributes 3 two-trit Hesse words plus one two-trit Pauli frame (8 trits):

$$16 + 48 + 8 = 72 \text{ trits.}$$

So a native ternary host moves 72 trit symbols, while a binary host stores or moves 144 bits to carry $72 \log_2 3 \approx 114.1$ bits of information. The missing ~ 29.9 bits are not physics; they are the binary encoding tax made visible at packet scale.

The live router. `docs/index.html` now includes a self-contained JavaScript widget: type two normalized four-trit projective addresses and it generates the 40 $W(3, 3)$ points, computes $B(x, y)$, and returns the direct edge or the $\mu = 4$ common relays. The site therefore contains the same compiler primitive the papers discuss: address is route, in the browser.

Honest scope: the assembler certifies the classical router/compiler layer only; quantum advantage still requires the photonic/magic layer. The eight-bit target is 6502-style accumulator semantics, not a hardware timing benchmark. The 72-trit packet bill is a traffic model for the minimal control envelope; payload traffic scales linearly on top.

C.67 Pass 49 — exported retro targets and golden traces

Pass 48 showed that the router has a compiler semantics. Pass 49 makes that semantics into files a reader can inspect and diff. The witness `analysis/w33_holonet_retro_export.py` writes a small artifact bundle under `artifacts/holonet_asm/`:

4004-style	22 instructions	<code>router_4004_style.asm</code>
6502-style	130 instructions	<code>router_6502_style.asm</code>
Z80-style	111 instructions	<code>router_z80_style.asm</code>

Each listing computes the same symplectic routing bit

$$B(x, y) = x_0y_1 - x_1y_0 + x_2y_3 - x_3y_2 \pmod{3},$$

and the artifact bundle includes golden traces for the sample route

$$1000 \longrightarrow 0100, \quad B = 1.$$

The Z80-style target is an independent check, not a re-label of the 6502 program: it is assembled to its own accumulator/control-flow instruction set and verified on all 1600 ordered $W(3,3)$ address pairs, with no MUL and no MOD₃ opcode. Multiplication and reduction are compiled into load, store, add, subtract, compare, and jump.

The architectural interpretation is important. The Holonet router has crossed from “there exists an emulator” to “there exists a bootstrapping compiler surface.” The smallest layer is the ternary algebra {ADD₃, SUB₃, MUL₃}; above it, ordinary binary-era control machines can synthesize the same route decision; above that, the browser widget executes the same primitive interactively; and above that, the photonic packet layer supplies the quantum timing and magic. The network stack is therefore fractal: the same $B(x,y)$ form appears as machine code, routing rule, packet admission test, and projective geometry.

Honest scope: the exported files are canonical Holonet target listings, not vendor assembler promises. They prove deterministic compilation and traceability of the classical control layer; they do not claim cycle-accurate performance on historical 4004, MOS 6502, or Z80 silicon.

C.68 Pass 50 — firmware-to-fabric accounting

The last missing practical bridge was not another emulator. It was an accounting identity from firmware all the way up to the global tomosphere/Witting fabric. Pass 48 gave the minimal packet:

$$72 = 16 \text{ route trits} + 48 \text{ body trits} + 8 \text{ epilogue trits.}$$

Pass 49 gave the exported route-decision programs. The typed packet ABI already contains the global mirror bus:

$$2160 = 45 \cdot 48, \quad 51840 = 24 \cdot 2160.$$

`analysis/w33_holonet_firmware_fabric_profile.py` joins those facts and verifies the cross-layer closure

$$2160 = 30 \cdot 72, \quad 51840 = 720 \cdot 72.$$

A complete mirror atlas is exactly thirty minimal packet frames, and a full $Sp(4,3)$ runtime supercycle is exactly seven hundred twenty packet frames. This is the system-level meaning of the Boerdijk–Coxeter/E8 beat: the global bus does not float above the packet machine; it is an integral multiple of the packet clock.

The same witness pushes the firmware counts through that accounting. A worst-case packet asks for two route decisions, so the exported targets cost

target	listing instructions/decision	static instructions/packet	static instructions/mirror atlas
4004-style	22	44	1320
6502-style	130	260	7800
Z80-style	111	222	6660

with dynamic accumulator-step upper bounds recorded separately for the eight-bit targets. The same packet on a binary host moves 144 bits; a filled mirror atlas therefore corresponds to $30 \cdot 144 = 4320$ binary host bits, while the ternary information content is only about 3423.5 bits. That gap is precisely the encoding tax, now visible at the bus epoch rather than at one digit.

The Witting admission desk also becomes a traffic statement. The logical ROM splits

$$1600 = 520 \text{ accepted} + 1080 \text{ retry-shadow}, \quad \frac{520}{1600} = \frac{13}{40}.$$

Therefore a random ordered Witting query opens a full packet with probability $13/40$, and the expected full-frame traffic is

$$72 \cdot \frac{13}{40} = \frac{117}{5} = 23.4$$

trits per query. The retry-shadow sector is not ignored; it is now counted as the complementary $27/40$ admission shadow and paired with the symbolic BT1706 retry economics. At the nominal rates $p_{\text{loss}} = 10^{-5}$, $p_{\text{dark}} = 10^{-6}$, and $p_{\text{parity}} = 5 \times 10^{-6}$, the combined retry-exit probability is $64p_{\text{loss}} + 8p_{\text{dark}} + 24p_{\text{parity}} = 7.68 \times 10^{-4}$, below one tenth of one percent per packet.

Honest scope: this pass is deterministic accounting, not a measured latency benchmark. The retro listings are canonical target programs rather than vendor-cycle timing claims; the packed mirror-atlas capacity assumes disjoint typed slots; and the retry numbers use the nominal symbolic rates of the existing runtime model. What is proved is the architectural fact that the firmware router, 72-trit packet, 2160-slot mirror bus, and 51840-slot supercycle are one integer clockwork.

C.69 Pass 51 — the machine audits itself, and the live playground

Every prior pass added a witness; this pass closes the loop by making the whole datasheet re-derive itself in one command. The Holonet’s signature claim is that the device specification *is* its own audit: there is no separate trusted checker, because every architectural constant is forced by the geometry and can be recomputed from the single integer $q = 3$. Pass 51 makes that literal. The witness `analysis/w33_master_audit.py` recomputes — it does not store — the headline constant of every layer, in one pass, into a single pass/fail ledger:

network	SRG(40, 12, 2, 4); diameter 2; $\lambda_2 = 2$; bisection $100 = \frac{n}{4}(k - \lambda_2)$
processor	runtime $51840 = 24 \cdot 2160 = W(E_6) $; 40 totally-isotropic line-contexts
contextuality	max partial ovoid 7; max satisfiable contexts $36/40 \Rightarrow \text{CF} = 1/10$; CSW $\chi = 10 > 7$
magic	robustness $R = 3$ (LP); mana $\ln(5/3)$
fault tolerance	distance-3 break-even at $p = 0.05$; Byzantine $t = 5 = \min(\frac{n-1}{3}, \frac{k-1}{2})$
I/O & minimal	Holevo $\log_2 3$; forwarding $= 7 \bmod 3$ ops, zero table; ternary tax $2/\log_2 3$

Sixteen checks, all passing. Crucially, none of these is a stored number: the audit rebuilds $W(3, 3)$ from $q = 3$, recomputes the spectrum, solves the contextuality ILP and the magic-robustness LP, runs a small distance-3 Monte-Carlo, and only then compares. Running the audit is therefore exactly the act of trusting the machine — if the geometry is what the specification says it is, every layer’s headline constant follows. It is exposed as a command, `holonetaudit`, and wired into continuous integration, so the entire datasheet is re-verified on every push.

The same pass makes the architecture a no-install playground. The page `docs/holonet.html` runs the classical machine live in a browser in vanilla JavaScript: route a packet (the address *is* the route, two hops, $\mu = 4$ multipath), drive the qutrit Clifford register (Fourier and SUM gates with a symplectic commutation check), run the Cabello–Severini–Winter contextuality witness as a shot count converging on the 7 versus 10 gap, and splice a self-reproducing child. The JavaScript uses the identical symplectic form

$$B(x, y) = x_0y_1 - x_1y_0 + x_2y_3 - x_3y_2 \pmod{3}$$

that the Python witnesses verify against, so the browser demonstration and the audited specification are the same arithmetic.

Honest scope: fifteen of the sixteen audit checks are exact recomputations; the distance-3 break-even check is a small Monte-Carlo and is therefore statistical, and the heavy $|\text{Sp}(4, 3)| =$

51840 closure is checked here through its arithmetic factorization $24 \cdot 2160 = |W(E_6)|$ rather than re-enumerated (the full closure remains in `analysis/w33_isa_encoding.py`). The physics identifications and the quantum-advantage values are re-stated, not re-measured, by the audit; they keep their own witnesses and the priced 9^t dial. What Pass 51 proves is the self-auditing property: one command, the whole datasheet, re-derived and checked from $q = 3$ alone.

C.70 Pass 52 — the parity law, and the performance face

Pass 51 made the datasheet re-derive itself from $q = 3$. Pass 52 asks the two questions that audit invites: *is $q = 3$ forced*, and *how fast does the machine run*.

The parity law (is $q = 3$ forced?). The witness `analysis/w33_audit_qscan.py` runs the same audit across the sister geometries $W(q) = \text{GQ}(q, q)$ for $q = 2$ and $q = 3$ and verifies that every layer constant moves with q exactly as its closed form predicts:

$$\begin{array}{lll} n = (q+1)(q^2+1) & k = q(q+1) & \lambda = q-1 \\ \mu = q+1 & \lambda_2 = q-1 & \text{points/line} = q+1 \\ |\text{Sp}(4, q)| = q^4(q^2-1)(q^4-1) & \text{Hoffman} = q^2+1 & \text{lines} = n \end{array}$$

giving $(n, k, \lambda, \mu) = (15, 6, 1, 3)$ at $q = 2$ and $(40, 12, 2, 4)$ at $q = 3$, and $|\text{Sp}(4, 2)| = 720$ against $|\text{Sp}(4, 3)| = 51840$. The headline is the contextuality. By Thas's theorem $W(q)$ possesses an ovoid — a set of $q^2 + 1$ pairwise non-collinear points meeting every line once — *if and only if q is even*. An ovoid is precisely a Kochen–Specker 0/1 assignment satisfying all n contexts, so the audit finds, computationally:

$$q = 2 \text{ (even)} : \text{ovoid of 5 exists} \Rightarrow \text{CF} = 0 \text{ (classical),}$$

$$q = 3 \text{ (odd)} : \text{no ovoid, max partial ovoid } 7 < 10 \Rightarrow \text{CF} = \frac{1}{10} \text{ (contextual).}$$

The contextual power that drives the whole machine is therefore not inserted by hand: it is forced by the parity of the field, and $q = 3$ is the smallest odd prime that supplies it. The audit exhibits, side by side, the $q = 2$ ovoid that kills contextuality and the $q = 3$ obstruction that creates it.

The performance face. Where `holonetaudit` certifies the machine is *correct*, `analysis/holonet_bench.py` (the command `holonetbench`) certifies it is *cheap*, and it separates two kinds of number honestly. The deterministic operation counts — 7 mod-3 ALU operations per routing decision, ≤ 2 hops per packet, $\mu = 4$ multipath, four syndromes per error-correction cycle — are properties of the geometry and are identical on any computer. The wall-clock timings — route latency, correction-cycle time, teleport time, audit time, and the derived throughputs — are reported as medians over fixed, seeded workloads, but are host-relative. On ordinary Python on ordinary hardware the classical layer routes on the order of 10^5 packets per second and applies Clifford gates at a similar rate; these are a floor, not a ceiling, since a compiled or ASIC realization of the same op counts would be far faster.

The audit, live in the browser. Finally `docs/holonet.html` now renders the audit ledger client-side: thirteen of the sixteen rows are recomputed live in the reader's browser from $q = 3$ (the page rebuilds $W(3, 3)$, counts the fabric, derives λ_2 and the line-contexts, and checks the arithmetic), with the three solver rows — the contextual fraction, the magic robustness, and the code threshold — marked as verified by the Python witness. The same symplectic form $B(x, y)$ drives the browser page and the audited specification.

Honest scope: the q -scan is an exact recomputation per q (scipy MILP for the max-satisfiable count, NETWORKX for the independence number) and $q \in \{2, 3\}$ are prime, so F_q is integers mod q ; it is a property test of the audit, not a new physical claim. The bench times the *classical* emulation only — the priced 9^t magic dial is excluded by construction — and its absolute times are host-relative.

C.71 Pass 53 — parity not primality, a falsifiable prediction, and the table-free win

Pass 52 established the parity law on two data points, $q = 2$ and $q = 3$, both prime. Pass 53 closes the obvious loophole and turns the law into an experiment.

Parity, not primality. The scan in `analysis/w33_audit_qscan.py` now reaches $q = 4 = \text{GF}(4)$, built over the genuine field $F_2[x]/(x^2 + x + 1)$ rather than the integers modulo 4, and optionally $q = 5$. The geometry $W(4)$ is $\text{SRG}(85, 20, 3, 5)$ with $|\text{Sp}(4, 4)| = 979200$, and it possesses an ovoid of $17 = q^2 + 1$ points, so its contextual fraction is *zero*. Because 4 is even but *composite*, this rules out a “contextuality requires prime order” reading: the cause is the parity of q , full stop. Across the scan the contextual fraction takes the exact closed form

$$\text{CF}(q) = \begin{cases} 0, & q \text{ even,} \\ \frac{1}{q^2 + 1}, & q \text{ odd,} \end{cases}$$

measured as 0, $\frac{1}{10}$, 0, $\frac{1}{26}$ for $q = 2, 3, 4, 5$. The Holonet sits at $q = 3$, the smallest order with nonzero CF, and its value $1/10$ is the $q = 3$ instance of this one-line law.

A falsifiable prediction. This makes the construction testable by a contrast rather than a single number. The same photonic Kochen–Specker test, laid out on an *even*-order symplectic fabric — $W(2)$ with 15 rays or $W(4)$ with 85 rays — must return *zero* contextuality, because an explicit ovoid assignment reproduces every measurement outcome noncontextually. Observing contextuality on an even-order fabric, or failing to observe $\text{CF} = 1/10$ on the $q = 3$ fabric, refutes the model. The contextual fraction is not a tunable parameter; it is pinned to the field order, and the even-versus-odd contrast is the discriminating measurement.

The table-free win, measured. Finally `holonetbench--compare` quantifies the claim that “the address is the route”. It builds a classical all-pairs forwarding table explicitly: a next hop for each of the $n(n - 1) = 1560$ ordered destination pairs, $\lceil \log_2 40 \rceil = 6$ bits each, i.e. 1170 bytes of fabric-wide routing state, assembled by a breadth-first search from every node. It then verifies that the Holonet routes the same sampled pairs with *zero* table — the destination address is the route, every adjacency is the seven-operation symplectic test, every pair is reached in at most two hops — and that the two routers agree on reachability. The win is therefore a measured fact: $O(n^2)$ bytes of routing state and a rebuild on every topology change, reduced to nothing, because routing is stateless and local.

Honest scope: the $q = 4$ arithmetic is the true $\text{GF}(4)$; the $q = 5$ check is exact but its integer program is heavy (minutes), so it is opt-in and excluded from the fast path. The routing comparison is an architectural state-and-setup comparison, not a wall-clock race: a table lookup is a cheap operation, but it presupposes the very table the Holonet never builds.

C.72 Pass 54 — the control arm: an explicit even- q ovoid, and the scaling win

Pass 53 made the parity law falsifiable in principle; Pass 54 supplies the *positive control* that makes it falsifiable in practice, and shows the routing win is a scaling law rather than a single number.

The explicit noncontextual model. The demonstrator protocol measures the $q = 3$ fabric and expects $CF = 1/10$; its natural control is an even-order fabric, which the parity law predicts must read $CF = 0$. That prediction is now a constructed object. The witness `analysis/w33_ovoid_construct.py` builds the ovoid of $W(2)$ and $W(4)$ — the five rays

$$\{0100, 1000, 1101, 1110, 1111\}$$

for $q = 2$, and seventeen rays for $q = 4$ — and verifies each three independent ways: it has size $q^2 + 1$, it meets every one of the n contexts in exactly one ray, and its rays are pairwise non-collinear. Assigning the value one to the ovoid rays and zero to all others is then a single context-independent truth assignment that satisfies every context at once: a working noncontextual hidden-variable model, so $CF = 0$ on the even fabric. The same witness confirms that $q = 3$ admits no such assignment (the best leaves four of forty contexts unsatisfiable), which is exactly the obstruction the positive arm detects. This converts the single-sided test into a two-arm discriminator, written up in `holonet_parity_control.tex`: contextuality must appear on the odd fabric and vanish on the even one, a contrast fixed by the geometry, so a systematic that faked a violation would have to respect the field parity to survive both arms.

The scaling win. Finally `holonetbench--compare--scale` shows that the table-free routing advantage grows without bound. For $W(q) = GQ(q, q)$ with $n = (q + 1)(q^2 + 1)$ nodes, a classical all-pairs forwarding table costs $n(n - 1)$ next-hop entries of $\lceil \log_2 n \rceil$ bits — routing state of order $n^2 \log n$ that diverges with q — while the Holonet stores zero bytes and reaches any node in a constant two hops at seven modular operations per decision, because the address is the route and the generalized quadrangle has diameter two at every order. Tabulated from $q = 2$ to $q = 13$, the baseline routing state climbs from about a hundred bytes to several megabytes while the Holonet column stays identically zero and the hop count stays two.

Honest scope: the ovoid construction is an exact finite computation over F_2 and the genuine field $GF(4)$; it is the predicted data for a physical control arm, not itself a run, and the even-order analyzer bank is a different optical object from the forty Witting tetrads. In the scaling table the forwarding-table sizes are exact closed forms, and the address-routing / diameter-two property is verified explicitly for $q \leq 5$ and holds for the whole $GQ(q, q)$ family by the diameter-two theorem.

C.73 Pass 55 — the two-arm discriminator, runnable end to end

Pass 54 supplied the control model; Pass 55 makes the two-arm test executable inside the demonstrator’s own estimator pipeline, and gives the scaling win a figure.

One estimator, two fabrics, opposite verdicts. The demonstrator ships two contextual-fraction estimators — a normal-window check and an exact two-sided binomial test — both written against the raw-shot schema and both targeting $1/10$. The witness `analysis/holonet_control_arm.py` runs those *unmodified* estimators on two fixtures. The positive fixture is the committed synthetic $q = 3$ run; the control fixture is generated from the explicit $W(2)$ ovoid, where, because the ovoid satisfies every context, the noncontextual model has zero contextual excess and the signal

therefore clicks only at the dark-count baseline. The same estimators then return opposite verdicts on the one-tenth hypothesis:

positive $W(3) : \widehat{CF} \approx 0.125$ (compatible with $1/10$), control $W(2) : \widehat{CF} \approx 0$ (incompatible with $1/10$),

with a one-sided upper bound on the control fraction below $1/10$. This is the parity law as a discriminating measurement that runs through the existing pipeline: a systematic faking a $1/10$ signal on the odd fabric would have to fake it on the even fabric as well, where an explicit ovoid forbids any contextual excess. The parity law, the explicit ovoid, and this discriminator are registered at tier E in the public claim-tier spine (`analysis/BT1907_photonic_holonet_claim_tier_refactor.md`).

The scaling win, drawn. Finally `holonetbench--compare--scale--plot` renders the routing comparison as a figure: across $q = 2$ to 13 the classical forwarding table climbs from about a hundred bytes to several megabytes on a logarithmic axis, while the Holonet line stays pinned at zero bytes and two hops, written to `docs/holonet_scale.png`.

Honest scope: the control fixture is a synthetic run exercising the estimator path, not physical data; the positive arm reuses the committed synthetic fixture; the even-order analyzer bank remains a different optical object from the forty Witting tetrads. What is demonstrated is that the discriminator executes through the unmodified estimators and returns the geometry-forced verdicts, and that the figure is regenerated from the same closed-form scaling law.

C.74 Pass 56 — two contextualities, separated; and why one photon in $\mathbb{C}^4$

Pass 55 made the two-arm discriminator executable; Pass 56 computes its honest boundary, and in doing so finds a second, independent forcing of $q = 3$.

Sign versus selection. The parity law says the even fabric $W(2)$ has contextual fraction zero — yet the doily is the textbook home of the Mermin–Peres magic square. Both are true, because they are different statistics, and `analysis/w33_doily_mermin.py` computes both exactly on the same fifteen points. The *selection* system — exactly one click per context, the statistic the contextual-fraction estimators measure — is satisfied perfectly by the five-ray ovoid, so $CF = 0$. The *sign* system — noncontextual ± 1 values reproducing the product of each line’s commuting two-qubit Pauli triple, with the fifteen signs computed as actual 4×4 matrix products (three lines multiply to $-I$) — is *unsatisfiable*: exact F_2 elimination proves it, and the minimal certificate is six lines covering every point an even number of times with an odd number of $-I$ lines — a Mermin–Peres square sitting inside the doily. So the even-order control fabric is sign-contextual yet selection-noncontextual, while the $q = 3$ fabric is contextual in *both* senses. The two-arm contrast is therefore a statement about one declared observable, and both arms must measure exactly that observable; this boundary is now stated in `holonet_parity_control.tex` and registered in the public theorem ledger.

The realization dimension bound. A complete-basis realization of $W(q)$ maps each context to an orthonormal basis, hence lives in $\mathbb{C}^{q+1}$. The witness `analysis/w33_realization_dimension.py` proves this is *impossible* for $q = 2$: a non-collinear pair of distinct rays spans two dimensions of $\mathbb{C}^3$, leaving a one-dimensional orthocomplement containing exactly one ray, while the geometry demands $\mu = 3$ distinct common neighbours there. The same count at $q = 3$ — four rays in a two-dimensional orthocomplement of $\mathbb{C}^4$ — poses no obstruction, and the Witting realization exists (cited corpus result); $q = 4$ passes this test with existence left open. Consequence: $q = 3$ is the

smallest symplectic quadrangle the one-photon $\mathbb{C}^4$ apparatus can realize at all, and by the parity law the smallest *contextual* one. Two independent forcings — one metrological, one statistical — select the same order.

Honest scope: both results are exact finite computations (matrix products and F_2 elimination; strongly-regular pair counts). The qubit-Pauli labelling of $W(2)$ is standard; the Witting existence at $q = 3$ is cited, not re-derived; $q = 4$ existence in $\mathbb{C}^5$ is open. Neither result weakens the discriminator — it sharpens what the discriminator’s contrast is a contrast *of*: the selection statistic, declared in advance, measured identically on both arms.

C.75 Pass 57 — the contextuality tax: the defect is one movable point-star

Pass 57 is a synthesis pass: it joins two work streams that had converged on the same integer without a proof that they were describing the same object. The contextuality arc established max-sat = 36 of 40 contexts ($CF = 1/10$) and the parity law $CF(q) = 1/(q^2 + 1)$ for odd q . Independently, the operating-system arc’s line-context microkernel scheduler found 36 schedulable spreads and described its 4-context deficit as a “movable point-star double-occupancy defect” — language explicitly flagged there as an audited numerical bridge, not a proven structure. The witness `analysis/w33_contextuality_tax.py` proves the structure.

The complete classification. Iterating the max-satisfiable integer program with no-good cuts — each round forbidding the exact four-line failure set just found and re-solving — enumerates *every* failure set achievable by an optimal (36-context) assignment; the enumeration terminates when the objective drops below 36, so the classification is exhaustive rather than sampled. The result: there are exactly 40 achievable failure sets, and each is a *point-star* — the four lines through a single point — with every one of the 40 points realized as a center. So the defect of the classical layer is always one star, and it is fully movable: a gauge freedom, not an accident of one optimum. The deficit law follows structurally: for odd q the deficit is exactly $q + 1$, one line’s worth of contexts (verified 4 at $q = 3$; 6 at $q = 5$ under `--deep`), and 0 for even q , which restates $CF = (q + 1)/n = 1/(q^2 + 1)$ as “the tax is one star.”

What the operating system pays. Operationally the classical scheduler can satisfy everything except the star of one point *of its choosing*, and may move that star to the least-loaded point. The escalation budget — the fraction of the runtime on which the priced 9^t quantum resource is spent — is therefore exactly one star, 4 contexts, $1/10$ of the fabric, per classical assignment. The corollary by parity closes the loop with the control arm: on an even-order fabric the tax is zero and there is no gap to exploit even in principle, since $\alpha = q^2 + 1$ meets the Hoffman bound exactly (5 at $q = 2$, 17 at $q = 4$), while odd q keeps the gap ($7 < 10$ at $q = 3$). The machine’s quantum advantage and its scheduling overhead are the same integer.

Honest scope: the $q = 3$ classification is exact and exhaustive; the deficit law is verified exactly at $q \in \{2, 3, 4\}$ (plus $q = 5$ opt-in) and stated as the observed closed form beyond; the identification with the scheduler arc’s $36 = 36$ bridge underwrites that bridge’s structural language but, like the bridge itself, claims no canonical bijection between optimal assignments and spreads; the runtime tie-in (twelve-tick admission words, the 2160-slot atlas, the 51840-tick supercycle) is accounting on committed constants.

C.76 Pass 58 — the spread side of the tax, and the anatomy of the optima

Pass 57 proved the tax theorem from the assignment side. Pass 58 (`analysis/w33_spread_star_anatomy.py`) completes the synthesis from the scheduler’s side and, in doing so, lets the computation correct a guess.

The clock, counted independently. The scheduler arc’s spread-clock rests on the spreads of $W(3,3)$ — sets of ten pairwise-disjoint lines partitioning the forty points — counted there by its own `find_spreads`. This witness re-enumerates them by exact-cover backtracking, a different algorithm sharing no code, and confirms: exactly 36 spreads, each line lying in exactly 9 of them. The bridge’s 36 is now verified from two independent directions.

The service-rate guarantee. Every spread covers every point exactly once, hence contains exactly one line of every point-star (the scheduler arc’s in-flight defect/spread tensor records the same complete incidence; it is re-verified here for all 36×40 pairs). Combined with the tax theorem — the failure set of any optimal assignment is one star — this yields the operational statement: *under any optimal classical assignment, every spread has exactly nine of its ten lines satisfied*. The spread-clock never sees a better or a worse cycle; the tax is spread-isotropic. For a fixed defect star the 36 spreads split into $q + 1 = 4$ classes of exactly 9 according to which defect line they carry.

The anatomy of the optima. Exhaustive enumeration (integer-program feasibility with no-good cuts, terminating) finds *exactly 20 optimal assignments per defect center* — 800 in all across the forty movable centers. The computation corrected the naive “double occupancy everywhere” guess and found a sharper structure, then pinned it: the center is *free* (optima come in center-lit/unlit pairs, exactly half each), and every optimum loads its four defect lines *uniformly* — the occupancy vector is always (c, c, c, c) , never mixed, with $c \in \{2, 3, 4\}$: center-unlit optima are $(2, 2, 2, 2)$ with 11 lit rays or $(3, 3, 3, 3)$ with 12; their center-lit partners are $(3, 3, 3, 3)$ with 12 or $(4, 4, 4, 4)$ with 13. A defect line is never empty and never accidentally satisfied. The scheduler arc’s “double-occupancy defect” is therefore literally the *minimal* class of optima — the ground state of the defect — while the full spectrum is uniform loading at two, three, or four rays per failed context.

Honest scope: all counts are exact finite enumerations; the optimum enumeration is exhaustive for the chosen center, with movability across centers supplied by Pass 57; the complete defect/spread incidence is the scheduler arc’s observation, independently re-verified here; as before, no canonical assignment-to-spread bijection is claimed — the service-rate lemma shows the relation is many-to-many and uniform.

C.77 Pass 59 — the tax orbits: one group moves the defect’s ground states

Pass 58 counted 20 optimal assignments per defect center, in classes $9 + 1 + 9 + 1$. Pass 59 (`analysis/w33_tax_orbits.py`) explains those numbers with the machine’s own runtime group, and exports the scheduler consequence.

The group, rebuilt — and a distinction enforced. Symplectic transvections $T_v: x \mapsto x + B(x, v)v$ generate the symplectic group; closing a handful of transvection permutations of the forty points by breadth-first search yields the faithful point action. The check itself caught the classic subtlety the claim-tier spine warns about: on *points* the order is $25920 = 51840/2$ — the projective group $\mathrm{PSp}(4,3)$ — because $-I$ acts trivially projectively; the 51840 matrix group is its double

cover. The witness now enforces the distinction rather than blurring it, and verifies transitivity on the forty points.

The orbit theorem. Acting with the group on the twenty optima of one center generates *all* 800 optima — twenty per center, every center reached — and they fall into exactly *four orbits*, of sizes 360, 40, 360, 40, with the pair (center lit?, uniform load c) a complete orbit invariant. The nine ground-state $(2, 2, 2, 2)$ optima per center form one 360-orbit (their center-lit flips another), while the special $(3, 3, 3, 3)$ optimum of each center forms a canonical 40-orbit: one distinguished twelve-ray configuration attached to each point of the fabric. The one-group-machine result thus extends to the tax at the projective level: the group that is the processor, the network, and the memory also moves the defect’s ground states transitively.

The queue invariant, and the family probe. Re-verified across all 800 optima globally: the defect star’s load vector is always uniform, with spectrum $\{2, 3, 4\}$ and ground state 2 — so a round-robin escalation queue has *zero* load imbalance in every optimal configuration, a statement exported machine-readably for the scheduler arc’s in-flight queue simulator. Under `--deep`, the same anatomy question is put to $W(5)$: in a capped sample of forty optima at one center, uniform loading persists on the six-line defect star, the incidence law $L = 25 + c$ holds, and the observed load spectrum is $\{4, 5, 6\}$ — i.e. $\{q - 1, q, q + 1\}$, matching the $q = 3$ spectrum $\{2, 3, 4\}$ under the same formula. A parity-family load law is suggested, honestly marked as a capped sample rather than an exhaustive classification.

Honest scope: the group closure, orbit decomposition, validity of all 800 optima, and the global uniform-loading law are exact finite computations; the $q = 5$ probe is exact but capped; the queue statement is a restatement of the loading theorem in scheduler language, with no claim about the in-flight simulator’s internals.

C.78 Pass 60 — the perp states: the canonical orbit named, the spectrum explained

Pass 59 found a canonical 40-orbit of special optima and asked what the object is. Pass 60 (`analysis/w33_perp_states.py`) answers: it is the oldest object in the geometry.

The identification. The special $(3, 3, 3, 3)$ optimum at defect center p is *exactly* the neighborhood $\Gamma(p)$ — the deleted perp $p^\perp \setminus \{p\}$ — verified unique and equal per center; its center-lit flip is the full perp $p^\perp$, a classical geometric hyperplane of the quadrangle. The optimality proof is the *defining axiom* of a generalized quadrangle: a point off a line is collinear with exactly one of its points. Lighting $\Gamma(p)$ therefore satisfies every non-star line exactly once (the axiom, verified for every point–line pair at $q = 2, 3, 4, 5$) and fails only the star, loaded uniformly with the three non-center points of each star line. The mysterious canonical configuration was the perp all along.

The spectrum, explained. Because the axiom holds in every generalized quadrangle, perp states exist at *every* order, loading the star uniformly at q (deleted) and $q + 1$ (full) — verified at $q = 2, 3, 4, 5$. This explains the tax load spectrum $\{q - 1, q, q + 1\}$: the top two loads are the deleted and full perp states, now *proven* present for all q (upgrading the $q = 5$ sample observation for those classes), while the bottom load $q - 1$ belongs to the ground states. Parity still decides optimality: for odd q one failed star is the proven minimum, so the perp state is optimal; for even q it is strictly dominated by the ovoid.

The ground states, anatomized. Each of the nine $(2, 2, 2, 2)$ ground states per center decomposes as 8 in-perp + 3 out-of-perp points. Verified exactly: the nine out-triples are pairwise non-collinear triads, each with common perp of size exactly 4 (centric triads), and they *partition* the 27 non-neighbors of the center — the second subconstituent, regular of degree 8 on 27 vertices, the same 27-point structure the corpus ties to the Schläfli/ E_6 story. The point stabilizer in $\mathrm{PSp}(4, 3)$ (order 648) is *transitive* on the nine ground states, so per center there is one ground state up to local symmetry, just as globally there is one up to the full group.

Honest scope: the perp and hyperplane are standard finite-geometry objects; the new content is the exact identification with the Pass 58/59 optima and the resulting explanation of the load spectrum. The Schläfli-adjacent observations are recorded as computed invariants (a partition of the 27 into nine centric triads; an 8-regular induced graph) with no further corpus identification claimed.

C.79 Operator-on-photon self-applied circuit packet

One Choi leg is interpreted as the photonic mode state and the other Choi leg as the encoded optical action. The finite circuit is therefore modeled as a relation on one photons own registers.

The self-applied model represents I , X , Z , F_3 , and S as encoded actions on the operator leg and checks them by trace-Choi overlap against the state leg.

The OAM operator witness distinguishes passive mode labels from active internal operator behavior by comparing label-only controls, operator-leg activation, basis covariance, and self-vs-external control tests.

C.80 Internal optic algebra, physical dictionary, and falsifier regimes

The internal operations I , X , Z , F_3 , and S generate the finite single-qutrit projective Clifford action of size $216 = 9 \cdot 24$ on the state/operator Choi legs.

The operations are mapped to optical analogues: reference/no-op, OAM shift or prism step, spiral phase plate, tritter or mode mixer, and lens or phase-curvature analogue. The lens row remains calibration-level.

The falsifier simulator separates passive OAM labels from active internal-operator behavior. Only the internal-operator scenario passes the symbolic criteria: trace-Choi pattern, low leakage, basis covariance, and operator activation.

C.81 Internal Clifford census, lens calibration, and passive-active protocol

The internal operations I , X , Z , F_3 , and S generate a 216-element single-qutrit projective Clifford action on the state/operator Choi legs. All 216 elements preserve the state/operator split in the internal-register model; only the identity stabilizes the identity Choi state itself. The 24 zero-translation frame changes preserve centered OAM opposition, while the 192 translated elements move the OAM origin and require recentering.

The lens calibration row is exact at the finite phase-signature level. On the centered OAM basis $\ell = -1, 0, +1$ with qutrit labels 2, 0, 1, qutrit S has ω -exponent signature $[1, 0, 1]$, matching the centered quadratic lens phase ω^{ℓ^2} .

The passive-vs-active protocol uses centered OAM preparation, axial time-bin support, operator off/on controls, $I/X/Z/F_3/S$ settings, leakage checks, basis covariance, and external-reference comparison. It is a protocol design, not experimental evidence.

C.82 OAM recentering, exact optical calibration, protocol table, and literature bridge

The 192 translated Clifford elements split into eight nonzero shift classes, each with 24 frame choices. The recentering correction is the inverse shift, with OAM-only, phase-only, and mixed shift classes.

The S and F_3 calibrations are placed on the same centered OAM basis. The S row has phase signature $[1, 0, 1]$, and the F_3 row is the finite three-mode mixer on the same qutrit labels.

The passive-vs-active protocol is converted into a paper-ready table covering preparation, passive controls, active gate settings, leakage, covariance, reference comparison, expected readouts, thresholds, failure modes, and claim tiers.

The OAM literature bridge supports high-dimensional qudit and same-photon register directions, but requires radial leakage and mode-overlap witnesses because radial and azimuthal spatial structure are coupled in realistic modes.

C.83 Recentered witnesses, leakage regimes, and appendix splicer

The calibrated I, X, Z, F_3, S witness applies directly to the centered frame class. Translated Clifford elements require inverse recentering before the same witness table is applied.

The radial leakage simulator turns symbolic shell leakage into pass/fail regimes, separating good core-gate behavior from radial leakage, visibility mismatch, and basis-covariance failures.

The operator/OAM appendix splicer follows the dry-run and idempotent splicer pattern. It records the appendix insert packets and changes the paper only when apply mode is used in check-out.

C.84 Appendix validation, recentered protocol table, and claim ledger

The appendix packet is validated by a dry-run artifact checker: all insert paths exist, the bounded splice block is idempotent, and the target paper admits the expected splice anchor or bounded fallback.

The protocol table is expanded by recentering class. The resulting table has four Clifford recentering classes times five calibrated witnesses, with correction operators, leakage thresholds, failure modes, and claim tiers.

The claim ledger separates exact finite claims from calibration-level optics, engineering leakage, protocol witnesses, literature motivation, and blocked physical overclaims.

C.85 Full operator/OAM appendix splice and recentered transaction ABI

The operator/OAM appendix is now part of the main Holonet paper through a bounded idempotent splice. The splice carries the operator-on-photon circuit model, the internal Clifford falsifier, the recentering protocol, the leakage and claim ledger, and the present synthesis block.

The architectural closure is the recenter-to-transaction ABI:

$$216 = 9 \cdot 24 = (1 + 2 + 2 + 4) \cdot 24.$$

The factor 24 is the centered BT1495 transaction-word kernel. The nine affine shifts are one centered class, two OAM-only shifts, two phase-only shifts, and four mixed shifts. Thus the class profile is

$$24, \quad 48, \quad 48, \quad 96,$$

and every translated internal Clifford action is executed by inverse recentering followed by one of the same 24 centered 72-tick transaction words.

One operation sweep therefore has $216 \cdot 72 = 15552$ ticks. The full five-witness I, X, Z, F_3, S protocol has 77760 ticks. This is an exact finite ABI statement: it certifies the centered word reuse and the recentering burden. It does not claim measured OAM leakage, optical loss, or a finished calibrated device.

Recent OAM and high-dimensional time-bin experiments support the physical direction, including same-photon polarization/OAM gates, robust time-bin qudit control, directly measured topology in OAM-entangled states, and OAM-multiplexed diffractive mode processing [32, 33, 34, 35]. BT1588 keeps those sources in the literature-motivation tier only: they motivate the mode interface, multiplexing optics, and leakage tests, while the exact theorem remains the local finite $216 = 9 \cdot 24$ recenter ABI.

C.86 LG radial covariance, full witness lane sheet, and OAM front end

The recenter ABI now has a radial-shell falsifier. Model the first three Laguerre–Gaussian radial shells by the stochastic channel family $L(\eta)$: diagonal entries $1 - \eta$, off-diagonal entries $\eta/2$. The BT1577 operation envelopes I, Z, S, X, F_3 all lie in this same channel family, so the radial recenter correction and the centered witness gate commute at shell level. For two envelopes,

$$\eta(a, b) = a + b - \frac{3}{2}ab.$$

Thus the $216 = 9 \cdot 24$ affine actions do not require 216 separate radial calibrations. They require one centered witness table plus the recenter tax. The symbolic worst case is a mixed OAM/phase recenter followed by F_3 ,

$$\eta_{\max} = 0.18296 < 0.20,$$

while the centered gate threshold remains the BT1577 value 0.10. This is a covariance simulator and falsifier, not measured beam leakage.

The complete witness replay is now compiled as a physical lane sheet:

$$5 \text{ gates} \times 9 \text{ recenter sectors} \times 24 \text{ centered words} \times 72 \text{ ticks} = 77760.$$

It is stored as 1080 exact 72-tick segments with the global tick formula

$$((((g \cdot 9) + s) \cdot 24 + w) \cdot 72 + t).$$

Every Hesse lane appears 12960 times, every detector slot appears 19440 times, native D_4 square-pulse words occupy 25920 ticks, and the remaining S_4 analyzer-relabel words occupy 51840 ticks.

The hardware interpretation is deliberately sharper. OAM-multiplexed diffractive optics are a proposed *front-end sectorizer*: a passive mode-processing layer can sort or fan out the nine affine recenter sectors before the exact 24-word transaction kernel runs. Recent OAM-MDNN work supports this as an optical mode-processing direction, but it does not prove a quantum gate, preserve coherence by itself, or calibrate loss. The promoted theorem remains local:

$9 \text{ recenter sectors} \times 24 \text{ centered words} = 216 \text{ exact finite ABI addresses,}$

with acceptance delegated to radial tomography, sector crosstalk measurements, and the 77760-tick witness replay.

C.87 Lab tomography, LG alphabet, and the Hesse/T witness loop

The OAM front end is now written in a lab-facing form. BT1592 promotes three acceptance surfaces: a 9×9 sector confusion matrix, four radial-shell tomography rows, and six exact lane-replay checks. The synthetic sector fixture uses diagonal probability 0.93, off-diagonal probability 0.00875, and acceptance thresholds

$$p_{\text{diag}} \geq 0.90, \quad p_{\text{off}} \leq 0.02.$$

The generated CSV ingest file has 91 rows, all passing, and is deliberately replaceable by bench data. This is a falsification harness, not measured OAM crosstalk.

BT1593 chooses the explicit Laguerre–Gaussian sector alphabet:

$$s = 3x + z, \quad \ell = s - 4, \quad p = x + 2z \pmod{3}.$$

Thus the nine recenter sectors occupy symmetric OAM charges

$$\ell = -4, -3, -2, -1, 0, 1, 2, 3, 4,$$

with three modes in each radial shell $p = 0, 1, 2$. The 24-word centered selector is a separate handoff layer:

$$\text{address} = 24s + w, \quad 0 \leq s < 9, \quad 0 \leq w < 24.$$

The architecture therefore uses 9 physical sector modes plus a 24-word selector, not 216 separate OAM channels. The 216 exact finite addresses split into 72 native D_4 square-pulse addresses and 144 S_4 analyzer-relabel addresses.

The universality port now lands in the same witness loop. BT1404 already made a Hesse/T microframe with

$$9 \text{ Hesse outcomes} \times 8 \text{ ticks} = 72 \text{ ticks}.$$

BT1590 already made each OAM witness segment 72 ticks. BT1594 overlays these rather than appending a second schedule:

$$1080 \text{ segments} \times 72 \text{ ticks} = 1080 \text{ Hesse/T microframes} = 77760 \text{ ticks}.$$

Each Hesse outcome appears once per segment, hence 1080 times total. The T-frame bit profile is 0 : 5400 and 1 : 4320, matching the parity rule

$$t = (r + p) \pmod{2}.$$

This closes the finite architecture loop: the Clifford/OAM witness, the sector-mode front end, the tomography acceptance harness, and the explicit Hesse/T non-Clifford port now inhabit one replayable 77760-tick object. The claim boundary remains unchanged: this is ABI and schedule compatibility, not a physical Hesse-SIC device, injection-fidelity measurement, or fault-tolerant magic-state-yield theorem.

C.88 The Witting reject shell as the universal fuel rail

The deepest closure in the current architecture is that the newest OAM/Hesse witness loop is not merely a validation sweep. It is the Witting matter shell itself. In the Witting delayed-query desk, each selected ray has

$$1 \text{ same} + 12 \text{ compatible} + 27 \text{ incompatible} = 40$$

possible queried rays. Therefore the incompatible ordered-pair shell has

$$40 \cdot 27 = 1080$$

records. BT1595 constructs an explicit bijection

$$40 \cdot 27 = 5 \cdot 9 \cdot 24$$

between those rejected Witting pairs and the BT1590/BT1594 witness segments. The per-gate refinement is even sharper:

$$9 \cdot 24 = 216 = 8 \cdot 27.$$

Thus each of the five witness gates consumes the full incompatible-target shell of eight Witting source rays. Each rejected Witting pair is one 72-tick contextual fuel segment, and each such segment carries the Hesse/T overlay.

BT1596 writes the resulting runtime economy:

$$520 \cdot 72 = 37440 \quad \text{accepted communication/control ticks,}$$

$$1080 \cdot 72 = 77760 \quad \text{contextual fuel ticks,}$$

and

$$1600 \cdot 72 = 115200 \quad \text{ticks for the complete Witting ordered-pair cycle.}$$

The ratio of accepted control to contextual fuel is 13 : 27, matching the Witting split 13/40 versus 27/40.

BT1597 packages this as a single universal transaction object. Accepted Witting pairs form the communication and control rail. Rejected Witting pairs are not discarded; they are exactly the OAM/Hesse contextual-fuel rail:

$$1080 = 5 \text{ gates} \cdot 9 \text{ OAM sectors} \cdot 24 \text{ words} = 40 \text{ rays} \cdot 27 \text{ incompatible targets.}$$

The Hesse/T microframe on that fuel rail supplies the explicit non-Clifford boundary required by the universal-computation contract. The claim remains a finite ABI and timing theorem: it does not assert optical coherence, loss tolerance, cryptographic security, or a fault-tolerant threshold.

C.89 Accepted control rail and the full Witting transaction cycle

The Witting fuel rail now has its complementary control rail. BT1598 compiles the accepted 13/40 Witting ordered pairs into concrete 72-tick control frames. The only ambiguity is the same-ray case: a ray lies in four Witting bases. BT1598 resolves this by a perfect matching between the 40 rays and the 40 bases. Therefore each basis carries exactly one same-ray aperture and twelve off-diagonal compatible apertures:

$$40 \text{ bases} \cdot 13 \text{ controls} = 520 \text{ accepted frames,}$$

$$520 \cdot 72 = 37440 \text{ accepted control ticks.}$$

The detector handoff is the existing one:

$$\text{Witting analyzer slot} \rightarrow \text{detector slot} \rightarrow \text{mirror slot mod } 4 \rightarrow \text{tomotope flag.}$$

Using the selected Witting basis as a tomotope block gives

$$\text{tomotope flag} = 4 \text{ basis} + \text{detector slot,}$$

using the first 40 of the 48 tomatope blocks and leaving 8 slack blocks for non-Witting packet traffic.

BT1599 explains the 120-record surplus in the basis-local Witting table. There are 160 same-ray basis-local apertures, but BT1598 promotes only 40 of them to selected control frames. The remaining records are

$$160 - 40 = 120 = 3 \text{ local tomatope sheets} \cdot 40 \text{ W33/Witting lines},$$

exactly the BT1365 selector phase-sheet bundle. Thus the same-ray surplus is a qutrit phase sidecar rather than extra communication throughput.

BT1600 then compiles the full ordered-pair desk:

$$40 \text{ source rays} \cdot 40 \text{ target rays} = 1600 \text{ frames},$$

$$1600 \cdot 72 = 115200 \text{ ticks}.$$

Every source ray has

$$13 \text{ accepted control frames} + 27 \text{ contextual fuel frames}.$$

The complete cycle is therefore not two unrelated protocols. It is one finite transaction machine: accepted pairs run the Witting analyzer/mirror control rail, rejected pairs run the OAM/Hesse contextual fuel rail, and the unused same-ray contexts are exactly the ternary selector phase sheets.

C.90 Physical automaton, Fano detector bins, and the universal ABI

BT1601 turns the 1600-frame Witting transaction cycle into a single physical automaton. The carrier is one single-photon transaction packet. Each ordered pair frame keeps the common 72-tick shape:

$$1600 \cdot 72 = 115200 \text{ ticks}.$$

The frame template is now explicit:

ticks	role
0...7	source switch
8...31	program/delay placeholder
32...47	analyzer or OAM/Hesse body
48...63	detector or Hesse/T handoff
64...71	dark-reference placeholder

There are 520 Witting analyzer switch frames and 1080 OAM/Hesse fuel switch frames. The four detector residues are balanced across the whole automaton:

$$400, 400, 400, 400.$$

Every frame has a symbolic loss channel and a dark-reference bin. These fields are intentionally placeholders: BT1601 compiles the interface, not a measured loss or dark-count model.

BT1602 identifies the old 168 active detector-bin count with the Fano plane. The active bus factors as

$$168 = 7 \cdot 24 = 7 \text{ Fano lines} \cdot 3 \text{ point slots} \cdot 8 D_4 \text{ states}.$$

The Witting source side factors as

$$40 = 5 \text{ witness gates} \cdot 8 D_4 \text{ source states}.$$

The five witness gates occupy five Fano lines. On those lines,

$$27 = 3 \text{ point slots} \cdot 9 \text{ Hesse/OAM residues}$$

is the contextual fuel shell. The two remaining Fano lines carry the accepted compatible controls:

$$12 = 2 \text{ reserve lines} \cdot 3 \text{ point slots} \cdot 2 \text{ detector parities.}$$

The same-ray control is the source anchor on the gate line. Thus every source ray has

$$27 \text{ fuel} + 12 \text{ compatible control} + 1 \text{ same-ray control} = 40$$

targets, and every one of the 168 active detector bins is touched. The usage profile across the 1600 frames is

$$80 \text{ bins} \times 9 + 88 \text{ bins} \times 10 = 1600.$$

BT1603 packages the architecture as a finite universal-computation ABI theorem. The accepted Witting frames provide finite Clifford transport through the analyzer/mirror/tomotope handoff. The rejected Witting frames provide the contextual OAM/Hesse fuel. The Hesse/T overlay supplies the explicit non-Clifford port required by the BT1377 universal-computation contract. The detector and row values then hand into the BT1486 retwined CSS syndrome layer and the BT1493 physical row-pulse compiler. In short,

$$\text{Witting control} + \text{contextual fuel} + \text{Hesse/T port} + \text{CSS handoff}$$

is now one finite programmable photonic-computation ABI. The theorem remains inside the claim firewall: thresholds, measured loss, detector efficiency, and magic-state yield are downstream physical requirements, not consequences of the finite compiler alone.

C.91 Fano charge, time-bin envelope, and guard-page closure

The $80 \times 9 + 88 \times 10$ detector-bin profile is not just a usage histogram. BT1648 shows that it is a conserved Fano charge. The 168 active detector bins split into three exact classes:

$$80 \cdot 9 + 40 \cdot (9 + 1) + 48 \cdot 10 = 1600.$$

The 80 bins are unanchored contextual-fuel bins. The 40 middle terms are fuel bins with one same-ray control anchor. The 48 final bins are compatible-control reserve bins. Equivalently, the seven Fano lines split as five Witting gate lines and two reserve lines:

$$5(16 \cdot 9 + 8 \cdot 10) + 2(24 \cdot 10) = 1600.$$

This explains why the high-usage side has size 88: it is exactly

$$40 \text{ same-ray anchors} + 48 \text{ compatible reserve bins.}$$

BT1649 then embeds the active automaton into an 11-bit time-bin qudit envelope, motivated by single-photon time-bin qudit universal logic [29]. The active Witting transaction uses the first 1600 time bins. The remaining addresses factor without remainder:

$$2^{11} = 2048, \quad 2048 - 1600 = 448 = 7 \cdot 64.$$

Thus every Fano point receives a 64-slot guard page. Each guard page is typed as

$$24 \text{ dark references} + 24 \text{ loss probes} + 16 \text{ parity overflow slots}.$$

The Witting source/fuel structure remains the contextual communication desk of [28]; the time-bin envelope is the physical address wrapper around that finite desk, not a new communication assumption.

BT1650 closes the calibration interface. The $7 \cdot 24 = 168$ dark-reference guards cover all active detector bins once, and the 168 loss-probe guards also cover all active detector bins once. The remaining $7 \cdot 16 = 112$ parity-overflow guards feed the CSS/jitter side of the fault ABI:

$$\text{dark} \mapsto \text{DARK_CLICK}, \quad \text{loss} \mapsto \text{MISSED_CLICK}, \quad \text{parity} \mapsto \text{CSS_SYNDROME}.$$

This is the new physical closure: the 448 time-bin slack addresses are not unused capacity. They are the calibrated guard shell required to run the 1600-frame Witting automaton as a bench-addressable single-photon qudit program.

References

- [1] W. Dahn, *The $W(3, 3)$ Substrate: Master Manuscript* (`w33_paper.tex`), Theory of Everything Project (2026) — master equation $q! = 2q$, parameter closure, Bose–Mesner field equations, and the full observable derivation chain.
- [2] W. Dahn, *Self-Entanglement Companion*, Theory of Everything Project preprint (2026).
- [3] M. Howard, J. Wallman, V. Veitch, J. Emerson, Contextuality supplies the ‘magic’ for quantum computation, *Nature* **510**, 351 (2014).
- [4] P. Zanardi, M. Rasetti, Holonomic quantum computation, *Phys. Lett. A* **264**, 94 (1999).
- [5] F. Wilczek, A. Zee, Appearance of gauge structure in simple dynamical systems, *Phys. Rev. Lett.* **52**, 2111 (1984).
- [6] P. T. Dumitrescu, R. Vasseur, A. C. Potter, Logarithmically slow relaxation in quasiperiodically driven random spin chains, *Phys. Rev. Lett.* **120**, 070602 (2018).
- [7] I. Martin, G. Refael, B. Halperin, Topological frequency conversion in strongly driven quantum systems, *Phys. Rev. X* **7**, 041008 (2017).
- [8] E. Witten, Three-dimensional gravity revisited, [arXiv:0706.3359](https://arxiv.org/abs/0706.3359) (2007).
- [9] A. Almheiri, X. Dong, D. Harlow, Bulk locality and quantum error correction in AdS/CFT, *J. High Energy Phys.* **04**, 163 (2015).
- [10] M. Howard, J. Wallman, V. Veitch, J. Emerson, Contextuality supplies the ‘magic’ for quantum computation, *Nature* **510**, 351 (2014), [arXiv:1401.4174](https://arxiv.org/abs/1401.4174).
- [11] P. T. Dumitrescu *et al.*, Dynamical topological phase realized in a trapped-ion quantum simulator, *Nature* **607**, 463 (2022).
- [12] D. Gottesman, I. Chuang, Demonstrating the viability of universal quantum computation using teleportation, *Nature* **402**, 390 (1999).

- [13] Y. Saad and M. H. Schultz, Topological properties of hypercubes, *IEEE Transactions on Computers* **37**, 867 (1988).
- [14] W. H. Kautz, Unit-distance error-checking codes, *IRE Transactions on Electronic Computers* **EC-7**, 179 (1958).
- [15] E. L. Wilmer and M. D. Ernst, Graphs induced by Gray codes, *Discrete Mathematics* **257**, 585 (2002), [abstract and bibliographic data](#).
- [16] A. Smith, Universality of Wolfram’s 2, 3 Turing machine, *Complex Systems* **29**, 1 (2020).
- [17] S. Wolfram, The Wolfram 2,3 Turing Machine Research Prize, <https://www.wolframscience.com/prizes/tm23/>.
- [18] I. Niven, *Irrational Numbers*, MAA (1956).
- [19] J. A. Thas, Ovoids and spreads of finite classical polar spaces, *Geom. Dedicata* **10**, 135 (1981).
- [20] B. Monson, D. Pellicer, G. I. Williams, The tomatope, *Ars Mathematica Contemporanea* **5**, 355–370 (2012).
- [21] T. Breuer, *GAP Character Table Library: $U_4(2) \cong \text{PSp}(4, 3)$* , ATLAS-derived maximal-subgroup data, [online ATLAS table](#).
- [22] B. McMillan, Two inequalities implied by unique decipherability, *IRE Transactions on Information Theory* **2**(4), 115–116 (1956), [doi:10.1109/TIT.1956.1056818](#).
- [23] J. Conrad, J. Eisert, F. Hangleiter, Gottesman–Kitaev–Preskill codes: A lattice perspective, *Quantum* **6**, 648 (2022).
- [24] M. Lin, C. Chamberland, K. Noh *et al.*, Closest lattice point decoding for multimode Gottesman–Kitaev–Preskill codes, *PRX Quantum* **4**, 040334 (2023) — the D_4 lattice is optimal in the two-mode case.
- [25] S. Lloyd, S. L. Braunstein, Quantum computation over continuous variables, *Phys. Rev. Lett.* **82**, 1784 (1999) — Gaussians plus any single degree- ≥ 3 Hamiltonian are universal on the oscillator.
- [26] I. Frenkel, J. Lepowsky, A. Meurman, *Vertex Operator Algebras and the Monster*, Academic Press (1988) — the Monster module $V^\natural$ (24 free bosons on the Leech torus, $\mathbb{Z}/2$ -orbifolded), $c = 24$.
- [27] A. Dymarsky, A. Shapere, Quantum stabilizer codes, lattices, and CFTs, *J. High Energy Phys.* **03**, 160 (2021), [arXiv:2009.01244](#) — the code/Narain-CFT correspondence.
- [28] A. Yu. Vlasov, Scheme of quantum communications based on Witting polytope, *Moscow Univ. Phys.* **80**, 560 (2025), [arXiv:2503.18431](#).
- [29] A. Delteil, N-qubit universal quantum logic with a photonic qudit and $O(N)$ linear optics elements, *APL Quantum* **1**, 046101 (2024), [arXiv:2410.06334](#).
- [30] M. Reck, A. Zeilinger, H. J. Bernstein, P. Bertani, Experimental realization of any discrete unitary operator, *Phys. Rev. Lett.* **73**, 58 (1994), [doi:10.1103/PhysRevLett.73.58](#).

- [31] W. R. Clements, P. C. Humphreys, B. J. Metcalf, W. S. Kolthammer, I. A. Walmsley, Optimal design for universal multiport interferometers, *Optica* **3**, 1460 (2016), [arXiv:1603.08788](#).
- [32] Y. Zhang *et al.*, Polarization and orbital angular momentum encoded quantum Toffoli gate enabled by diffractive neural networks, *Phys. Rev. Lett.* **133**, 140601 (2024), [doi:10.1103/PhysRevLett.133.140601](#).
- [33] A. Delteil *et al.*, Robust approach for time-bin encoded photonic quantum states, *Phys. Rev. Lett.* **134**, 180802 (2025), [doi:10.1103/PhysRevLett.134.180802](#).
- [34] R. de Mello Koch, P. Ornelas, N. Gounden, B.-Q. Lu, I. Nape, A. Forbes *et al.*, Revealing the topological nature of entangled orbital angular momentum states of light, *Nature Communications* **16**, 11095 (2025), [nature.com/articles/s41467-025-66066-3](#).
- [35] C. He, X. Li, H. Wang, J. Li, Z. Zhang, S. Liu, Q. Jiang, N. Zhang, Y. Wang, L. Huang, Orbital angular momentum multiplexing diffractive neural networks for high-capacity optical inference, *Nature Photonics* (2026), [doi:10.1038/s41566-026-01930-2](#).
- [36] H. H. Otto, Golden Quartic Polynomial and Moebius-Ball Electron, *Journal of Applied Mathematics and Physics* **10**, 1785–1812 (2022), [doi:10.4236/jamp.2022.105124](#).
- [37] H. H. Otto, Can Artificial Intelligence Help to Verify the Most Probable Electron Structure Model? ResearchGate preprint (March 2025), [ResearchGate](#).
- [38] M. J. J. Fleury and O. Rousselle, Critical Review of Zitterbewegung Electron Models, *Symmetry* **17**, 360 (2025), [MDPI](#).




Figure 19: The whole substrate tower on one axis. **A:** every measured observable’s fractional agreement $|\text{pred} - \text{obs}|/\text{obs}$ (log scale). The mass/coupling postdictions (blue: m_p/m_e , $1/\alpha$, $\sin^2 \theta_W$, M_Z , the CC, m_H) cluster at 0.01–0.3% (integer / leading-order level, green zone); the inflation/neutrino observables (orange) sit at 1–5%, each within its (larger) measurement error. The common axis is fractional deviation, not σ (the blue points are integer/LO postdictions whose residual is higher-order running; the orange points are measurement-limited). **B:** the dated falsification frontier — each future test is one measurement that can kill the theory: m_{DM} (LZ ~ 2028), the first r hint (SO ~ 2027), the Δm^2 ratio (JUNO), Σm_ν (DESI/CMB-S4), $m_{\beta\beta}$ (nEXO/LEGEND), and the headline $r = 1/300$ (LiteBIRD ~ 2035). One $q = 3$ alphabet, ~ 20 numbers, all consistent — and decided by the mid-2030s.




Figure 20: The kill-shot dashboard: computed discovery dates. Forecast $\sigma(t)$ for the two clean discovery observables, each normalised to its own 3σ threshold ($\sigma_{3\sigma} = \text{pred}/3$), declining as the experiments improve. Where a curve crosses the green dashed line the substrate prediction is a 3σ detection: $\Sigma m_\nu = 58 \text{ meV} \sim 2028$ (DESI+CMB-S4) and $r = 1/300 \sim 2032$ (SO→CMB-S4→LiteBIRD), with 5σ (red dotted) by ~ 2030 and ~ 2034 . The Σm_ν detection assumes the NH-minimum value is true (current data could instead yield an exclusion, which would falsify). Precision/coverage kill-shots (not shown): JUNO pins $\Delta m_{31}^2/\Delta m_{21}^2 = 33$ by ~ 2030 ; LZ covers $m_{\text{DM}} = 22.8 \text{ GeV}$ by ~ 2028 .